

Practical Messaging

- A 101 guide to messaging
- Ian Cooper
- X and Hachyderm: ICooper

Who are you?

www.linkedin.com/in/ian-cooper-2b059b

I am a polyglot coding architect with over 20 years of experience delivering solutions in government, healthcare, and finance and ecommerce. During that time I have worked for the DTI, Reuters, Sungard, Misys, Beazley, Huddle and Just Eat Takeaway delivering everything from bespoke enterprise solutions, 'shrink-wrapped' products for thousands of customers, to SaaS applications for hundreds of thousands of customers.

I am an experienced systems architect with a strong knowledge of OO, TDD/BDD, DDD, EDA, CQRS/ES, REST, Messaging, Design Patterns, Architectural Styles, ATAM, and Agile Engineering Practices

I am frequent contributor to OSS, and I am the owner of: <https://github.com/BrighterCommand>. I speak regularly at user groups and conferences around the world on architecture and software craftsmanship. I run public workshops teaching messaging, event-driven and reactive architectures.

I have a strong background in C#. I spent years in the C++ trenches. I dabble in Go, Java, JavaScript and Python.

Day One Messaging



Distribution



Integration Styles



Messaging Patterns



Queues and Streams



Managing Asynchronous Architectures

Day Two Conversations



Conversations



Flow



Process Automation



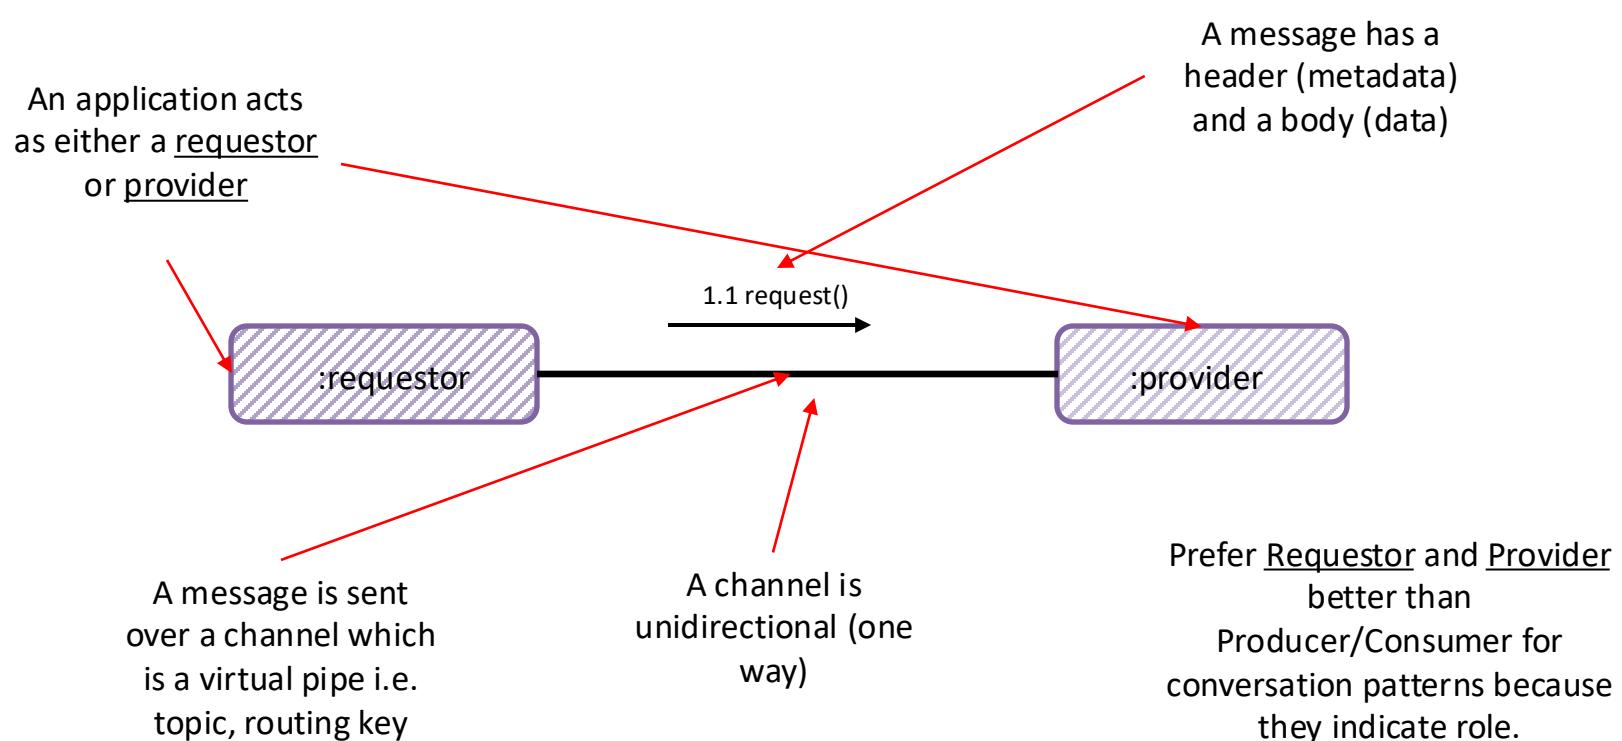
Orchestration and Choreography

Day Two

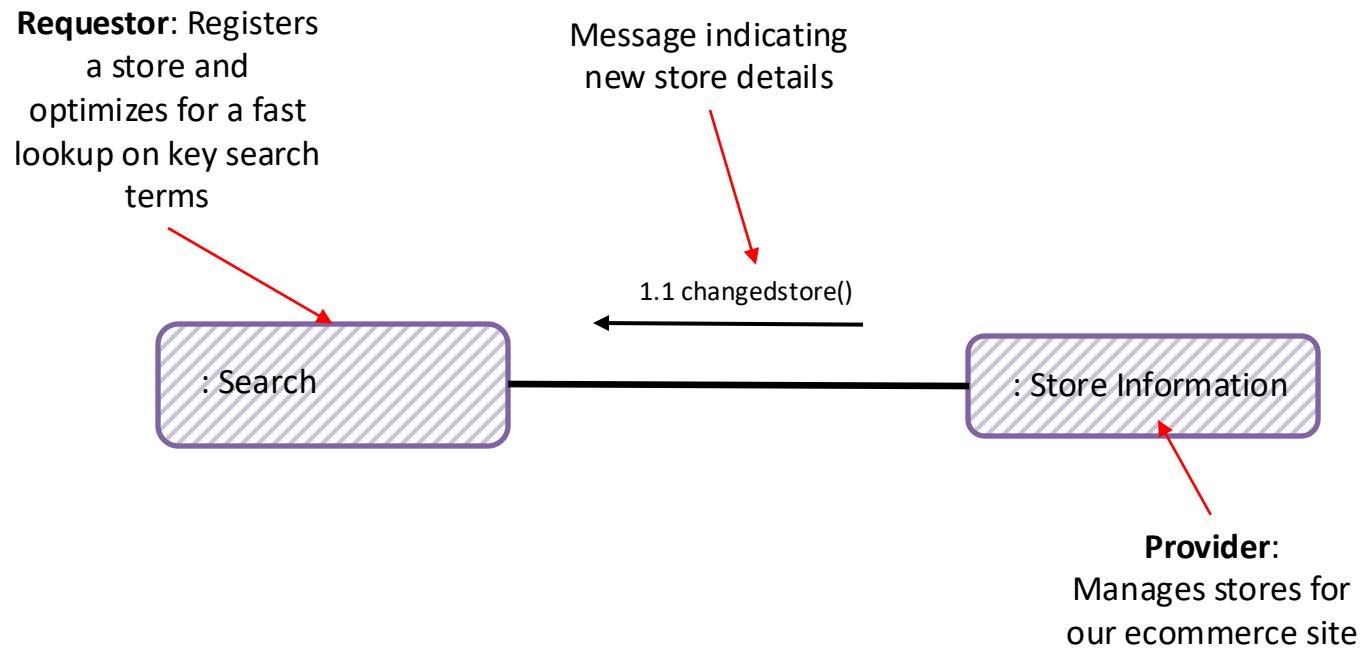
Conversations

In & Out

Messaging Participants



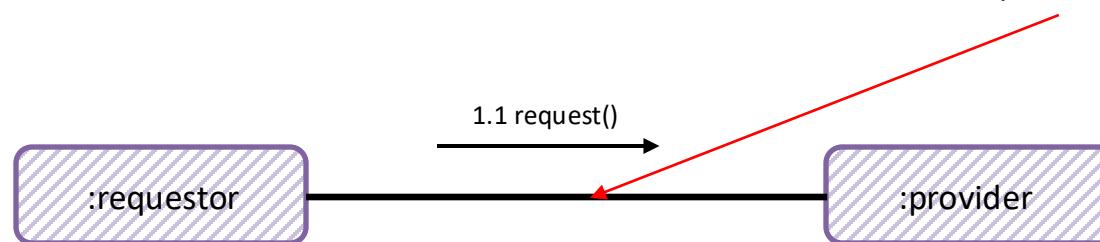
Messaging Participants



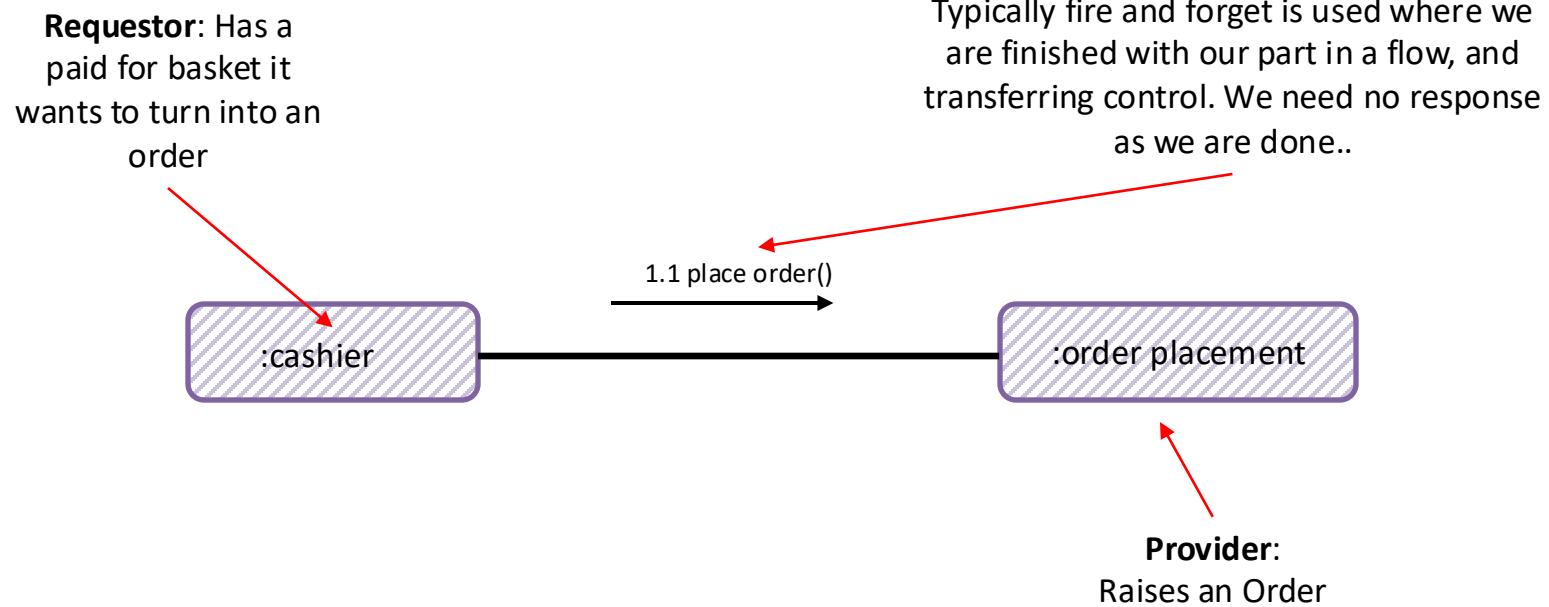
In-Only (Fire and Forget)

Typically we call this **fire-and-forget**.

Under a In-Only pattern the requestor sends a request to the provider, but does not seek an acknowledgment of completion of the requested operation

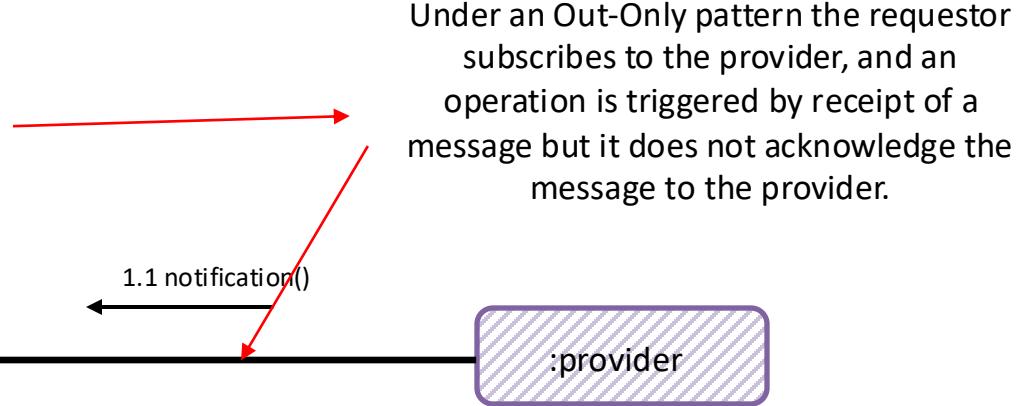


In-Only (Fire and Forget)

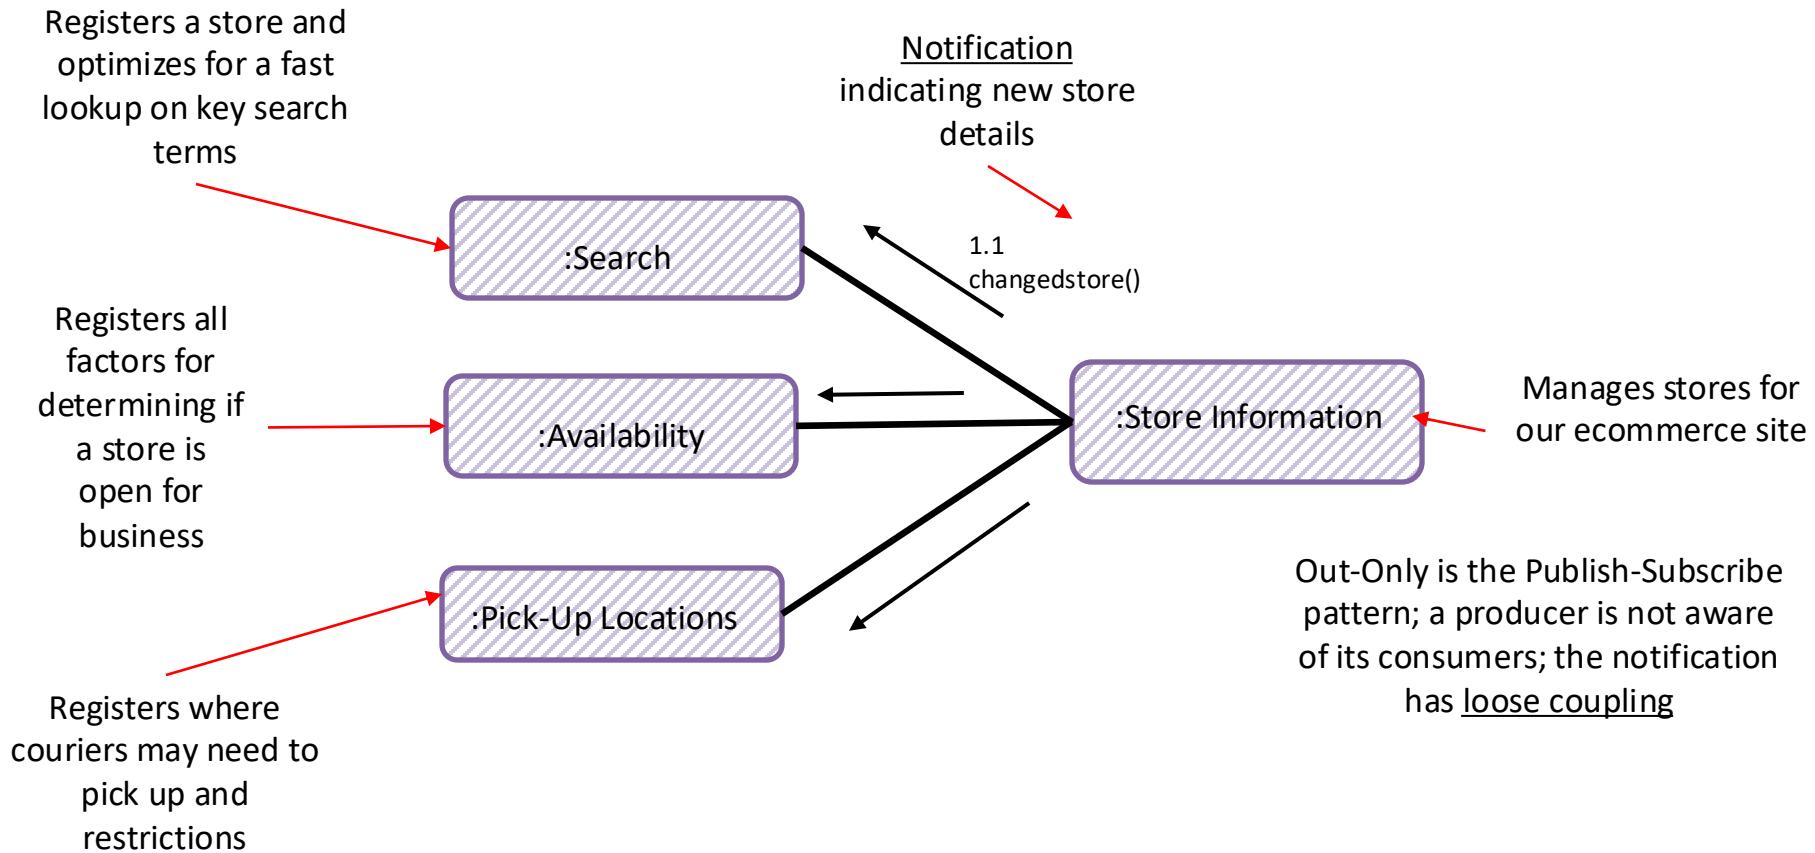


Out-Only (Notification)

Typically we call this a **notification**.

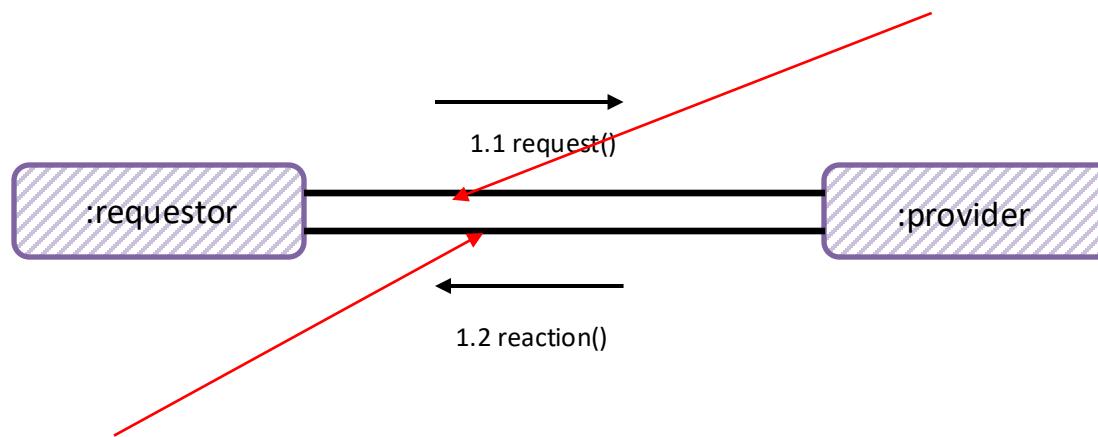


Out-Only (Publish-Subscribe)



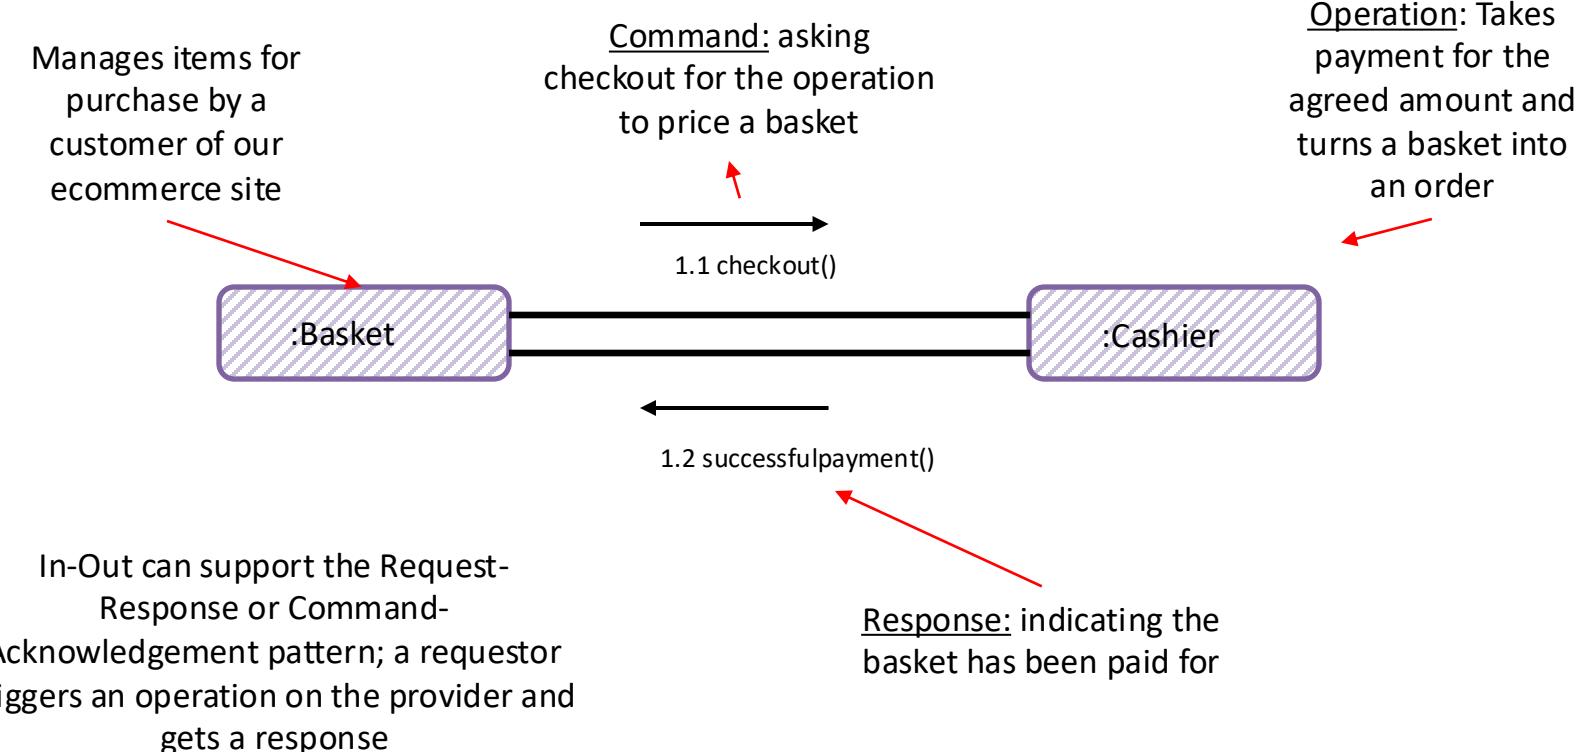
In-Out (Request-Reaction)

Under an In-Out pattern the provider receives a request from the requestor on one channel, and returns a response from the operation triggered by that message on a separate channel.



When using an In-Out pattern the provider sets a **correlation id** (conversation id) in the header of the request and the provider returns that id in the header of the response so that we can correlate replies sent over a separate channel.

Request-Reaction (Command-Acknowledgement)



In-Out (Query-Result)

Manages items for purchase by a customer of our ecommerce site



Query asking for delivery fees for this basket

1.1 getdeliveryfees()



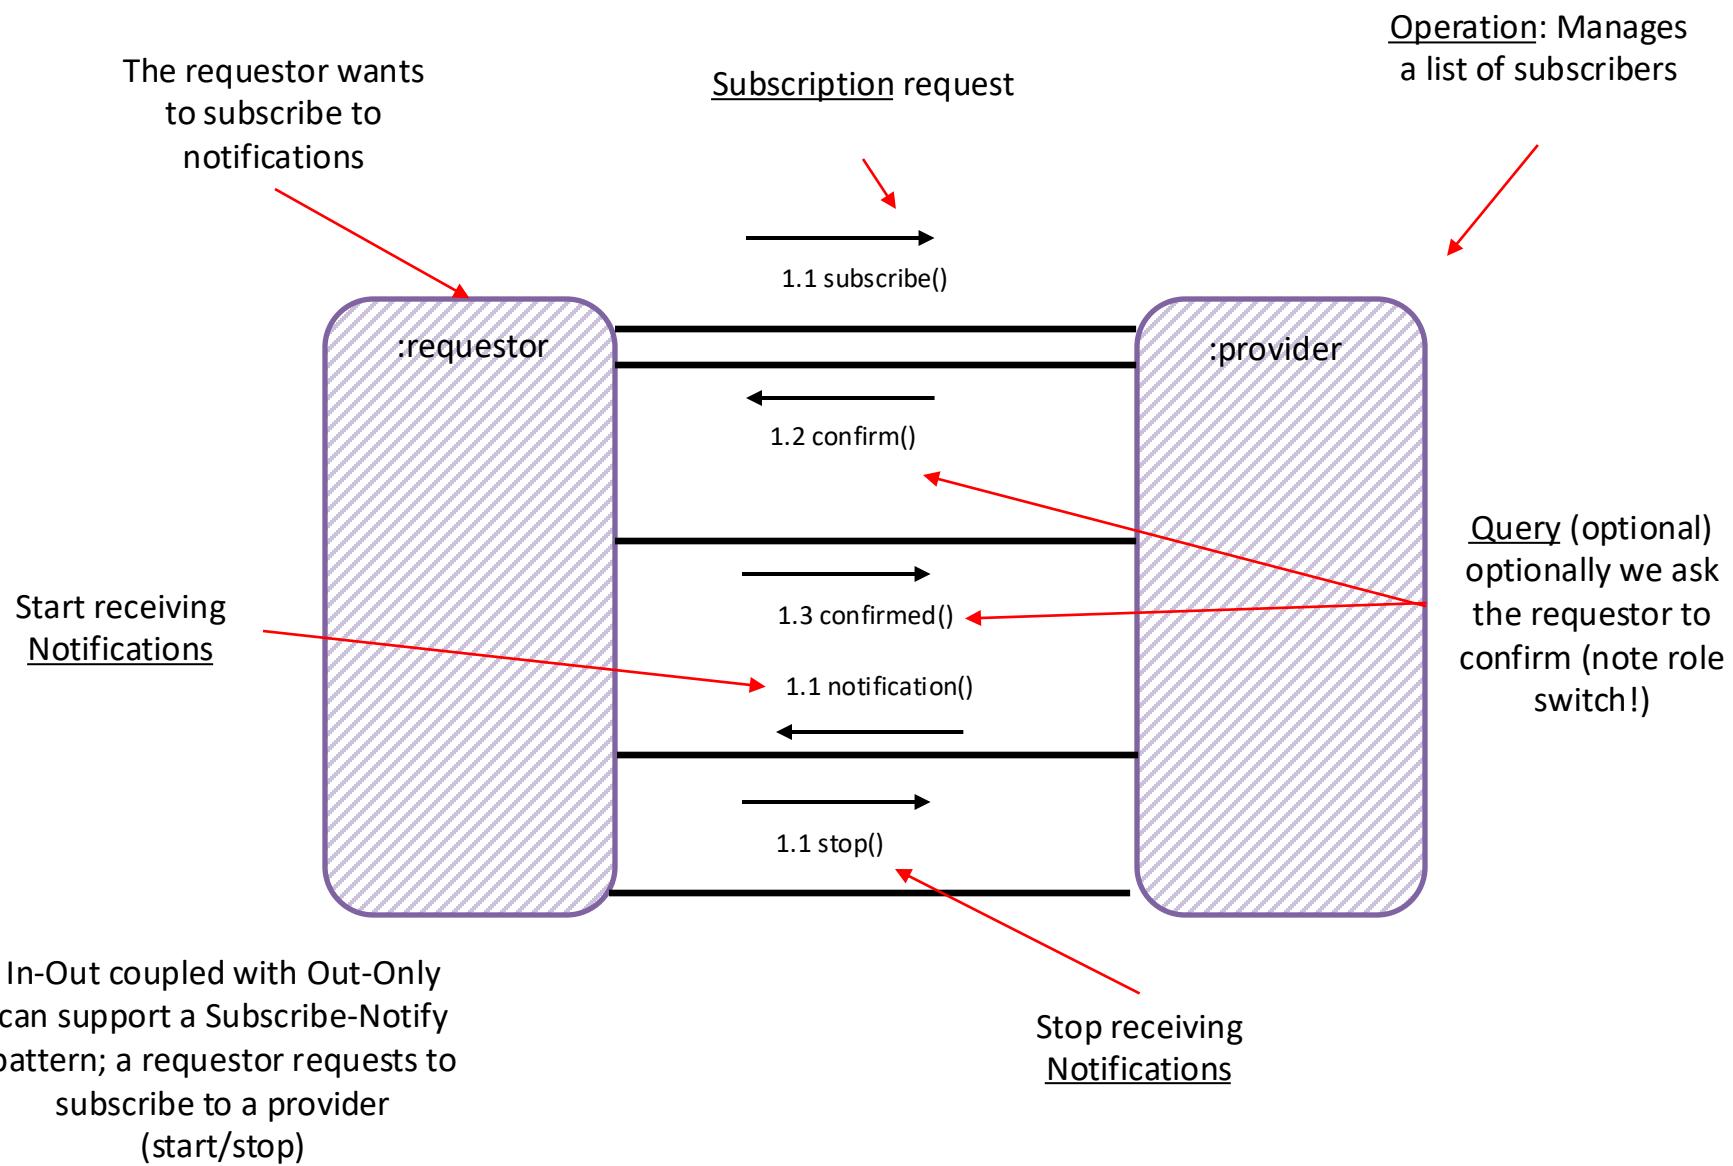
Operation:
Determines the delivery fee based on store and items in the basket

1.2 deliverfee()

In-Out can support the Query-Result pattern; a requestor triggers a query on the provider and gets a result

Result indicating the fee for the current state of the basket

In-Out (Subscribe-Notify)

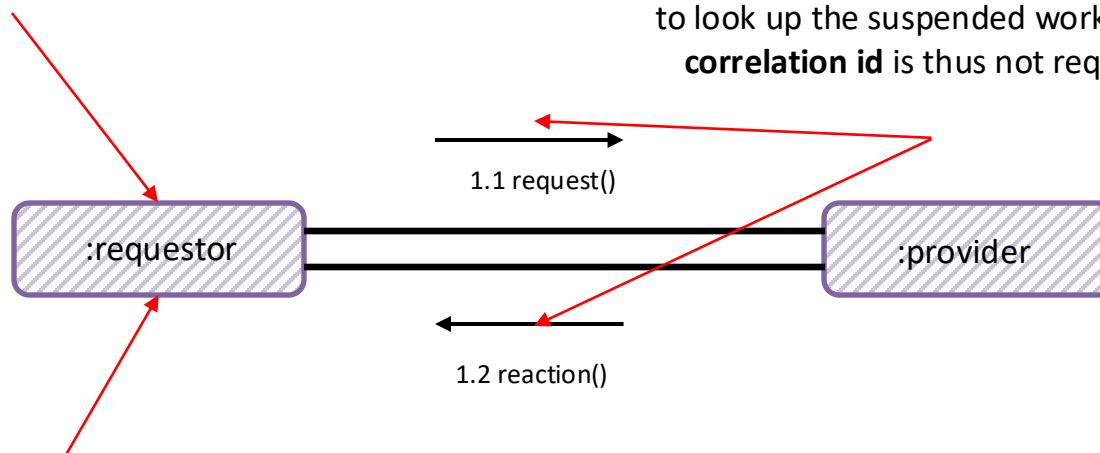


In-Out coupled with Out-Only can support a Subscribe-Notify pattern; a requestor requests to subscribe to a provider (start/stop)

Blocking In-Out (Request-Reply)

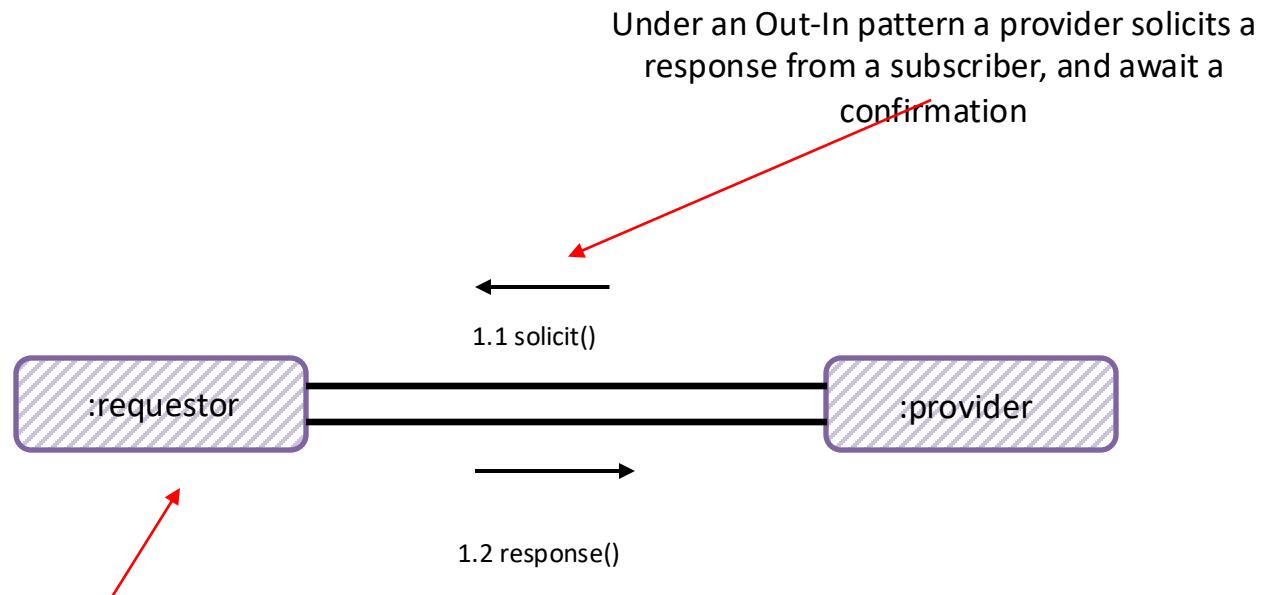
In order to resume, we need a way to identify the suspended workflow.

The requestor provides a unique **Reply To** channel (in the request message header/metadata) and the provider uses that channel for the response, allowing the sender to look up the suspended workflow. A **correlation id** is thus not required



The requestor may block on the Reply To channel whilst awaiting the reply.

Out-In (Solicit-Response)



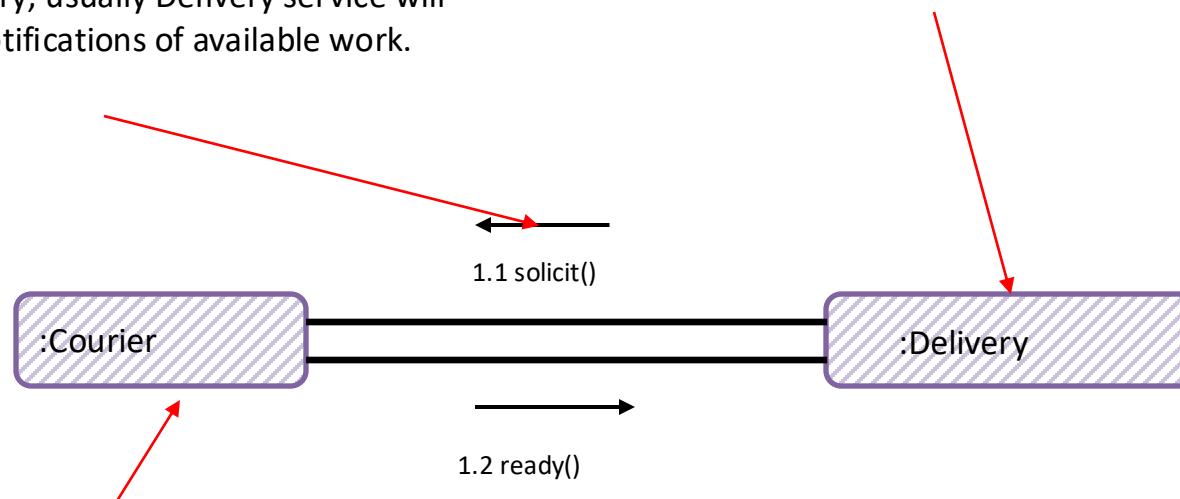
Under an Out-In pattern the subscriber confirms receipt of the provider's solicitation

Under an Out-In pattern a provider solicits a response from a subscriber, and await a confirmation

Out-In (Solicit-Response)

Query to see if a courier is available to receive orders for delivery; usually Delivery service will follow with notifications of available work.

Delivery service assigns delivery request to drivers



Courier offers jobs depending on location and busyness

In-Only (fire and forget)

Messaging

Has Intent

Request An Answer (Query)

Transfer of Control

(Command)

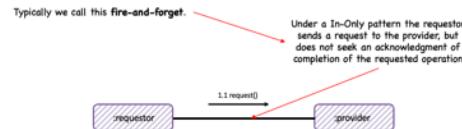
Transfer of Value

Part of a Workflow

Part of a Conversation

Concerned with the Future

In-Only (Fire and Forget)

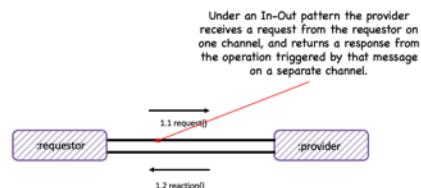


Ian Cooper

25

In-Out (request-reaction)

In-Out (Request-Reaction)



Ian Cooper

29

Eventing

Provides Facts

Things you Report On

No Expectations

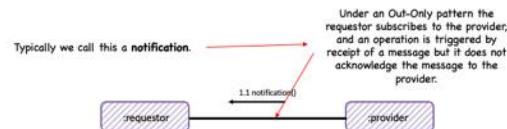
History

Context

Concerned with the Past

Out-Only (notification)

Out-Only (Notification)

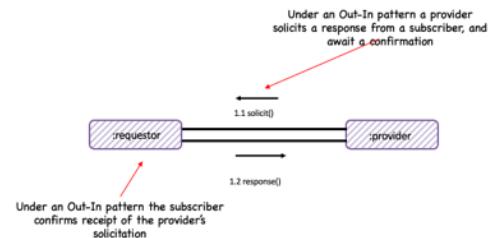


Ian Cooper

27

Out-In (solicit-response)

Out-In (Solicit-Response)

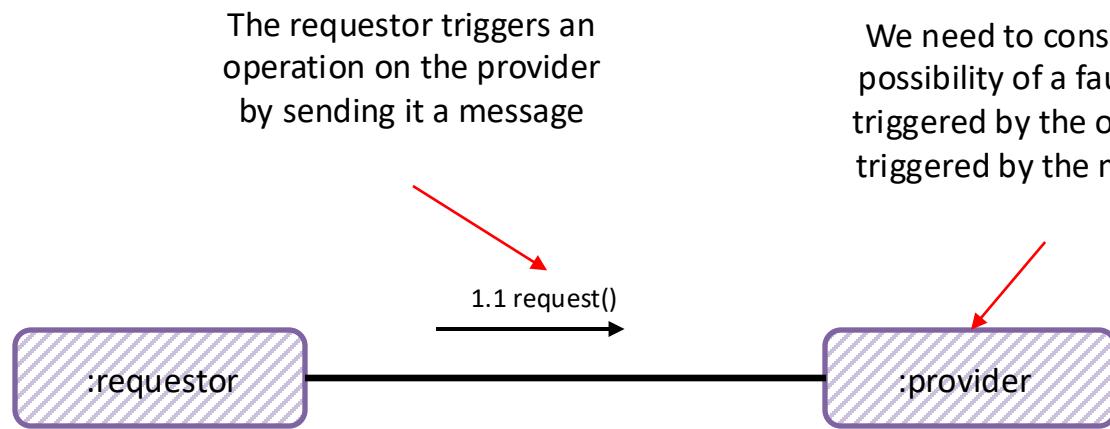


Ian Cooper

34

Repair and Clarification

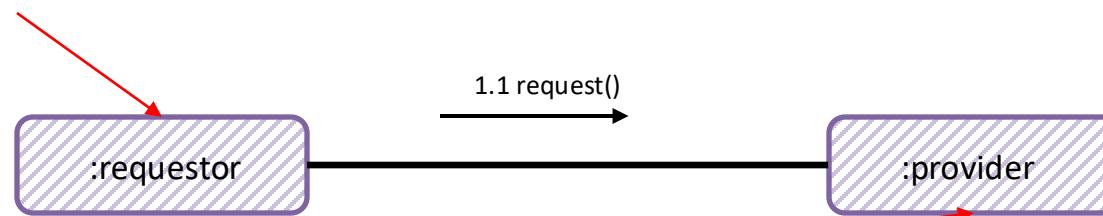
Repair and Clarification



What guarantees does the provider make to the requester about communicating any faults that are triggered?

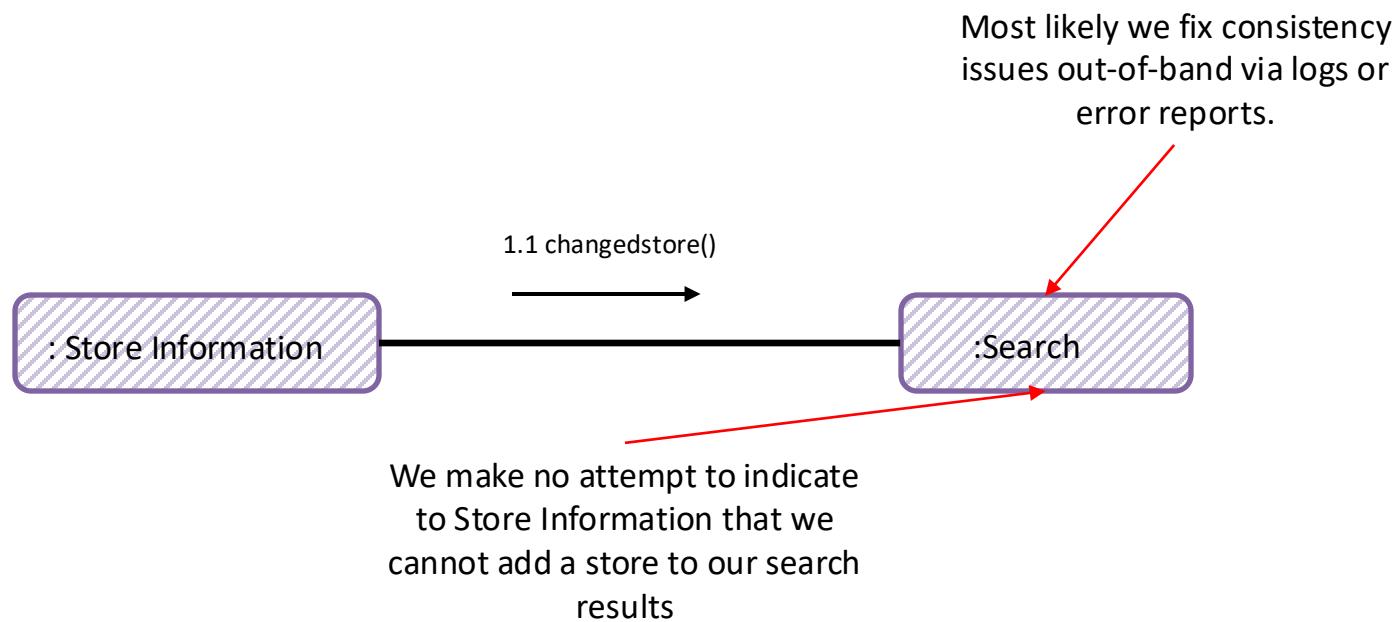
No Fault

From the requestor's perspective, faults are an application issue for the provider, and not its concern



Under a No Fault pattern the provider makes no attempt to communicate triggered faults from the operation to the requestor

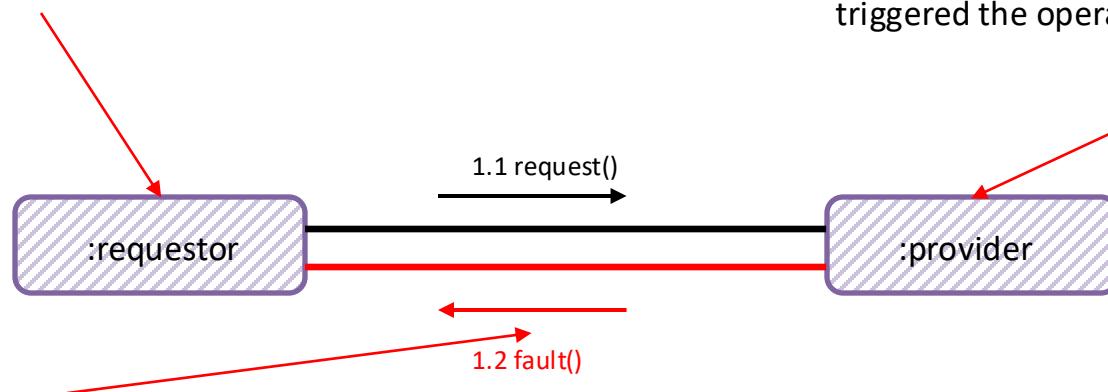
No Fault



Message Triggers Fault (Robust In-Only)

Assumption here is that the requestor can take action on receipt of the fault; but does not normally take a response.

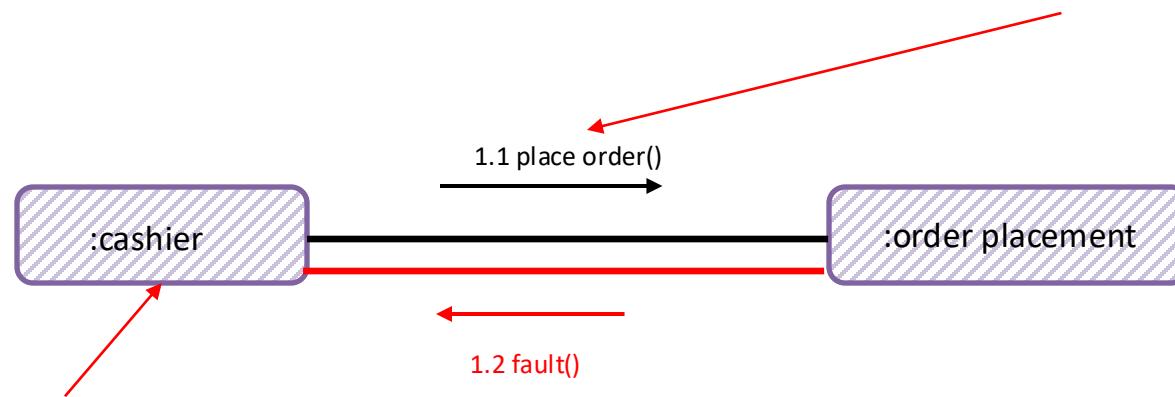
Under Message Triggers Fault pattern a **provider** propagates triggered faults from the operation back to the **requestor** that triggered the operation via a message



The message must have the opposite direction and go back to the requestor

Message Triggers Fault

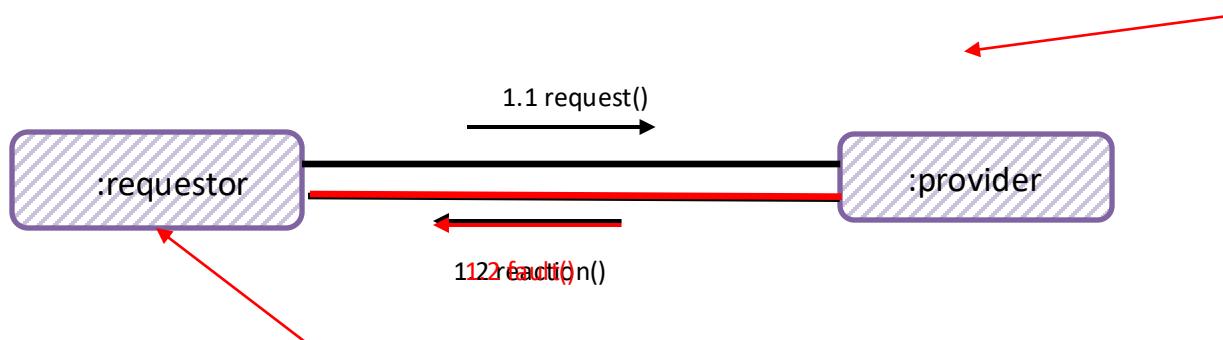
If Order Placement cannot place the order due to a fault we may raise an error



The requestor does not need to acknowledge success, but on a fault may need to take other action such as a refund

Fault Replaces Message

Under the Fault Replaces Message pattern a provider propagates faults triggered by an operation by switching to a fault flow, replacing any message subsequent to the first with a fault.

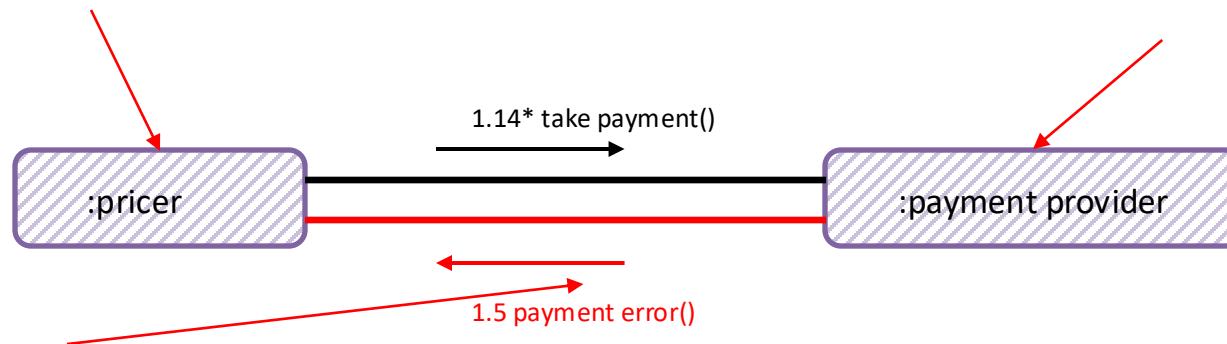


The requestor can handle the error; the error replaces the existing response

Fault Replaces Message (Robust In-Out)

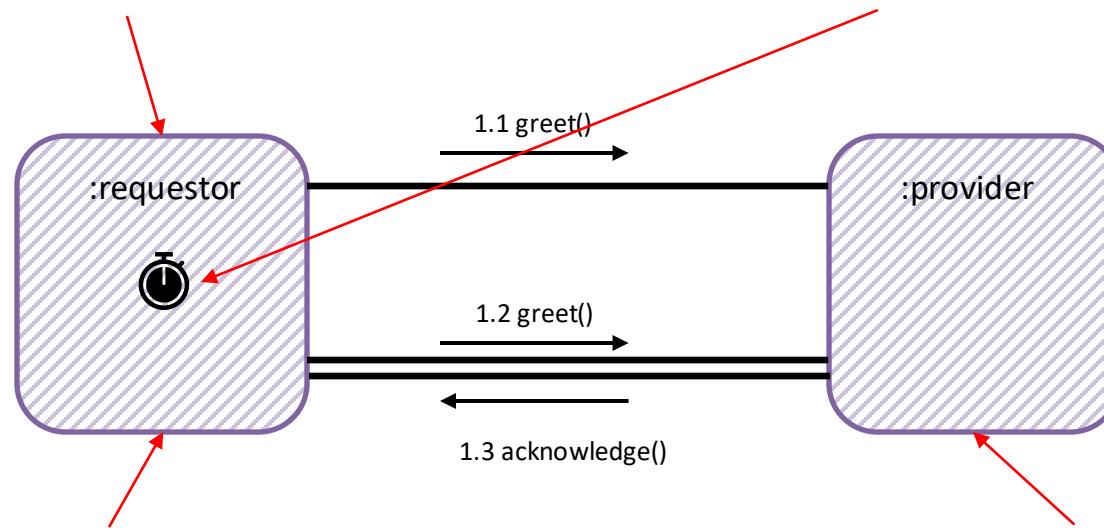
Assumption here is that the pricer can orchestrate a new flow, such as asking for an alternate payment method, or cancel the order.

Under Fault Replaces Message pattern the payment **provider** would signal errors taking a card payment back to the pricer



The message should indicate why the payment failed. This might be an issue with the payment provider but it also might be an issue like an invalid card or insufficient funds

The requestor may not receive an expected response from a provider. What can it do?



The requestor can choose to retry if it does not receive a response within that time window.

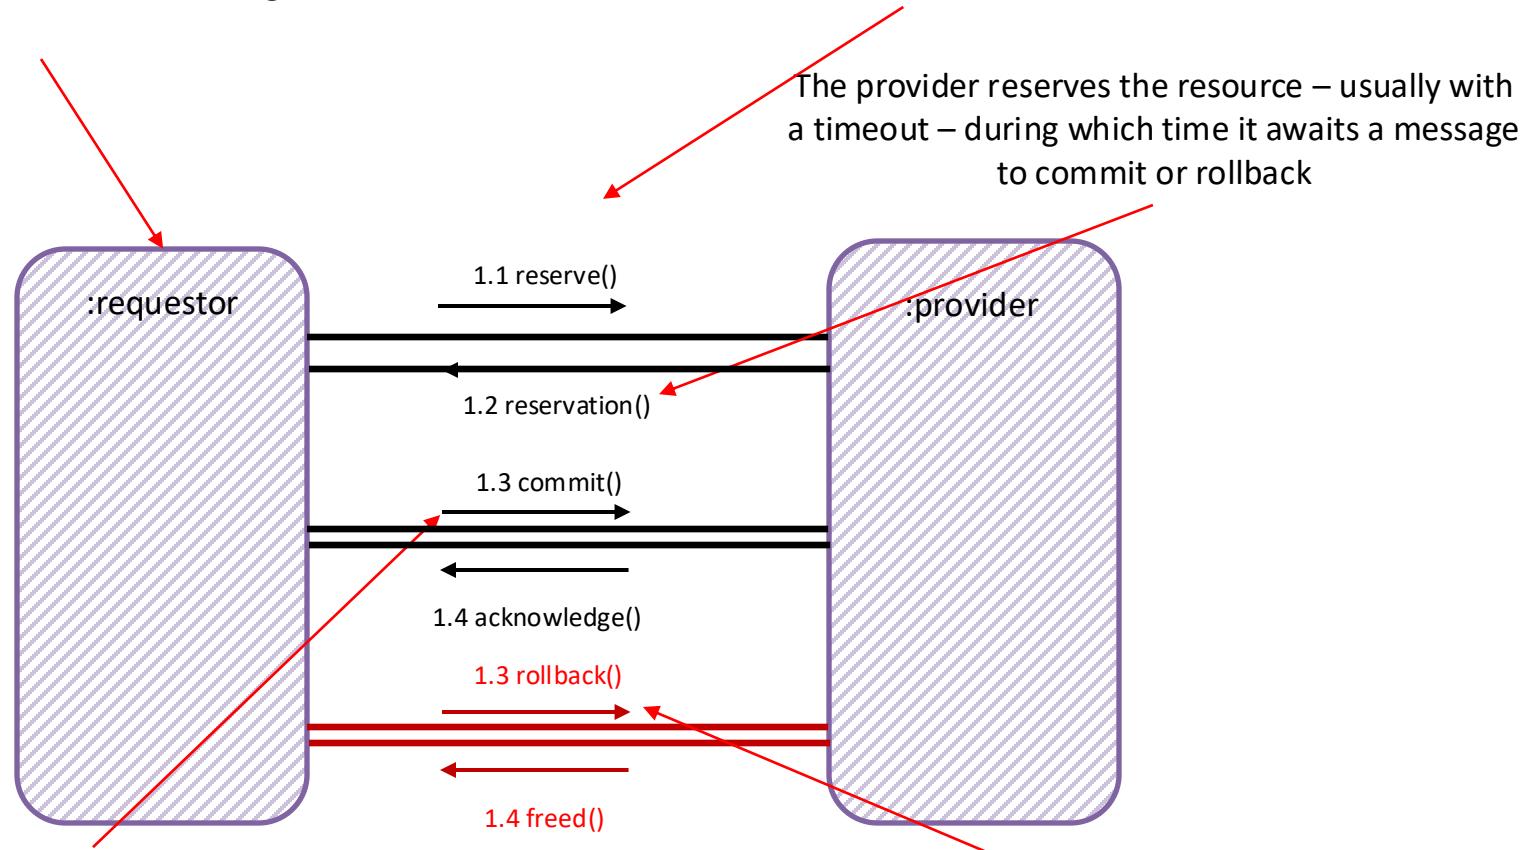
The requestor can set a timeout within which to receive a response.

Because we might send a message twice, the operation on the consumer must be idempotent, or the consumer must de-duplicate already seen messages

Tentative Operations

The requestor may not know if the provider will be able to succeed, and might not want to proceed without knowing that

The requestor asks if the operation is possible, and reserves the resources to perform it.



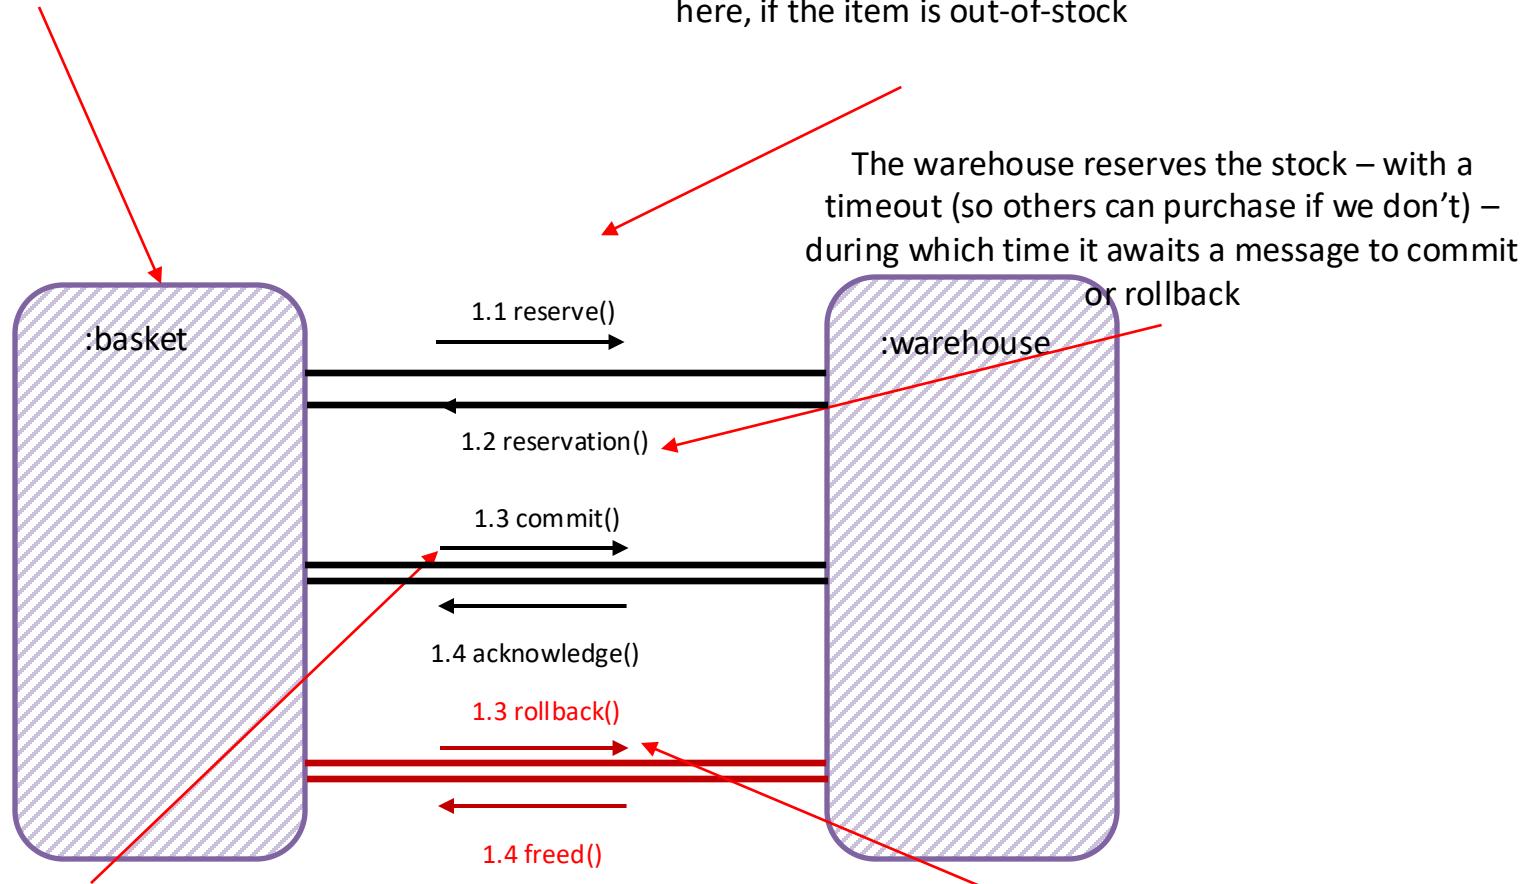
If the requestor completes its work, it then commits to ask the provider to allocate the reserved capacity.

If the requestor fails, it then rolls back to ask the provider to free the reserved capacity.

Tentative Operations

We don't want to purchase the item if the stock is not available

The basket asks the warehouse if there is stock, and if so reserves it. Fault could replace message here, if the item is out-of-stock



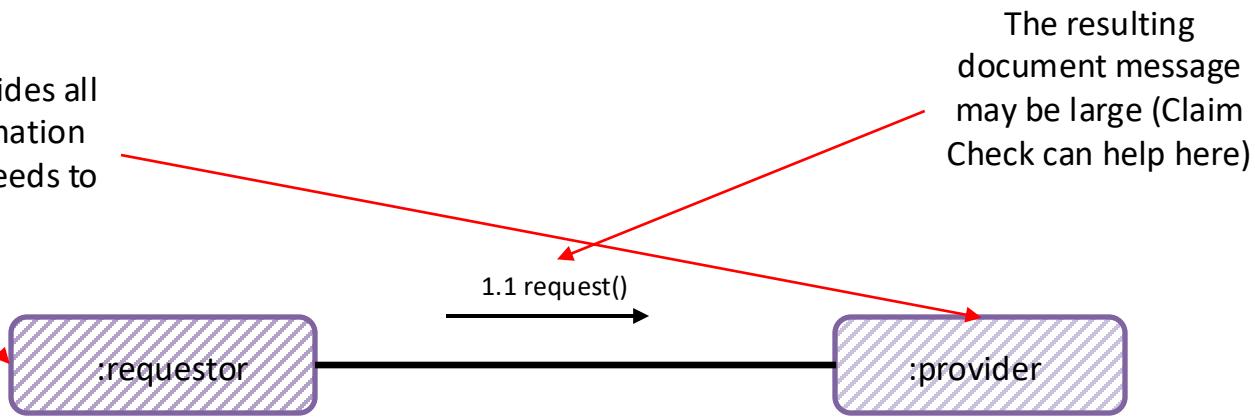
If the customer finishes shopping and pays for the basket before the time limit, then it commits to ask the warehouse to allocate the stock.

If the basket does not complete, then it rolls back to ask the warehouse to free the stock.

Fat & Skinny

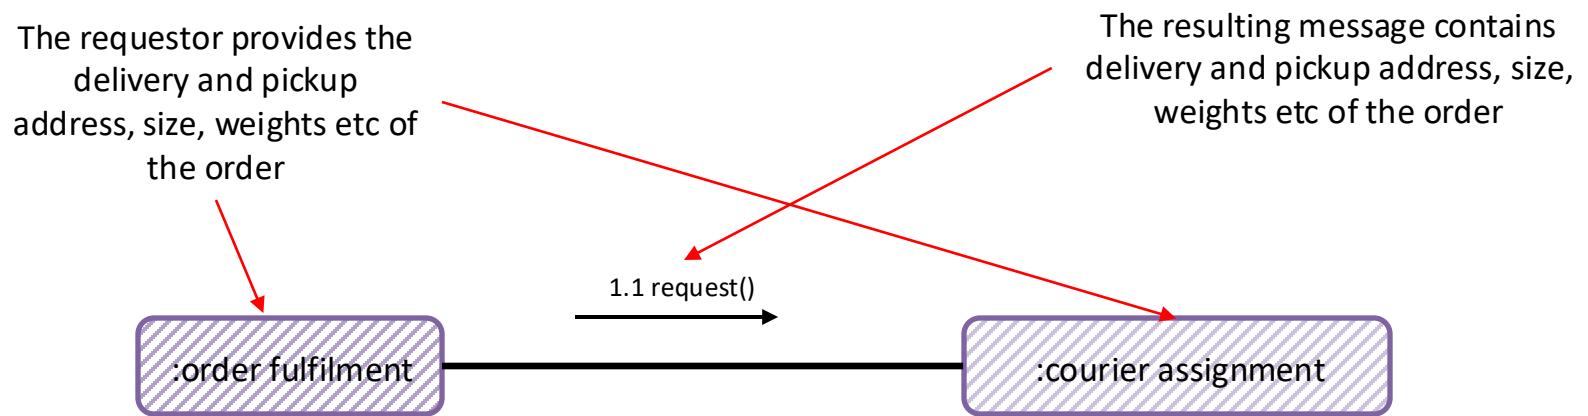
Fat Message

The requestor provides all external the information that the provider needs to act



With a **Fat Message** the requestor sends across all the external information a provider may need to perform the operation.

Fat Message Example



Fat Message, Transitive Dependencies

Purchase Order Message

OrderId	Customer First Name	Customer Last Name	Customer Post Code	Restaurant Name	Restaurant Post Code	Order Amount	Order items
12345	Jo	Doe	SW17 3NJ	Pizza 'R Us	SW17 5HK	1038p	...

This data has a lifetime of the message
i.e. the purchase order and their
schema changes if the purchase order
changes

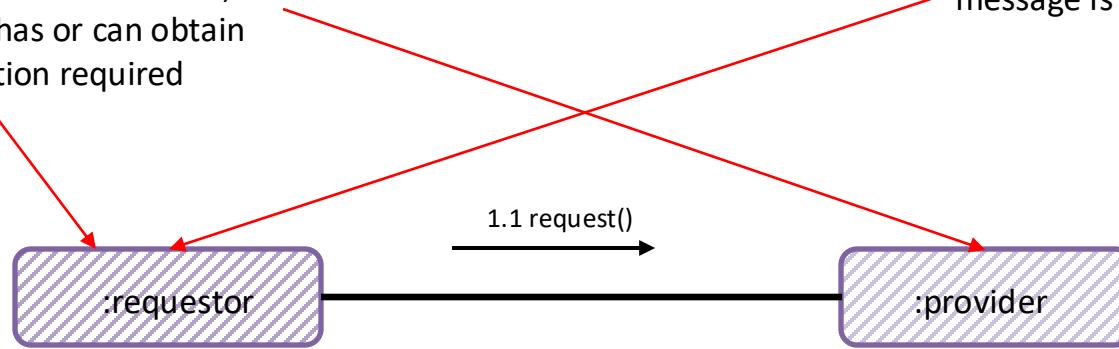
This data is a **transitive dependency**, it has
a lifetime of the Customer and their
schema changes if the Customer schema
changes – which may force us to change
the message for Customer changes

This data is a **transitive dependency**, it has
a lifetime of the Restaurant and their
schema changes if the Restaurant schema
changes – which may force us to change
the message for Restaurant changes

Skinny Message

The requestor provides only information unique to that event, assumes provider has or can obtain other information required

The resulting notification message is normally skinny

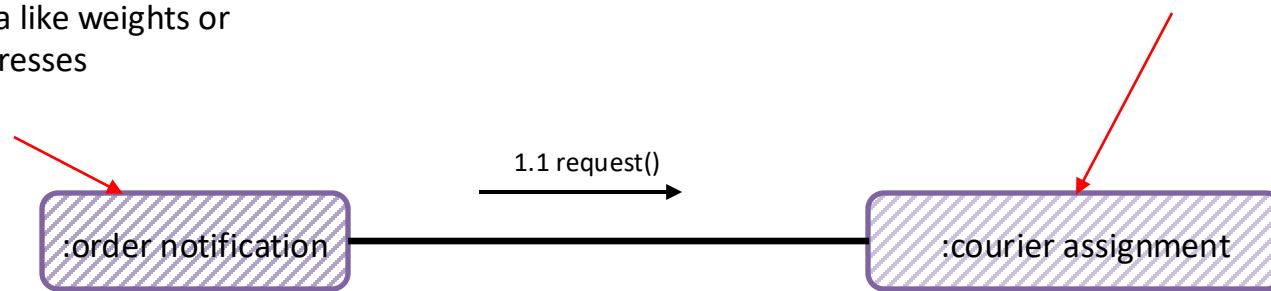


With a **Skinny Message** the requestor assumes the provider has necessary external information to perform the operation.

Skinny Request, Example

The requestor provides a notification that there is an order, but only inlines data that is unique to the order not reference data like weights or addresses

Courier assignment needs information that is not on the order: weights, addresses. It has to source that from elsewhere



Skinny Message, Normalized

Purchase Order Message

OrderId	Order Amount	Order items	CustomerId	RestaurantId
12345	1038p	...		

Customer

Customer First Name	Customer Last Name	Customer Post Code
Jo	Doe	SW17 3NJ

We use an Id where the data does not share the lifetime of the message; we assume the requestor obtains the data out-of-band

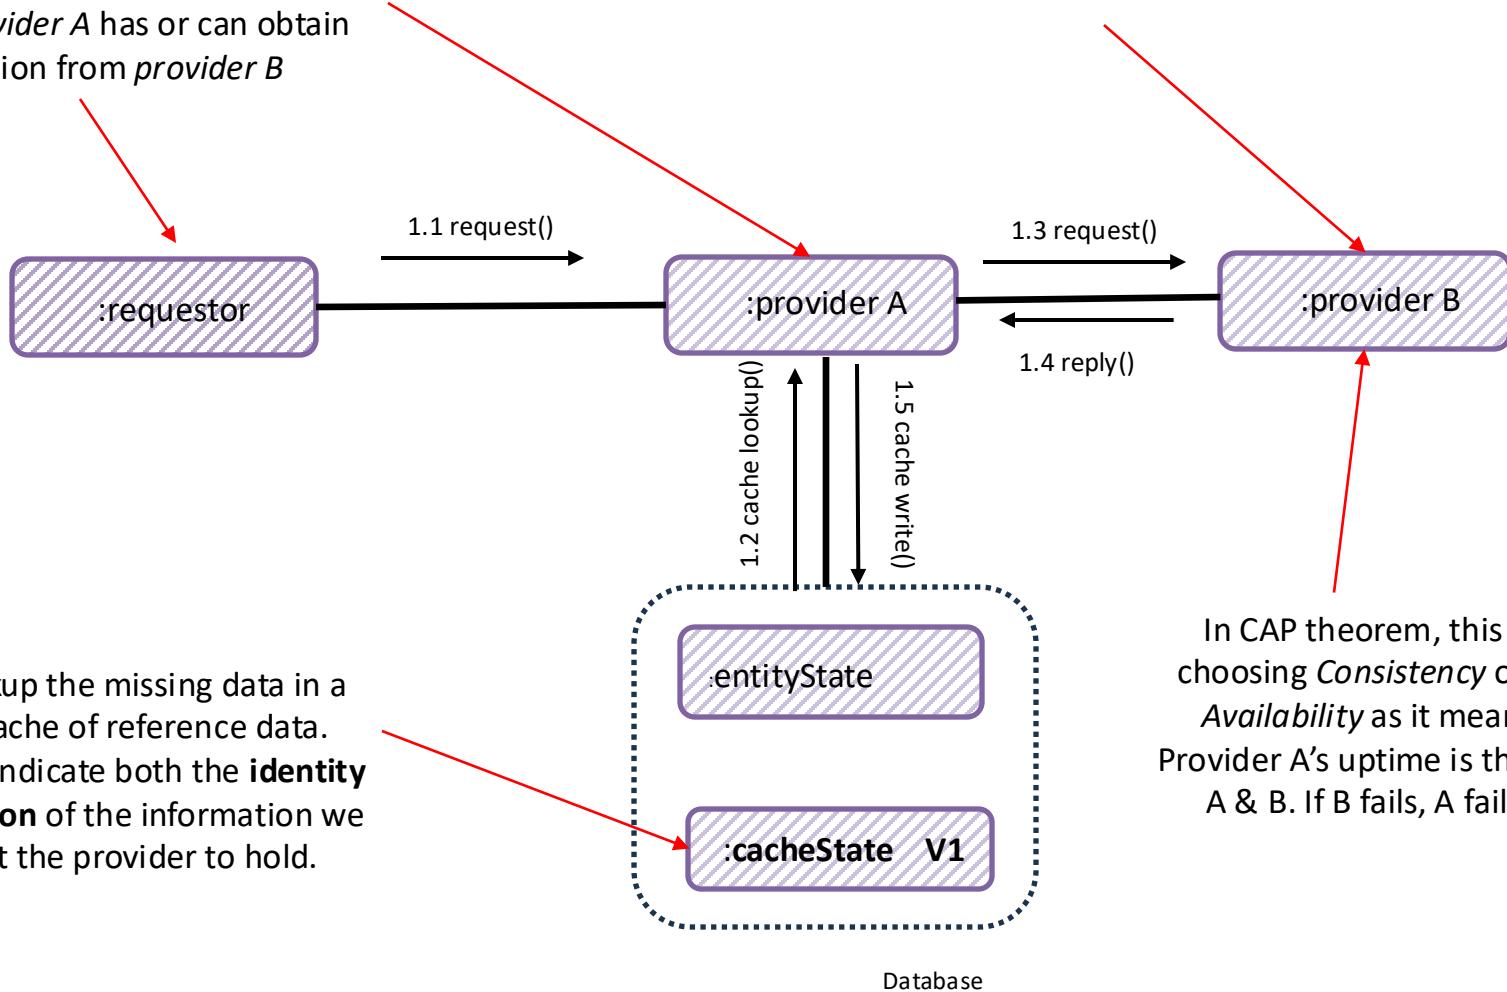
We need to lookup this id with other providers

Restaurant

Restaurant Name	Restaurant Post Code
Pizza 'R Us	SW17 5HK

Skinny Message via REST/RPC

The requestor provides only information unique to that event, assumes *provider A* has or can obtain information from *provider B*



We lookup the missing data in a local cache of reference data.
We may indicate both the **identity** and **version** of the information we expect the provider to hold.

Skinny Message using Reference Data

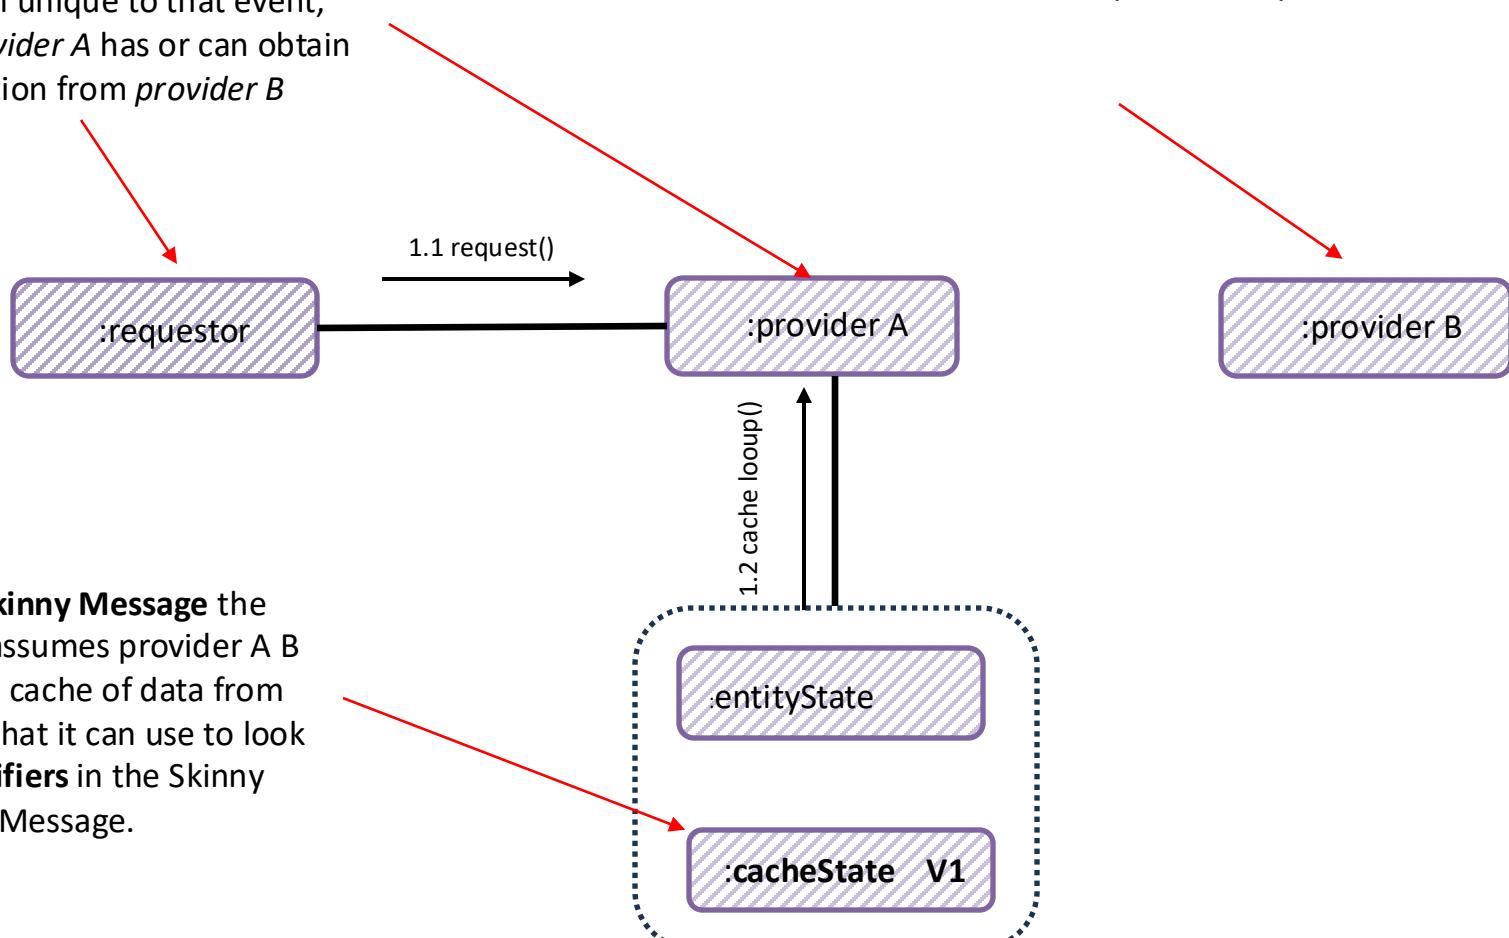
Data that leaves Provider B via an API is **Reference Data**.

It is:

- *Immutable*
- *Versioned*
- *Stale*

The requestor provides only information unique to that event, assumes *provider A* has or can obtain information from *provider B*

It can be cached locally to provider A to prevent the need for frequent lookups

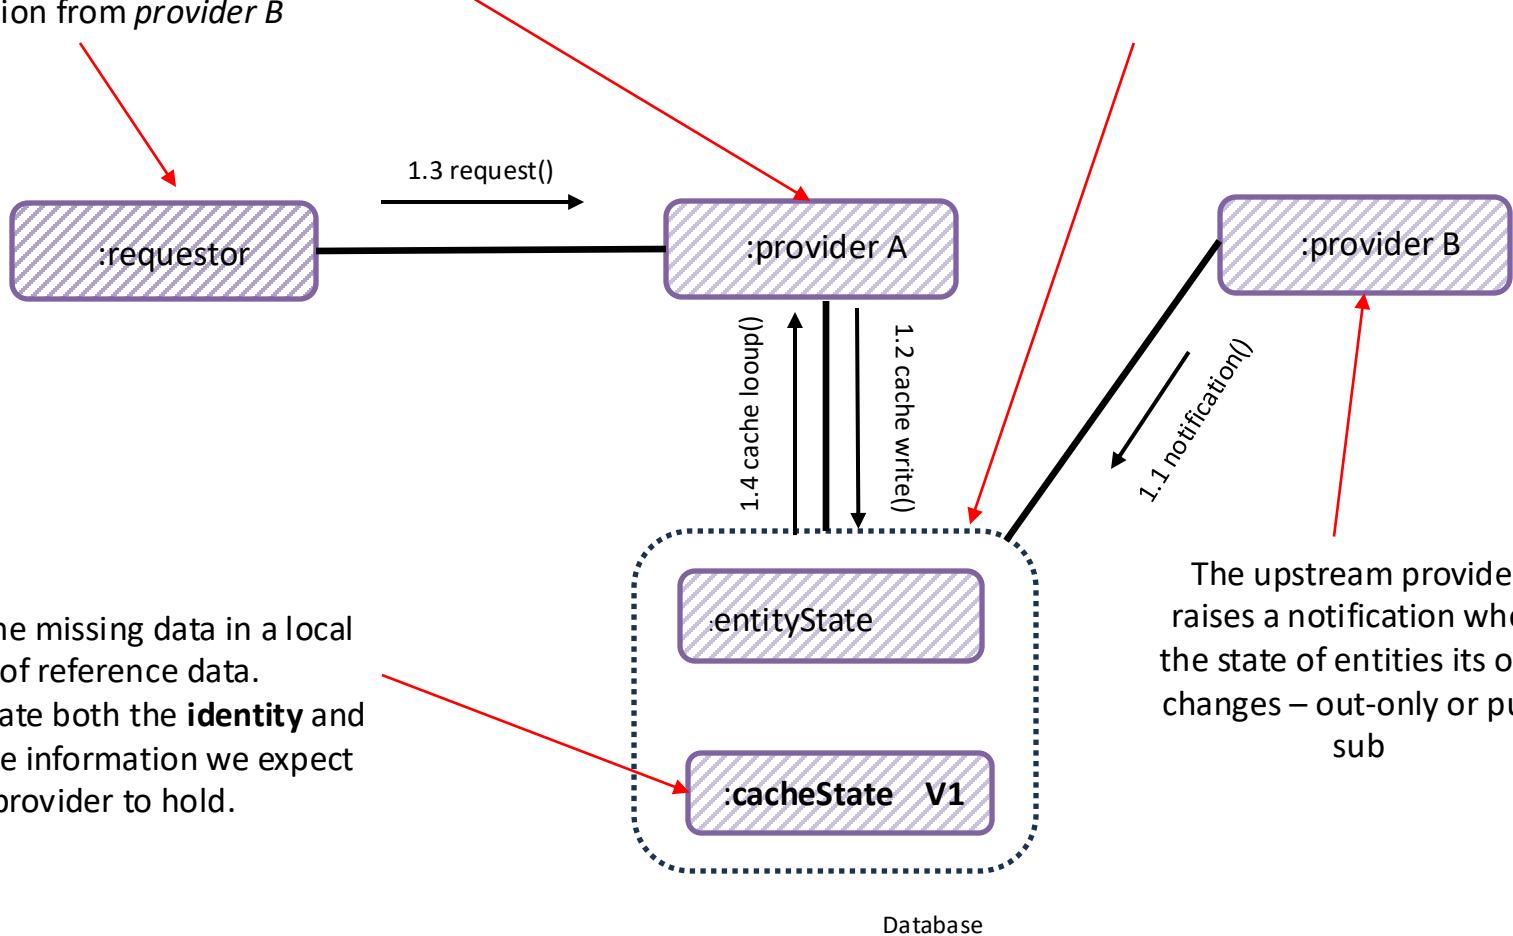


With a **Skinny Message** the requestor assumes provider A B has a local cache of data from provider B that it can use to look up **identifiers** in the Skinny Message.

Skinny Message, Reference Data via ECST (Event Carried State Transfer)

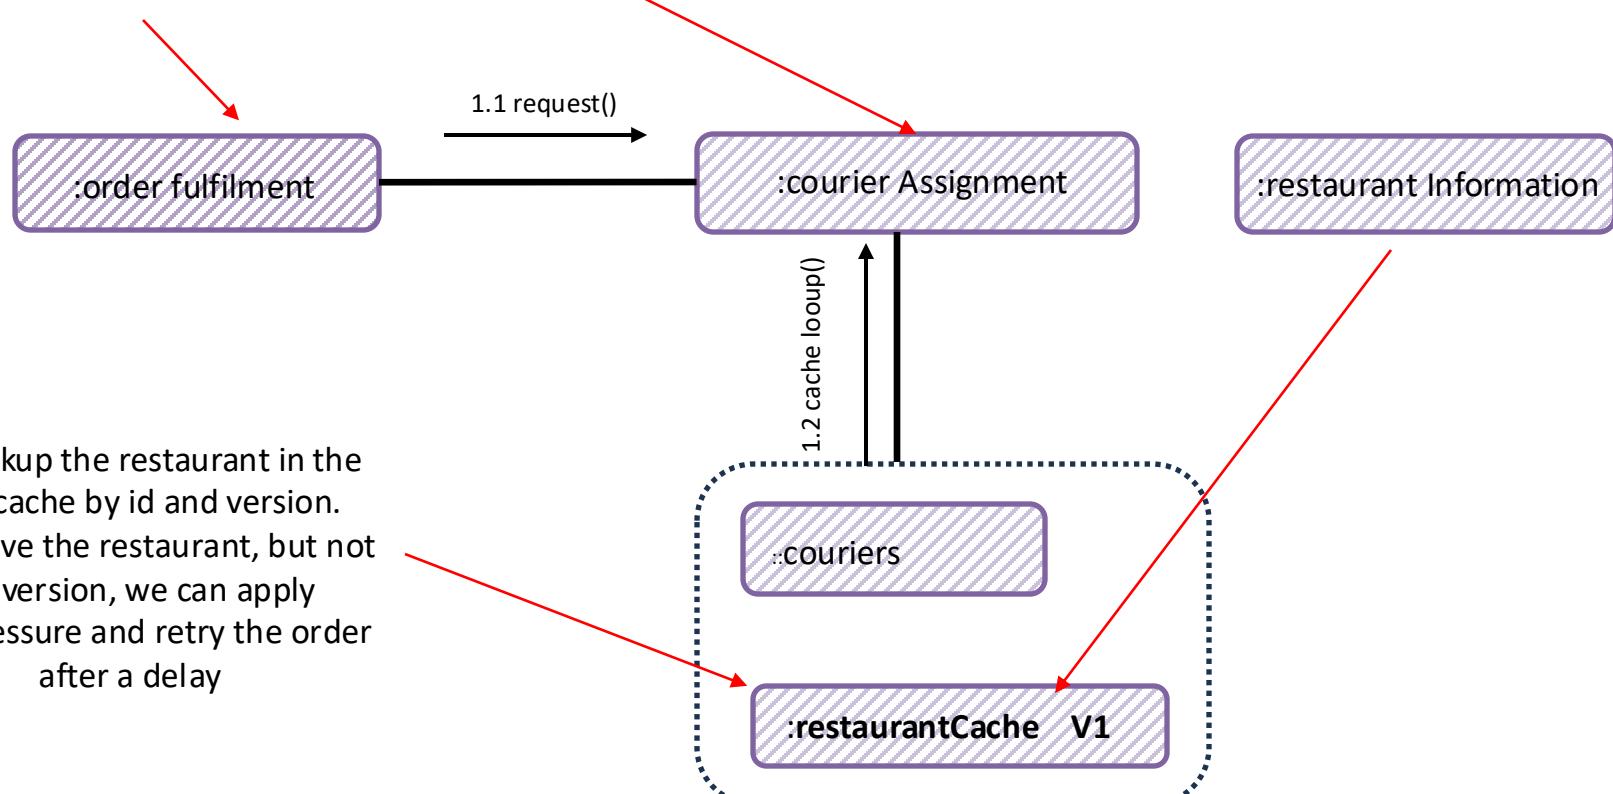
The requestor provides only information unique to that event, assumes *provider A* has or can obtain information from *provider B*

The downstream provider subscribes to these notifications and writes these to the local cache.
In CAP Theorem we choose *Availability* over *Consistency* – we choose to accept stale data over accepting the risk of failure due to a *Partition*.



Skinny Message using Reference Data

Order Fulfilment does not include address details for the restaurant for pickup, instead we assume that courier assignment has obtained it from restaurant information

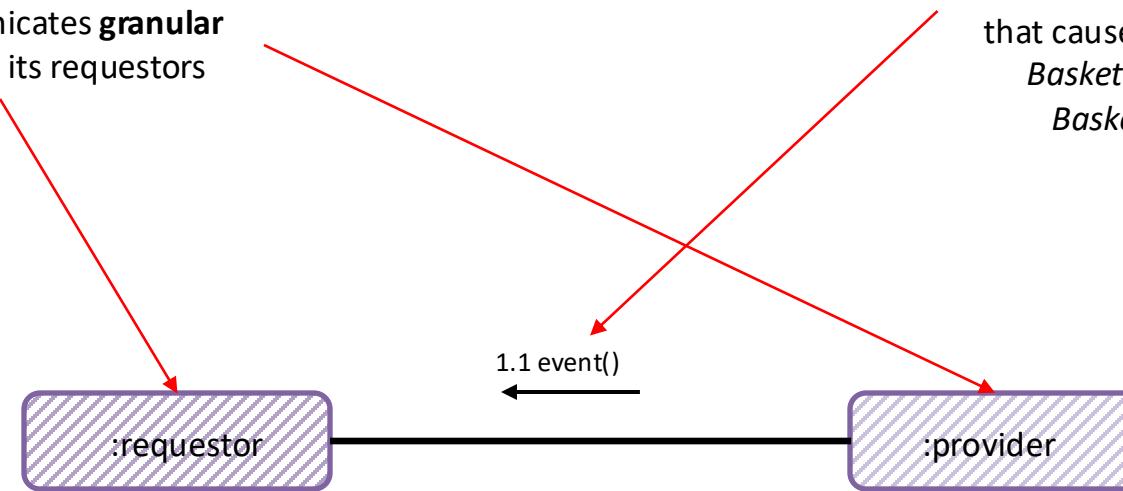


We lookup the restaurant in the local cache by id and version. If we have the restaurant, but not the version, we can apply backpressure and retry the order after a delay

Domain, Summary & Sequence

Domain or Delta Event

An approach to Pub-Sub where the provider communicates **granular** state changes to its requestors



The resulting message is usually a past participle named after the command that caused the change i.e.
BasketItemRemoved,
BasketItemAdded

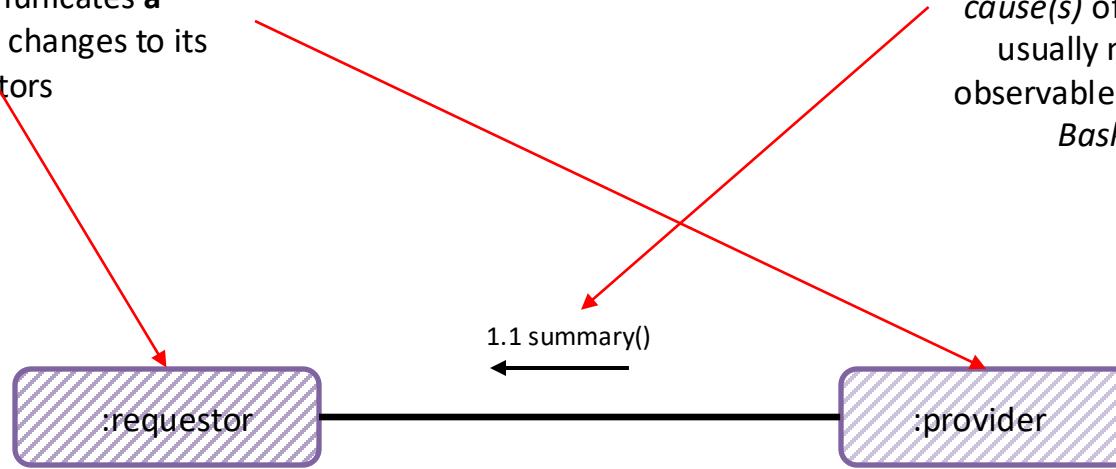
With a **Domain or Delta Event**, the provider cannot use a *DataType Channel* as the domain events for the observable must be ordered on the same channel, and have a different schema.

This puts an additional burden on the requestor to multiplex handling of the event

Summary or Snapshot Event

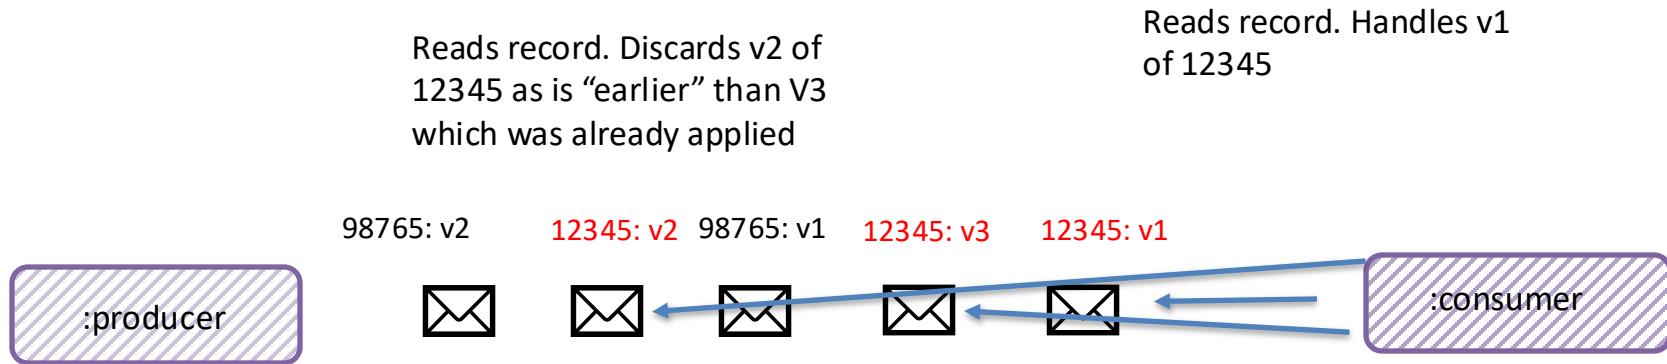
An approach to Pub-Sub where the provider communicates a **summary** of state changes to its requestors

The resulting message is **versioned** and contains *metadata describing the cause(s)* of any changes. It is usually named after the observable in the pub-sub i.e. *BasketChanged*



With a **Summary or Snapshot Event** the provider can use a *DataType Channel*, as there is just one schema used to communicate a snapshot of the observable.

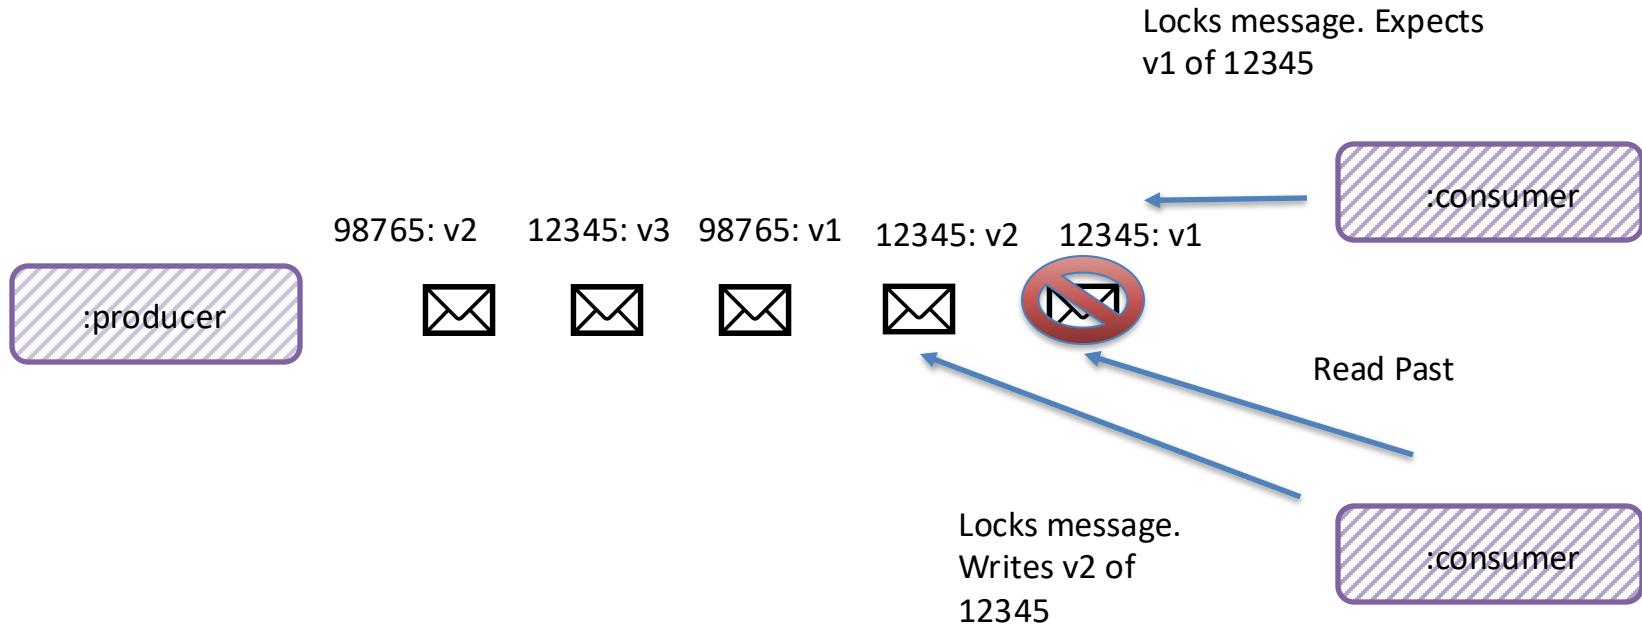
If Later, Stream



If Later allows us to use a versioned **Summary Event** to ignore ordering errors. Even on a stream, a non-blocking retry or guaranteed delivery via an outbox may result in out-of-order messages. We can also shed load

We cannot use If Later with a **Domain Event**, as they all must be applied. We can only use a blocking retry with a Domain Event and cannot shed load.

If Later, Queue

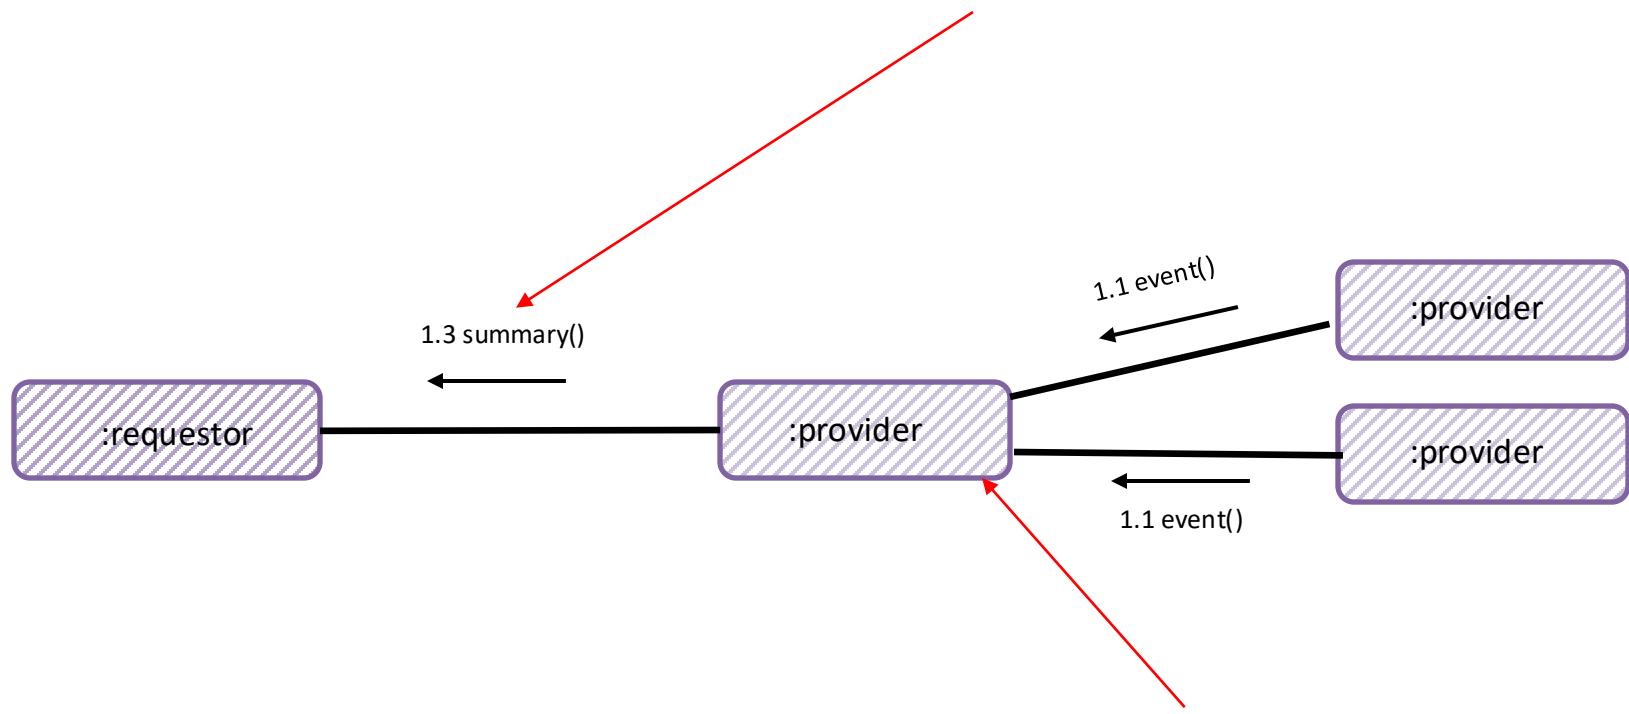


If later allows message to be processed out-of-order with a queue and competing consumers. A queue normally processes messages, and not events, so this only applies where we are using a queue.

If a message must be ordered, such as a series of commands, we must use requeue with delay, or a sequencer to re-order.

Public and Private Providers

The resulting message is **versioned** and contains *metadata describing the cause(s) of any changes.*



An approach to Pub-Sub where a **public** provider communicates with collaborators in other domains via a *summary event of domain events* from **private** providers

FLOW

Paper Workflows

“My life looked good on paper - where, in fact, almost all of it was being lived.”
- Martin Amis



Messaging/Discrete Event

Discrete, Immediately Actioned

Skinny



Series Event (Document)

Series, Supports Action

Fat





OE DIVISION
OEPI ROUTING SHEET

TITLE: EID-E15 (23 July 84)

REQD legend: I = Information
A = Action
R = Retain

Upon Completion, Return to **ETI Barmag**

**Destroy
File**

DPS/OE Form No. XXXXXXXXXX



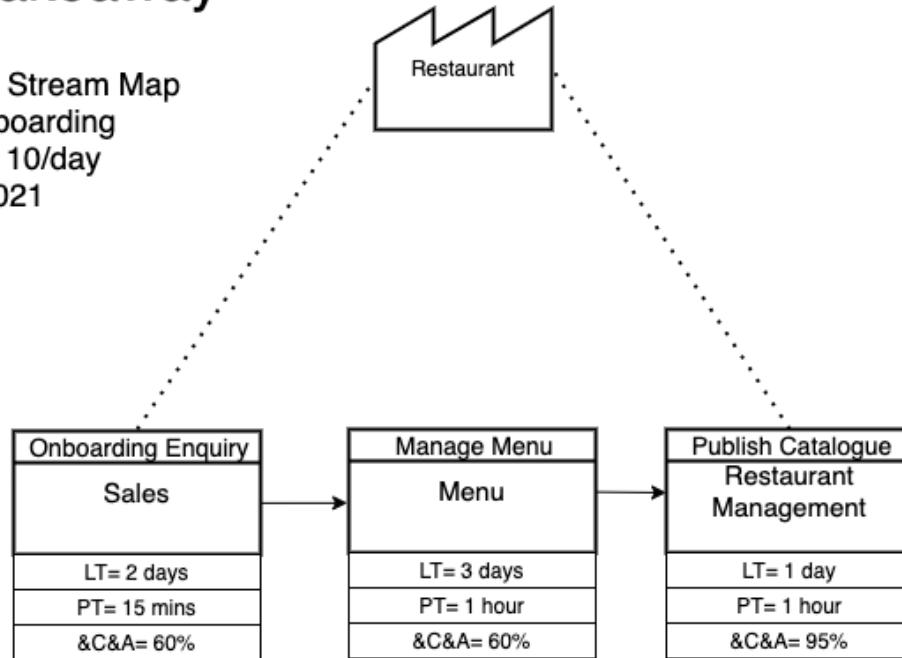
What is a microservice?

SOA is focused on business *processes*. These *processes* are performed in different steps (also called *activities* or *tasks*) on different systems. The primary goal of a **service** is to represent a “natural” step of business functionality. That is, according to the domain for which it’s provided, *a service should represent a self-contained functionality that corresponds to a real-world business activity*.

Josuttis, Nicolai M.. SOA in Practice: The Art of Distributed System Design .
O'Reilly Media. Kindle Edition.

Just Paper Takeaway

Current State Value Stream Map
Restaurant Onboarding
Demand Rate 10/day
29 SEP 2021



1

Restaurant
Owner



Restaurant Onboarding



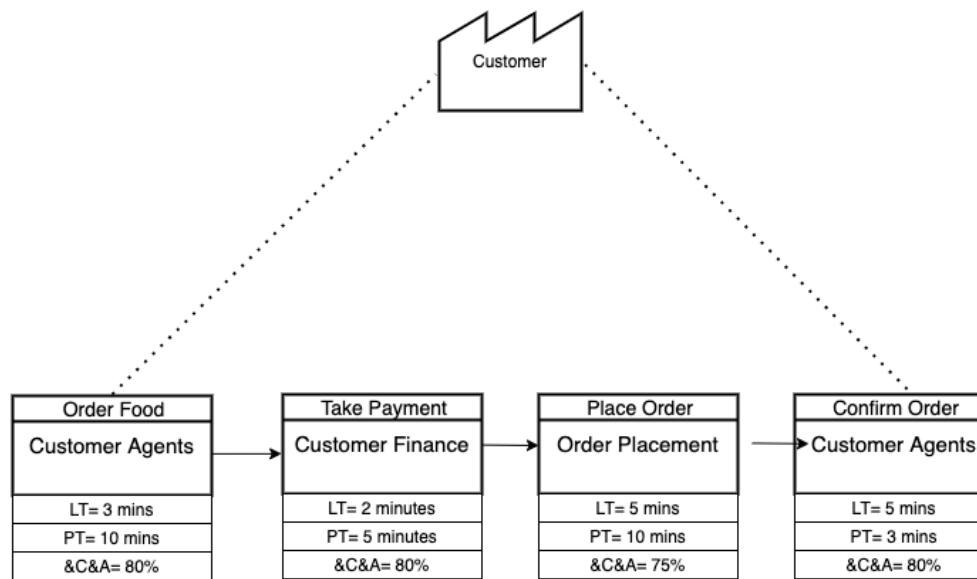
Just Paper Takeaway

Current State Value Stream Map

Order Flow

Demand Rate 10/day

29 SEP 2021



1



New Catalogue

To make a selection, we use the menu, so that we don't need to phone the restaurant to ask what food is available

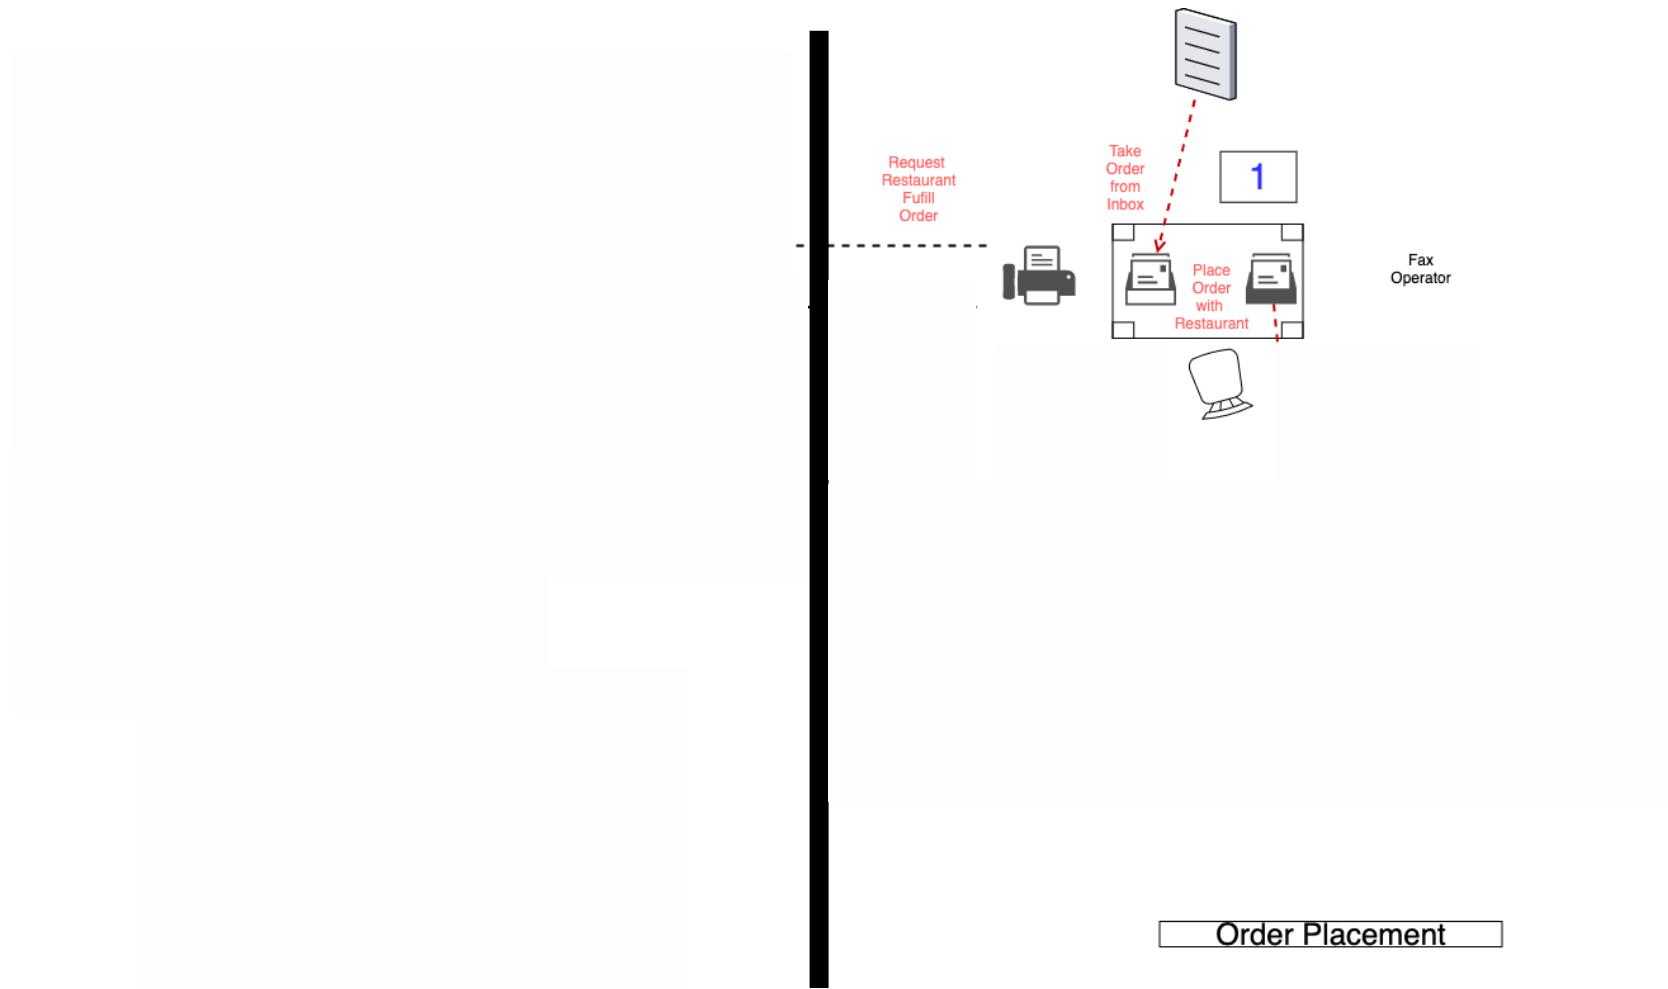
Hungry Customer

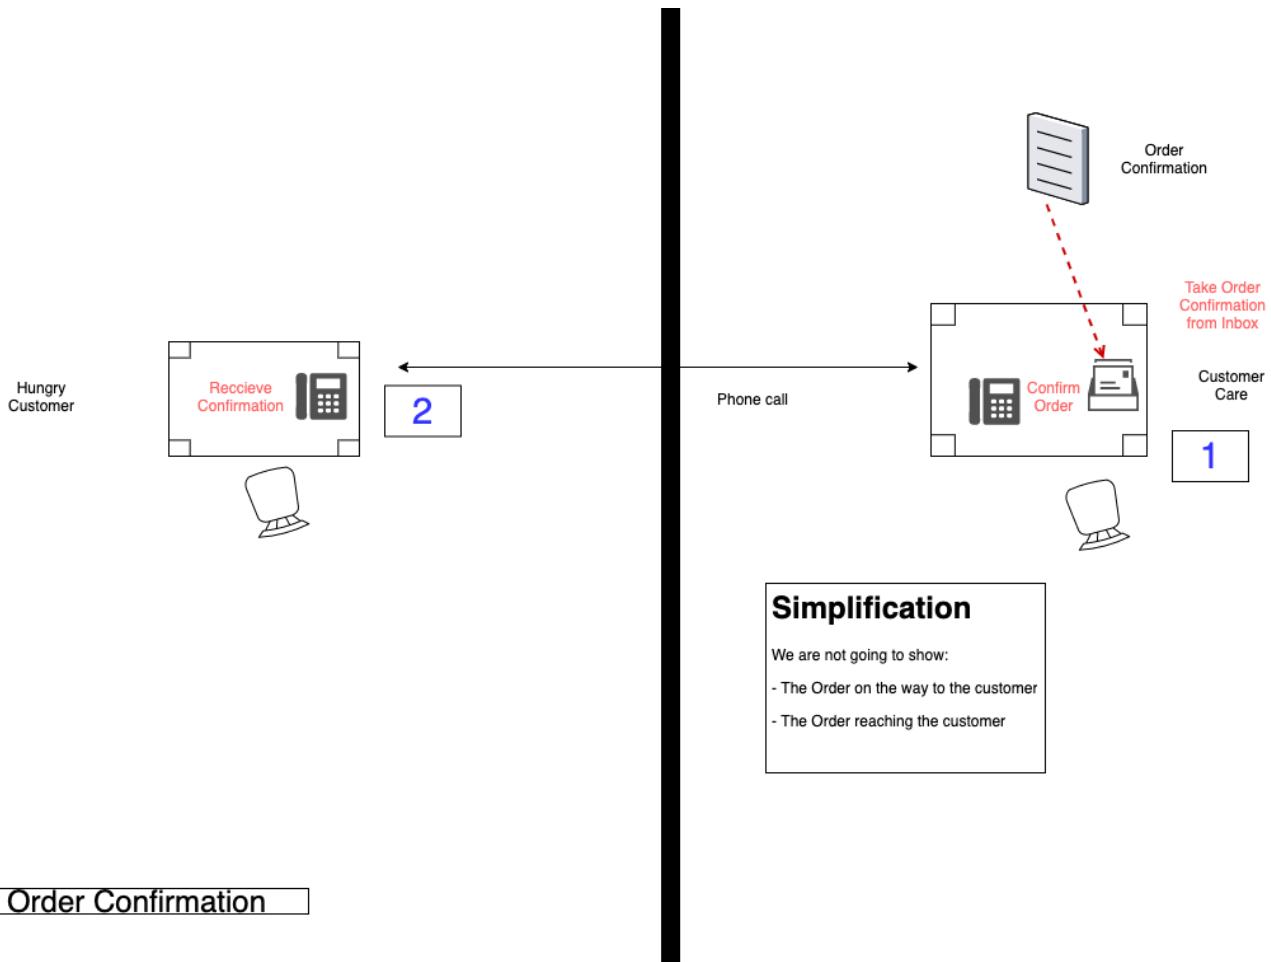


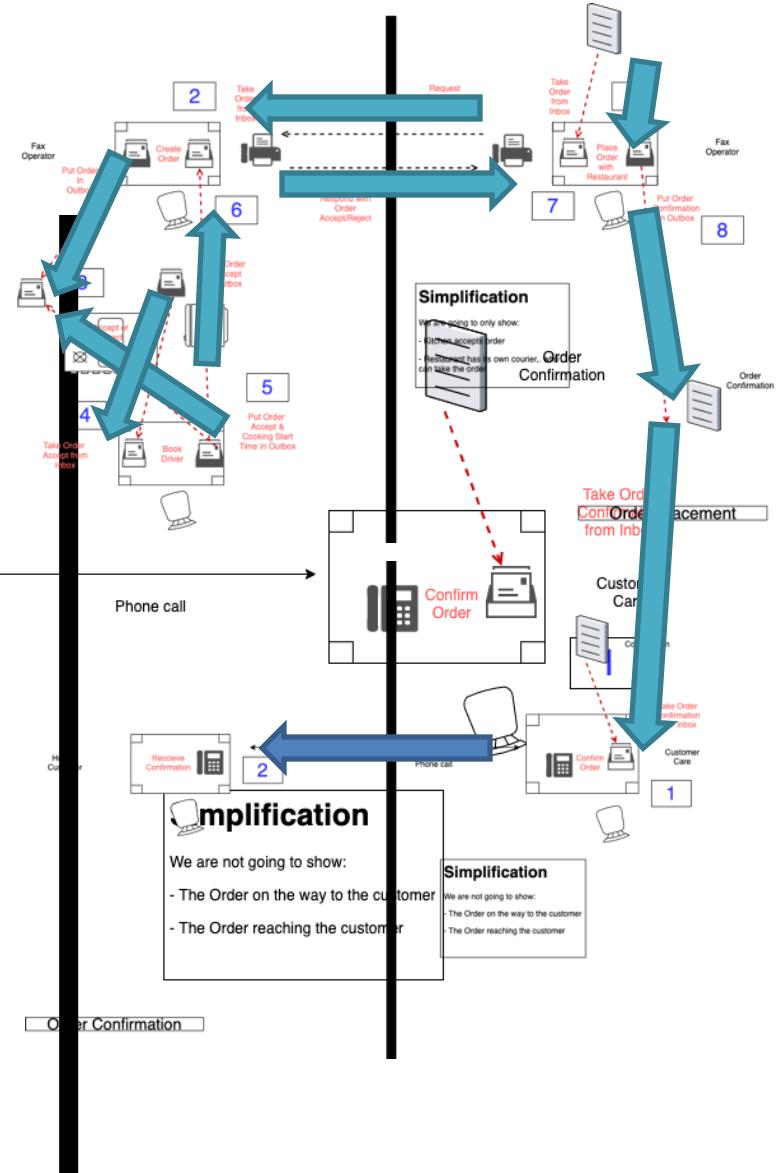
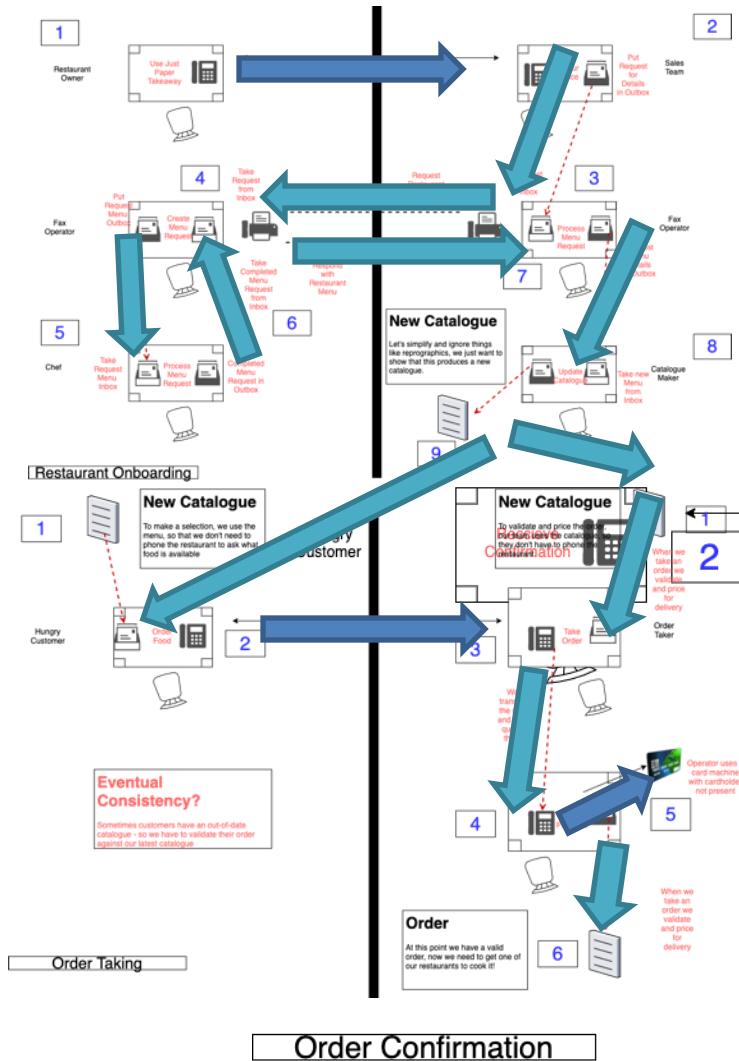
Order Taking

Order Wheel

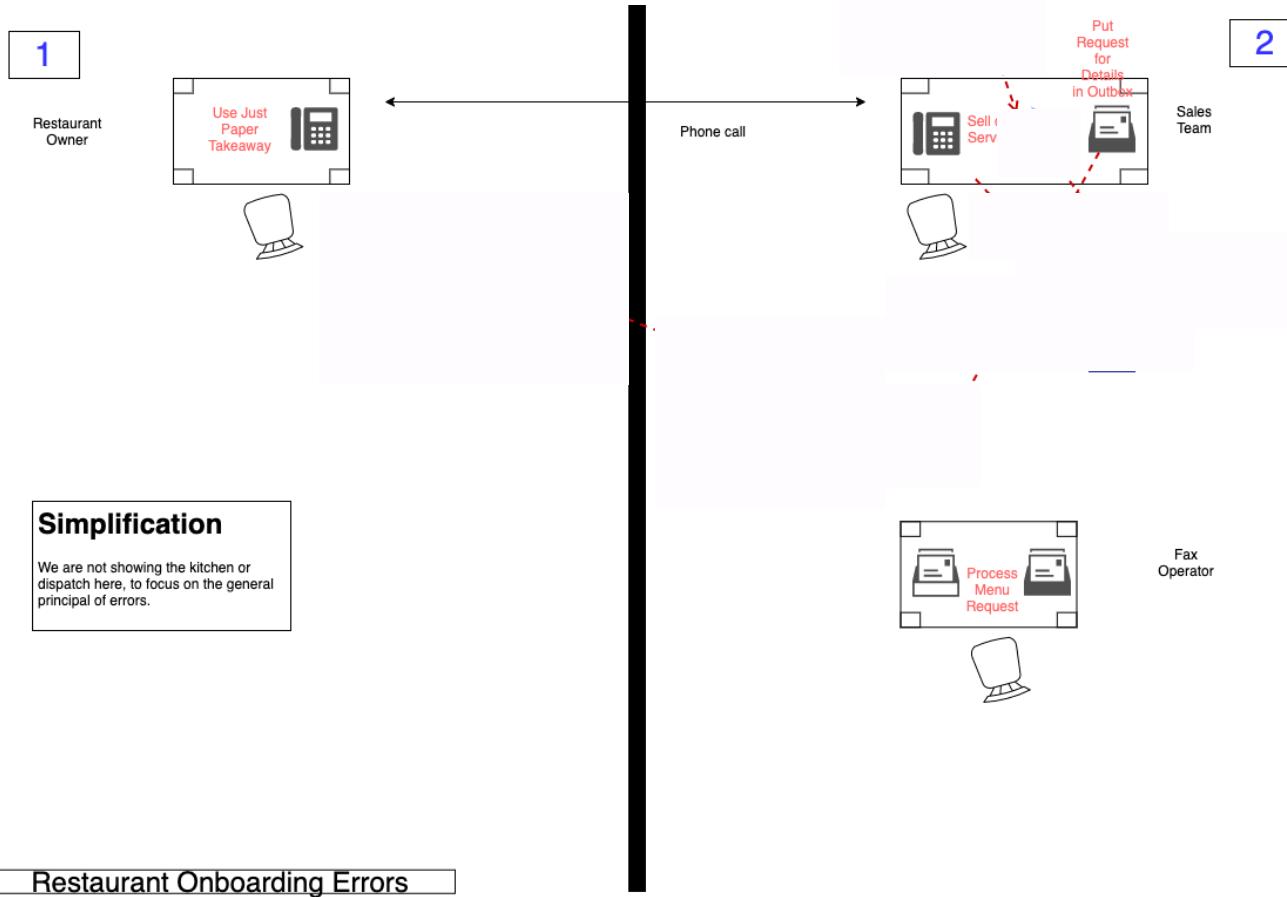








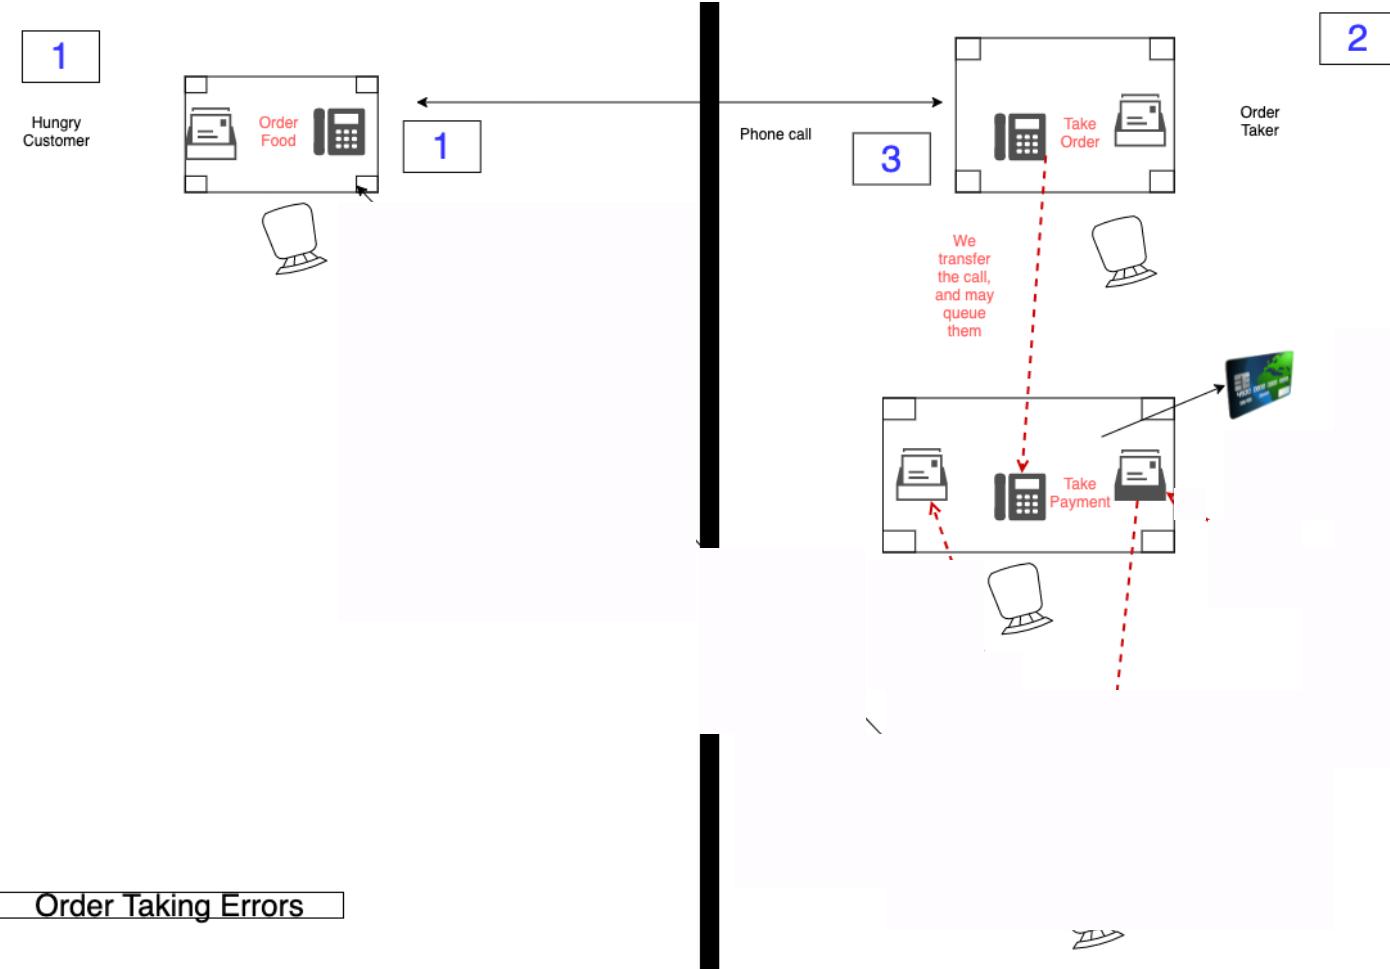


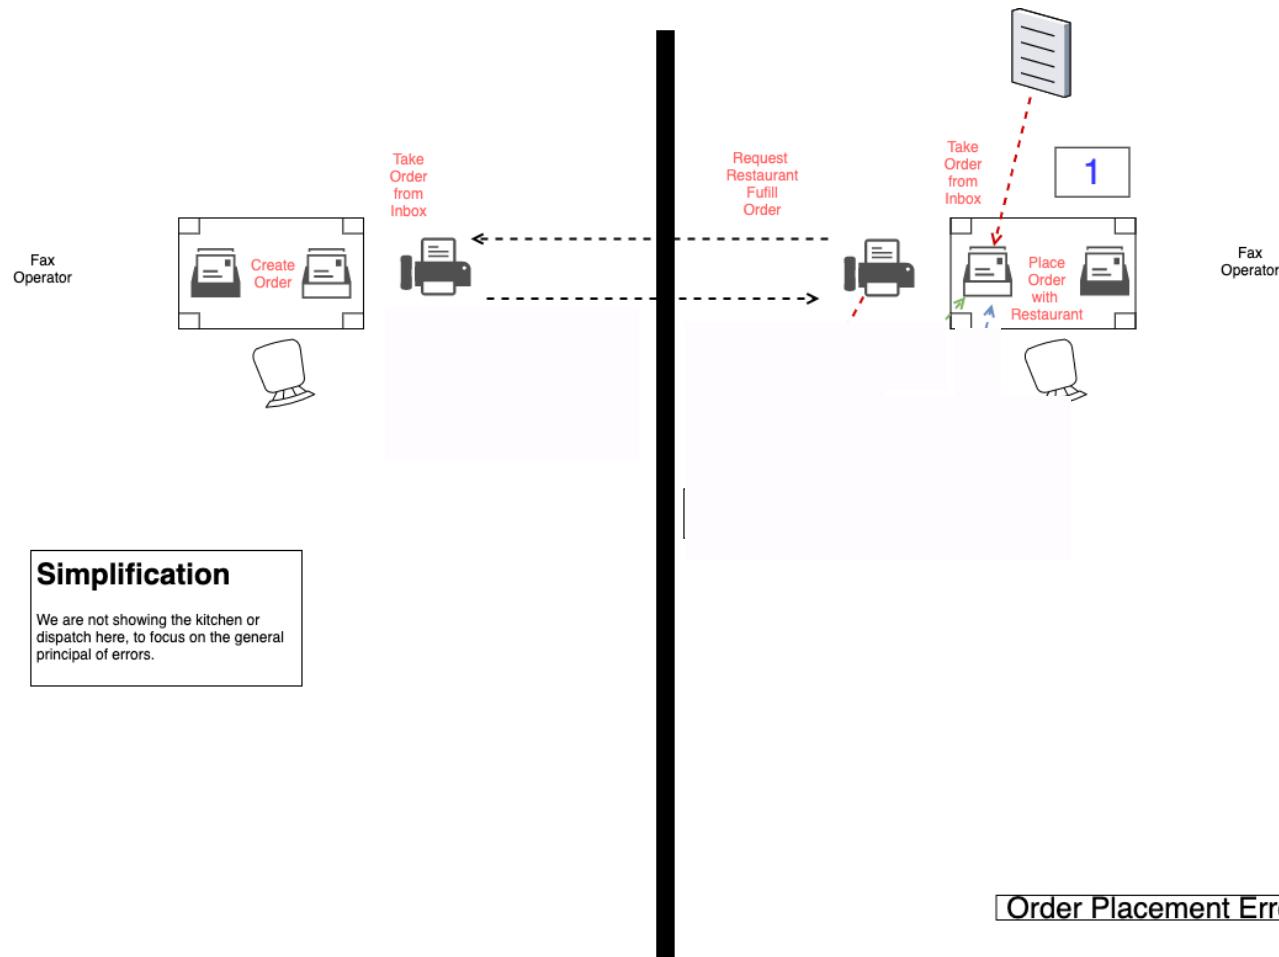


Fax Call Log

Monday, 2010-11-08 11:19

Date	Time	Type	Job #	Length	Speed	Station Name/Number	Pages	Status
2010-06-18	08:31	SCAN	92	0:25	28800	[REDACTED]	0	E-705 V.34 1M31
2010-03-22	11:39	RECV	59	0:20	26400	[REDACTED]	1	OK -- V.34 BM31
2010-03-22	11:45	RECV	60	0:20	26400	[REDACTED]	1	OK -- V.34 BM31
2010-03-22	12:29	RECV	61	0:21	26400	[REDACTED]	1	OK -- V.34 BM31
2010-09-14	14:46	SCAN	129	1:40	9600	[REDACTED]	2	E-606 V.29 AR30

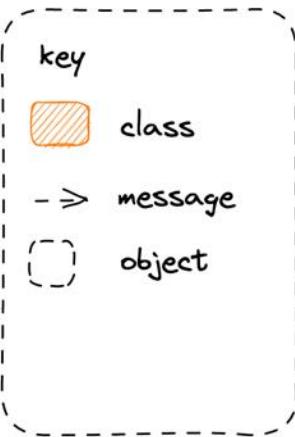




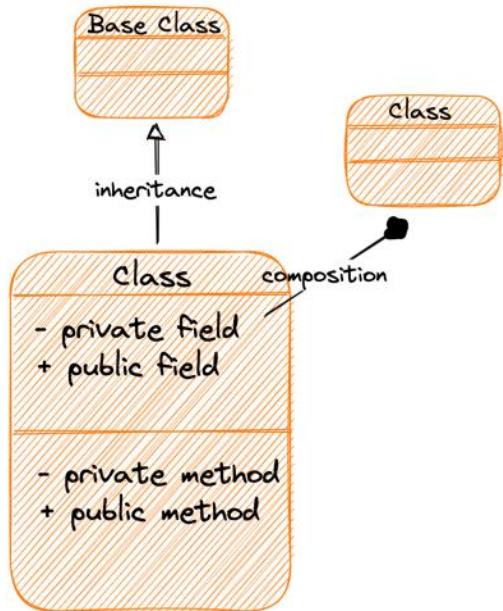
REACTIVE PROGRAMMING

Object Oriented Programming





Object Orientation



Data and Behaviour

A class couples data and the behaviour that depends upon that data. This allows encapsulation: the data can be hidden and just behaviour exposed.

Dynamic Dispatch

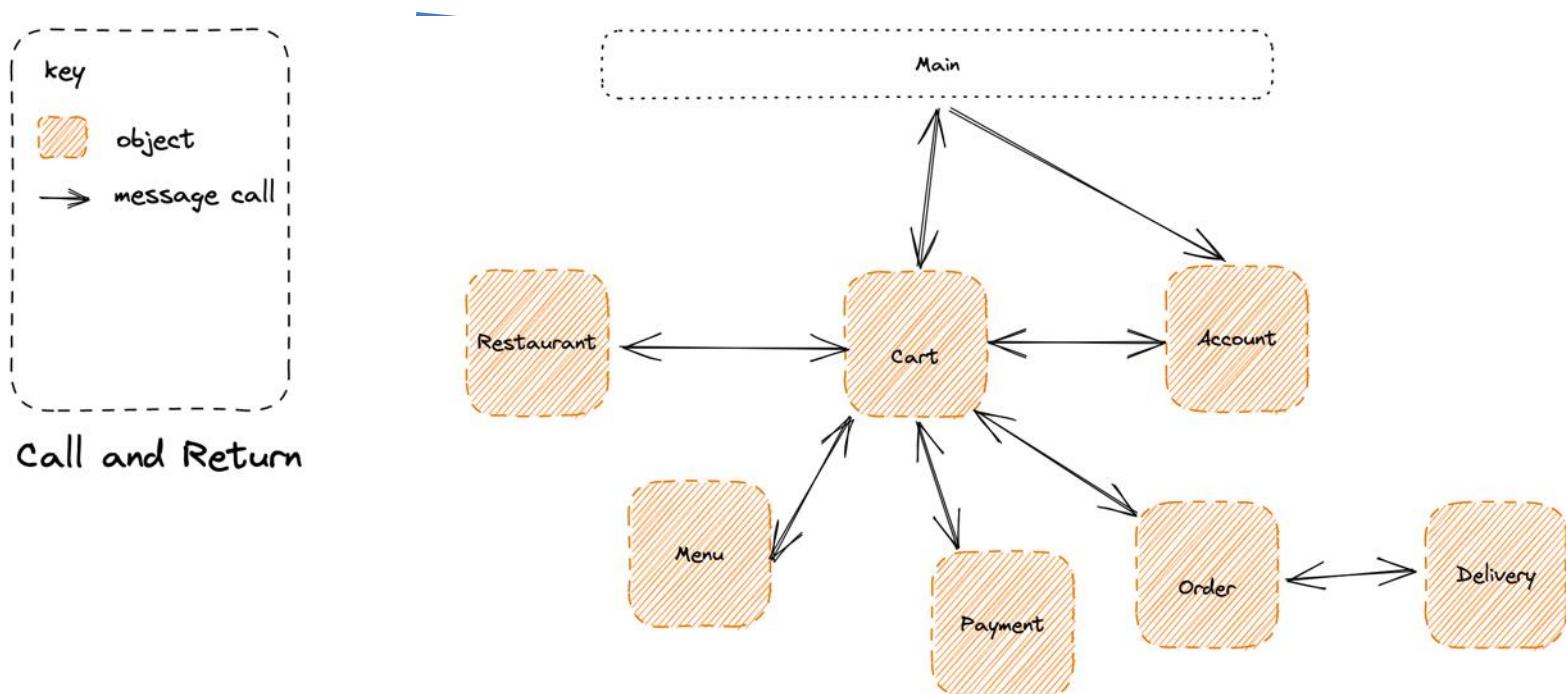
Due to inheritance the method chosen in response to a message may come from the base class of an object



Message Passing

Objects communicate by message passing. A message is the method to call, and the parameters to that method.

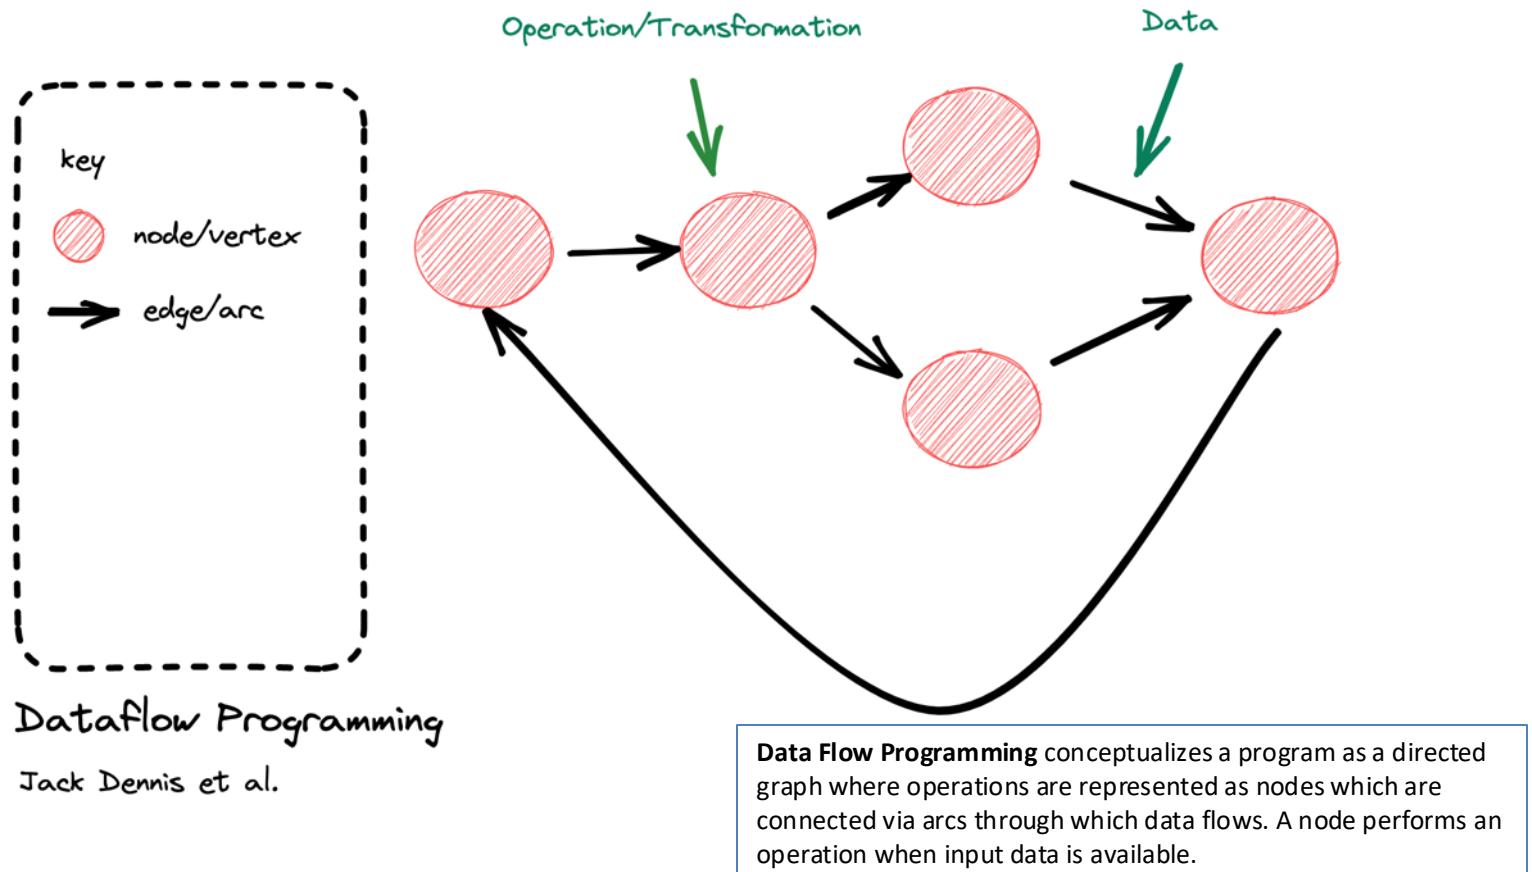
Call and Return: The main method represents our entry point and it uses message passing to invoke objects, which in turn invoke other objects, and return to their caller

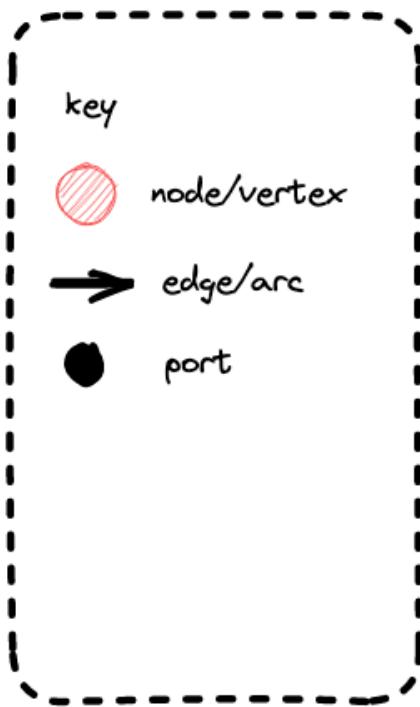


God Object: A danger here is that we end up with a "god" object, such as Cart here, that controls all the other objects. This is a high-degree of behavioral coupling

Data Flow Programming

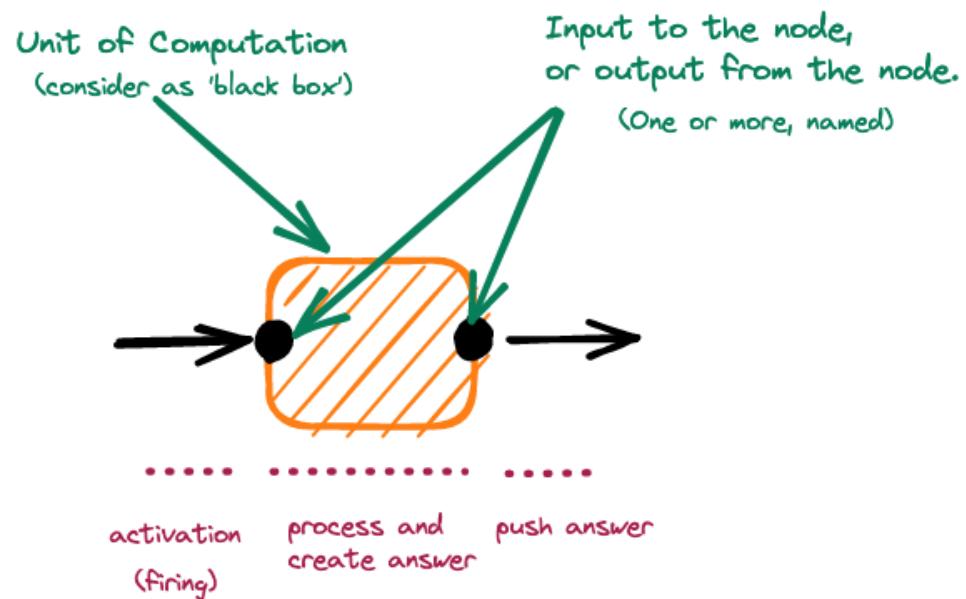






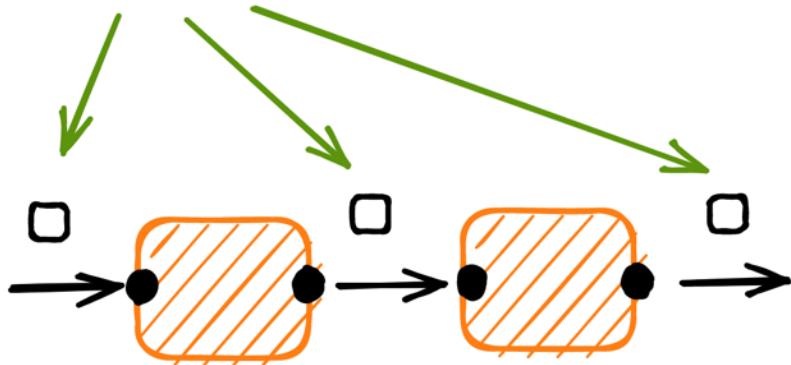
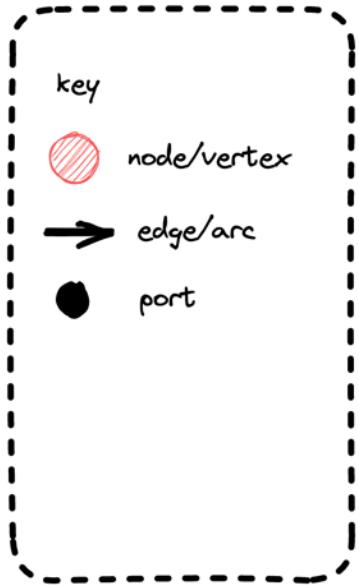
Dataflow Programming

Jack Dennis et al.



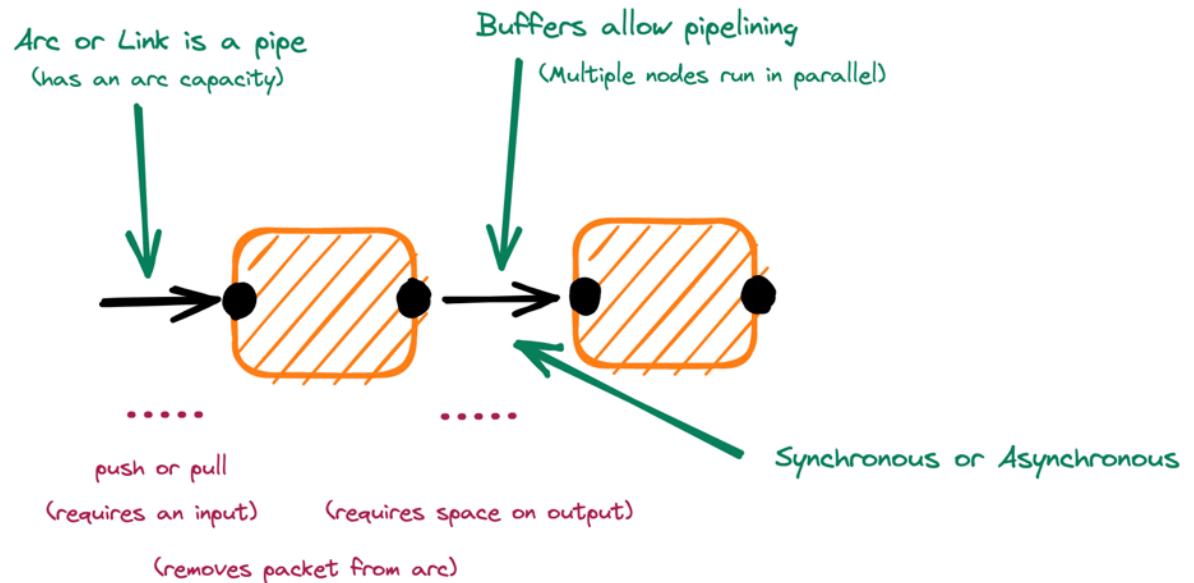
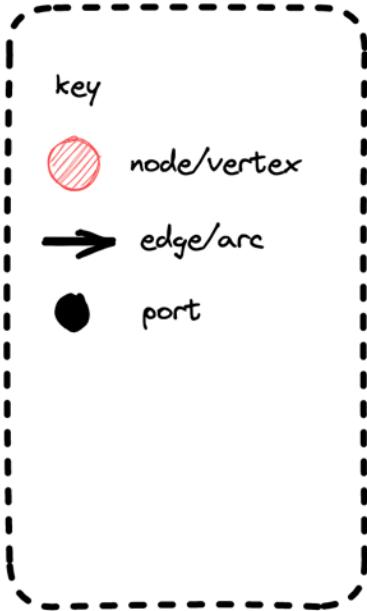
Data Packet

primitive/compound, structured/unstructured values, and even include other data packets



Dataflow Programming

Jack Dennis et al.



Dataflow Programming

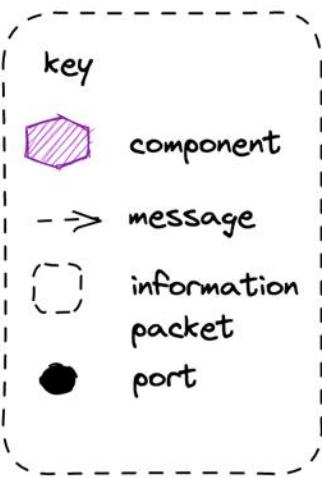
Jack Dennis et al.

Node Lifetime: In 'classic' dataflow programming the lifetime of a node is from activation when there is work to do (push) or work is requested (pull) until it has pushed the answer.

Capacity: In dataflow programming we do not activate a node to read from the input, if there is no space on the output. (Infinite capacity links exist only in theory). This creates **backpressure**.

Flow Based Programming

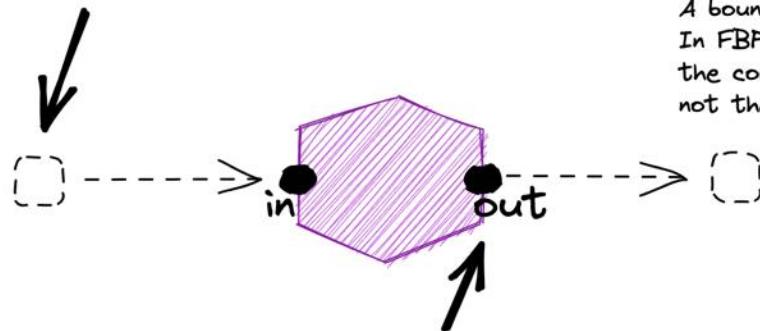
“Everything Flows
and Nothing
Stays.” -
Heraclitus



Flow Based Programming (J. Paul Morrison)

Information Packet (IP)

An independent structured piece of information with a defined lifetime.



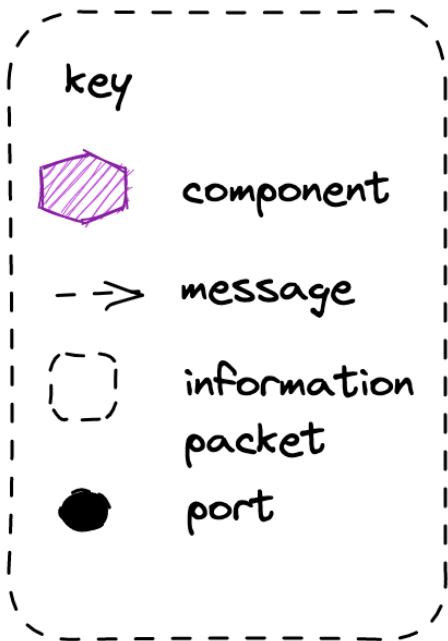
Port

The point where a connection makes contact with a process. A port is addressed by name.

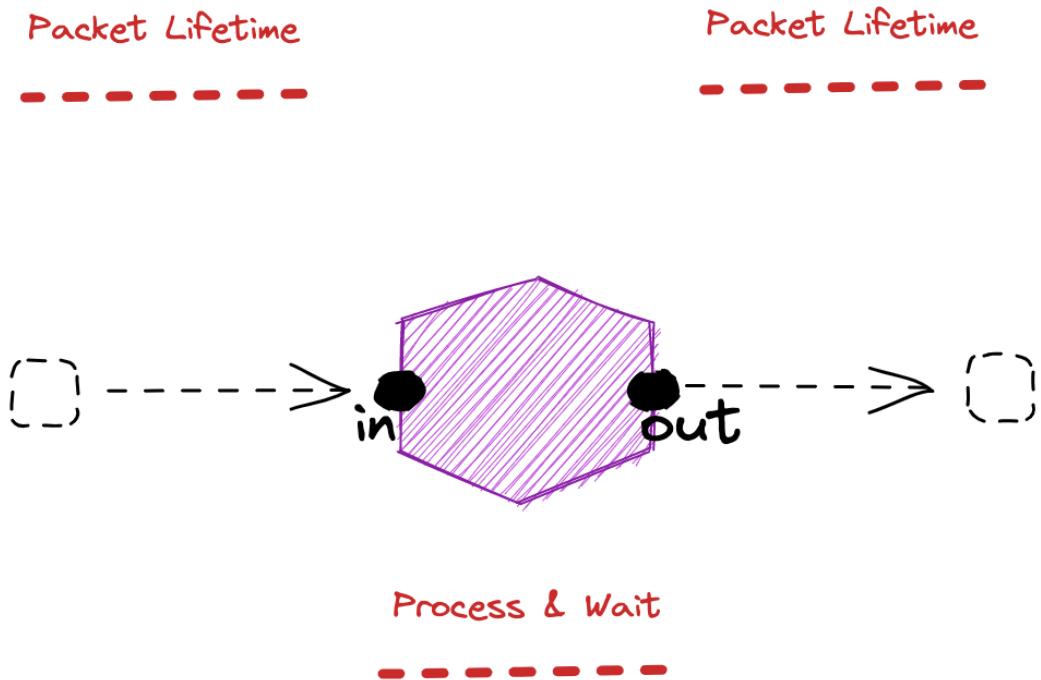
Flow-Based Programming (FBP) is a subclass of DFP. While DFP can be synchronous or asynchronous, FBP is always asynchronous. FBP allows multiple input ports, has bounded buffers and applies back pressure when buffers fill up.

Connector

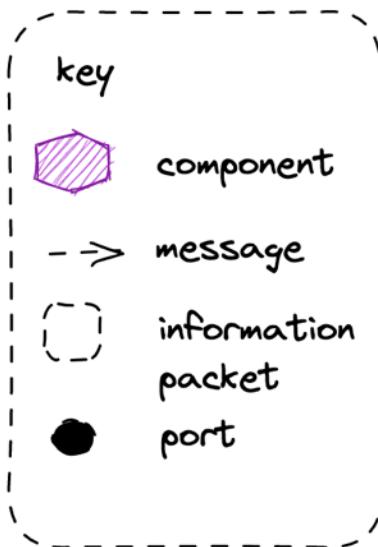
A bounded pipe of information packets.
In FBP a connector is external - that is the component is only aware of the named port not the connector, so it could be swapped out



Flow Based Programming (J. Paul Morrison)



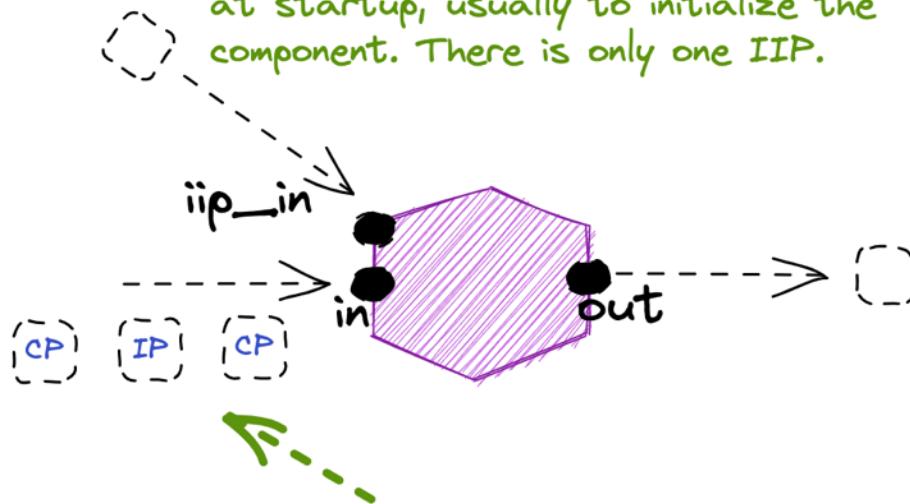
Node Lifetime: In flow base programming the a node can remain running whilst there is work on an input queue, can can suspend instead of terminate if there is no work on its connector.



Flow Based Programming (J. Paul Morrison)

Initial Information Packet

A packet that the component reads at startup, usually to initialize the component. There is only one IIP.

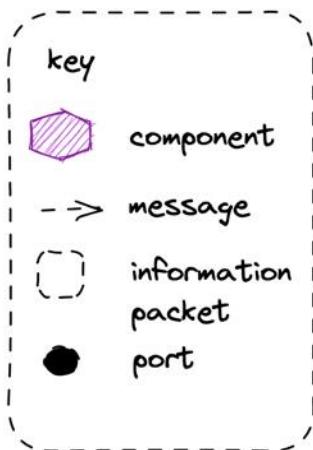


Control Packets

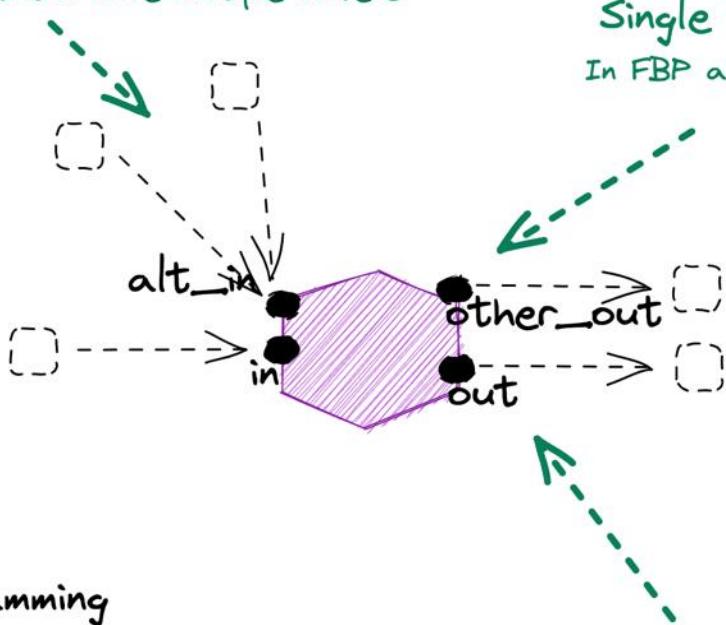
Control Packets can be used to 'bracket' groups

Multiple Writers

In FBP an in port can have multiple writers



Flow Based Programming (J. Paul Morrison)



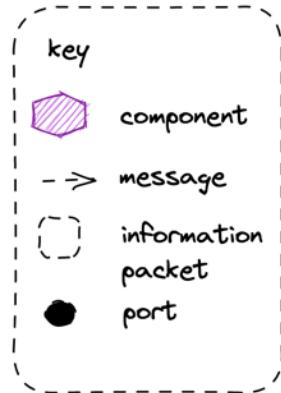
Single Writer

In FBP an out port is single writer

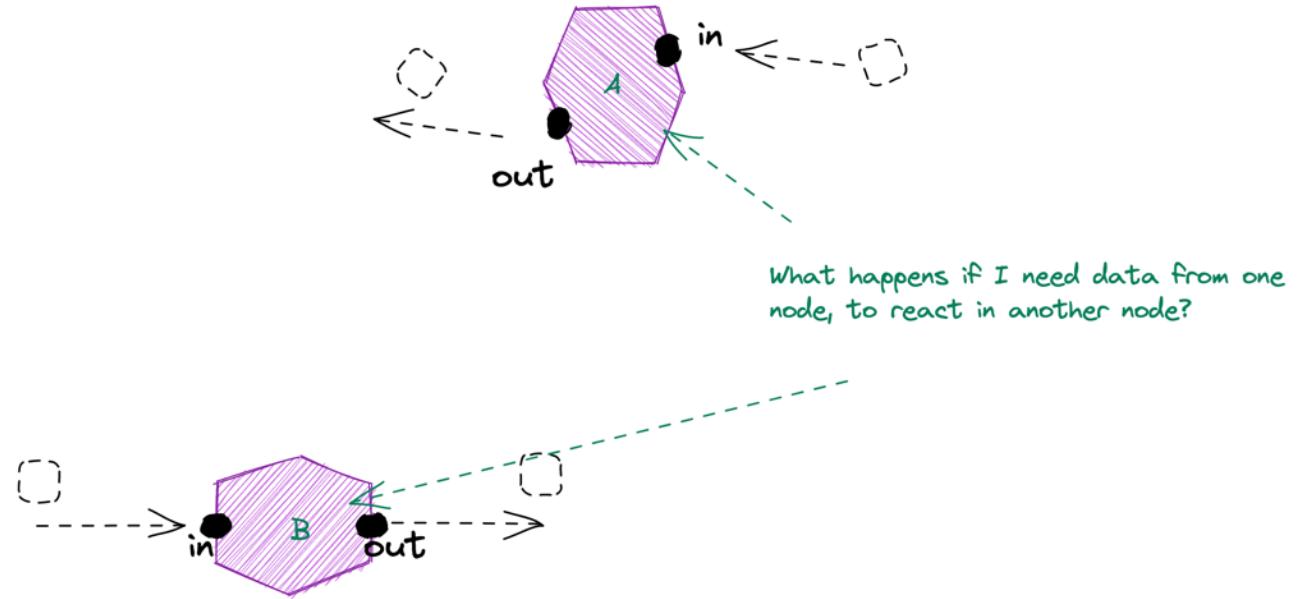
Multiple In/Out Ports

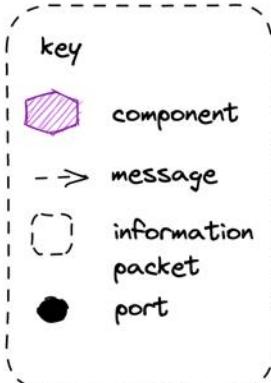
In FBP we can have multiple in or out ports

Capacity: In flow-based programming a component can test to see if it can send to out, and if not decide whether to halt or ignore.



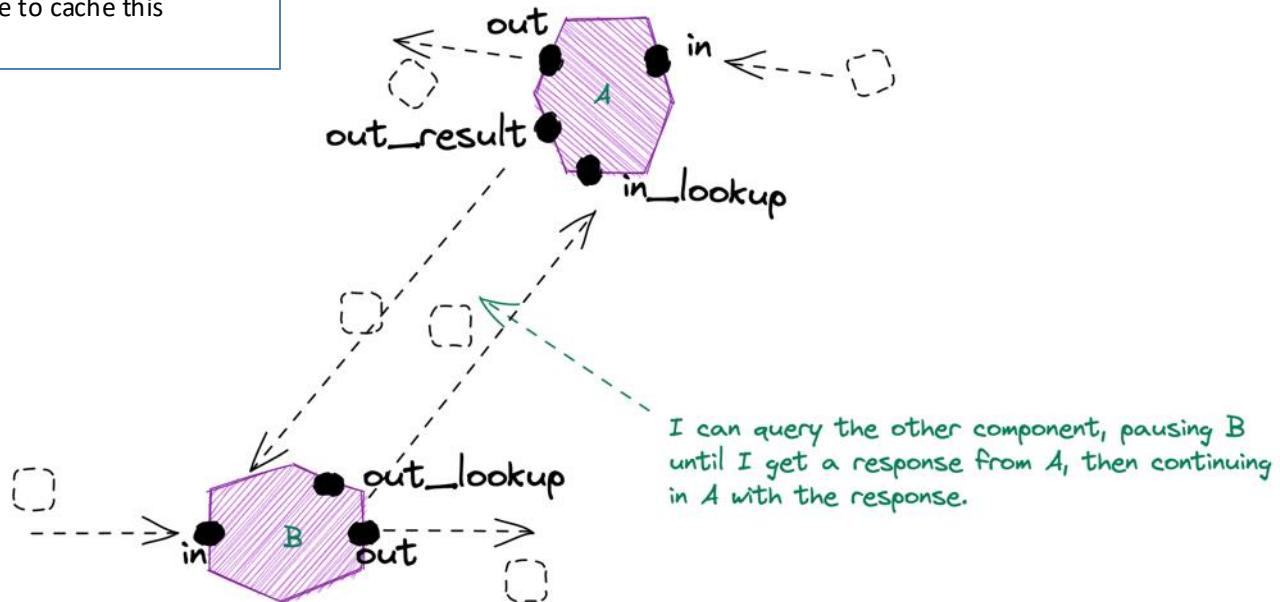
Flow Based Programming (J. Paul Morrison)

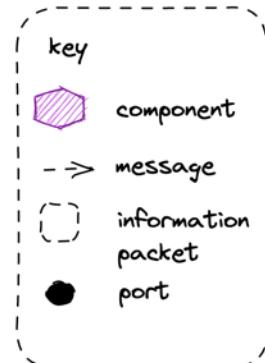




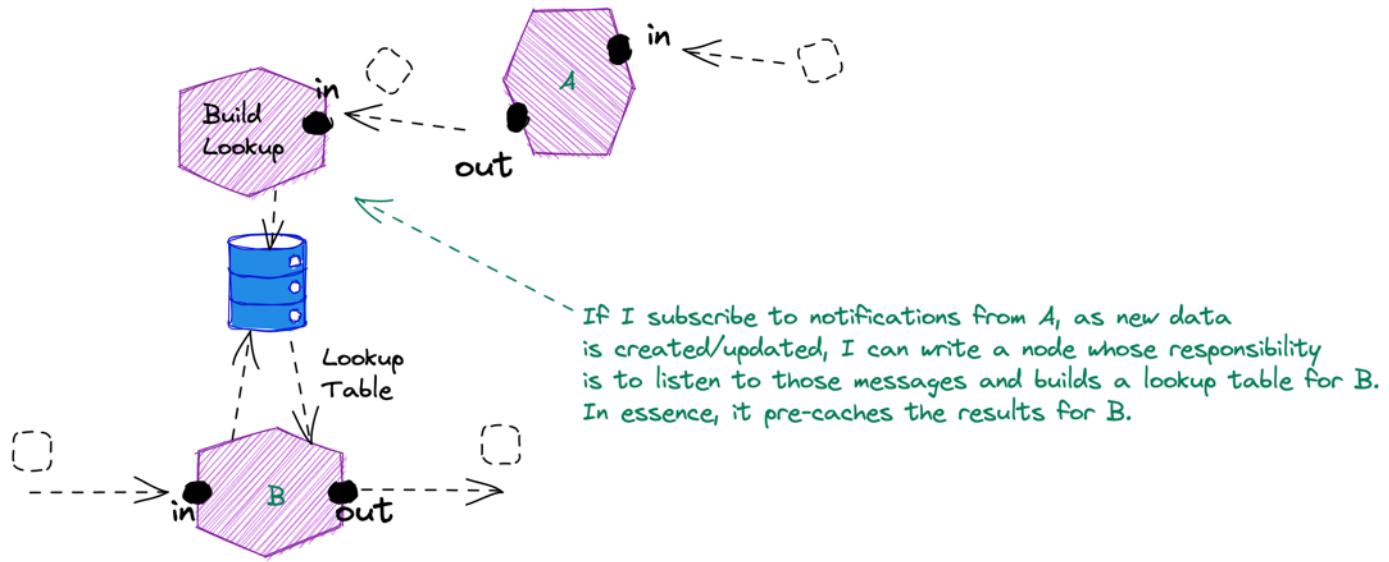
Flow Based Programming (J. Paul Morrison)

Walk of Shame: The trouble is that we may keep asking, so it makes sense to cache this somehow.



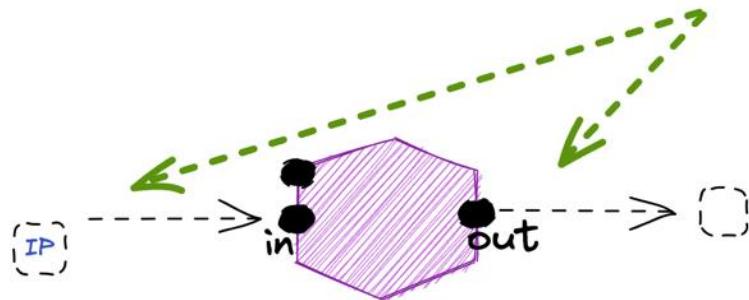
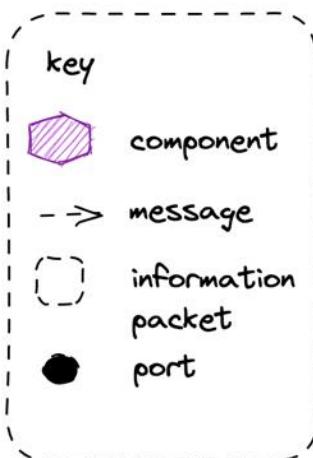


Flow Based Programming (J. Paul Morrison)



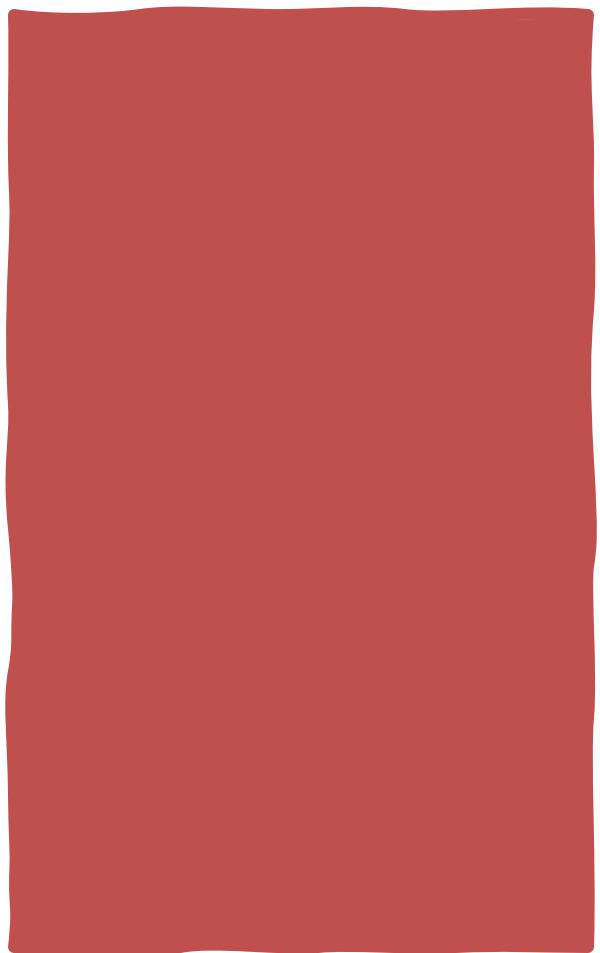
Message Oriented Middleware (MoM)

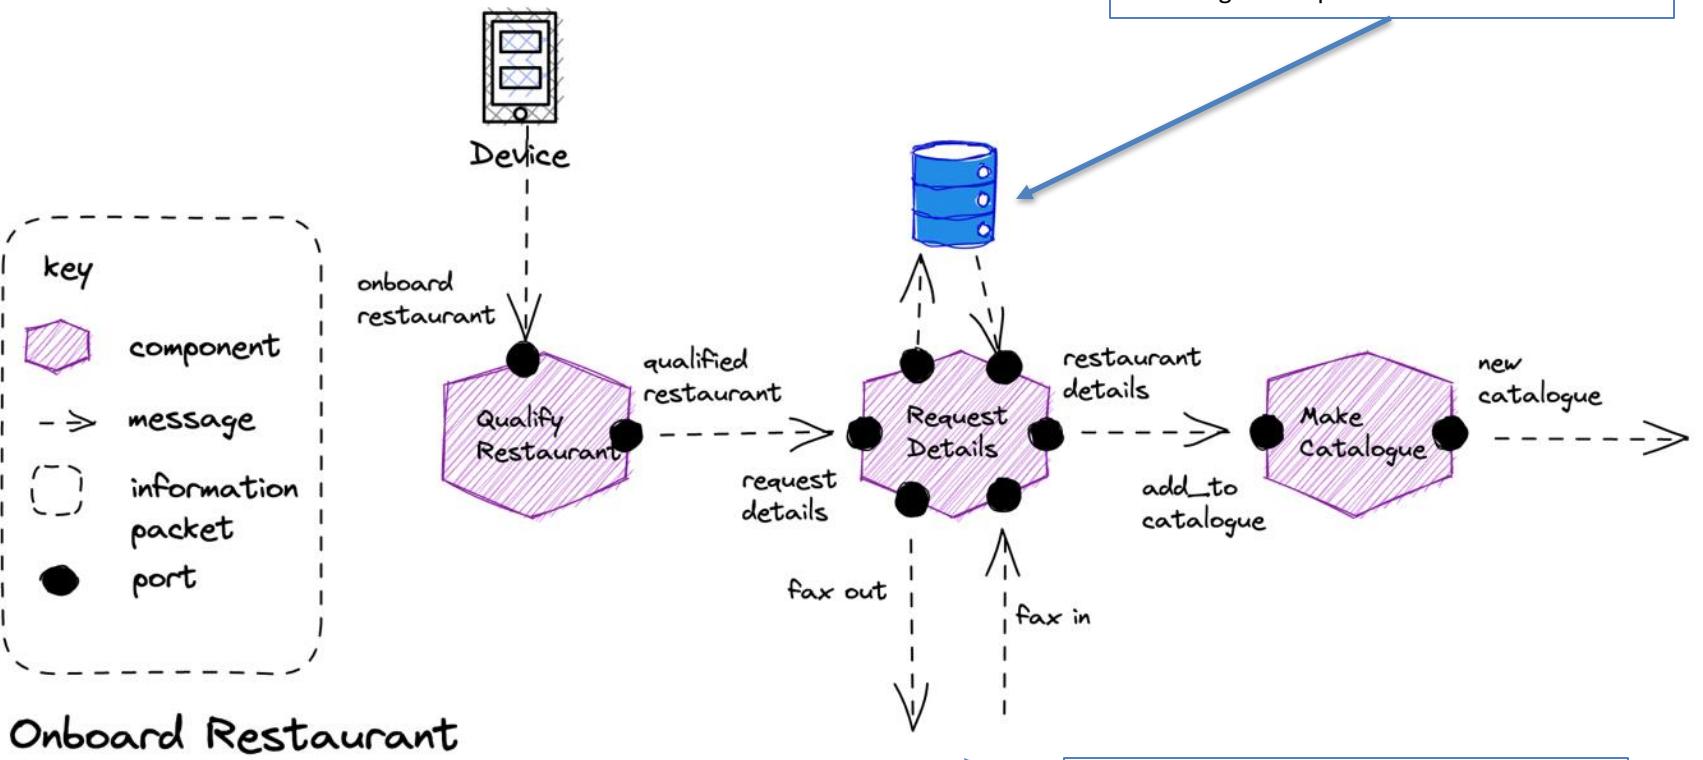
Because FBP allows for externally defined connectors in the 2010 update JPM outlined that it could be used in distributed systems. Nodes become processes and connectors use MoM.



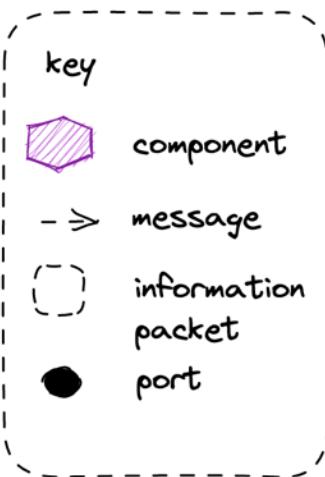
Flow Based Programming - 2010 Update
(J. Paul Morrison)

Example

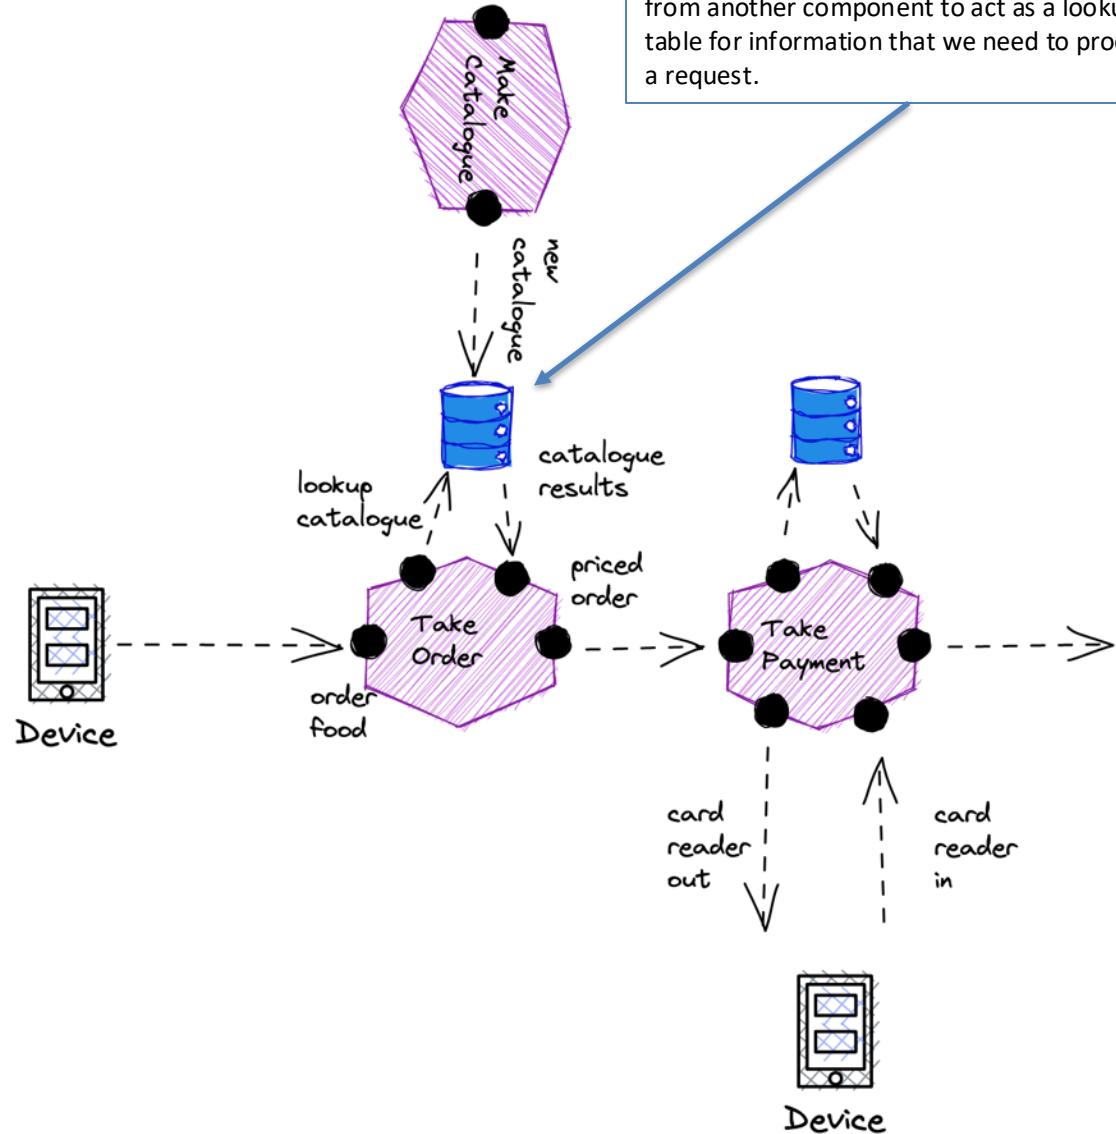


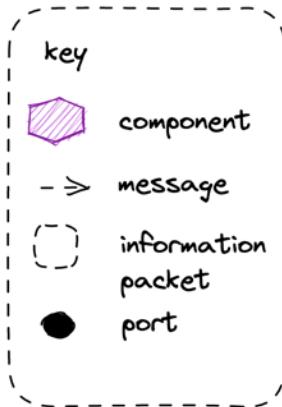


Correlation: To allow us to match the response to the saved state we use a *correlation id*. We send it on the ip to **fax_out**, and the restaurant includes it on the ip from **fax_in**. We use the ip to look up the workflow we stored.

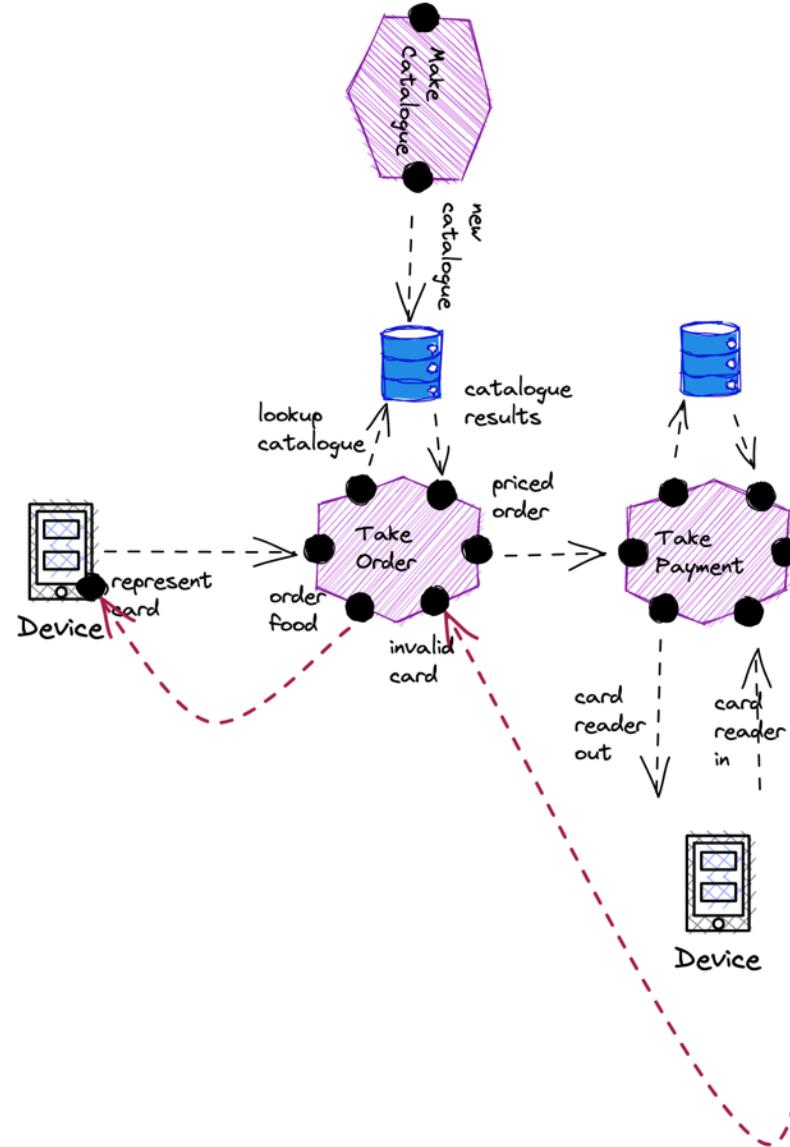


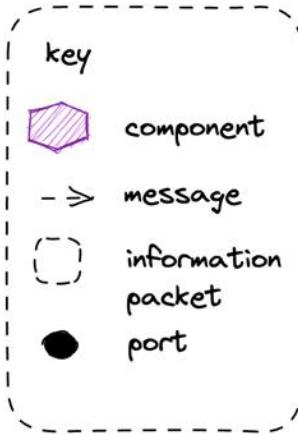
Order Food



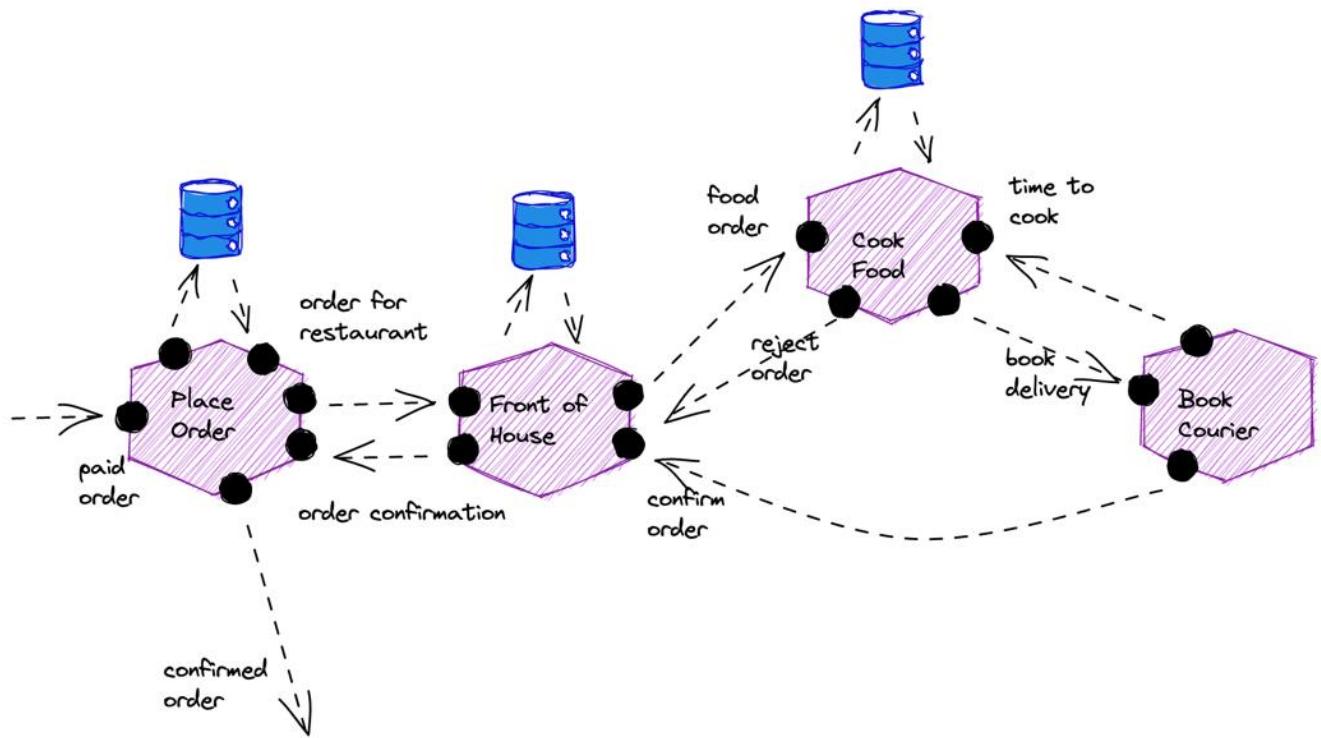


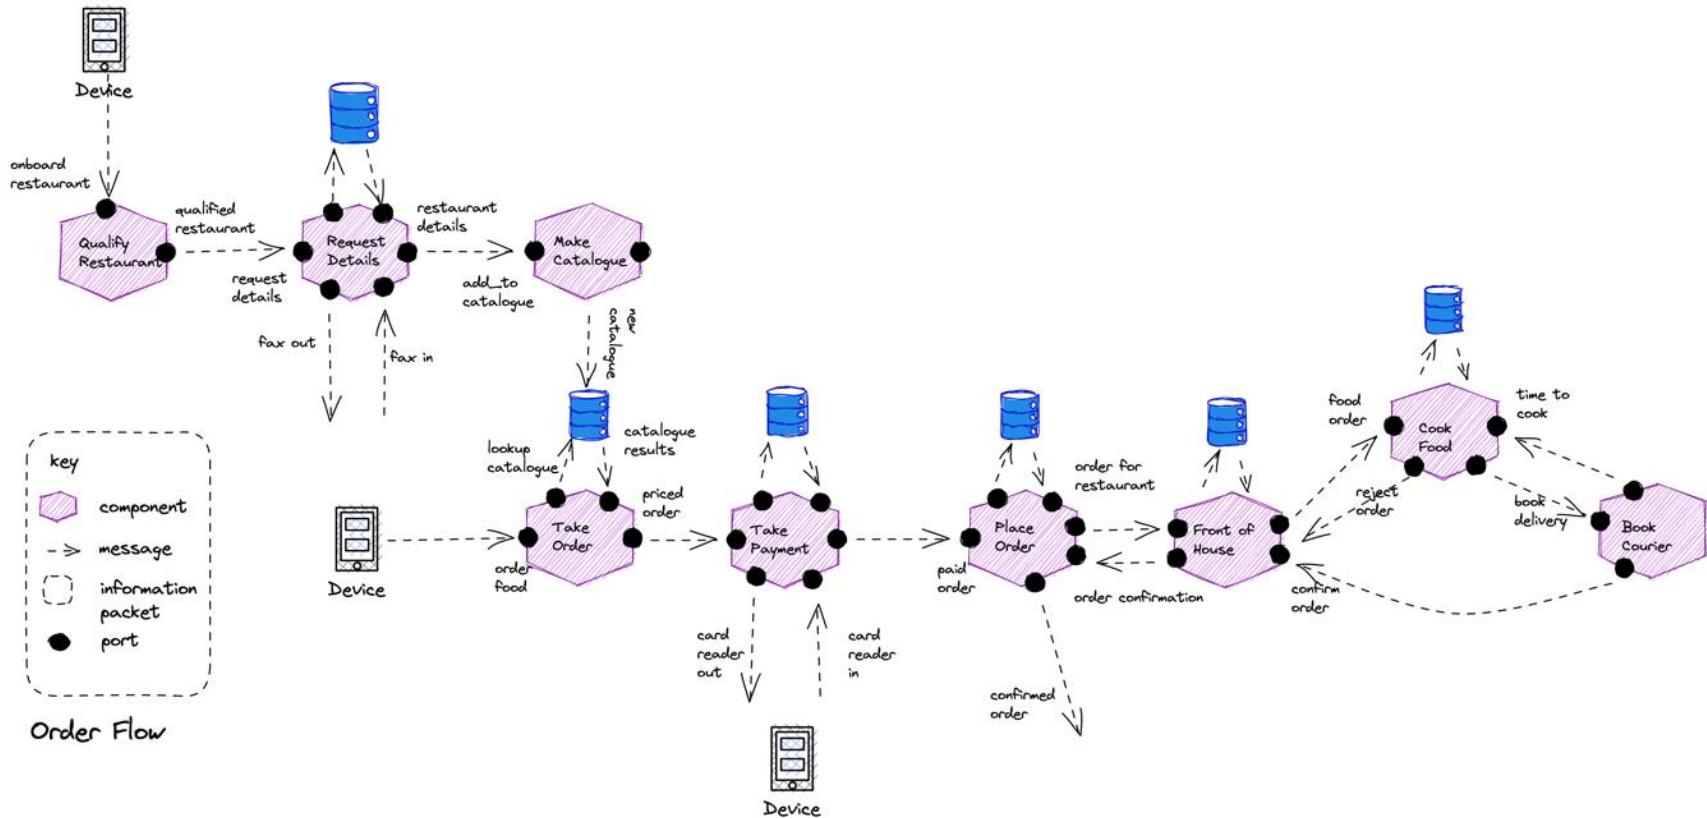
Order Errors





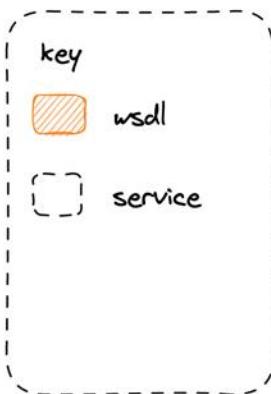
Order Placement



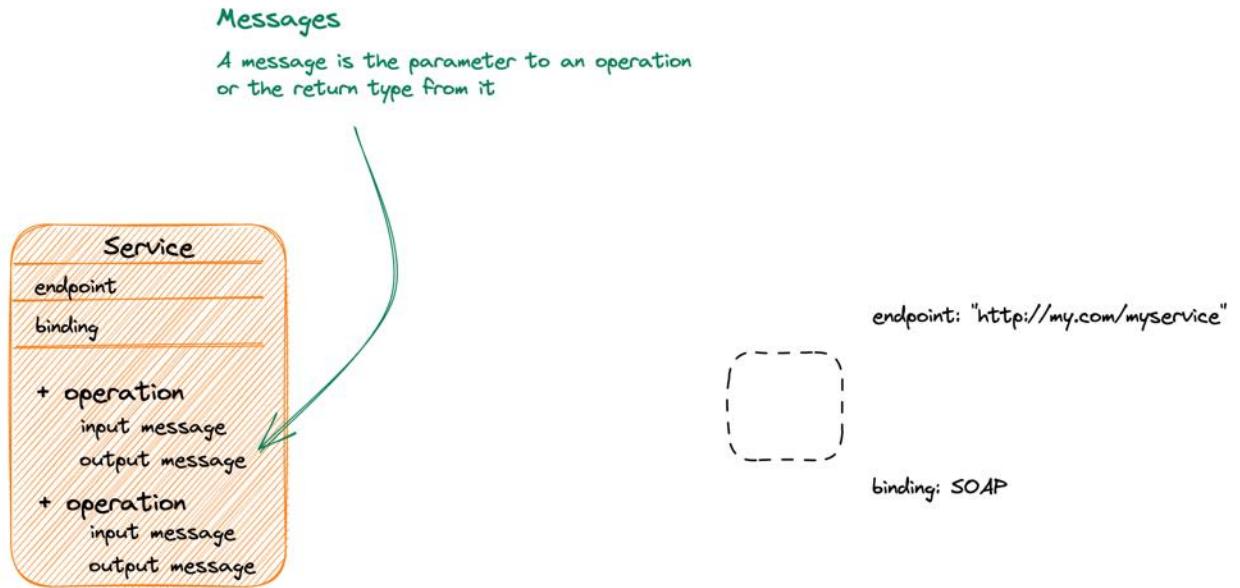


SOA Architectures





Service Orientation

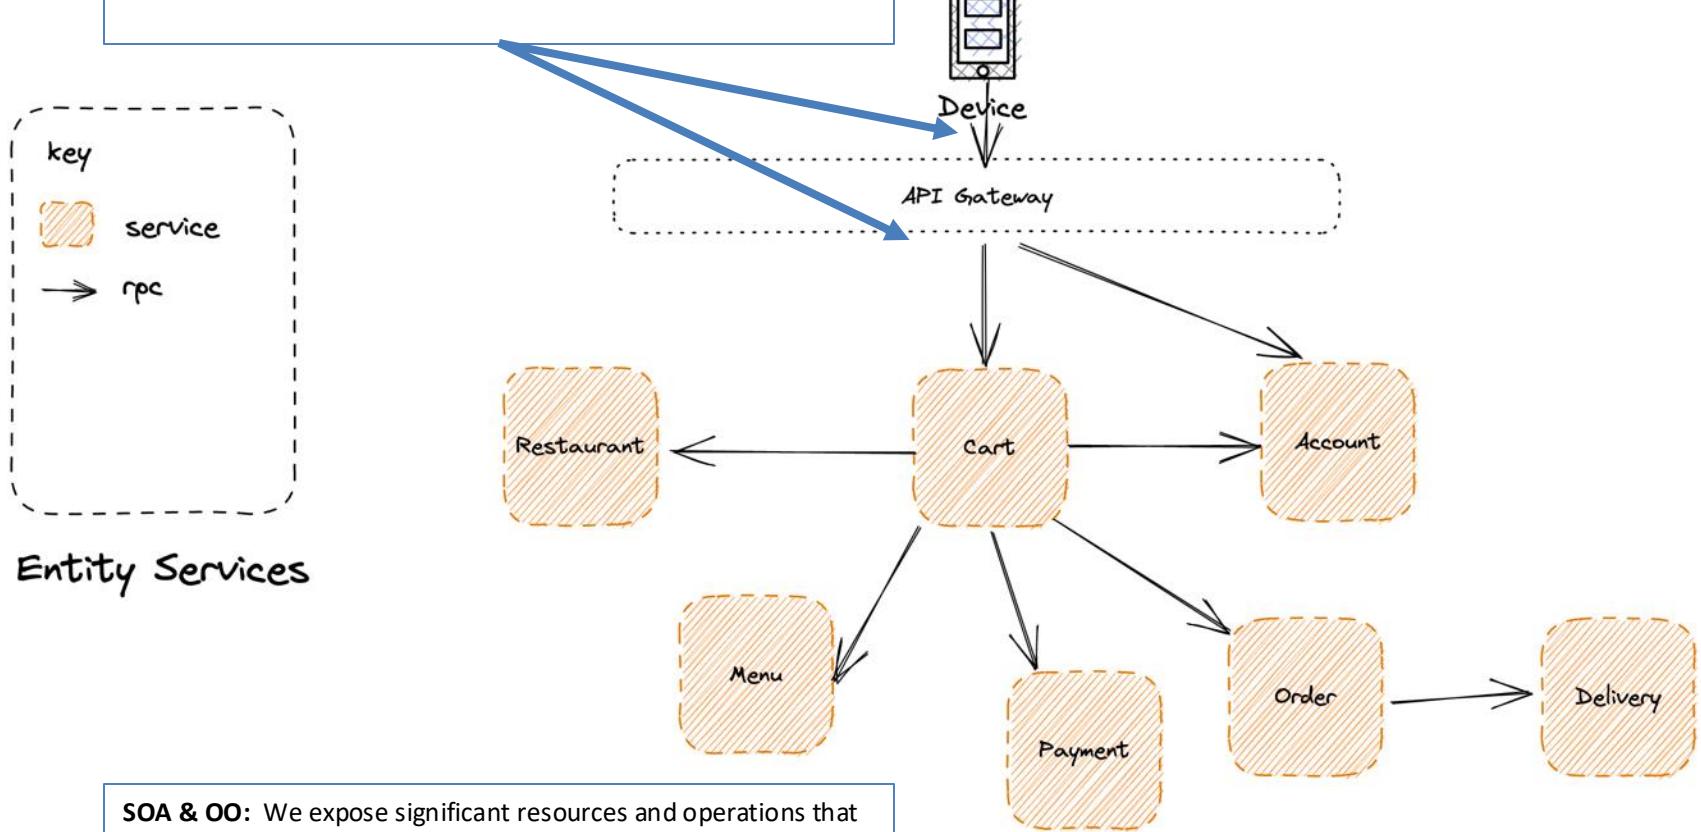


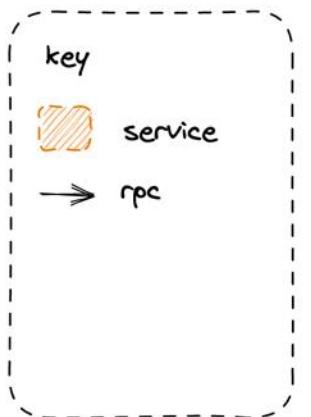
Data and Behaviour

A service couples data and the behaviour that depends upon that data. This allows encapsulation the data can be hidden and just behaviour exposed

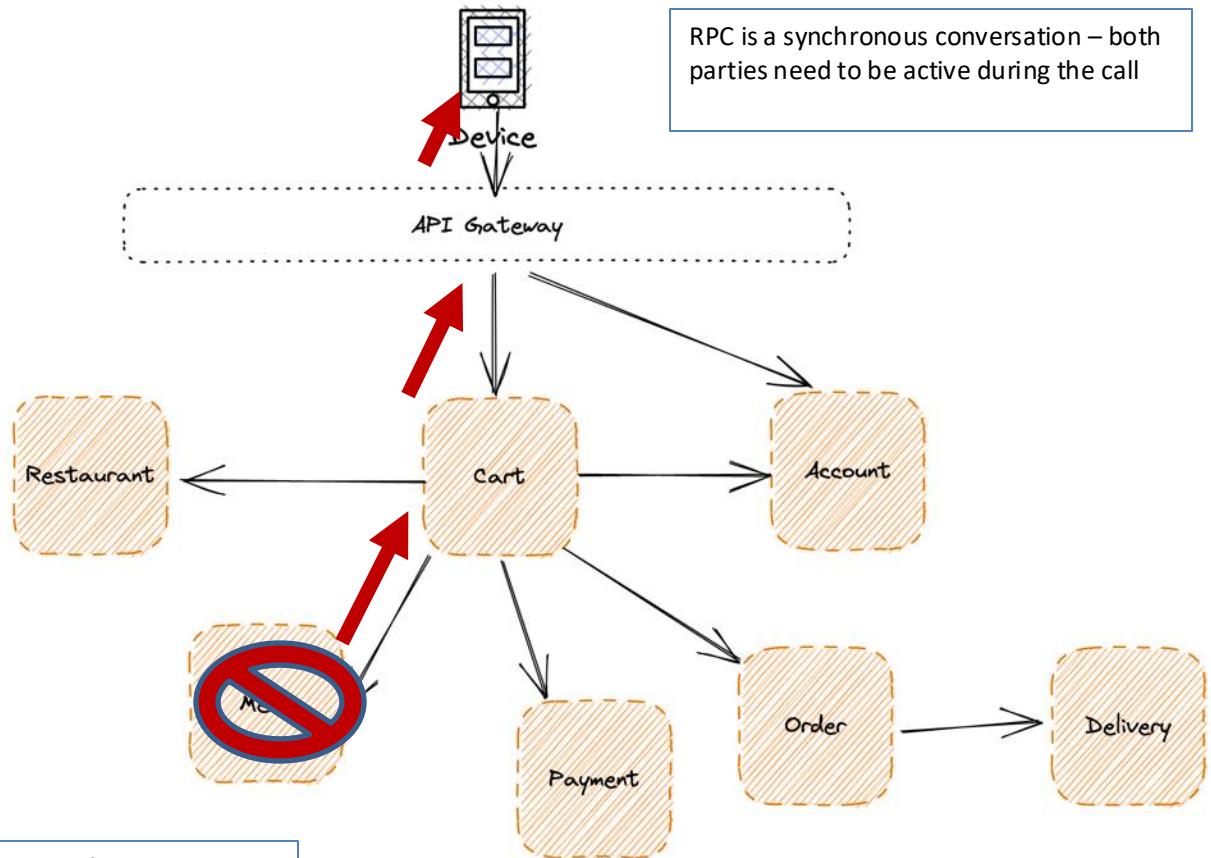
SOA & OO: SOA creates OO-like components, they encapsulate their data and expose behavior that is coupled to that data. This is the Web Services approach to services.

Feature Envy: Because domain logic needs to co-ordinate across these resources, it ends up either on the client, the API Gateway, or the Cart service and not in individual services. An *anti-pattern*.





Entity Services



Faults Propagate: With a synchronous service, a fault propagates

Reactive Architectures

The Reactive Manifesto

Published in September 2014

Author(s): Jonas Bonér (Erik Meijer, Martin Odersky, Greg Young, Martin Thompson, Roland Kuhn, James Ward and Guillaume Bort)

Defines an architectural style: **Reactive Applications**

Write applications that:

react to events: their event-driven nature enables the following qualities

react to load: focus on scalability rather than single-user performance

react to failure: resilient systems with the ability to recover at all levels

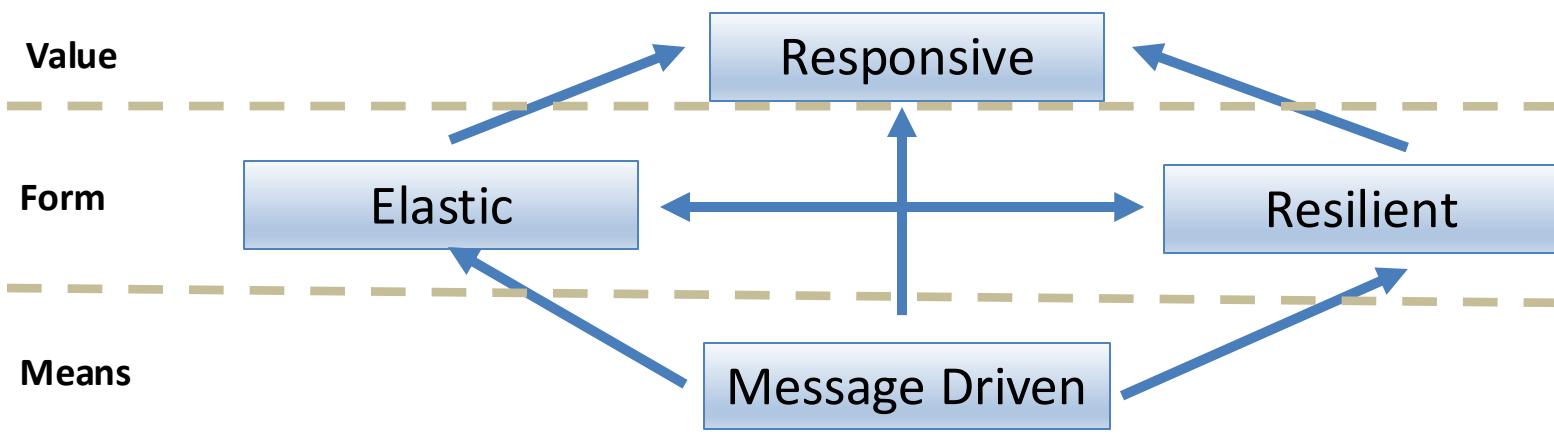
react to users: combine the above for an interactive user experience

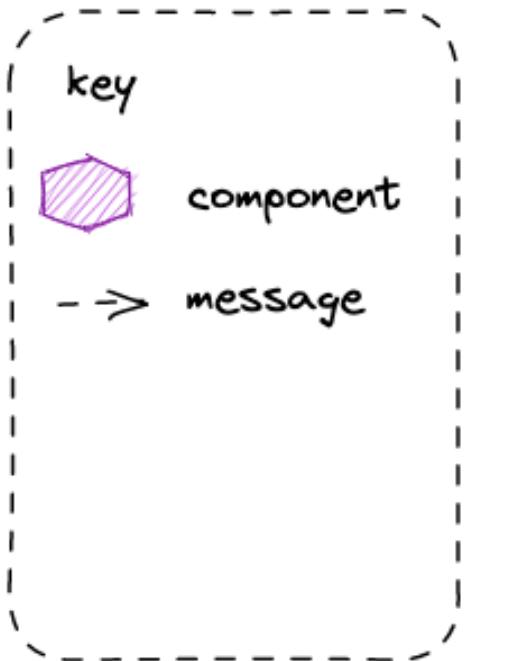
Responsive: The system responds in a timely manner.

Resilient: The system stays responsive in the presence of failure.

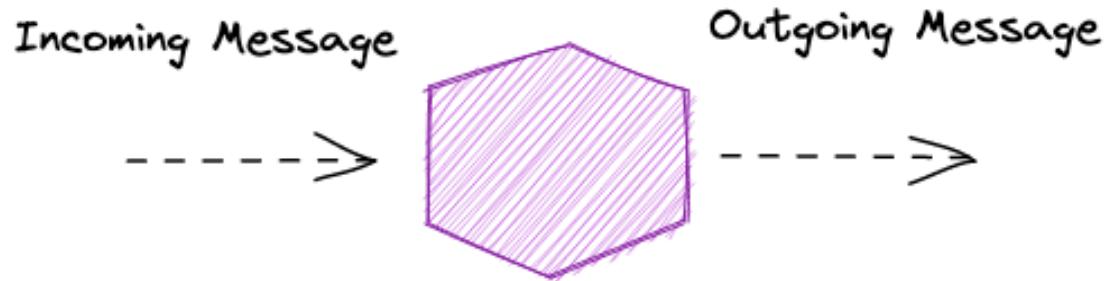
Elastic: The system stays responsive under varying workload.

Message Driven: Rely on asynchronous message passing

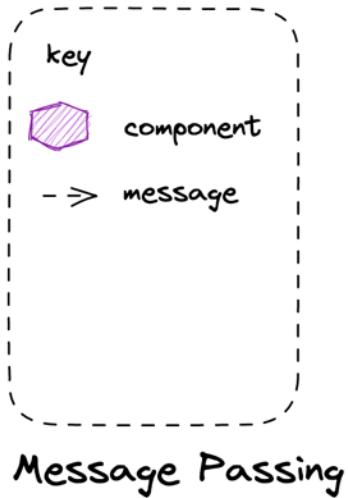




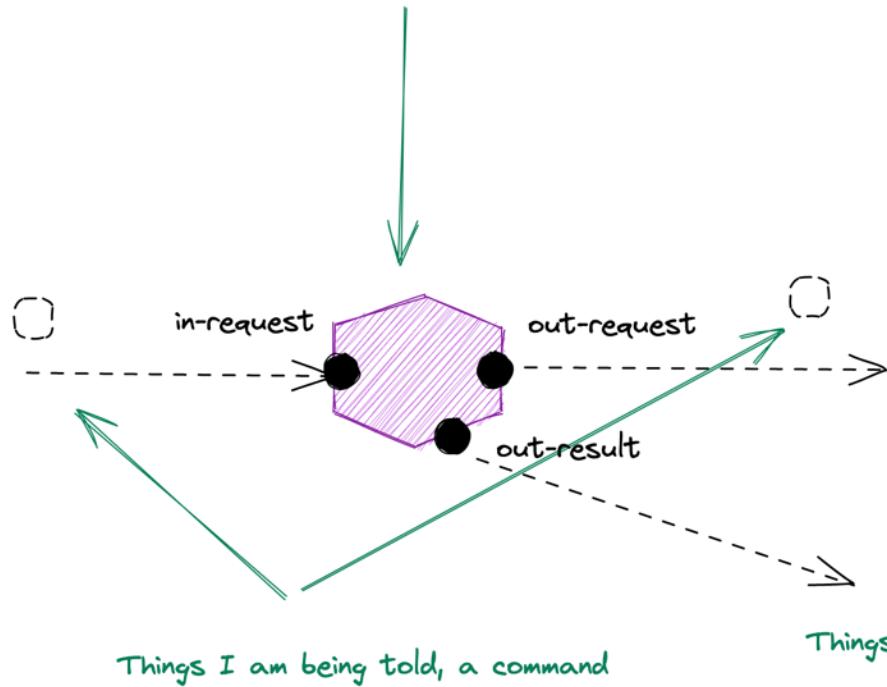
Message Passing

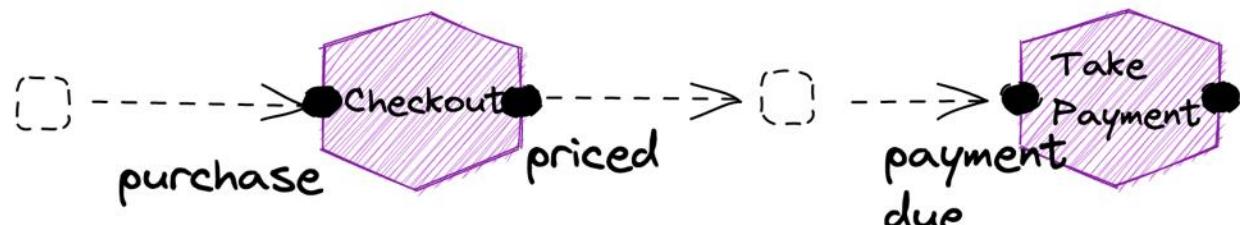
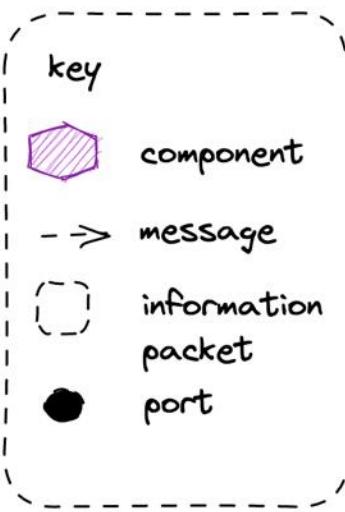


Message Passing is an asynchronous method of communication – both parties do not have to be simultaneously present for communication to occur, instead mail is delivered to ‘mailbox’ of some form for later retrieval.



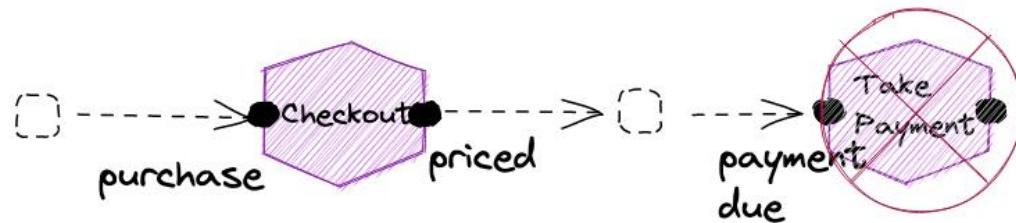
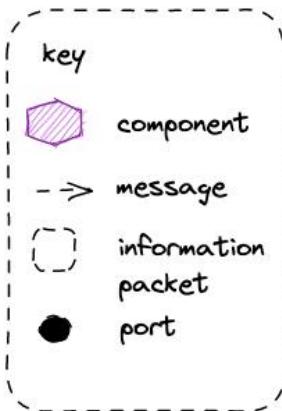
Service should focus on a transformation to the data
 Usually a verb + noun combination
 Place Order or Onboard Restaurant





Coordinate Dataflow A reactive *principle* where we “orchestrate a continuous steady flow of information” focusing on division by behavior, not structure

Partitioning Partition to take advantage of parallelism in the system, tasks that could be done concurrently with each other

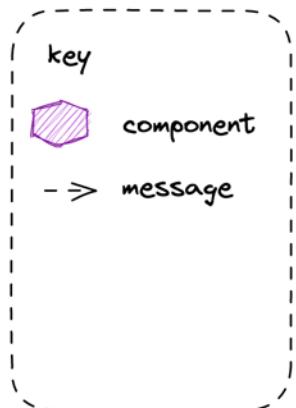


Work Queues on a Fault

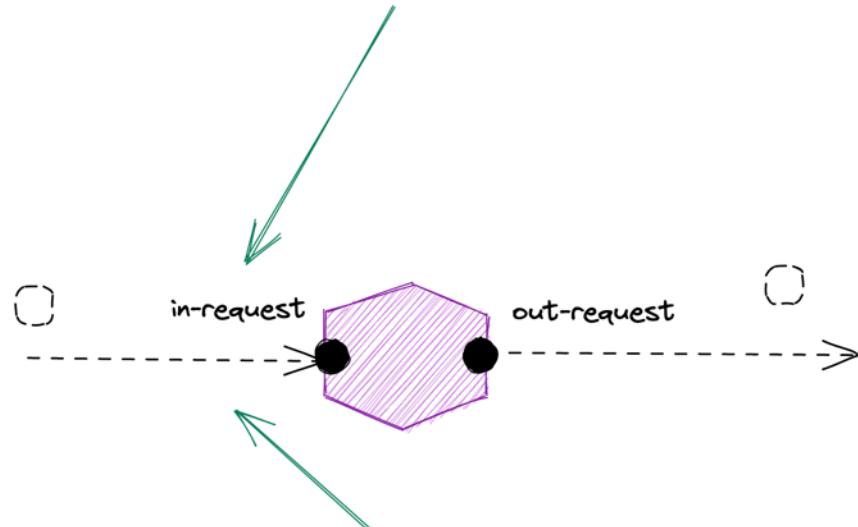
If Take Payment is not available, we can queue requests. When it starts up, Take Payment can process those requests and respond as normal.

Bulkheads: In an asynchronous conversation a fault does not propagate back up the chain – we have a bulkhead that protects us against failure

Push => open socket, middleware calls us as messages arrive
Pull => We poll for messages



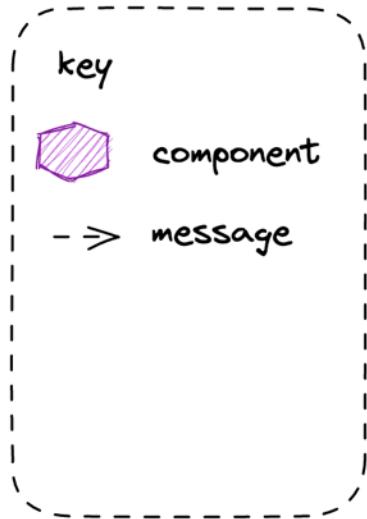
Message Passing



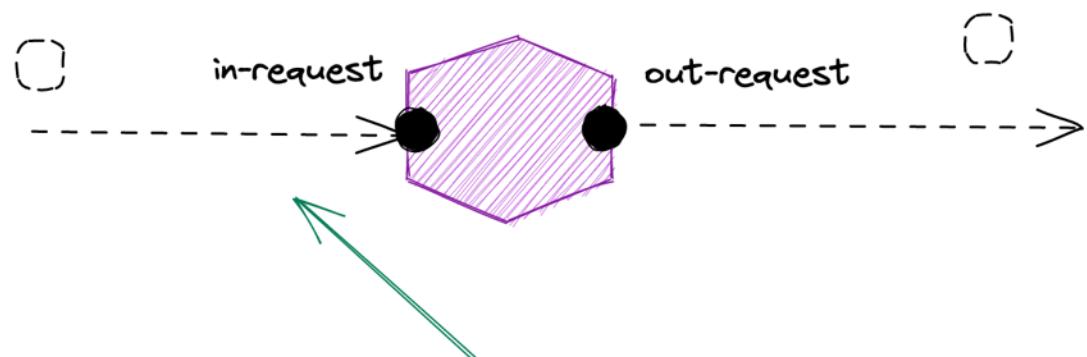
Backpressure

Push => Need to control how many messages are "in flight" with us
Pull => We control the rate at which we poll

Blocking Retry: Blocking Retry – when we keep trying to process a message, such as when we cannot connect to a DB and retry, creates backpressure as we slow consumption.

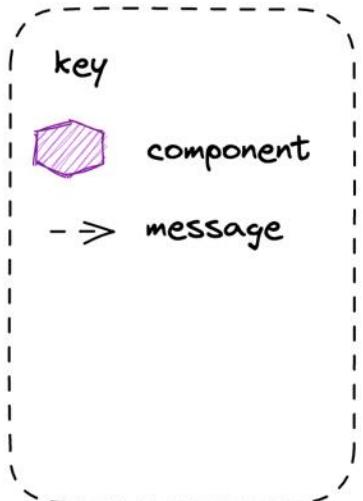


Message Passing

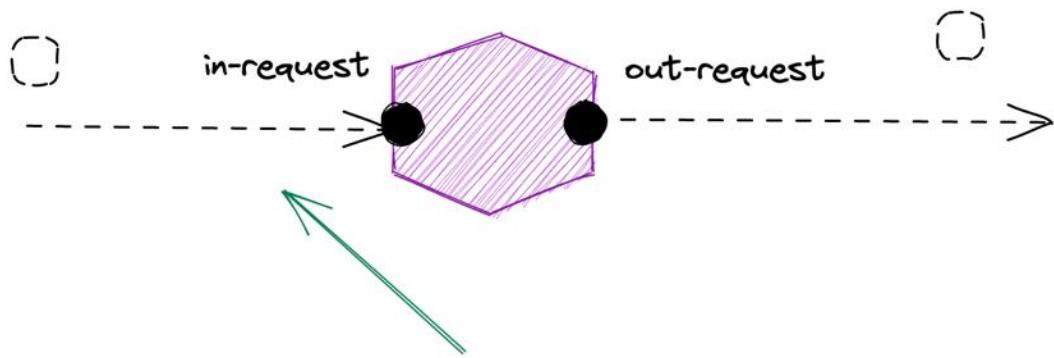


Load Shedding

Sometimes we may prefer to simply drop messages or shed load, when there is a fault.

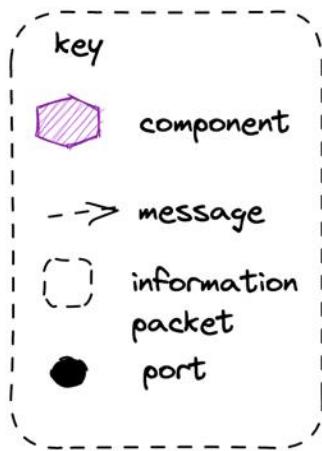


Message Passing



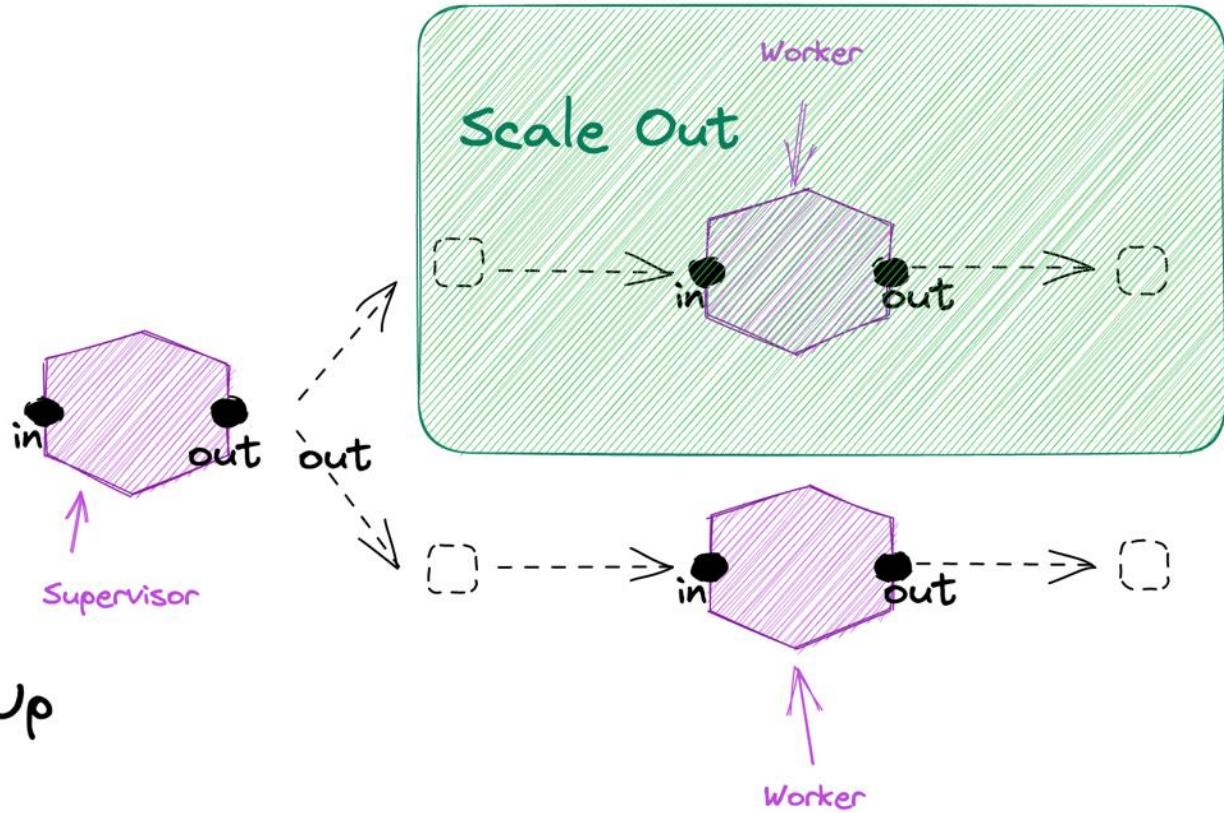
Circuit Breaker

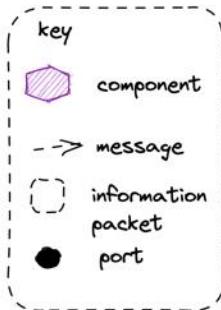
We may want to stop consuming, due to a fault.
We can periodically let a message through to see if the fault has cleared.



Scale Out, not Up

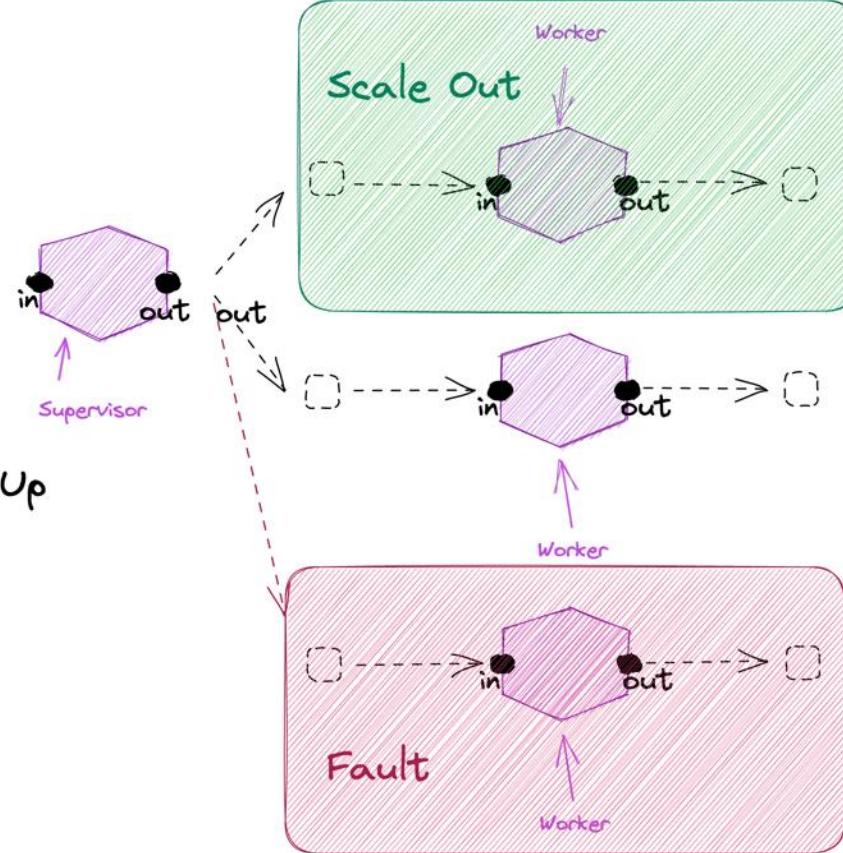
12-factor et al.



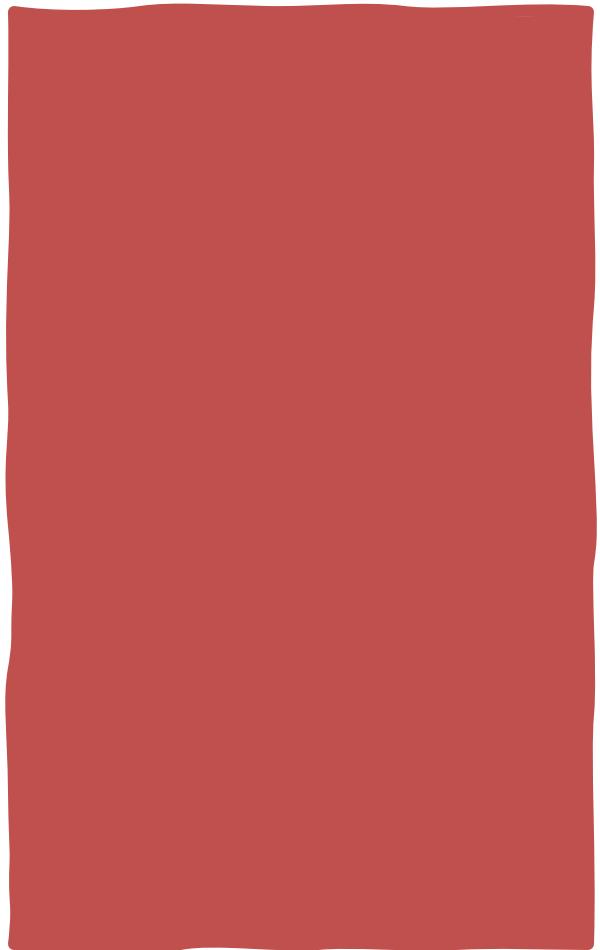


Scale Out, not Up

12-factor et al.



Putting It Together



EXERCISE MATERIAL

Flow

- Readme
- Slides



DON'T PANIC

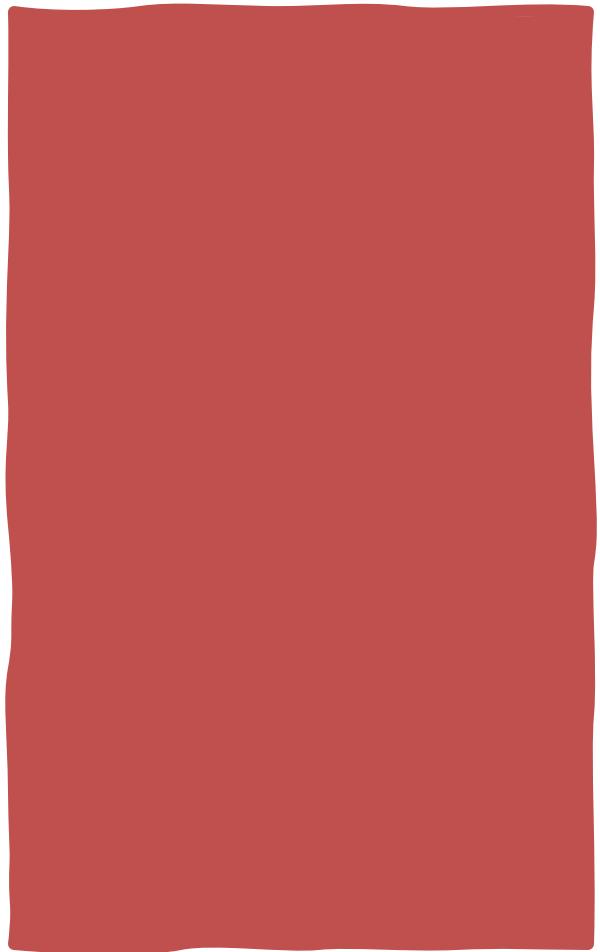
PROCESS AUTOMATION

What is a microservice?

SOA is focused on business *processes*. These *processes* are performed in different steps (also called *activities* or *tasks*) on different systems. The primary goal of a **service** is to represent a “natural” step of business functionality. That is, according to the domain for which it’s provided, *a service should represent a self-contained functionality that corresponds to a real-world business activity*.

Josuttis, Nicolai M.. SOA in Practice: The Art of Distributed System Design .
O'Reilly Media. Kindle Edition.

BPMN

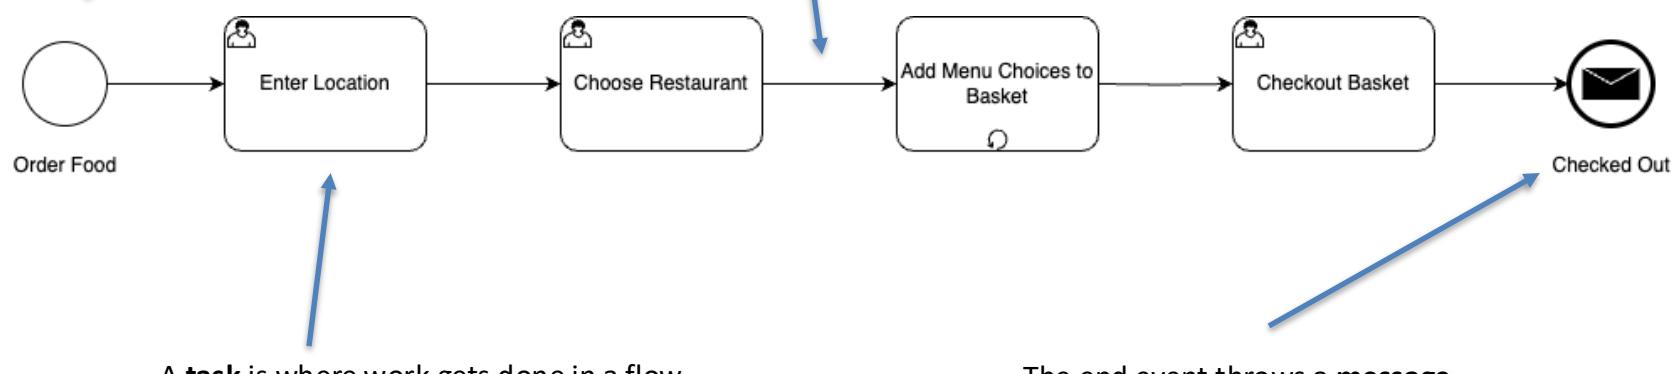


A event is something that begins, ends or interrupts a flow

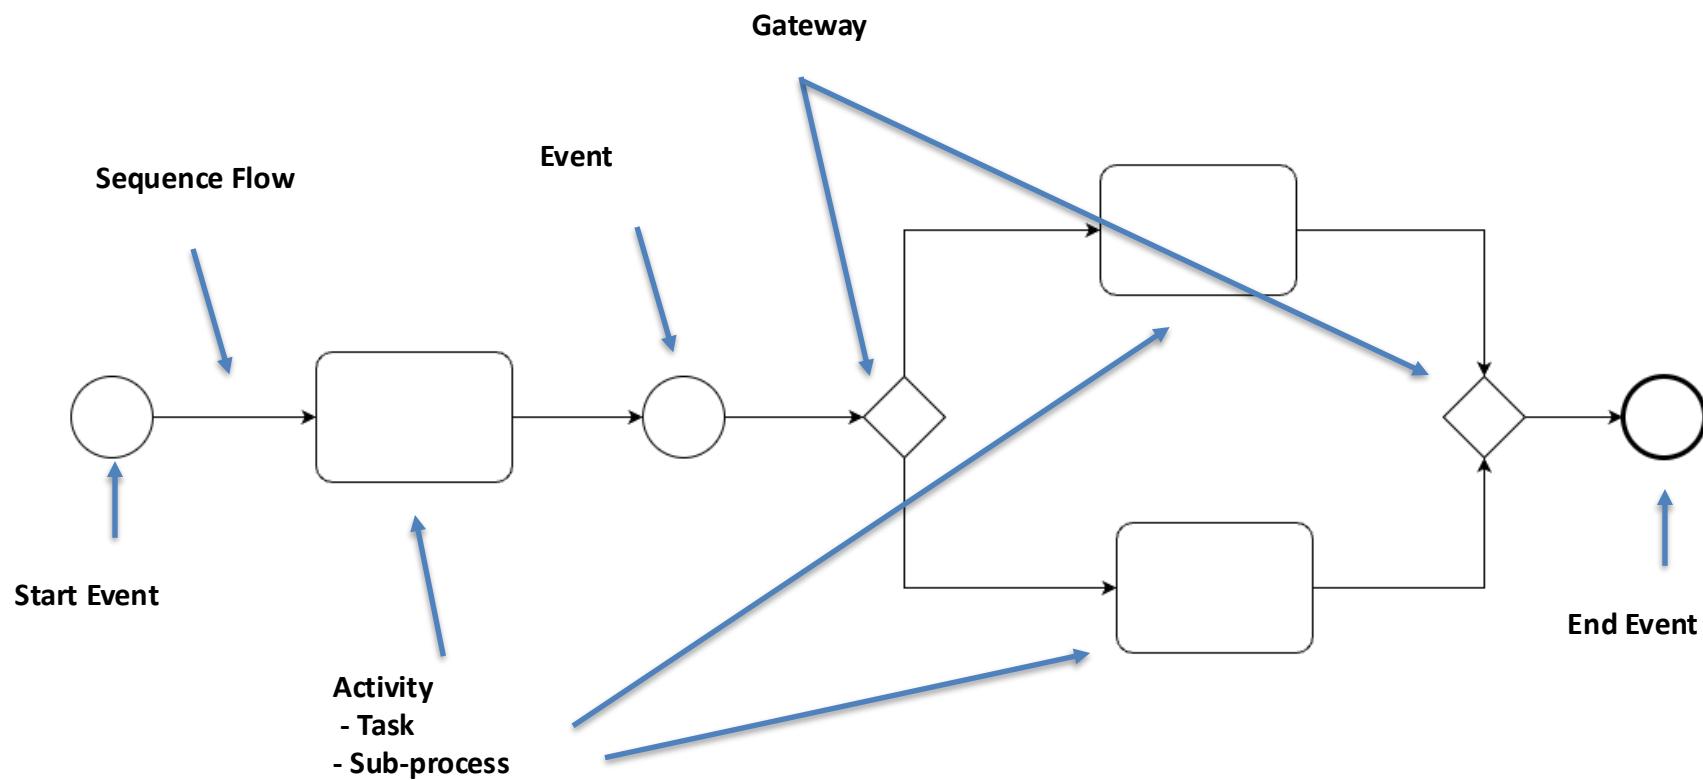
A **start** event begins a flow

An **end** event terminates a flow

In a **sequence** flow arrows show
The path of execution



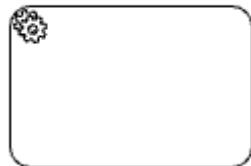
BPMN stands for *Business Development Management and Notation*. It is a visual language for diagramming a business processes in a clear, standardized and comprehensive way.



A BPMN diagram consists of a set of basic elements



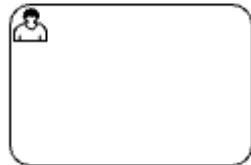
Generic Task: basic task



Service Task: uses a Service



Receive Task: waits to receive a message from another participant



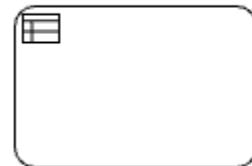
User Task: waits for human completion of task (via software)



Send Task: sends a Message to another participant



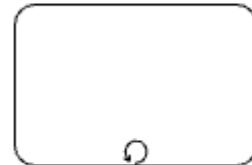
Manual Task: waits for human completion of task (not via software)



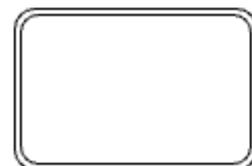
Business Rule Task: executes a business rules engine



Script Task: executes code

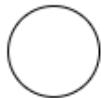


Loop: this task loops



Transaction: this task is a transaction

Task: an atomic activity. By implication, cannot be broken down further.



None Event: typically start or end,
no defined trigger



Message Event: a message is
received from another participant



Time Event: a date or time
triggers a process



Escalation Event: triggers a sub-
process to expedite completion of
an activity



Compensation Event: triggers a
compensation sub-process



Conditional Event: triggers when
a conditional expression becomes
true



Signal Event: a broadcast event
triggers flows

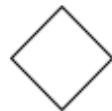


Parallel Event: multiple events
required to trigger a sub-process



Cancel Event: used to trigger a
sub-process with the abnormal
termination of a transaction

Event: something that happens within a flow: starts it, ends it, or interrupts it.



Exclusive: creates alternative paths based upon an expression, only one taken



Inclusive: creates alternative paths based upon an expression, may be taken in parallel



Parallel: split or join paths



Complex: a complex expression determines what paths are taken



Event: instead of an expression, an event determines what path is taken

Gateway: controls how sequence flow branches and converges

Sequence Flow 

Message Flow 

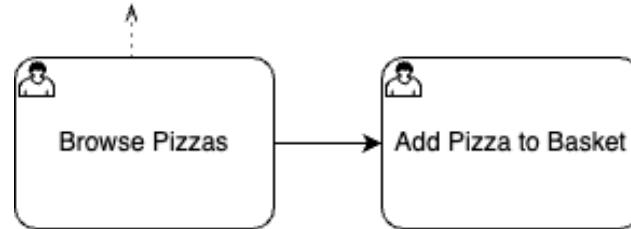
Connecting Object: connects activities

Workflow Patterns

- Workflow Patterns: *W.M.P. van der Aalst, A.H.M. ter Hofstede, B. Kiepuszewski, A.P. Barros, 2000*
 - **5 Basic Patterns**
 - 15 Advanced Patterns
- Workflow Control Flow Patterns: A Revised View: *W.M.P. van der Aalst, Nataliya Mulyar, Nick Russell, A.H.M. ter Hofstede, 2007*
 - 23 New Patterns

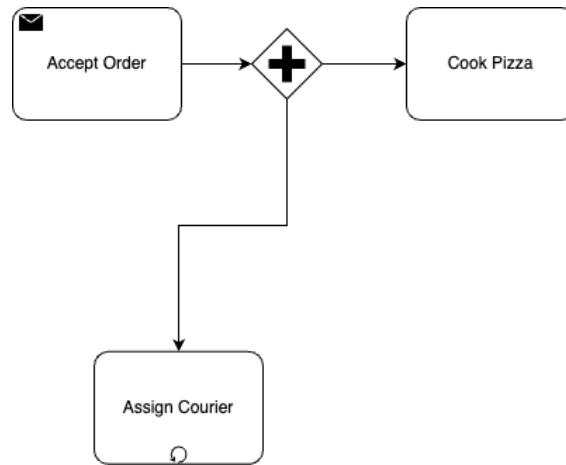
Pattern 1 – Sequence

- **Description:** One activity follows another.
- **BPMN:** Connect two tasks with a **sequence flow arrow**.
- Easy to model and execute.
 - **Example (Customer):** Browse Pizza → Add Pizza to Basket.



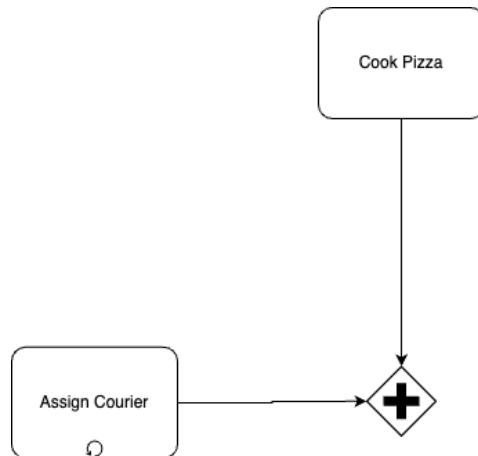
Pattern 2 – Parallel Split

- **Description:** One path splits into two or more **concurrent branches**.
 - **BPMN:** Use a **Parallel Gateway** (+) to fork.
- **Example (Pizza Shop):** Assign Courier + Cook Pizza.



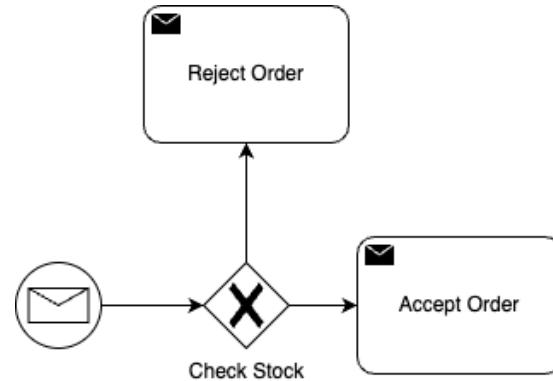
Pattern 3 – Synchronization (Join)

- **Description:** Wait until multiple concurrent branches complete.
- **BPMN:** Use a **Parallel Gateway** again to join.
 - **Example (Pizza Shop):** Assign Courier + Cook Pizza.



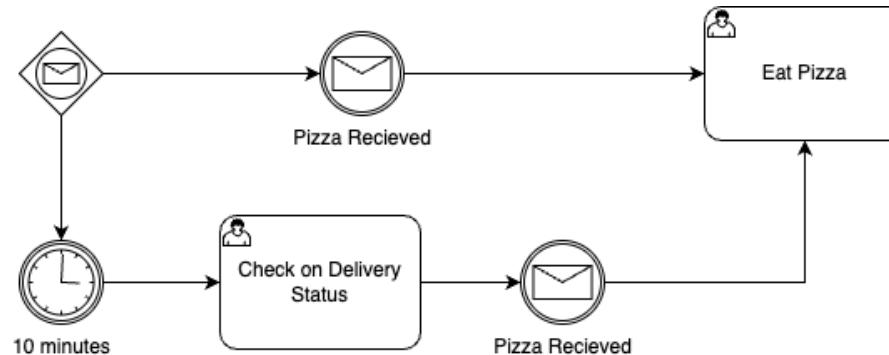
Pattern 4 – Exclusive Choice

- **Description:** Choose one path based on a condition.
- **BPMN:** Use an **Exclusive Gateway (X)** with condition expressions.
 - **Example (Pizza Shop):** Check Availability → Accept Order X Reject Order.

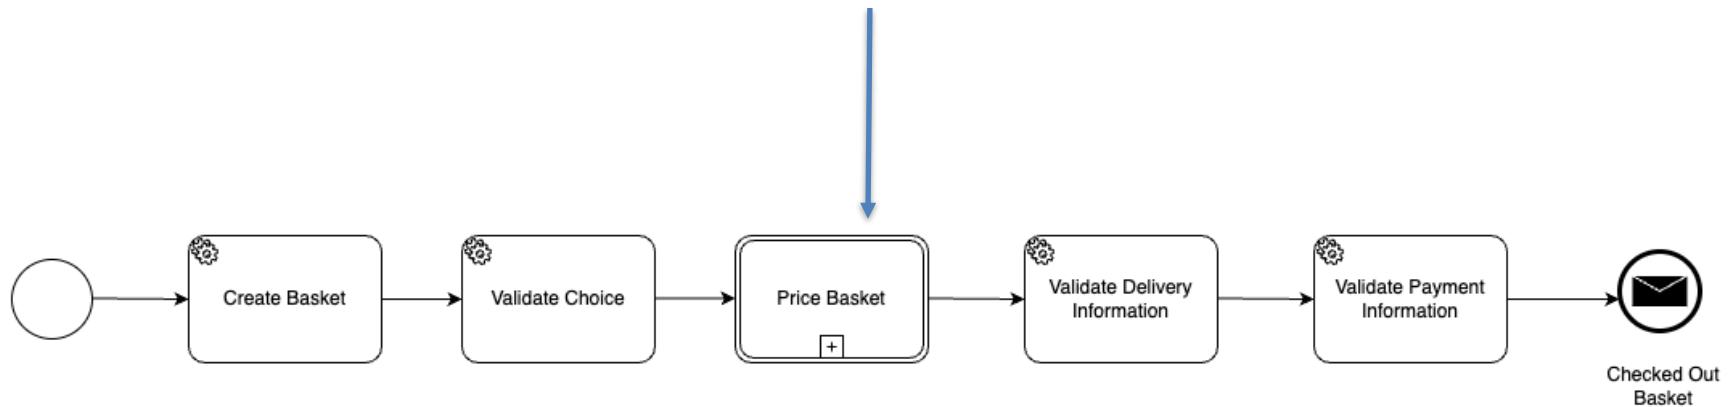


Pattern 5 – Simple Merge

- **Description:** Merge **non-concurrent** paths back into one.
- **BPMN:** Use a **converging sequence flow**; **no gateway is needed** if no synchronization is required.
 - **Example:** Multiple paths lead to Eat Pizza.



A **Process** describes a sequence or flow of **Activities** for carrying out work. In **BPMN** a **Process** is depicted as a graph of Flow Elements, a **Sequence Flow**, which is a set of **Activities**, **Events**, **Gateways**



A **Process** is an ***Orchestration***:

- The **Sequence Flow** represents control of a process
- The **token** represents state associated with an instance
- In implementation, generally a process orchestration lives within an address space i.e. is not distributed, either an embedded workflow or external process manager.

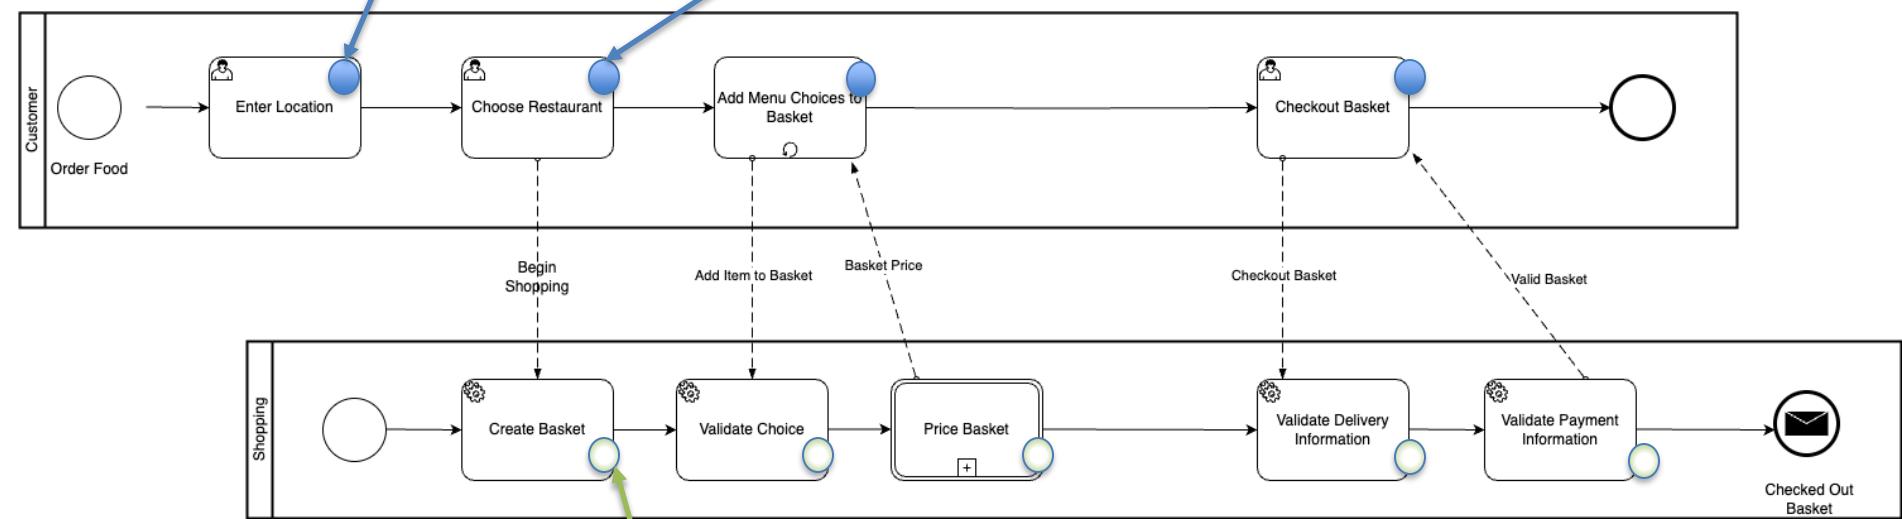
What is Orchestration?

Focused on a **single participant's perspective**

- Like writing *your own script*
- Includes control flow, state, and decisions
- All logic is local to the orchestrator
- Often implemented via state machines or workflow engines.

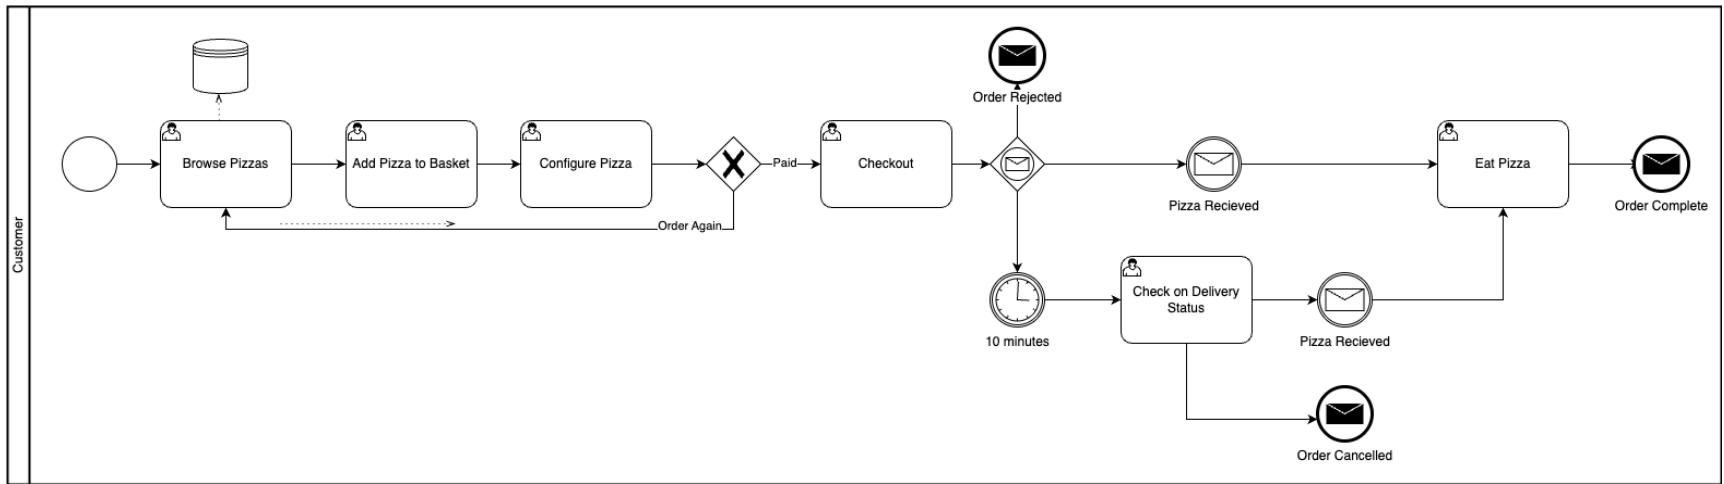
A **token** represents an instance of a process

A **token** flows down the process; it is the state of the process for that instance



A **token** flows down the process; it is the state of the process for that instance

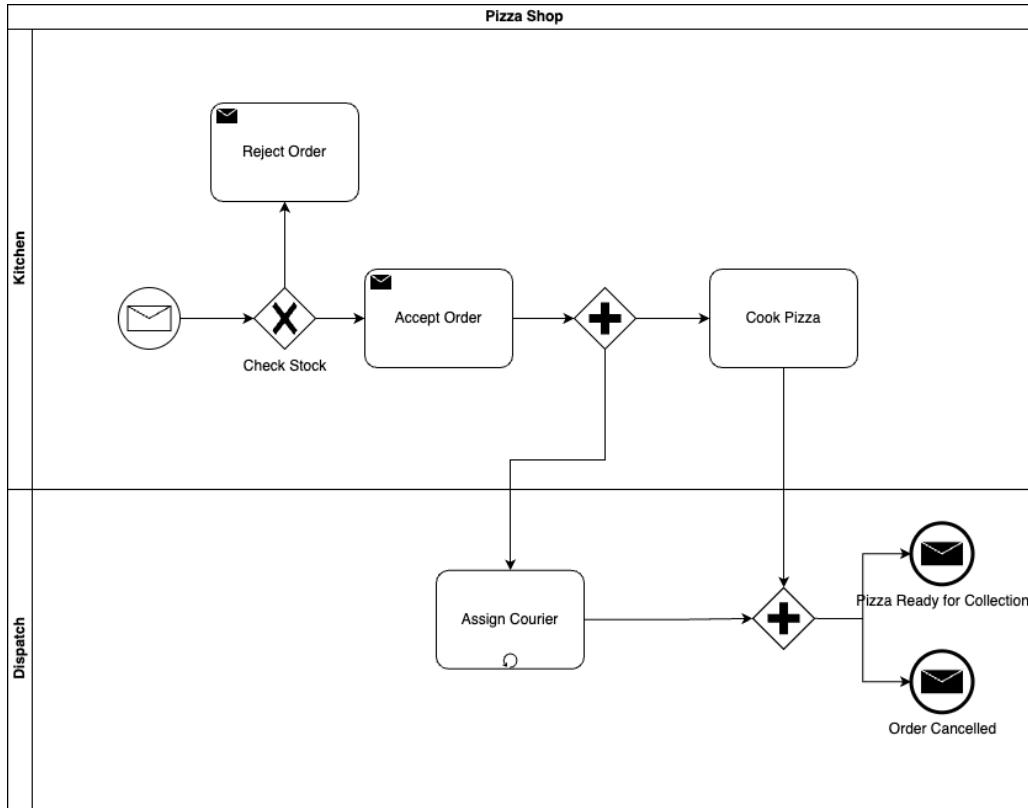
Pizza Example – BPMN Orchestration



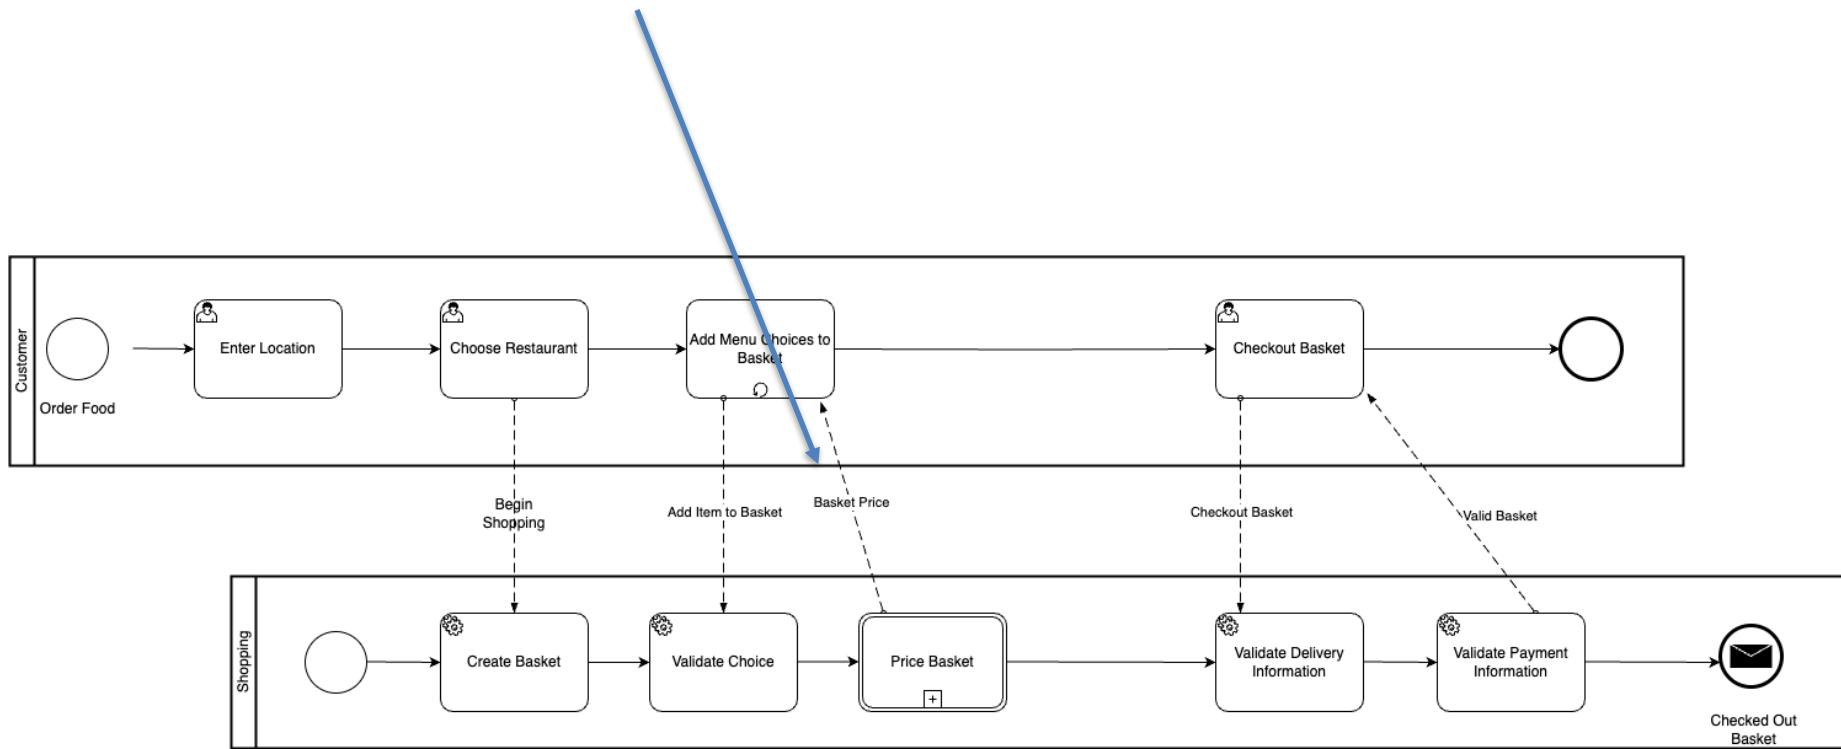
Customer Pool

Pizza Example – BPMN Orchestration

Pizza Shop Pool



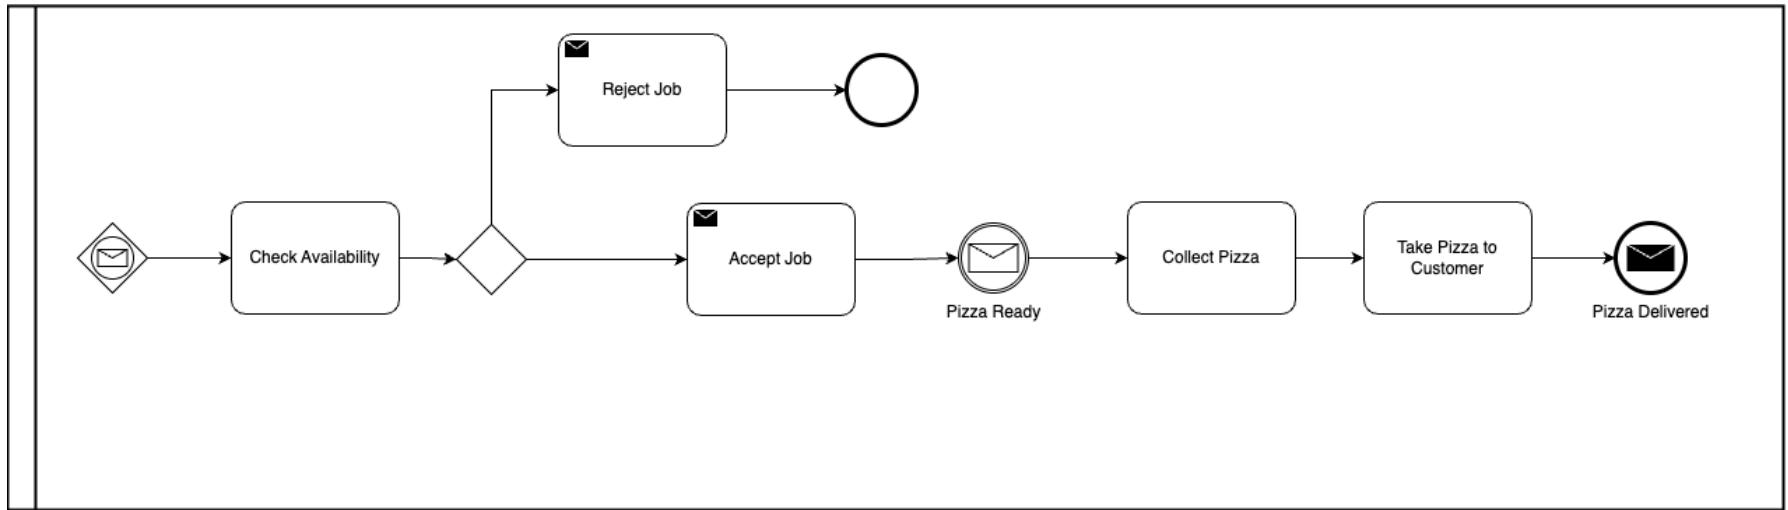
A **Collaboration** is where we have multiple participants.



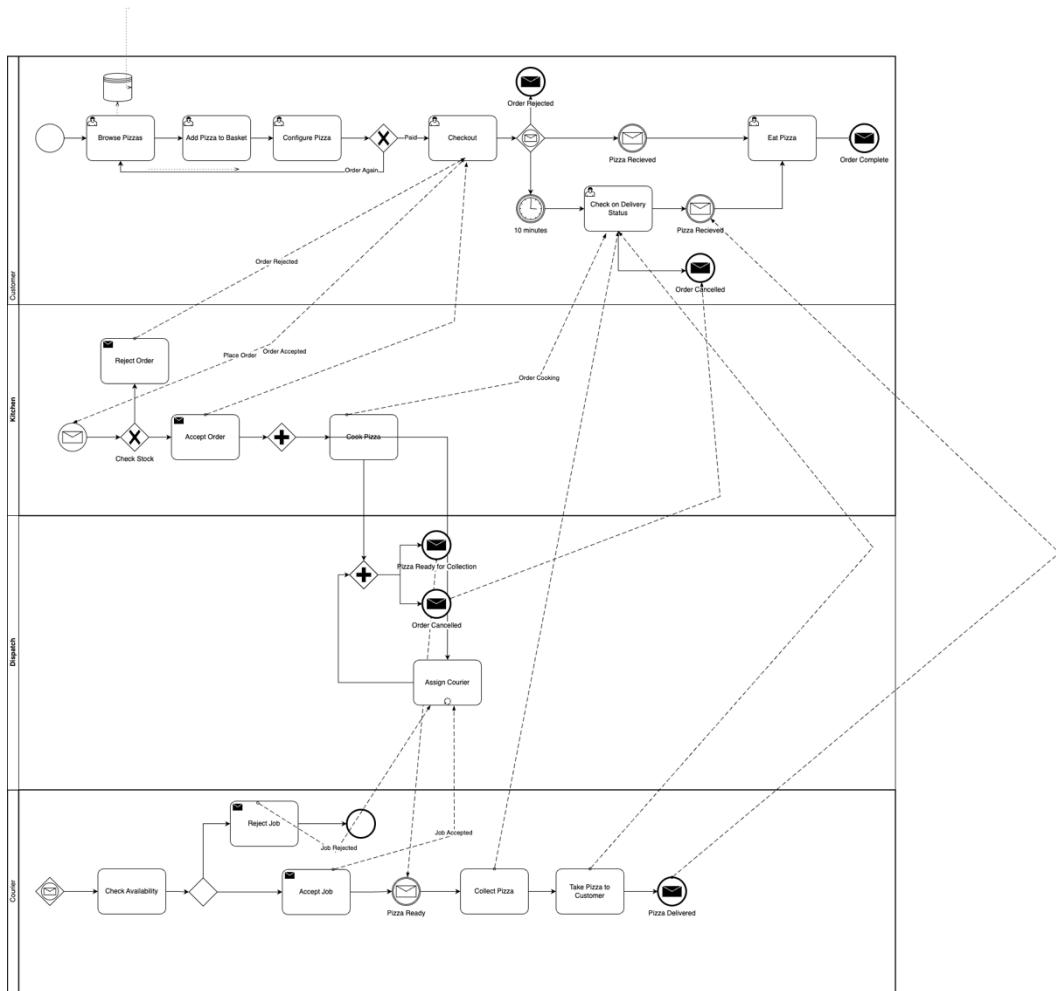
The message exchange between participants in a **Collaboration** is a **Choreography**

- In implementation, generally, a choreography is the flow of messages between participants.
- Within each orchestrated process, there are message events that are caught or raised
- But Processes react to what happens in the Choreography, no one owns it (it does not have its own tokens or state)

Pizza Example – BPMN Orchestration



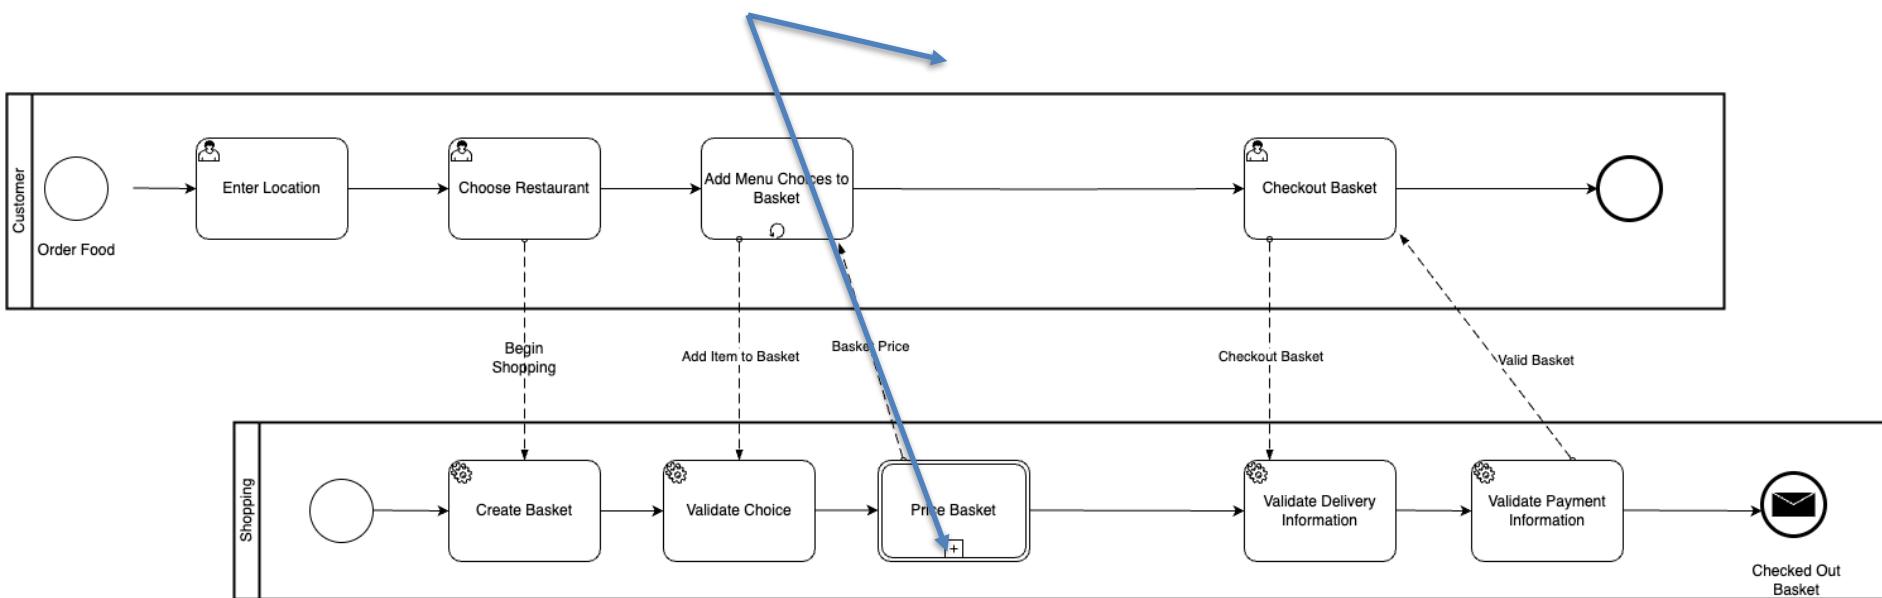
Courier Pool



Pizza Example – Pools and Lanes

A **Pool** is participant in a collaboration

- A participant can be a role or an organization
- Not required for the Main Internal Pool => modeler's own organization
- Required for other organization's processes
- A Lane distinguishes sequences of Activities in the Pool
- White Box - can see process; Black Box - cannot see process
- Activities within the Pool organised by sequence flow

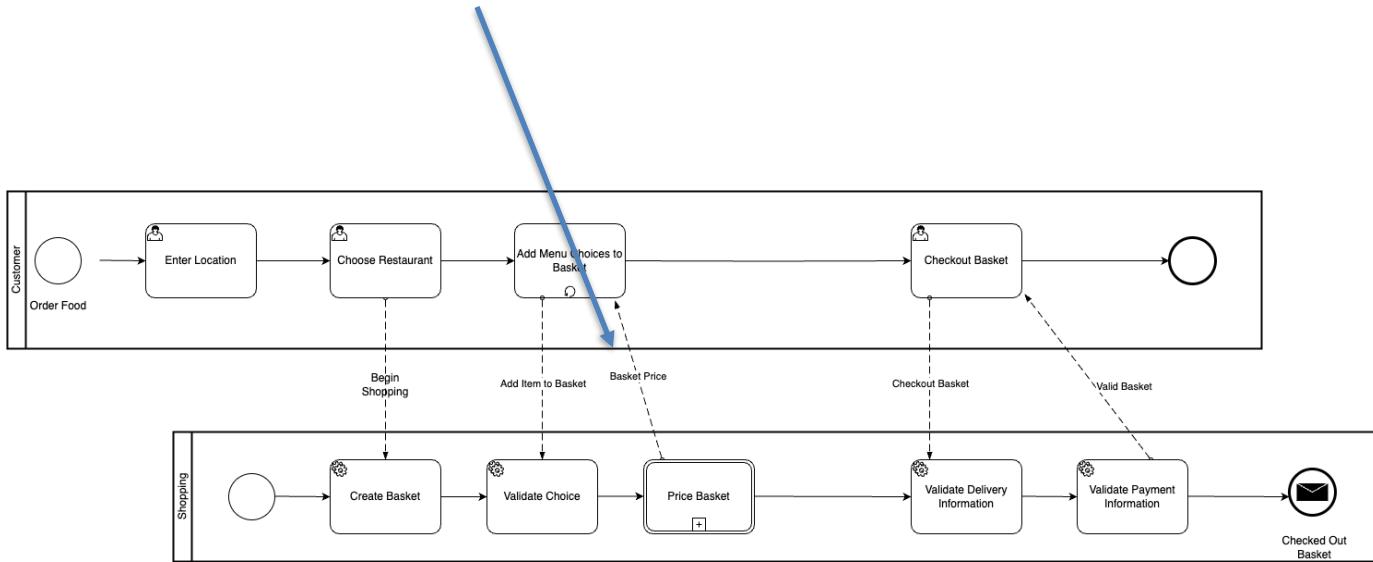


What is Collaboration

Focused on how **participants interact**

- Interaction might be within a workflow engine.
 - But this creates coupling, and both participants must run in the workflow engine.
- Interaction might be via an API
 - From an event-driven perspective, we are most interested in this model.

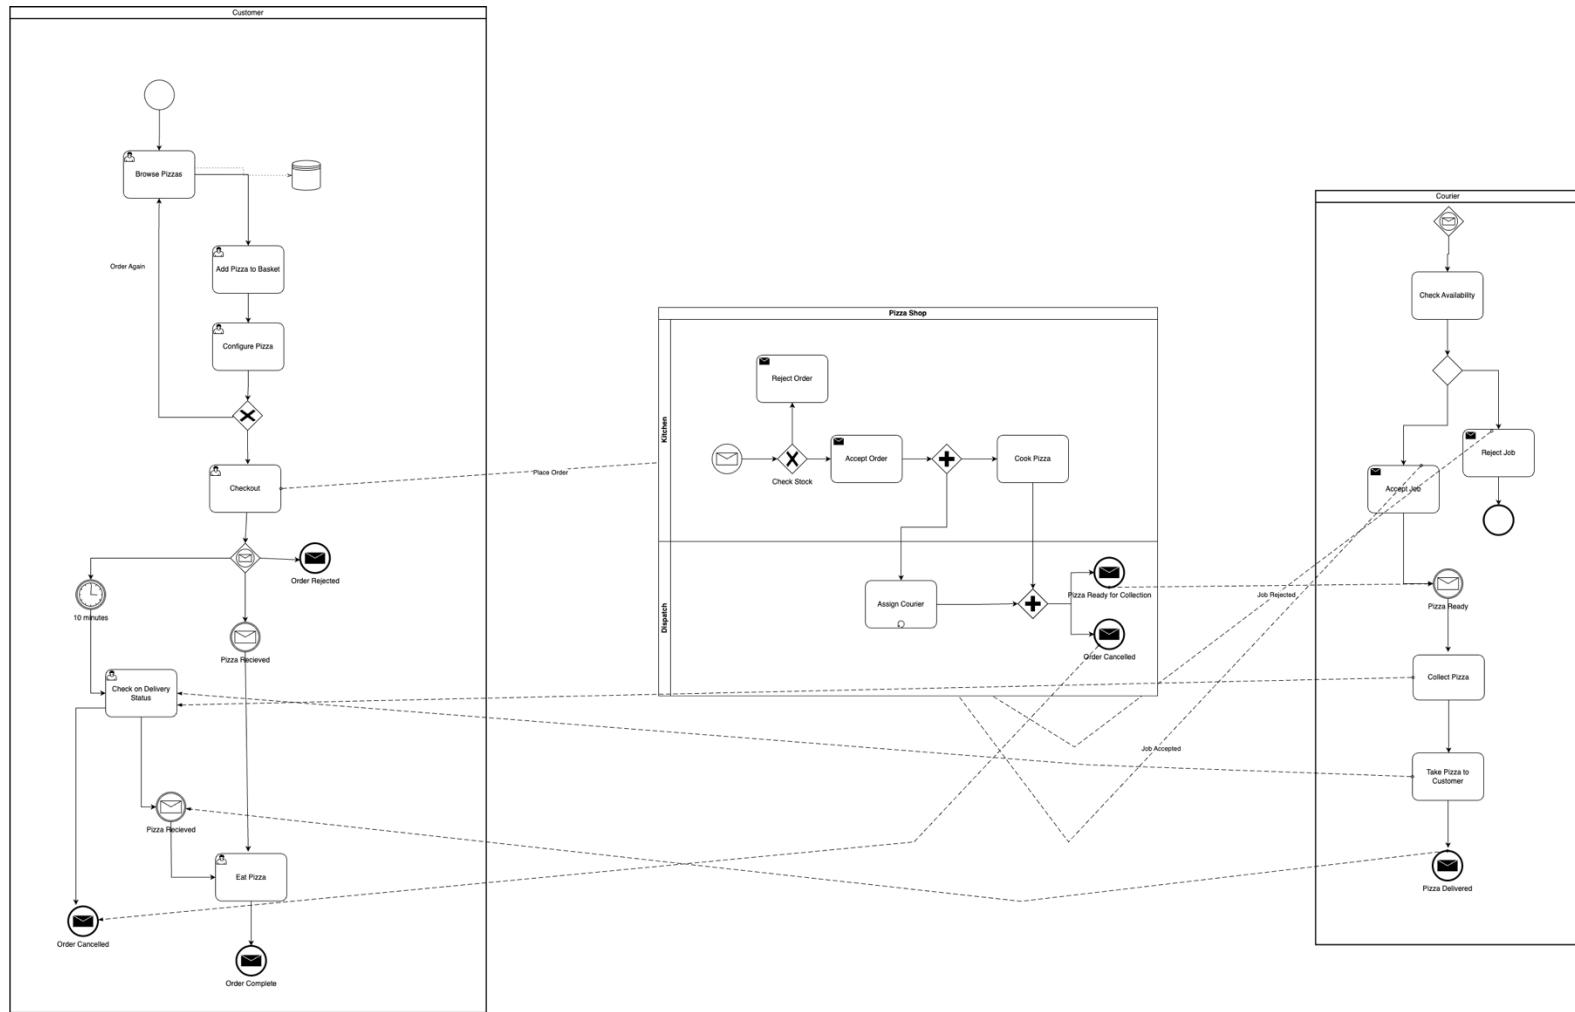
A **Collaboration** is where we have multiple participants.



The message exchange between participants in a **Collaboration** is a **Choreography**

- In implementation, generally a choreography is the flow of messages between participants.
Within each orchestrated process there are message events that are caught or raised
- But Processes react to what happens in the Choreography, no one owns it (it does not have its own tokens or state)

Pizza Shop Collaboration

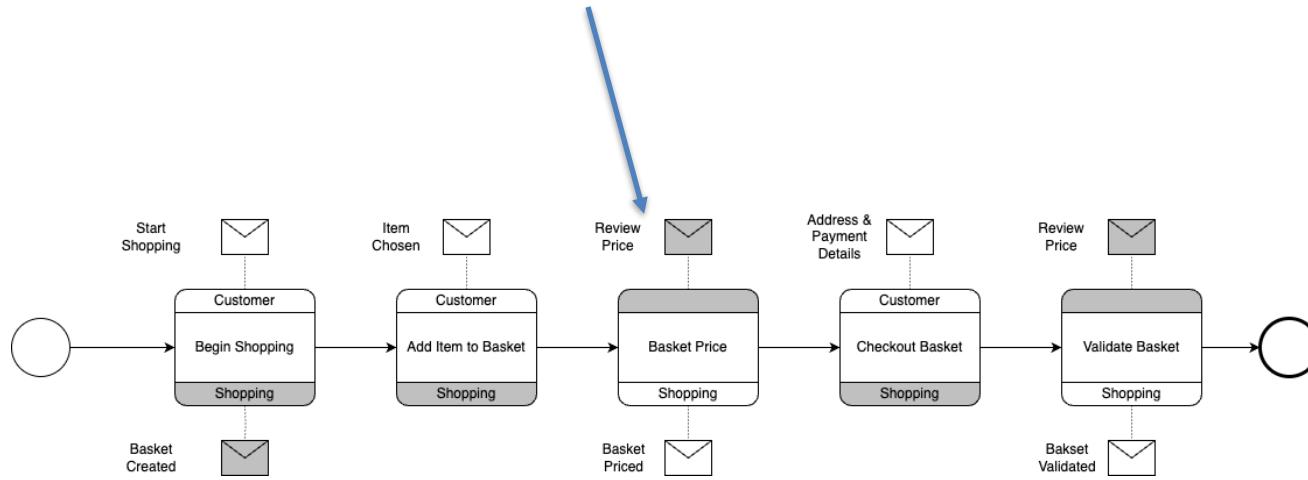


What is Choreography?

Focused on **interactions** across multiple participants

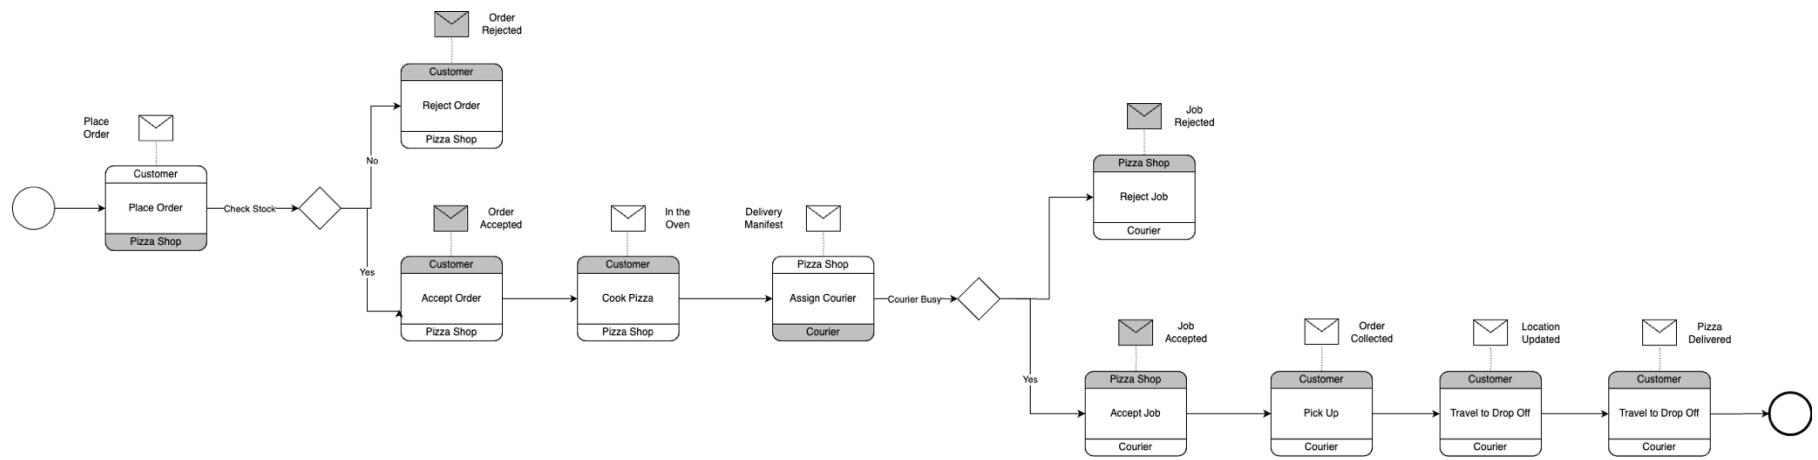
- Like describing a **dance**, no single owner of the flow.
- Defines *who talks to whom, in what order*
- No centralized control
- Cannot access shared internal data/state

A **Choreography** describes a sequence or flow of **Activities** between two different participants. In BPMN a **Choreography** is depicted as a graph of Flow Elements, a **Message Flow**, which is a set of **Activities**, **Events**, **Gateways**



A Conversation is a logical association of messages that can all be correlated
 - Correlation Keys are used to associate messages that form part of the same Conversation
 - This may be existing message data
 - The first task in a conversation MUST populate the conversation id

Pizza Example - BPMN Choreography



Pizza Order Choreography

Messaging and Eventing

A common notion within EDA:

- Messaging (commands) is orchestration; Eventing is choreography
- It is not a helpful rule.
- An orchestration is a sequence, with attendant session state [token]
 - Frequently, it is not distributed and *is within* an executable process
 - A routing slip is a distributed orchestration that utilizes message attributes as the engine for application state.
- Choreography is the flow of control and information between orchestrations.
 - Frequently, it is out of process.
 - Can use messaging or eventing as required.

Event Chaining

A sensible balance of orchestrated bounded contexts and choreography between will reduce event chaining.

- However, Event Pinball can still be a problem between orchestrated services.
- In-Out (or Robust In-Only)
 - Use of In-* patterns with explicit transfer of control reduces the “pinball effect” in choreography
 - Creates **causality** in workflows by awaiting a response; you don’t send an event, and wait for a response, you send a command and wait for a response.
- Out-Only
 - Often better at the termination of a workflow, when the work is done, or to broadcast intermediate states with no expectation of response.

Allocating Responsibility for Collaborations

Commands/Messaging – Transfer of Control/Data.

- Typically, this involves delegating responsibility for part of a workflow.
- In-Out. Expected to Return
- In-Only. Not Expected to Return. But May on Error (Robust In-Only).

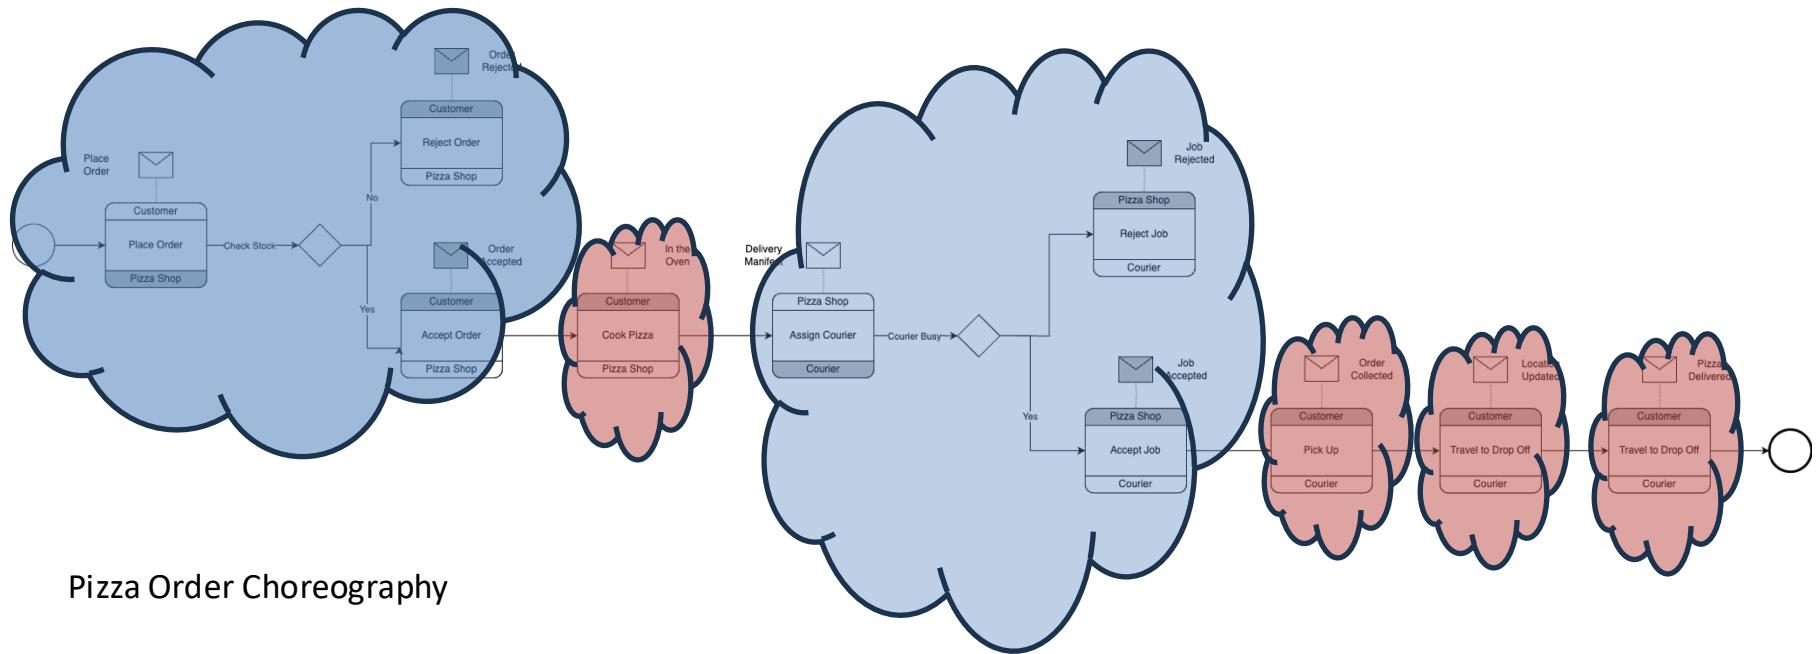
Events – Communicate the Outcome of a Workflow

- Typically, a workflow is complete or a milestone is broadcast for observation.
- Out-Only. No Expected Return.

Workflow code is Glue.

- Glue does not execute the domain; it is not the behaviour
- Instead, glue calls something that has the behaviour

Pizza Example – Messaging and Eventing



Aspect	Orchestration	Choreography
Control	Centralized (one participant)	Decentralized (shared)
BPMN Shape	Inside a Pool	Between Pools
Perspective	One party's process	Multi-party message exchange
Data Access	Internal	Shared messages
Execution Tool	API calls, Handlers, Workflow engine	Messaging protocol
Example	Pizza Shop coordinates cook + courier	Customer + Shop + Courier interact

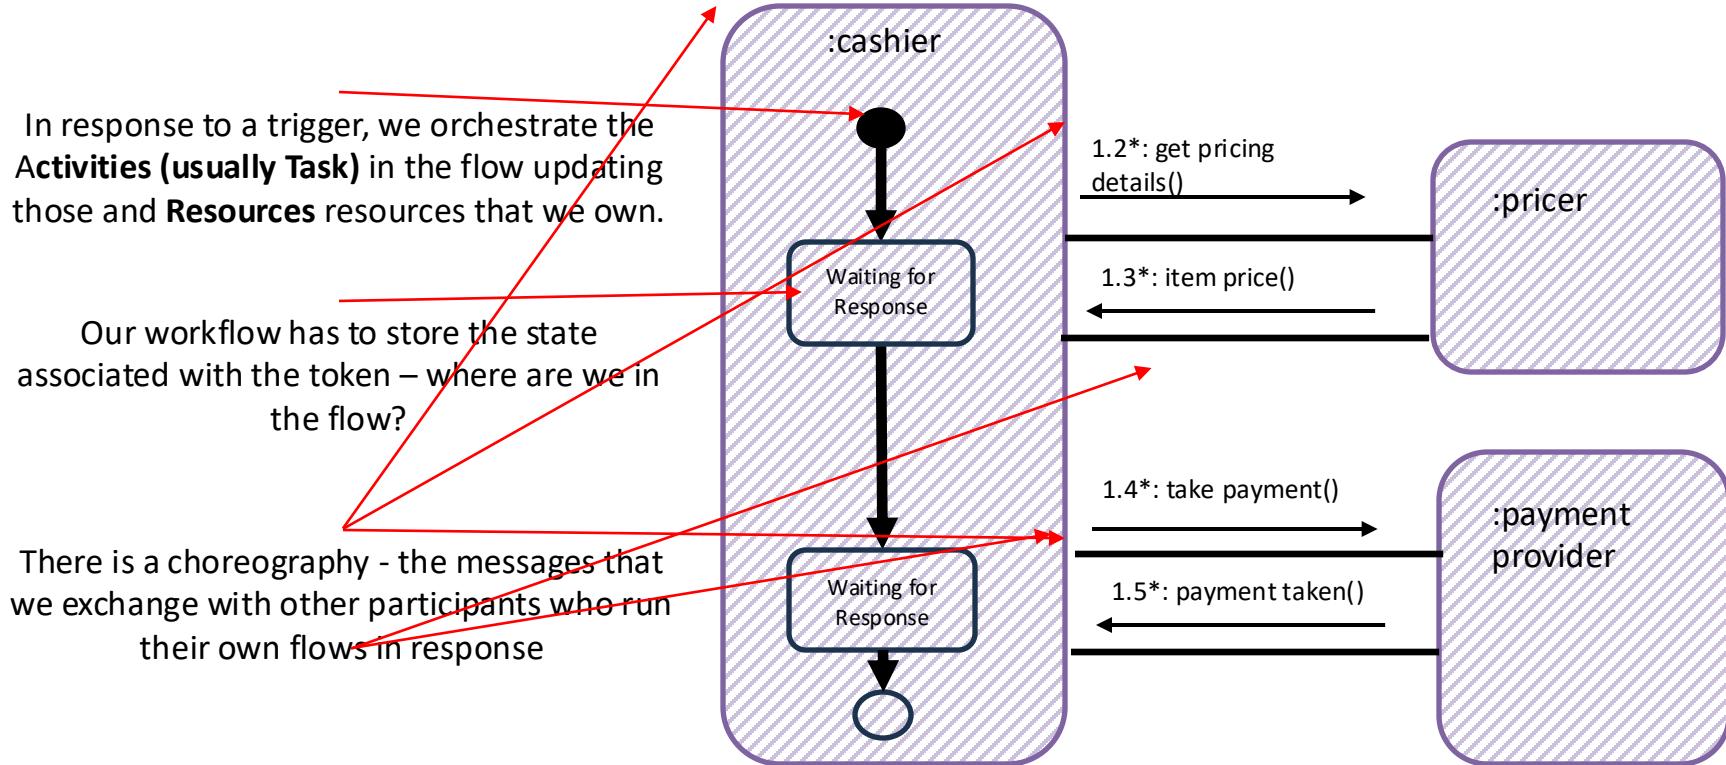
Orchestration vs. Choreography Summary

Activities and Resources

Broadly, a software component manages *activities* and *resources* [Pat Helland].

- **Resources:** domain concepts we manage, like restaurants, couriers, and customers.
- **Activities:** one or more sequences for our interaction with our resources.
- In response to the *token* moving through the sequence we update resources, and the activity to indicate our progress through the sequence.

The implementation question is: where does our activity state live, and who updates it?



Activity. Usually the task we need to perform. Most frequently a service task – our workflow engine just contains glue code that calls the domain logic internally, or messages another application. Might be a framework, might be bespoke state machine or pipes and filters.

Resources. Code and data used to manage shared items coordinated across multiple activities such as widgets in inventory or space on a truck needed for a shipment. Our domain model.

Durable Execution

Long-running processes must:

- Survive restarts
- Handle retries and failures
- Resume after waiting (e.g., waiting for payment or a courier)

Durable execution = we **persist activity state** to indicate which steps in a sequence have been completed, when we are awaiting an event etc.

Handlers + Activity State Updates

Handlers process events (change resources) and update activity state.

- Activity state is often implicit, such as within an aggregate.
- Durable because the state is stored (e.g., database), but **the control flow is implicit**
- Great for simple workflows, but logic becomes scattered across handlers
- Relies on guaranteed delivery to store work for retry unless ack'd
- Relies on Transactional Messaging (Outbox)
- Becomes complex once we introduce retry, circuit breakers & compensation
- Becomes complex once we have split, join, choice, merge (not a sequence)

```
1 public async Task<Order> HandleAsync(Order order, CancellationToken cancellationToken)
2 {
3     var posts = new List<Guid>();
4
5     var tx = await _uow.BeginOrGetTransactionAsync(cancellationToken);
6     try
7     {
8
9         //should save the order
10        orderRepository.Received(_uow, order);
11
12        //then raise accept/reject
13        bool canMake = await stock.Check(order);
14        if (canMake)
15            posts.Add(await _postBox.DepositPostAsync(new OrderAccepted(order.Id), cancellationToken:cancellationToken));
16        else
17            posts.Add(await _postBox.DepositPostAsync(new OrderRejected(order.Id), cancellationToken:cancellationToken));
18
19        await tx.CommitAsync(cancellationToken);
20
21        //kick off cook and assign courier tasks
22        await cookRequests.Writer.WriteAsync(new CookRequest(order), cancellationToken);
23        await deliveryRequests.Writer.WriteAsync(new DeliveryRequest(AssignCourier(), order), cancellationToken);
24
25    }
26    catch (Exception e)
27    {
28        _logger.LogError(e, "Exception thrown handling Add Greeting request");
29        //it went wrong, rollback the entity change and the downstream message
30        await tx.RollbackAsync(cancellationToken);
31        return await base.HandleAsync(addGreeting, cancellationToken);
32    }
33
34
35    await _postBox.ClearOutboxAsync(posts, cancellationToken:cancellationToken);
36
37    return await base.HandleAsync(addGreeting, cancellationToken);
38 }
```

(Fault) Handlers + Activity State Updates

How do we handle compensation flows?

- May retry the handler, and only if it fails X times, run the fallback.
- On fault, choose to run a “fallback handler” – initiates compensating action.
- Typically, needs to reverse actions that have occurred so far; write an “undo” for each “do”.
- Typically, a straightforward approach is to send a fault message back upstream and undo it upon receipt.
- Fault Message
 - For reliable In-Only, a fault channel is required, which is only used for compensation.
 - For In-Out, the message sent to the response channel should indicate fault, not success.

State Machine + Activity State Updates

When the interaction of Handlers becomes complex, we can make the activity explicit.

- A State Machine represents the activity and triggers actions in response to transitions.
- Transitions are caused by receiving a message.
- The handler receives a message and loads the corresponding state machine associated with the conversation ID.
- This triggers the transition denoted by the receipt of the message; the state machine runs any code associated with that transition.
- Save the new state of the state machine and ack the message.



```
1  public OrderStateMachine	ILogger<OrderStateMachine> logger)
2  {
3      InstanceState(x => x.CurrentState);
4      ConfigureCorrelationIds();
5
6      Initially(
7          When(OnOrderSubmitted) // Receive the request
8              .Then(x => logger.LogInformation($"Order submitted"))
9              .SendAsync(new Uri("queue:stock-check"), context => context.Init<DebitAccount>(new
10             {
11                 context.Message.OrderId,
12                 //... details
13
14             })) // Calls debit
15             .TransitionTo(PendingStock)
16         );
17
18     During(PendingStock, // While PendingStock, if...
19         When(OnHasStock) // Check has been succeeded?
20             .Then(x => logger.LogInformation($"Stock available"))
21             .SendAsync(new Uri("queue:order-accepted"), context => context.Init<CallPartner>(new
22             {
23                 context.Message.CorrelationId,
24                 context.Message.OrderId,
25                 //..details
26
27             })) // Calls the kitchen
28             .TransitionTo(AwaitingKitchen),
29         When(OnLackStock) // Stock check has been failed?
30             .Then(x => logger.LogInformation($"Stock failed"))
31             .TransitionTo(RejectOrder)
32         );
33
34     //...
35 }
```

(Fault) State Machine + Activity State Updates

How do we handle compensation flows?

- May retry the transition, and only if it fails X times, run the compensation.
- On fault, transition to a faulted state and initiate compensating flow to the “safe” state.
- Typically, this involves reversing actions that have occurred so far; write an “undo” for each “do”.
- Requires *durable execution* – if we fault partway through, on restart, we can resume.
- The “**saga**” pattern, when used by messaging frameworks, is typically a state machine.
 - In essence, guarantees that on a fault, the state machine will transition back to “safe”
 - Implementations run all fault transitions to safe before ack fault message.

Routing Slip + Activity State Updates

If we want to distribute steps in the sequence, we need to transport the activity state in the message.

- A routing slip is an envelope that indicates the activity state (next step, context, etc.).
- Think of it as “Message As the Engine of Application State”
- The handler updates the state when work is done and forwards according to the routing slip. Then acks.
- Complex if “next” is dynamic routing, as the code must understand how to execute.
- Complex because you are distributing the steps, hard to observe [Otel]

(Fault) Routing Slip + Activity State Updates

How do we handle compensation flows?

- Next is now the fault next, from this step; what do we need to reverse.
- Context is “faulted”
- The handler updates the state when work is done and forwards according to the routing slip. Then acks.

Workflow Engines + Activity State Updates

When we want more than just state transitions (split, join, merge, choice, etc.), use a workflow engine.

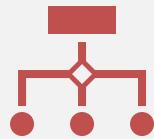
- A workflow engine explicitly models a sequence of activities and gateways.
- May directly support BPMN (or proprietary language); allows visual modelling.
- A step receives the token. The step is glue code: it invokes the action in the domain, calls another service, and sends a message.
- A step may wait for an event or a timer to signal that execution should resume. It may be the handler equivalent, or it may be triggered by the handler calling the engine.
- A job scheduler provides durable execution. Typically, any instance can resume a paused job – distributed.

(Fault) Workflow Engines + Activity State Updates

How do we handle compensation flows?

- Is the error a business error or a technical error?
- A known business error (i.e. card declined) is usually modelled explicitly as a branch via a gateway.
- A retry frequently handles technical faults. If the retry fails and a catch is present, then a fallback runs. If there is no catch, then either the workflow rolls back to the last wait state or terminates.

Embedded vs. External Workflow Engines



Embedded: The workflow runs in process and calls your code directly.



External: The workflow runs in its own process; it can execute code directly, but should favour calling services to get work done.

Workflow Engines: Lessons from SOA



SOA era favoured external workflow engines, Enterprise Service Buses (ESBs)



Business logic moved into the bus — tightly coupled, hard to test



Led to **Guerilla SOA**: reclaim domain logic into services, embrace simplicity

Smart Endpoints, Dumb Pipes

ESB products often include sophisticated facilities for message routing, choreography, transformation, and applying business rules. The microservice community favours an alternative approach: **smart endpoints and dumb pipes**.

<https://martinfowler.com/articles/microservices.html#SmartEndpointsAndDumbPipes>

Embedded vs. External Workflow Engines



Favor embedded Workflow engines when the workflow is within a bounded context; everything between bounded contexts is choreography

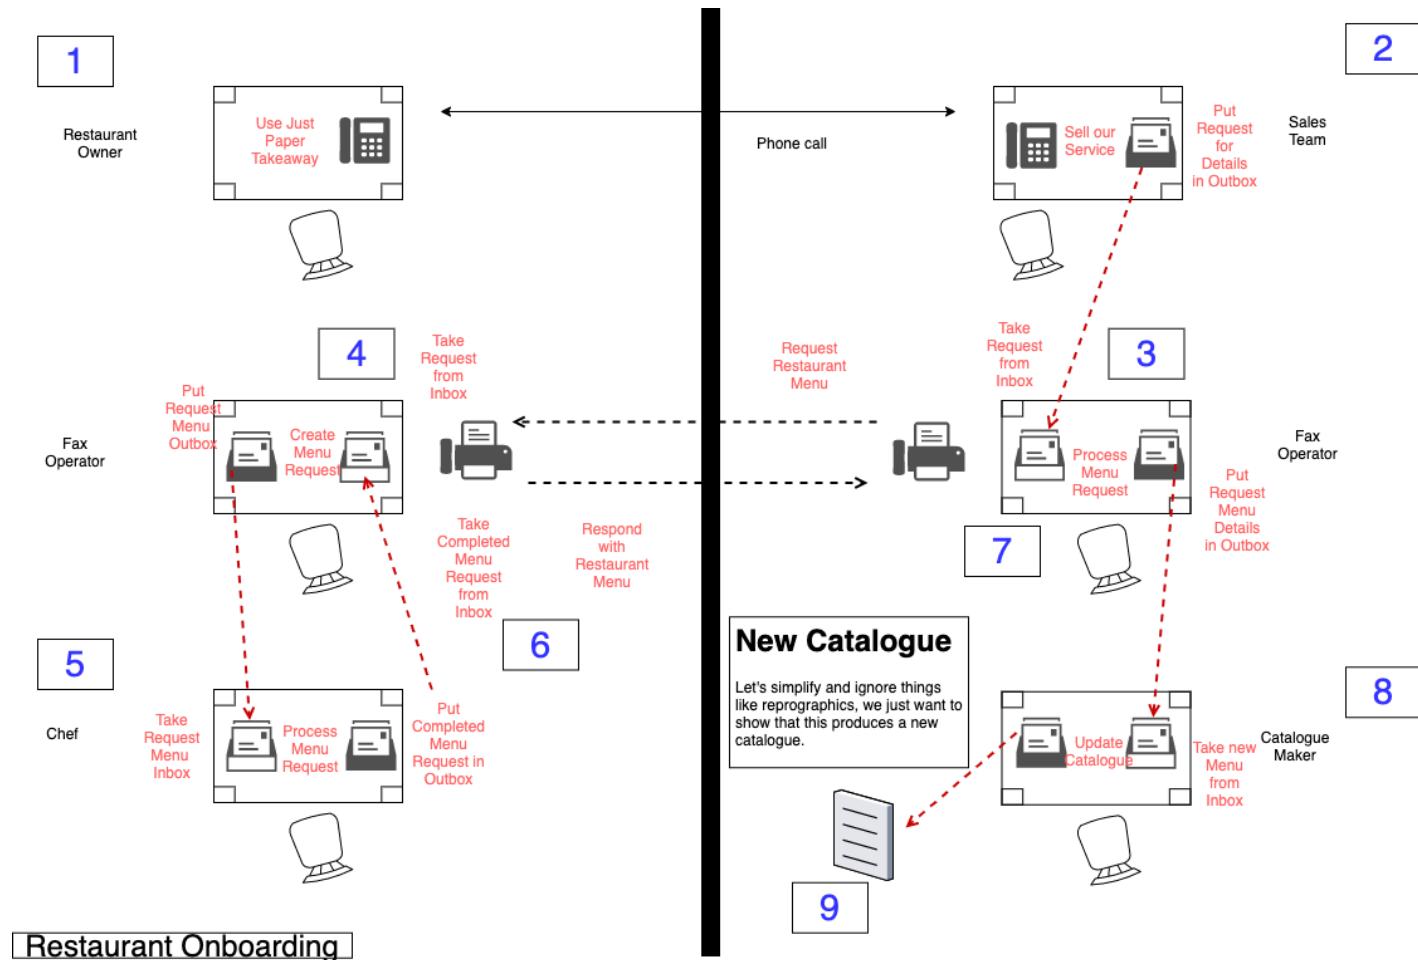


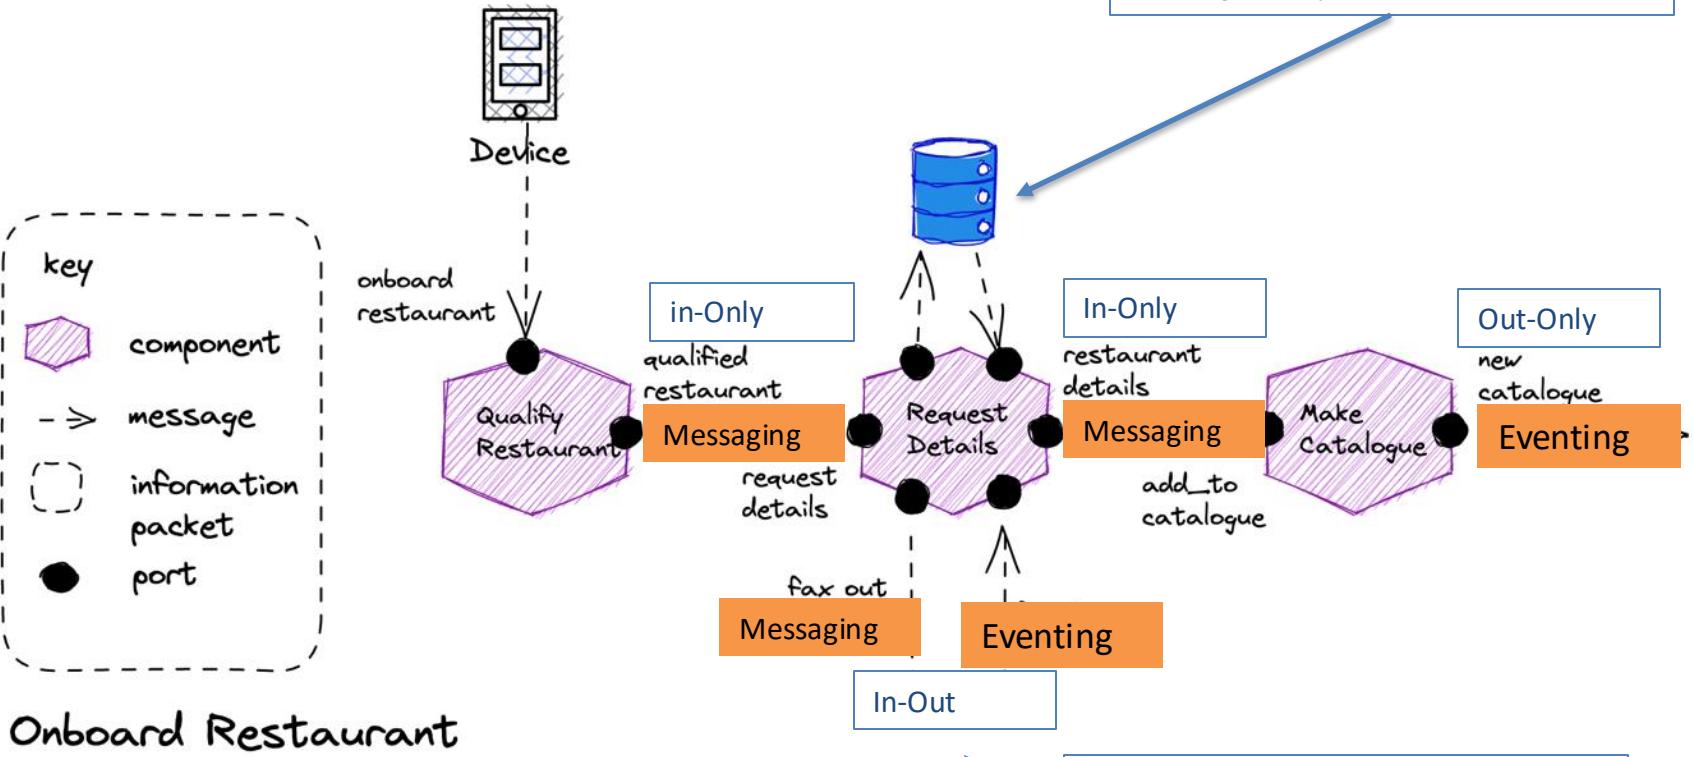
Reserve external Workflow engines for orchestrating the flow between bounded contexts, when relying on choreography is too complex.

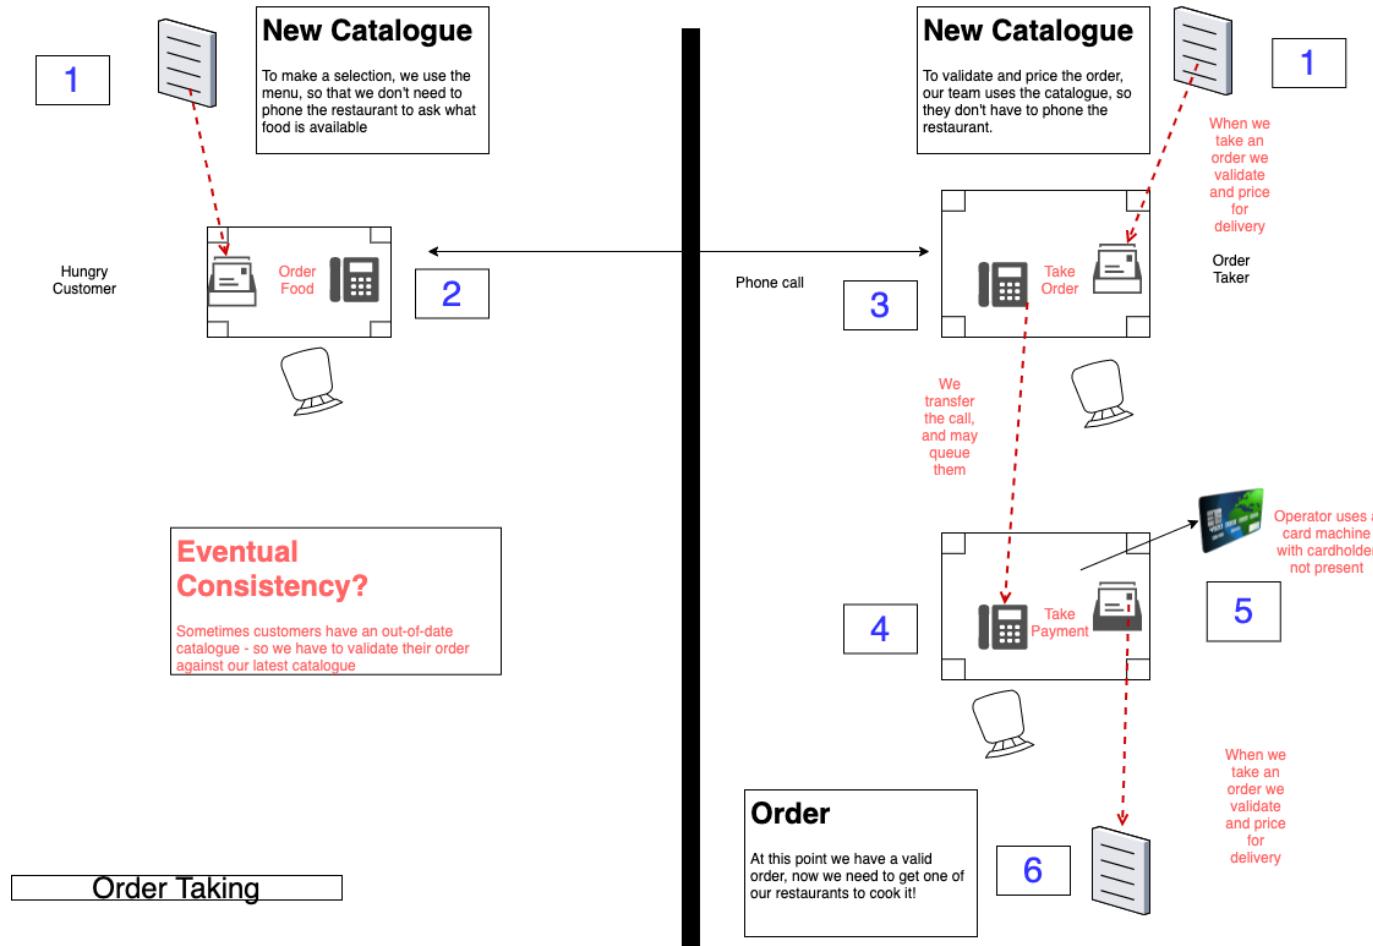
Implementing Workflow Patterns

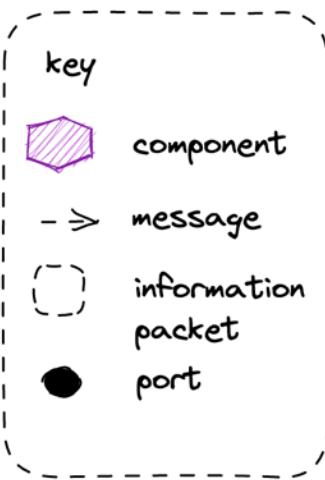
Pattern	Handlers Only	State Machine	Workflow Engine
Sequence	Chain message handlers	State transitions (A→B→C)	Sequential steps
Parallel Split	Trigger multiple handlers	 Complex coordination manually	Parallel branches with fork/join
Synchronization	Wait for multiple states	 Hard to do without orchestration	Join after parallel branches
Compensation	Fallback Handler	Saga	Implicit
Exclusive Choice	Conditional in handler	State transition based on input	Conditional branching
Simple Merge	Trigger single follow-up	Re-entrant transition or rehydration	Multiple incoming flows, OR gateway

PUTTING IT TOGETHER

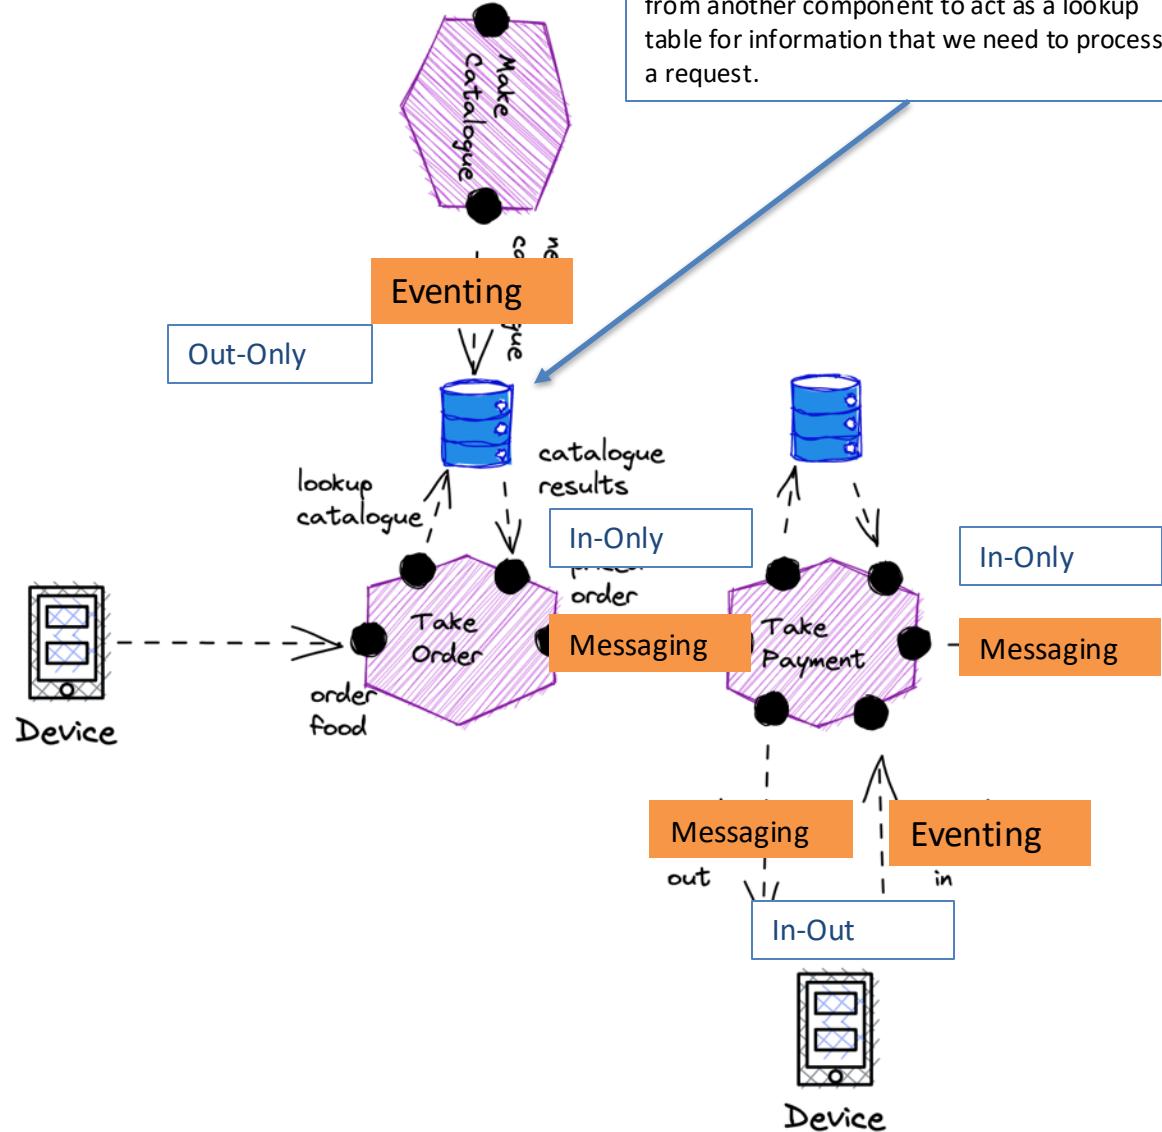


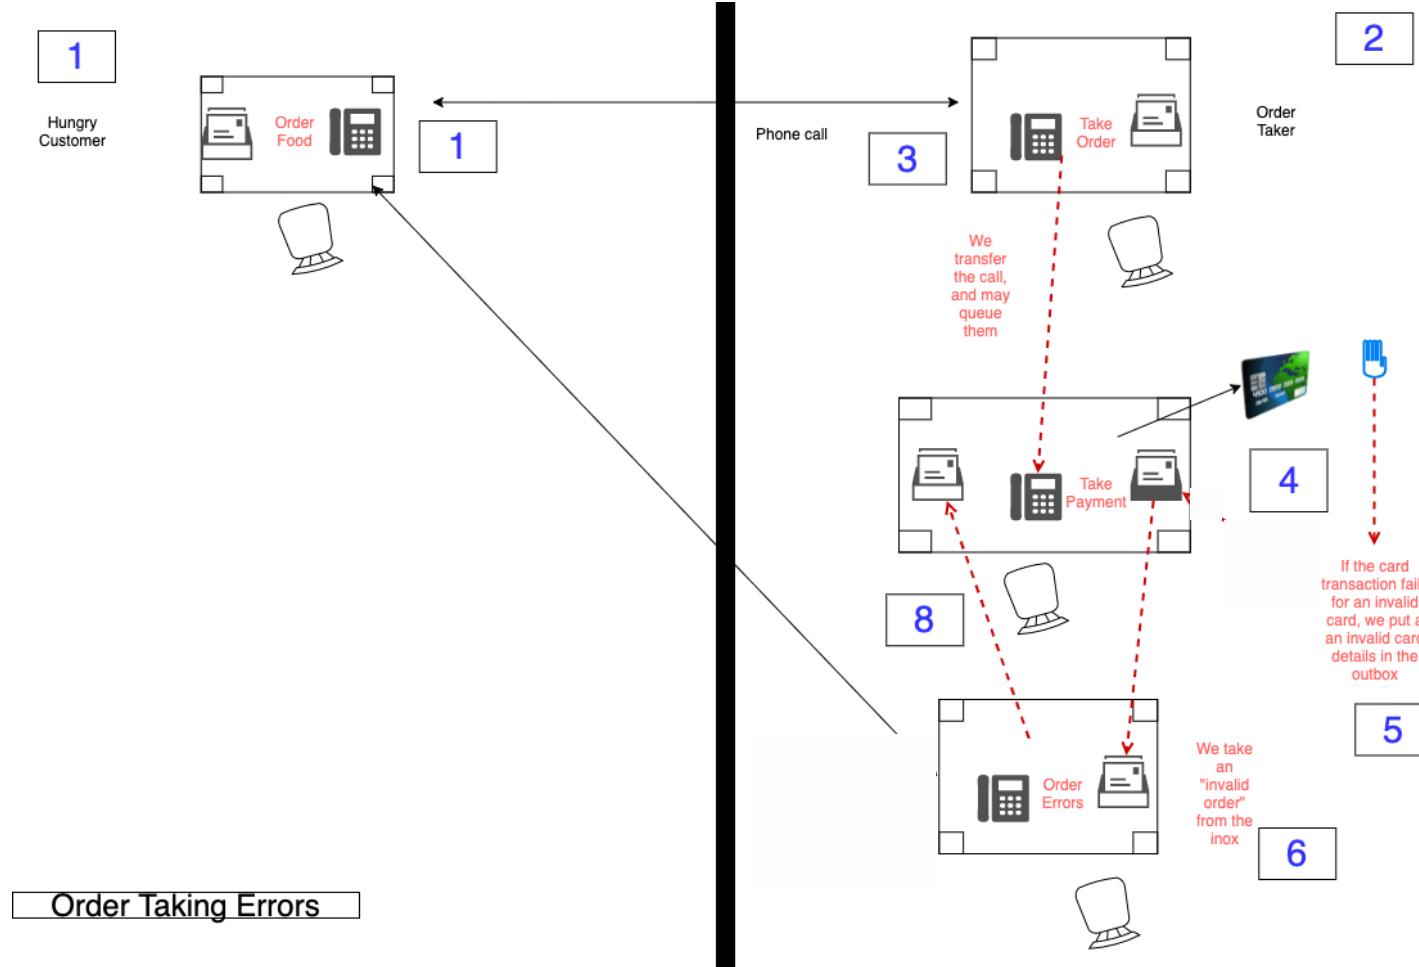


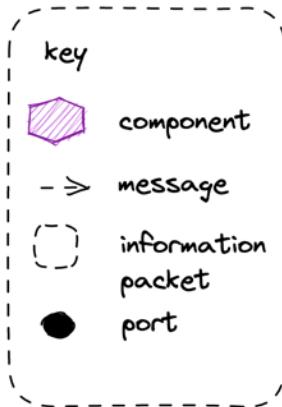




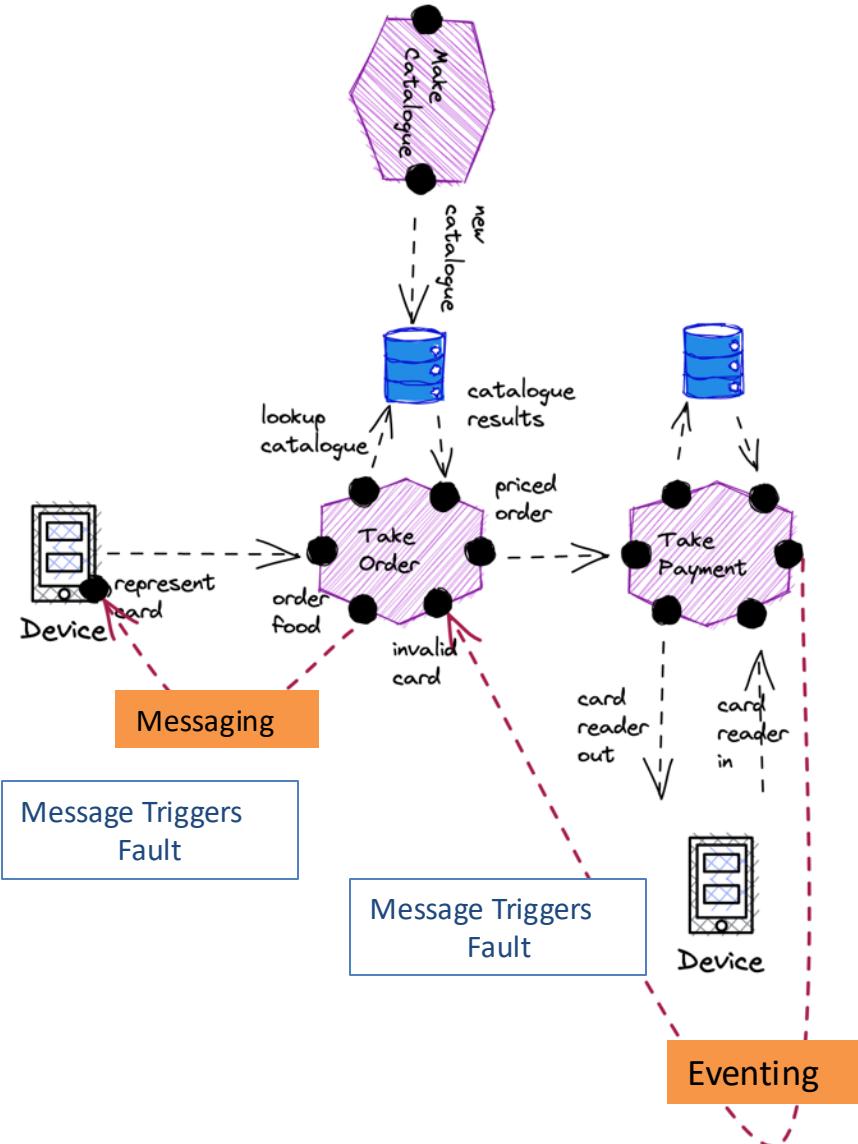
Order Food

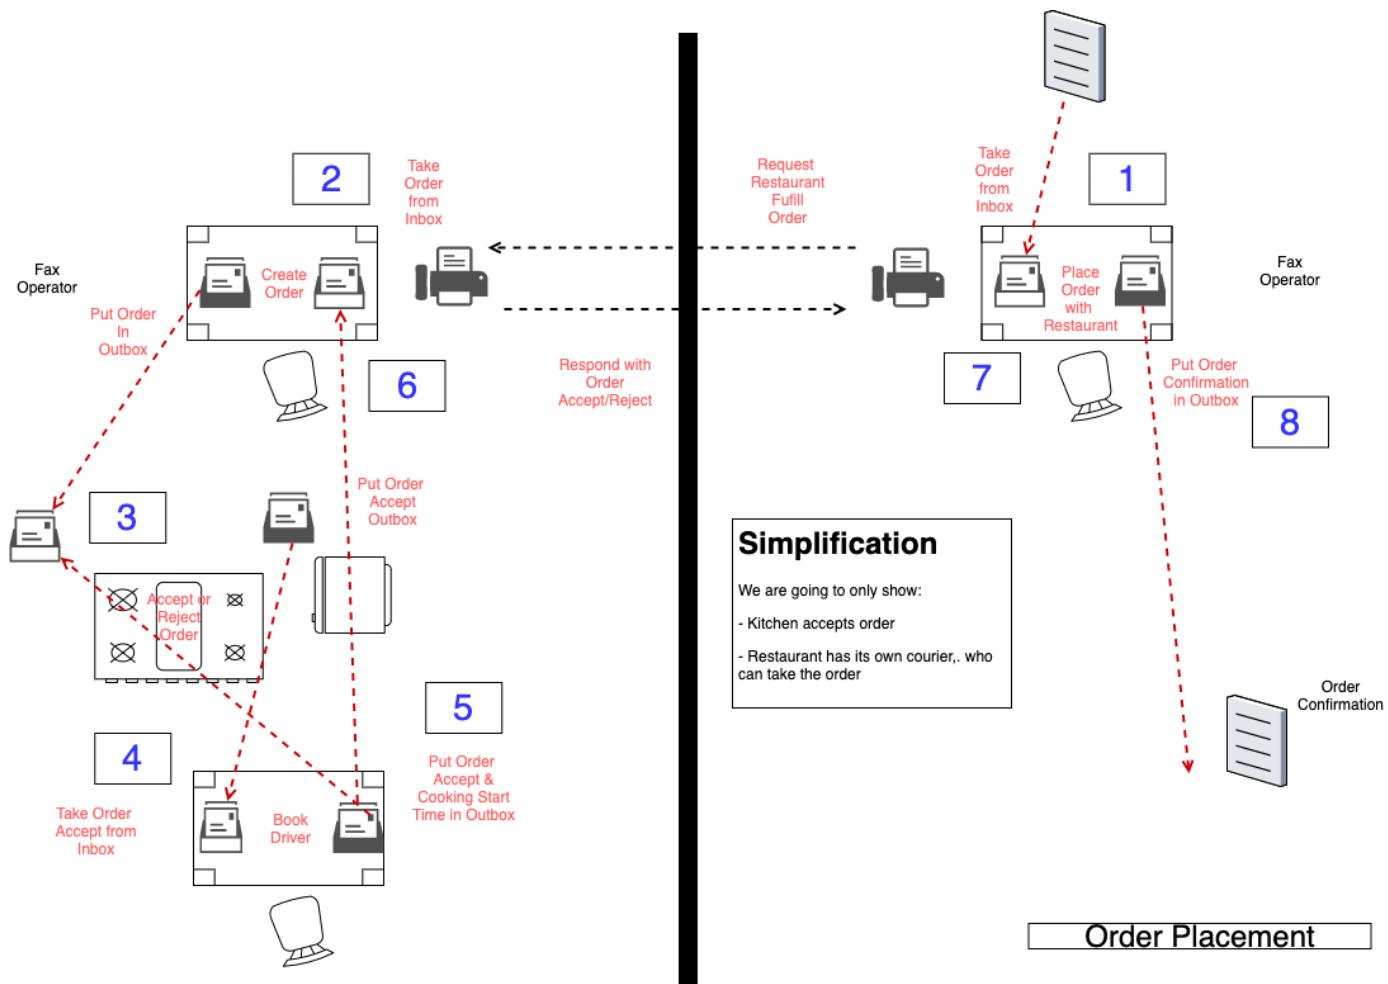


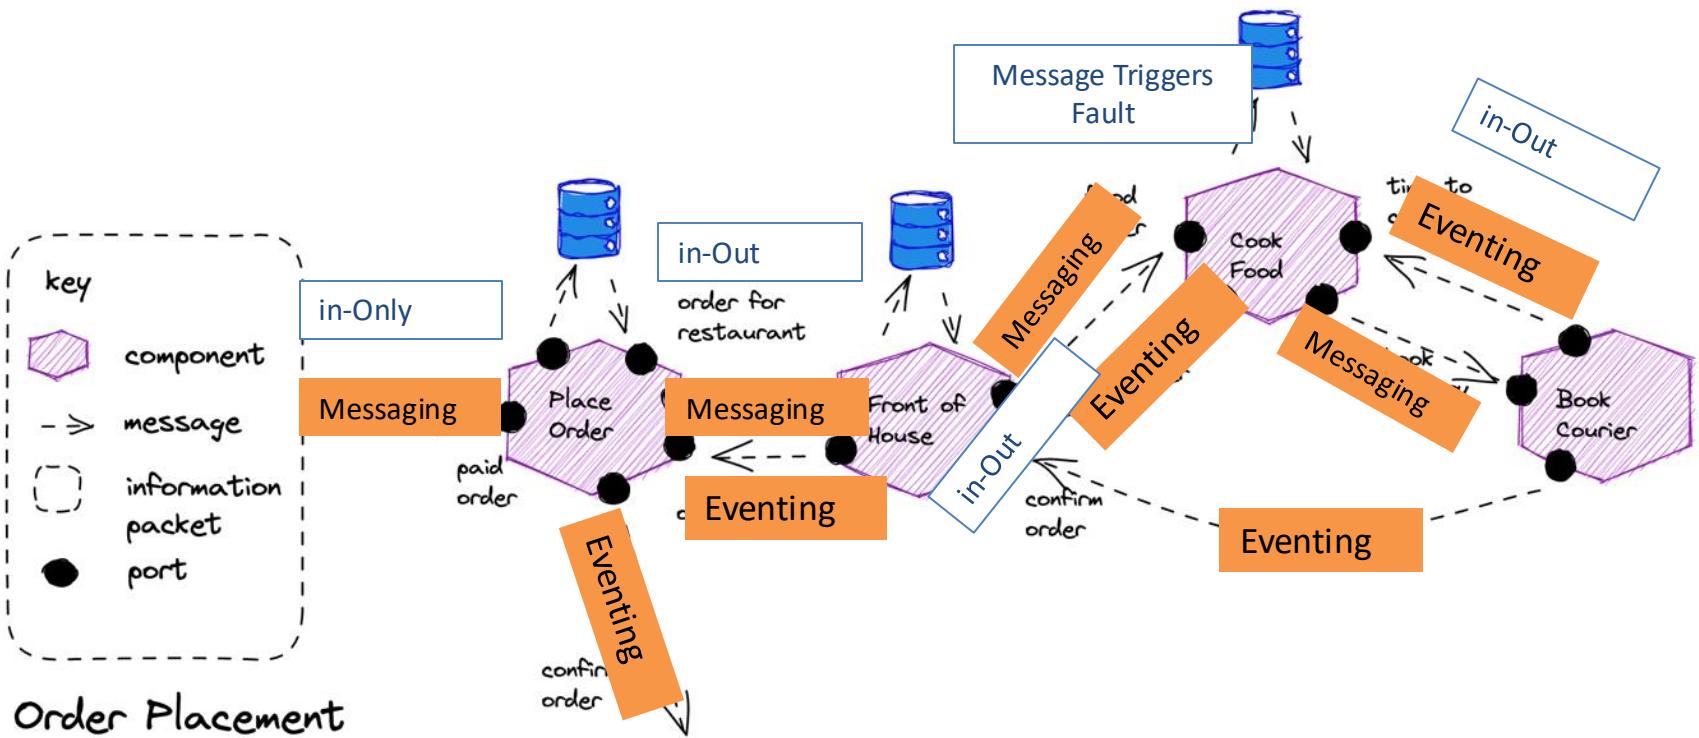


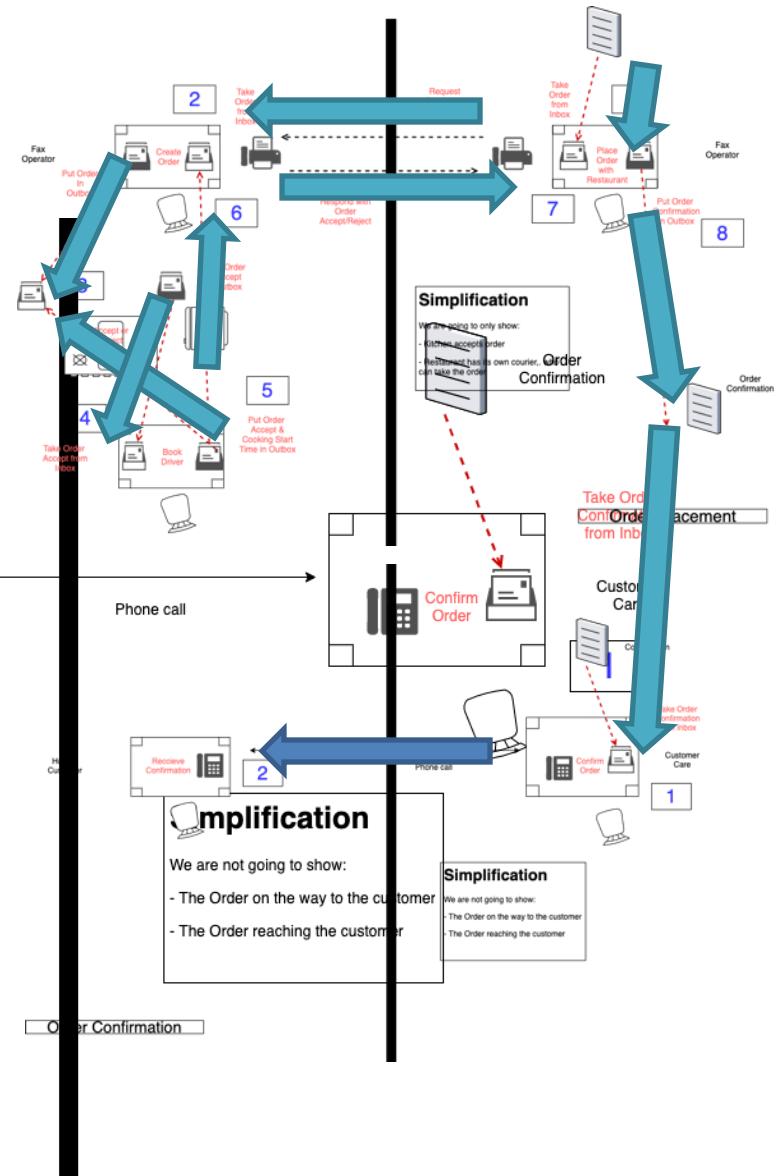
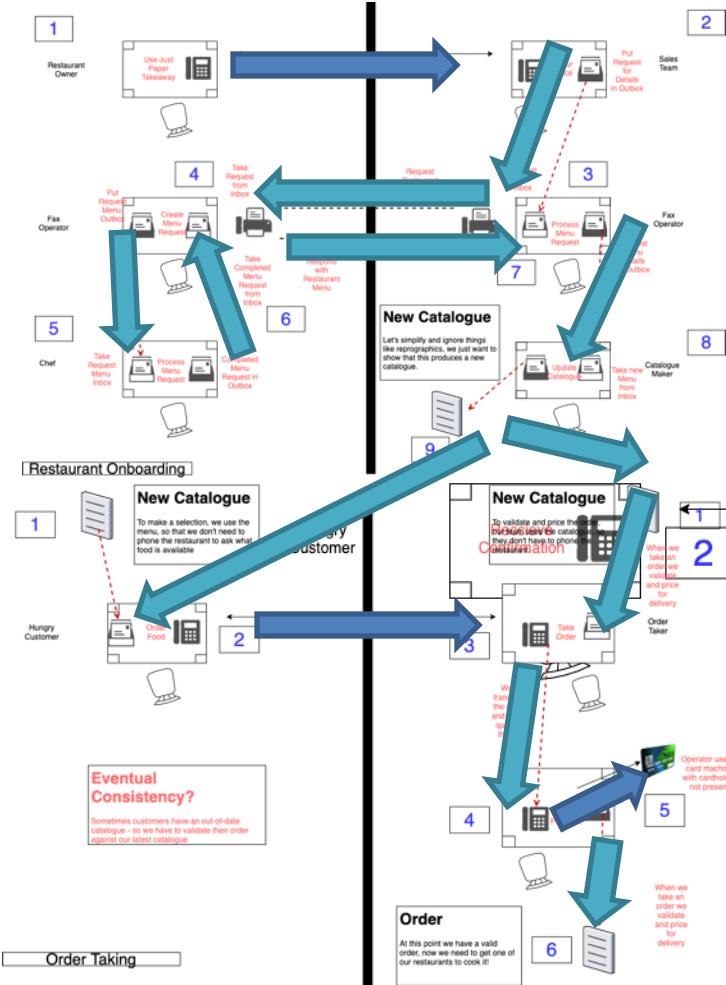


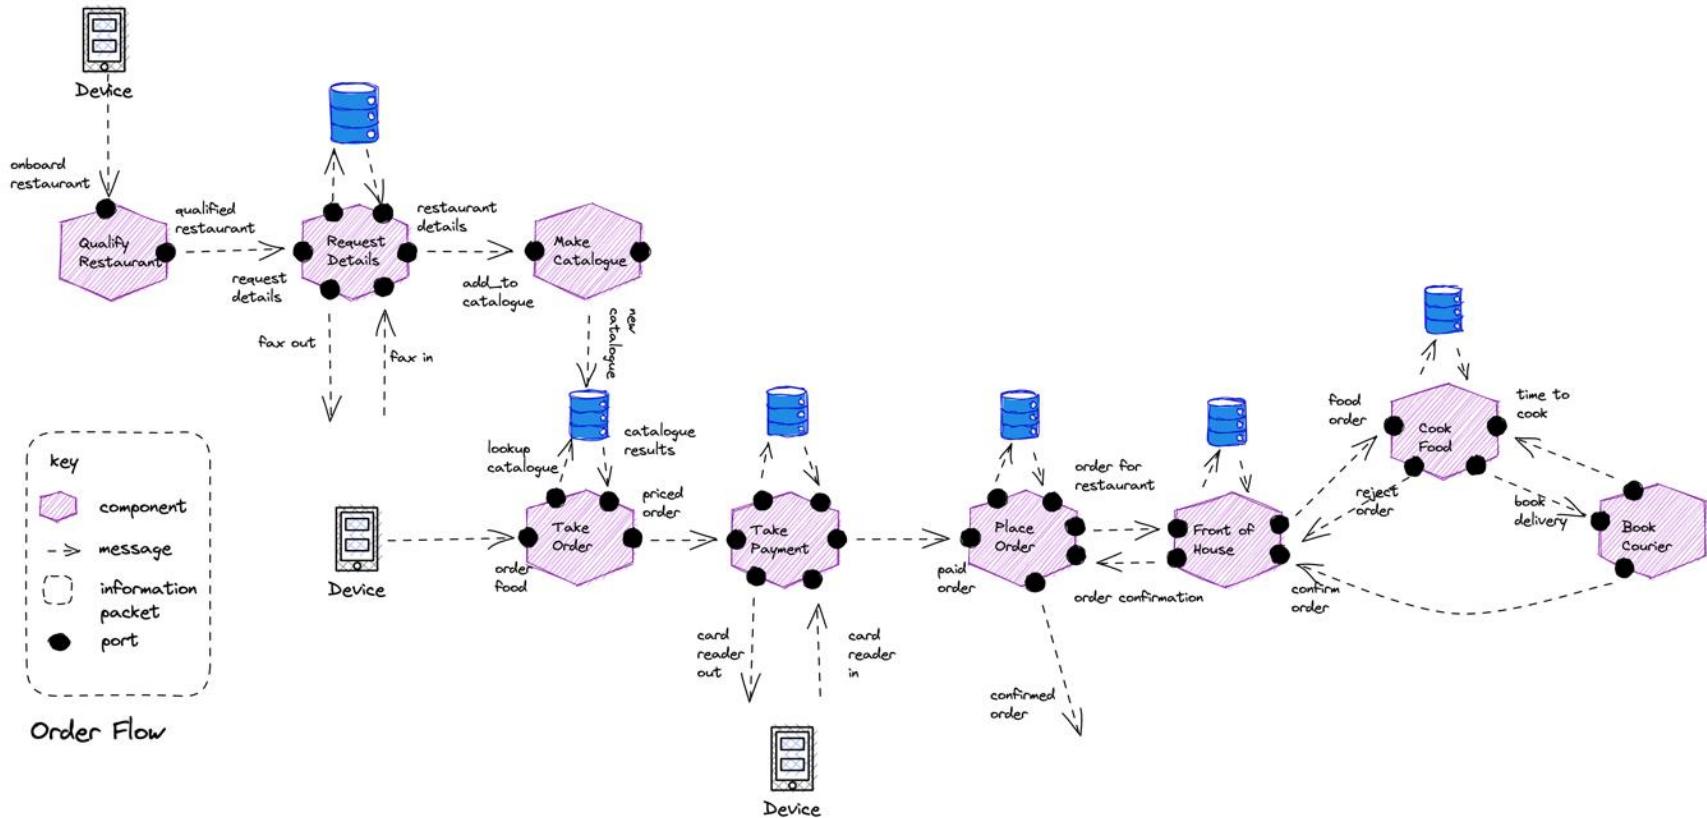
Order Errors







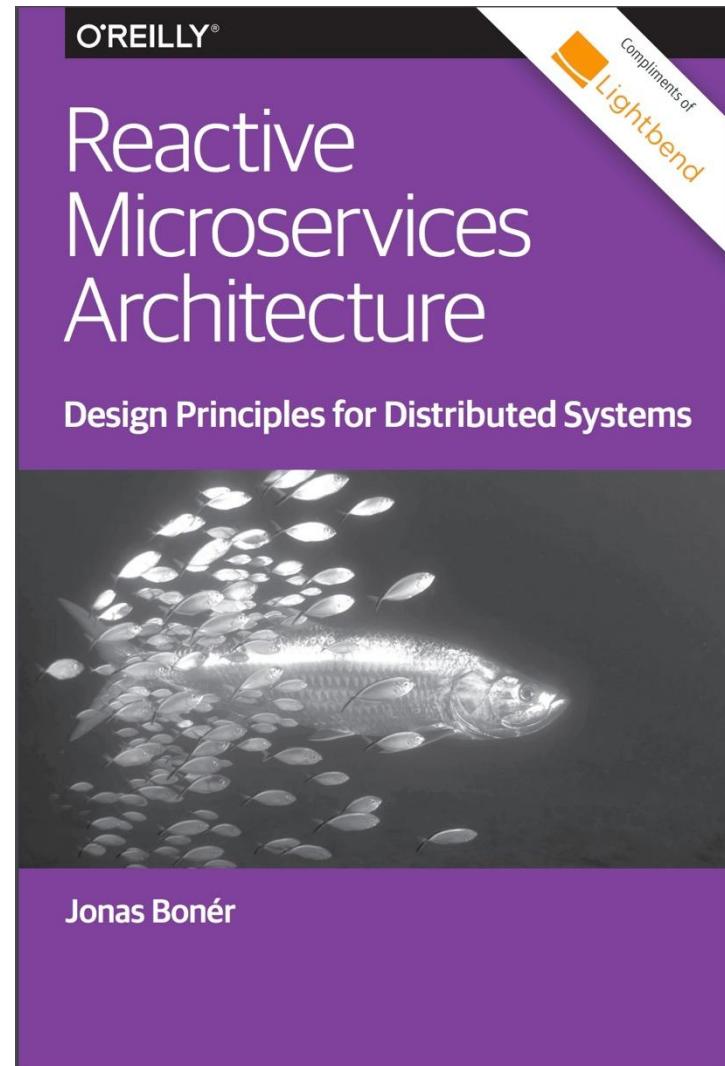




NEXT STEPS

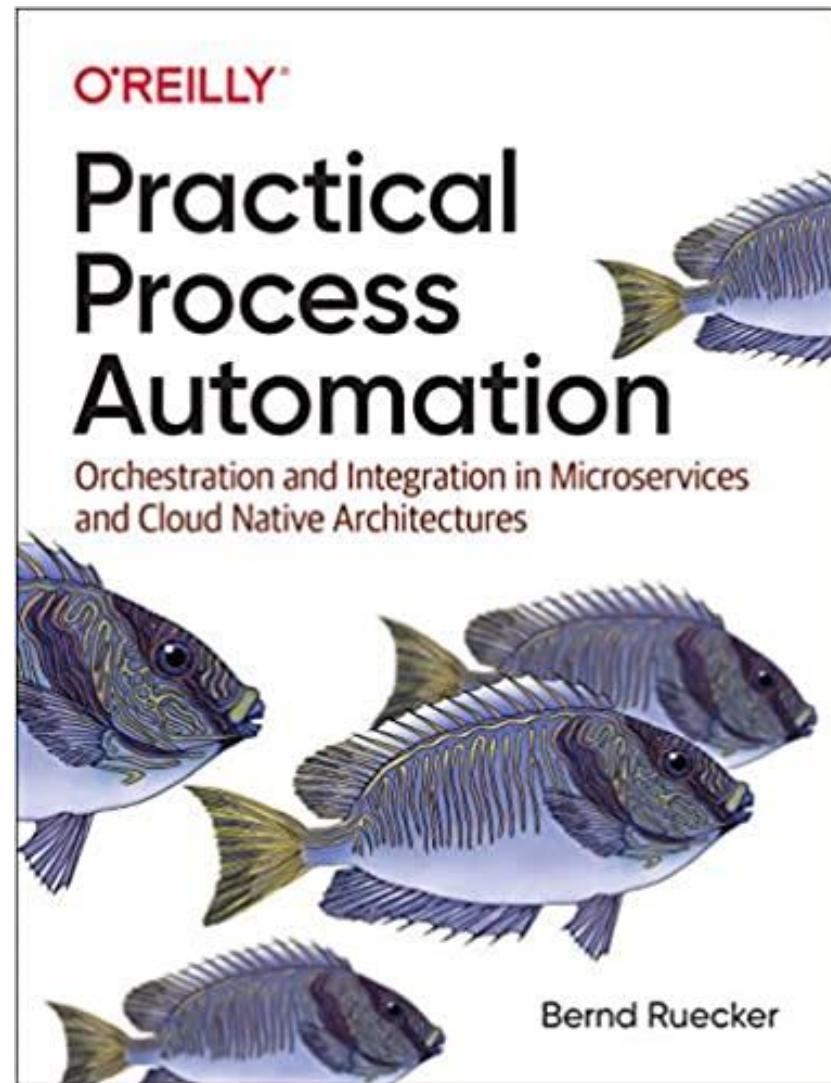
Further Reading

https://jonasboner.com/resources/Reactive_Microservices_Architecture.pdf



Further Reading

<https://processautomationbook.com/>



Further Reading

Serverless Land Content ▾ Learn EDA ▾ Search Search

EDA VISUALS

Small bite sized visuals about event-driven architectures

Designs and thoughts from [@bogmey123](#)

Avoiding the big ball of mud in event-driven architectures

Bite sized visual to help you understand the big ball of mud and why you might want to avoid it.

Local cache copy vs requesting data

Bite sized visual to help you understand local cache vs requests with event-driven architecture.

What are events in event-driven architectures?

Bite sized visual to help you understand events.

What is Event Sourcing?

Bite sized visual to help you understand event sourcing.

Queues vs Streams vs Pub/Sub

Bite sized visual to help understand the differences.

Reducing team cognitive load with event-driven architecture

Bite sized visual to help reduce cognitive load on teams and reuse code.

What are Events?

Bite sized visual to help you understand events.

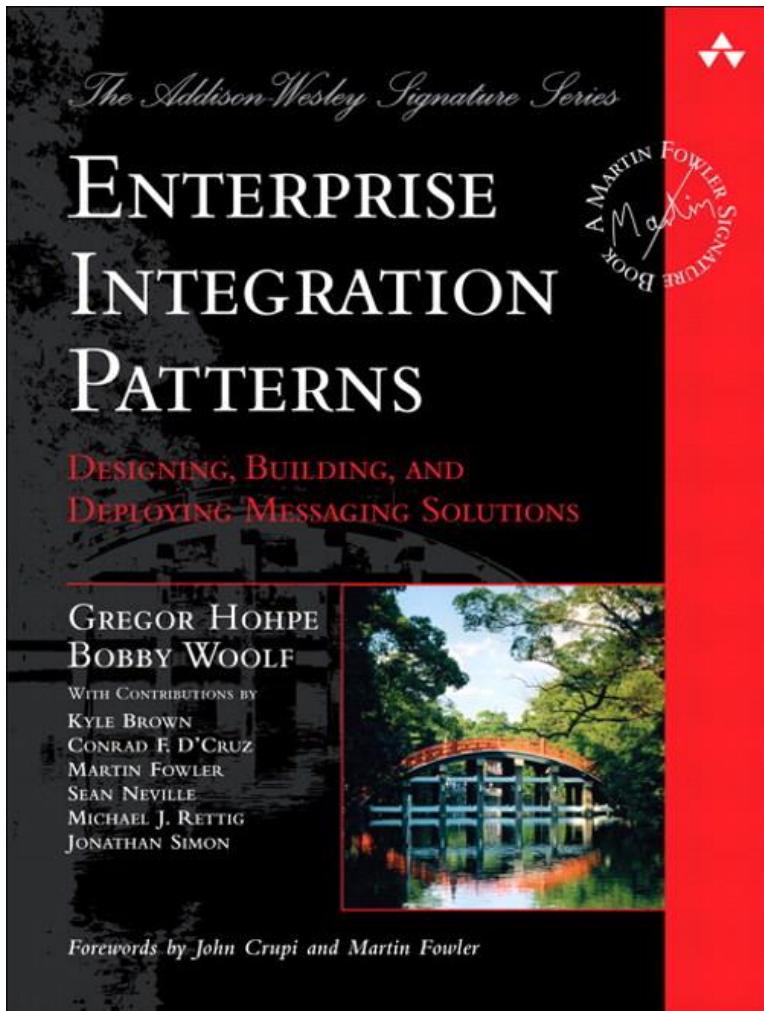
Reducing team cognitive load with event-driven architectures

Bite sized visual to help reduce cognitive load on teams and reuse code.

<https://serverlessland.com/event-driven-architecture/visuals>

Further Reading

<https://www.enterpriseintegrationpatterns.com/gregor.html>



Q&A