

Practical Messaging

A 101 guide to messaging

Ian Cooper

X/Hachyderm: ICooper

Who are you?

I am a polyglot coding architect with over 20 years of experience delivering solutions in government, healthcare, and finance and ecommerce. During that time I have worked for the DTI, Reuters, Sungard, Misys, Beazley, Huddle and Just Eat Takeaway delivering everything from bespoke enterprise solutions, 'shrink-wrapped' products for thousands of customers, to SaaS applications for hundreds of thousands of customers.

I am an experienced systems architect with a strong knowledge of OO, TDD/BDD, DDD, EDA, CQRS/ES, REST, Messaging, Design Patterns, Architectural Styles, ATAM, and Agile Engineering Practices

I am frequent contributor to OSS, and I am the owner of: <https://github.com/BrighterCommand>. I speak regularly at user groups and conferences around the world on architecture and software craftsmanship. I run public workshops teaching messaging, event-driven and reactive architectures.

I have a strong background in C#. I spent years in the C++ trenches. I dabble in Go, Java, JavaScript and Python.

www.linkedin.com/in/ian-cooper-2b059b



Welcome to Brighter

This project is a Command Processor & Dispatcher implementation with support for task queues that can be used as a lightweight library.

It can be used for implementing [Ports and Adapters](#) and [CQRS \(PDF\)](#) architectural styles in .NET.

It can also be used in microservices architectures for decoupled communication between the services

[GET STARTED](#)

Day One Messaging

- Distribution
- Integration Styles
- Messaging Patterns
- Queues and Streams
- Managing Asynchronous Architectures

Day Two Conversations

- Conversation Patterns
 - Activity and Correlation
 - Repair and Clarification
 - Reliable Messaging
 - Fat and Skinny
 - Conversations
- Reactive Architectures
 - Message Passing
 - Paper Based Flows
 - Flow Based Programming
- Next Steps

Prerequisites

We will use Rabbit MQ and Kafka for examples. You should have Docker (or an equivalent) installed on your machine, as exercises provide a Docker Compose file to spin up RMQ and Kafka.

You will need to be able to write code with an editor/IDE of your choice.

You can choose from: C#; Java; Python; Go; JavaScript

Course Content

<https://github.com/iancooper/practical-messaging>

Exercise Code

<https://github.com/iancooper/Practical-Messaging-Sharp>

<https://github.com/iancooper/Practical-Messaging-Python>

<https://github.com/iancooper/Practical-Messaging-JavaScript>

<https://github.com/iancooper/Practical-Messaging-Go>

<https://github.com/iancooper/Practical-Messaging-Java>

Day One

What is driving messaging

DISTRIBUTED SYSTEMS

Why Distribute?

Performance and Scalability

Availability

Maintainability

Inherent Distribution

Example: Task Queues

What if the work is time consuming?

What happens if the Db is not available?



Web Server

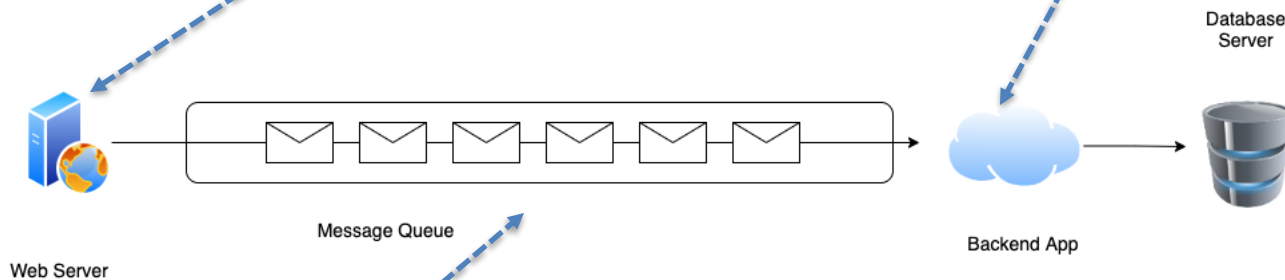
Database
Server

What if the work experience a surge?

What if the work is CPU intensive?

The web server puts the work on a queue

A backend application can perform long-running or CPU intensive work, allowing the web server to service new requests.

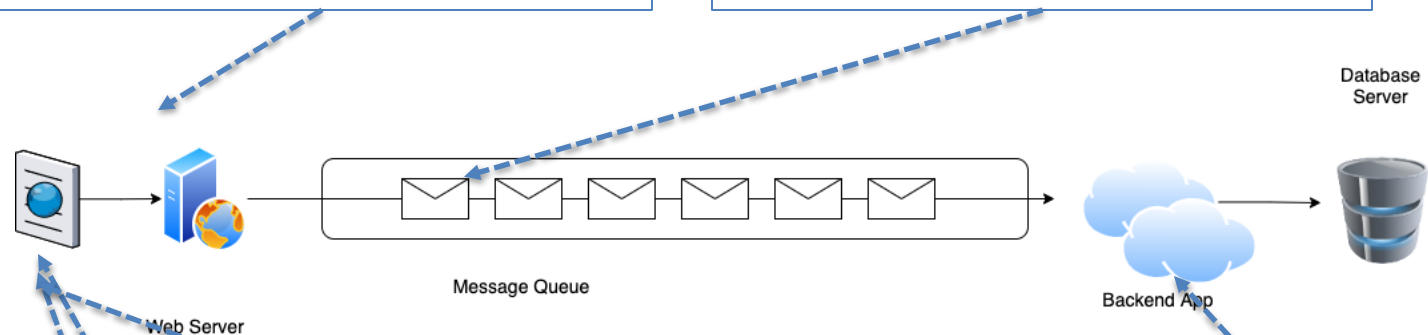


The queue stores work until we are ready to consume it. We can *throttle* to prevent surges.

We can scale out the backend services using a competing consumers approach, to ensure the queue does not backup.

We return 202 Accepted – we have your work request, and won't lost it.

We enqueue a work item for the request.



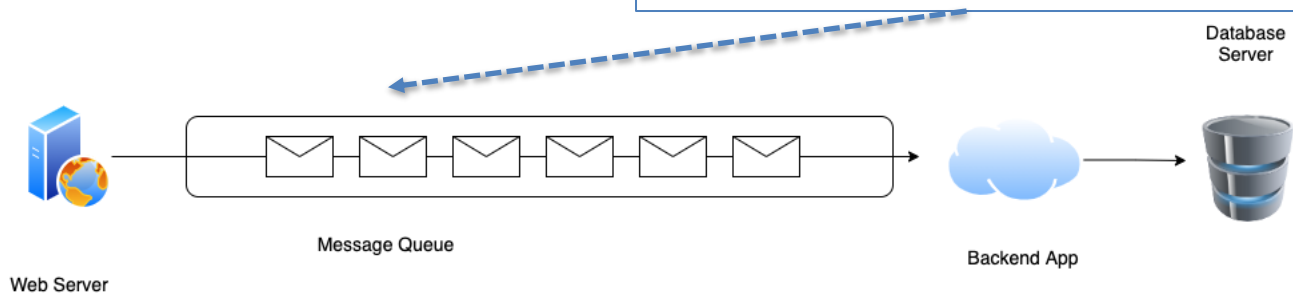
We return a Location header, so you can monitor progress.

Link to Resource – 404 until created

Backend app does work at sustainable pace, and updates KV store if required.

Link to a progress page, backed by a KV store where we note progress

This general technique is known as Decoupled Invocation – we separate building the command from executing it.



Decoupled Invocation Pattern

Use Decoupled Invocation. A producer puts a message onto a queue at the service endpoint. A consumer reads messages from the queue.

The queue stores messages for eventual processing.

If the rate of arrival at the endpoint is unpredictable, the queue acts as a buffer that makes it possible to predict the rate of consumption.

This makes it simpler to do capacity planning because peaks of requests are smoothed out by the queue.

The consumer must be able to control the rate of processing, otherwise a spike is simply passed down the wire.

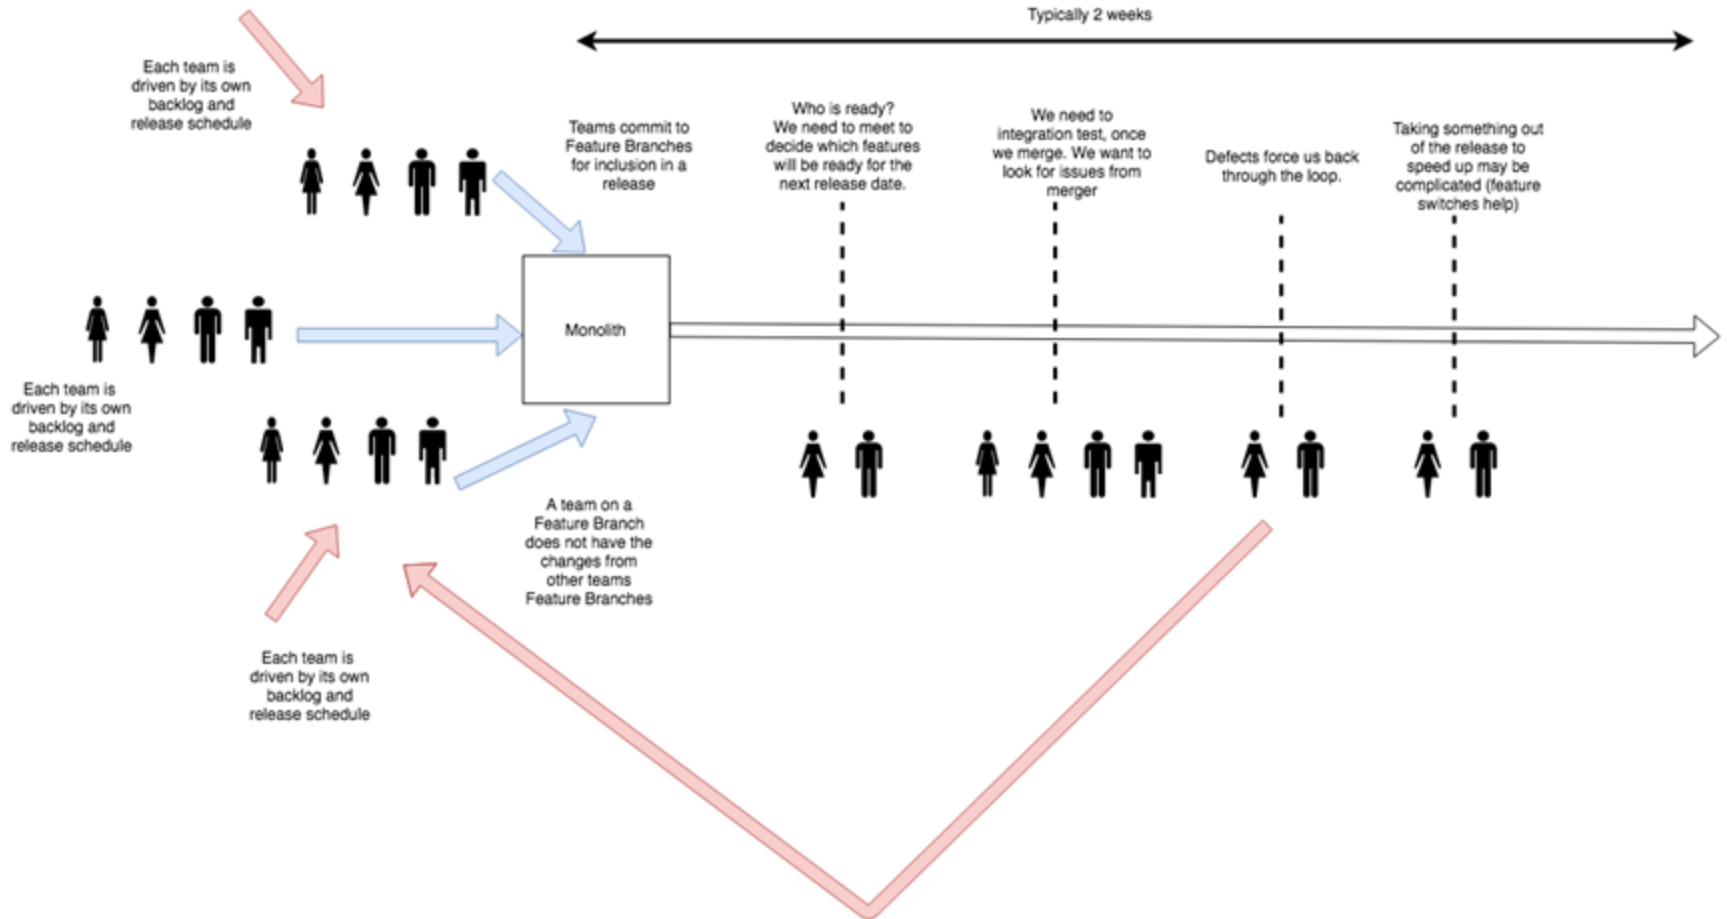
Example: Microservices

It's all about velocity!!!

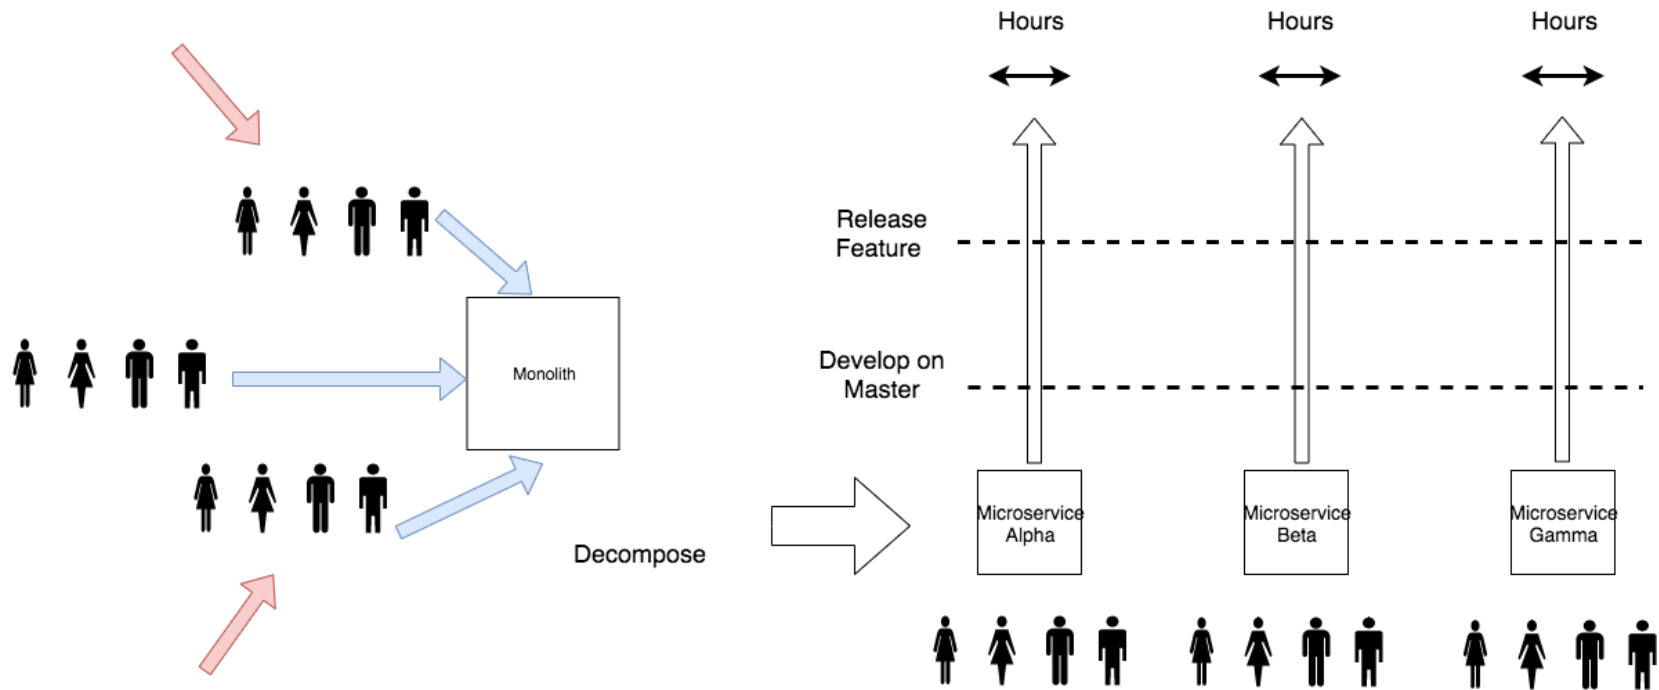
“Speed wins in the marketplace”

Adrian Cockcroft, former lead architect at Netflix

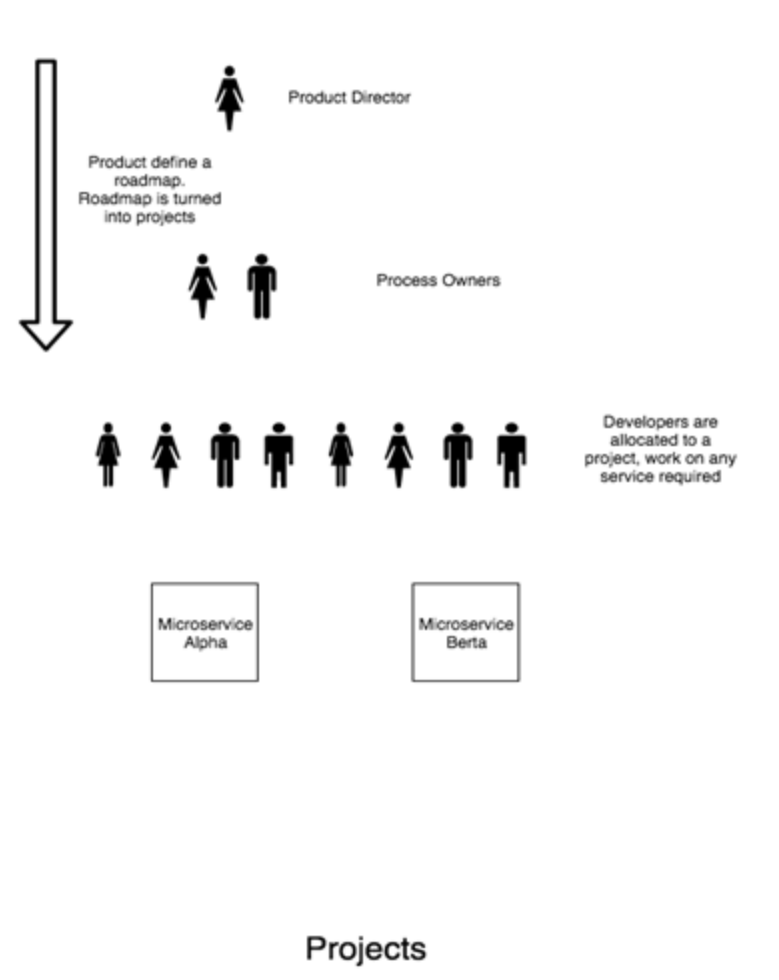
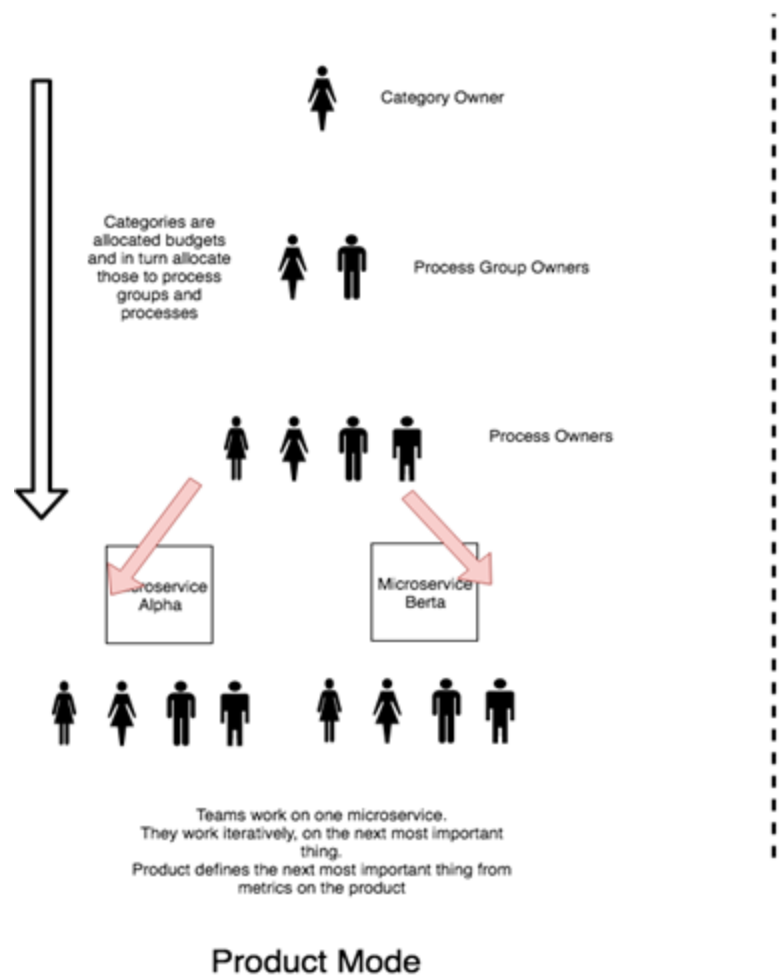
Monoliths Do Not Scale To Many Teams!



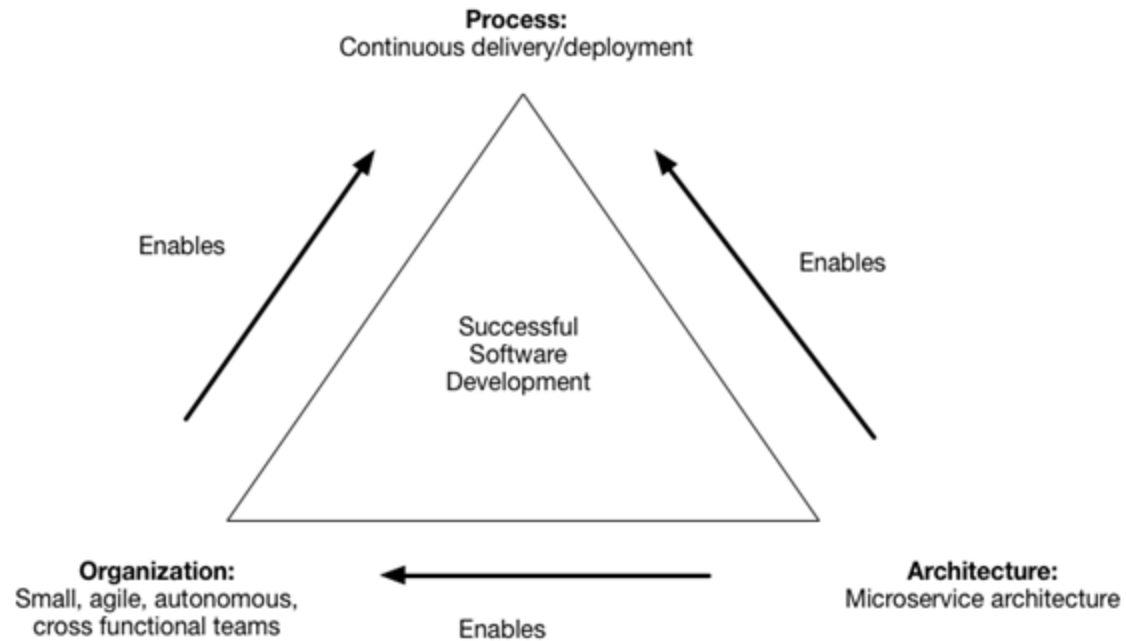
Microservices let us scale an organisation



Product Mode



Microservices Enable Agility



<https://microservices.io/patterns/decomposition/decompose-by-business-capability.html>

The Price of Distribution

Fallacies of Distributed Computing

The network is reliable.

Latency is zero.

Bandwidth is infinite.

The network is secure.

Topology doesn't change.

There is one administrator.

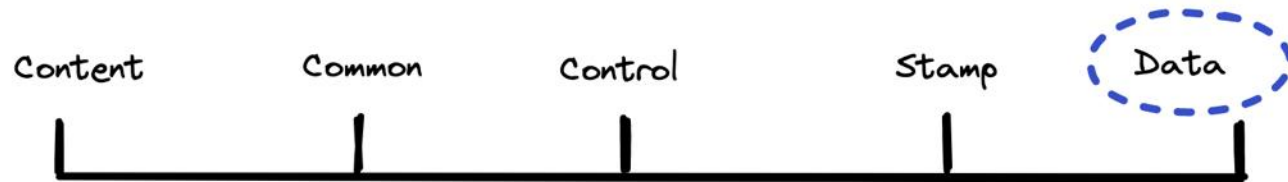
Transport cost is zero.

The network is homogeneous.

How do we communicate between microservices?

INTEGRATION STYLES

Coupling



Tight

More Interdependency

More Coordination

More Information Flow

Loose

Less Interdependency

Less Coordination

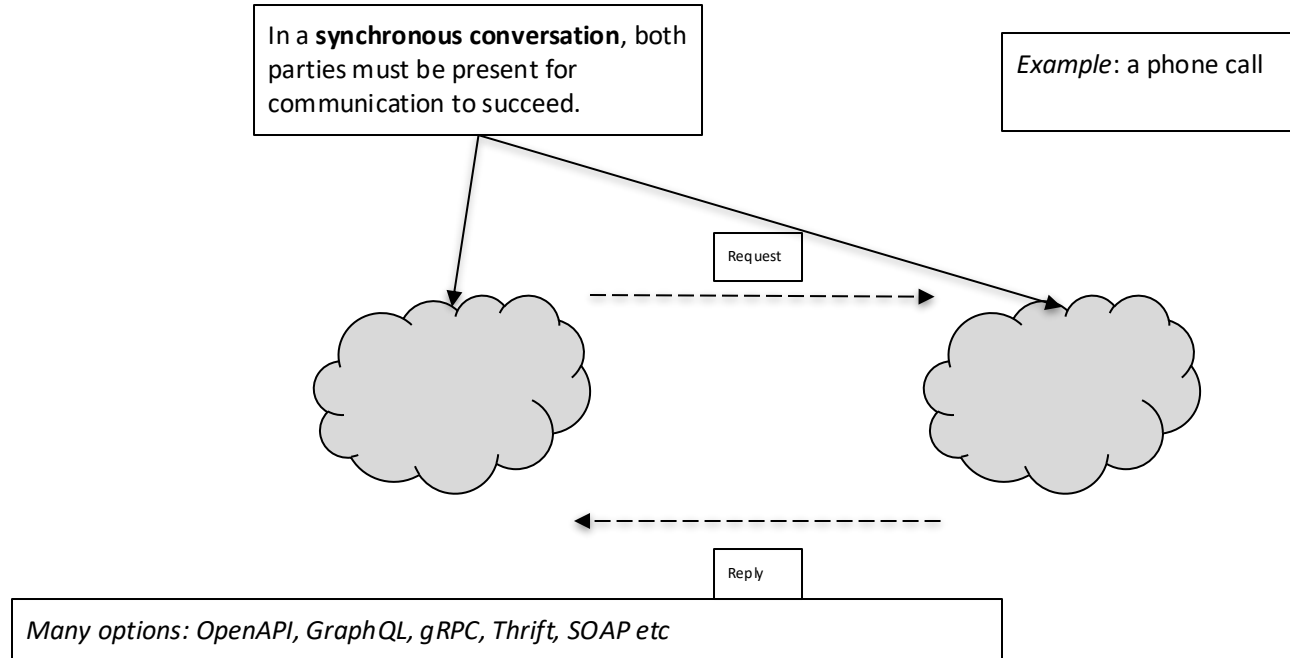
Less Information Flow

Behavioral Coupling: a form of control coupling where we exchange a sequence of calls to complete work

Behavioral Coupling

Temporal Coupling

Synchronous Conversation

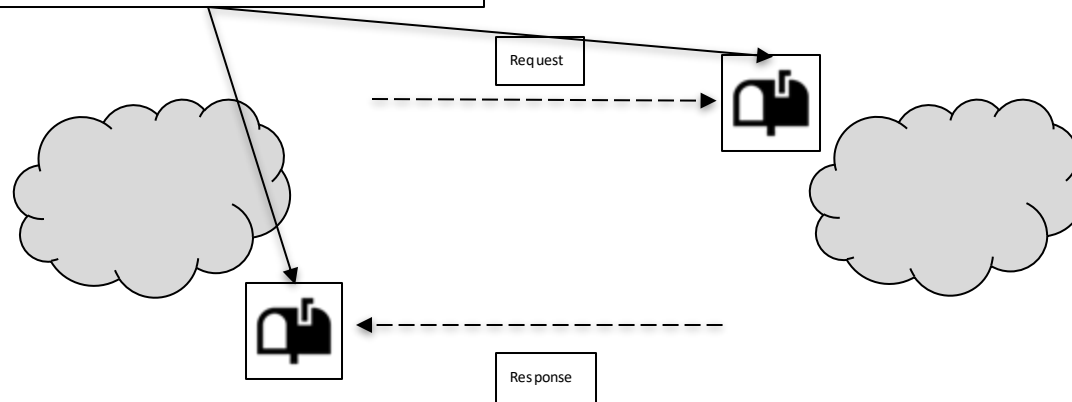


Temporal Coupling

Asynchronous Conversation

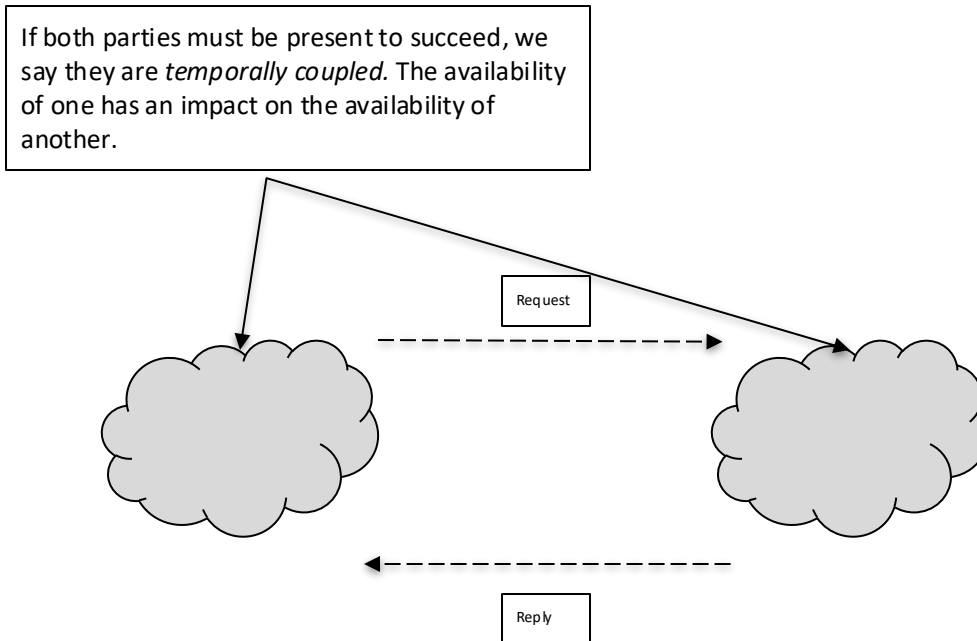
In an **asynchronous conversation**, the *receiver* does not need to be present at the time the *sender* communicates with them, using **store and forward** to pick up the message later.

Example: snail mail

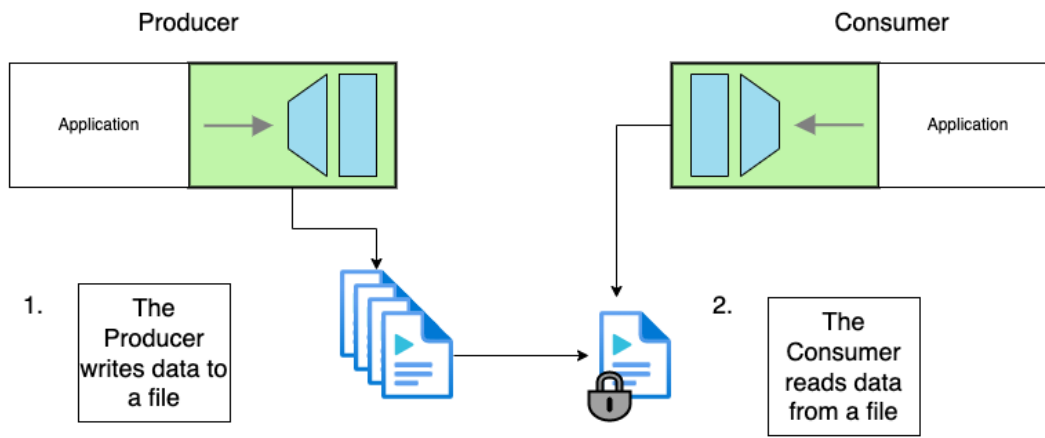


Many options: SQS, Kafka, AMQP 0-9-1 (RMQ), AMQP 1-0, MQTT, S3

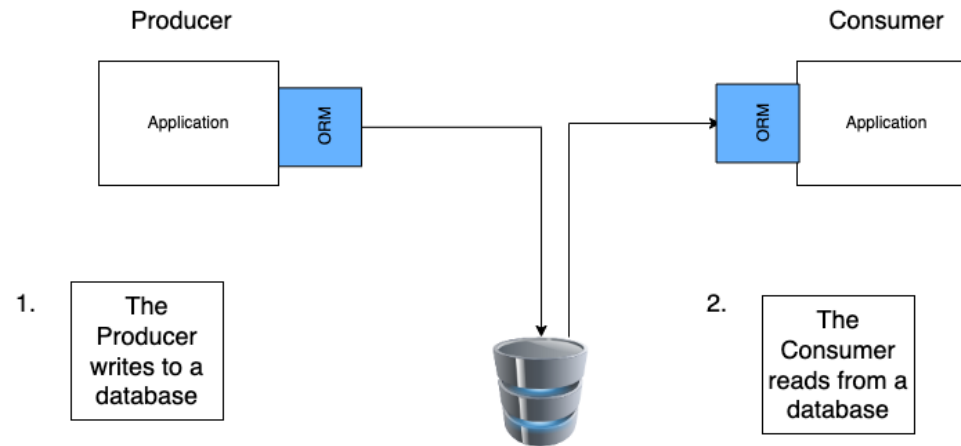
Temporal Coupling



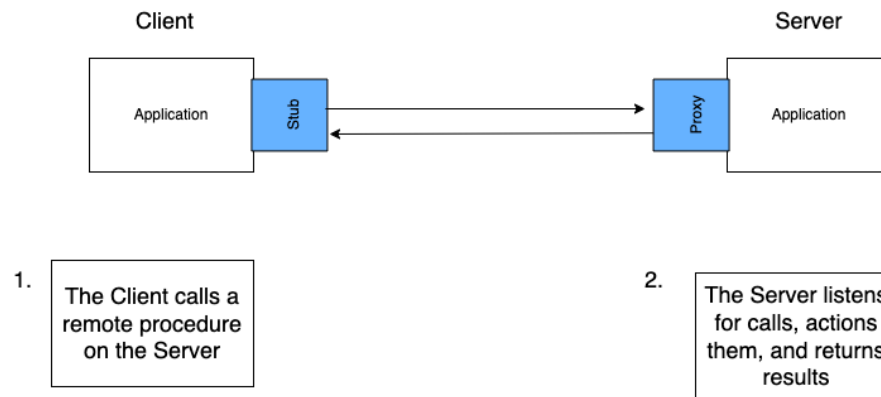
File Transfer



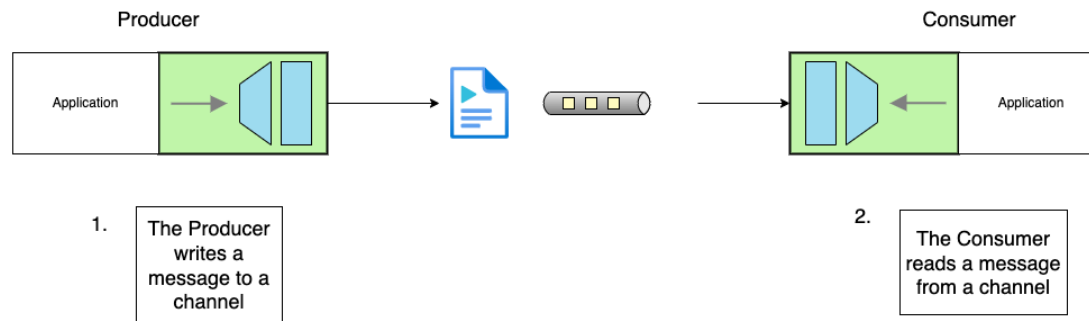
Shared Database

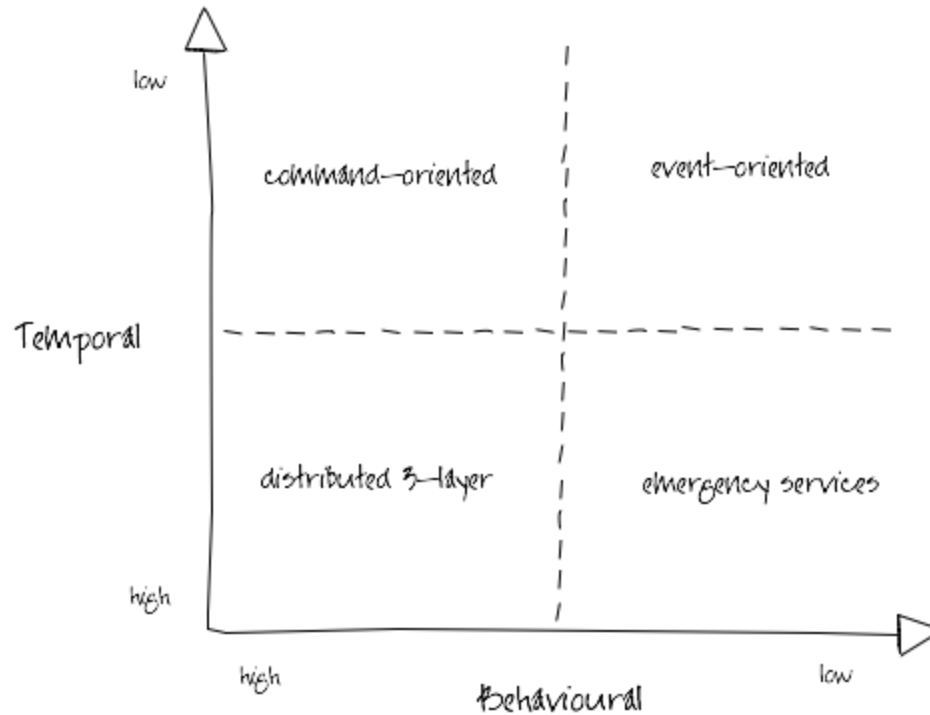


Remote Procedure Call



Messaging

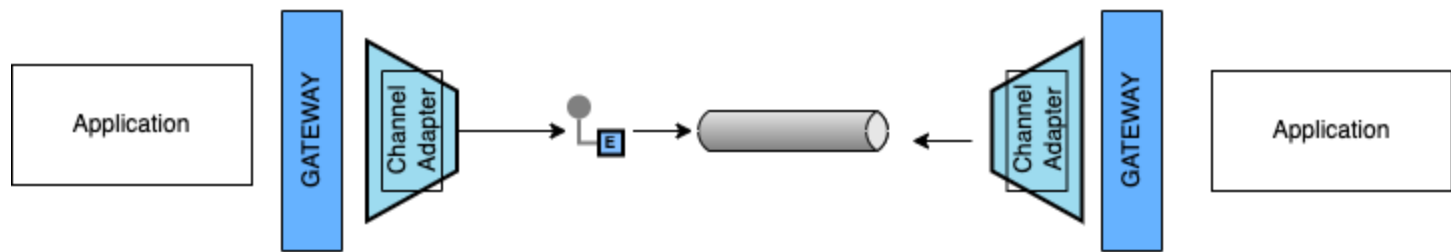




Ian Robinson: <http://iansrobinson.com/2009/04/27/temporal-and-behavioural-coupling/>

Integrating using events

MESSAGING PATTERNS



What is a message?

A MESSAGE

Message Construction

A message has a header and body

The body contains data for the consumer

The header contains metadata for any *filter* in the pipeline.

The header should indicate the format of the body

Break a large message into pieces as a Message Sequence or use a Claim Check

MESSAGING AND EVENTS

Message Types

Messaging

Has Intent

Request An Answer
(Query)
Transfer of Control
(Command)
Transfer of Value

Part of a Workflow
Part of a Conversation

Concerned with
the Future

Eventing

Provides Facts

Things you Report On

No Expectations

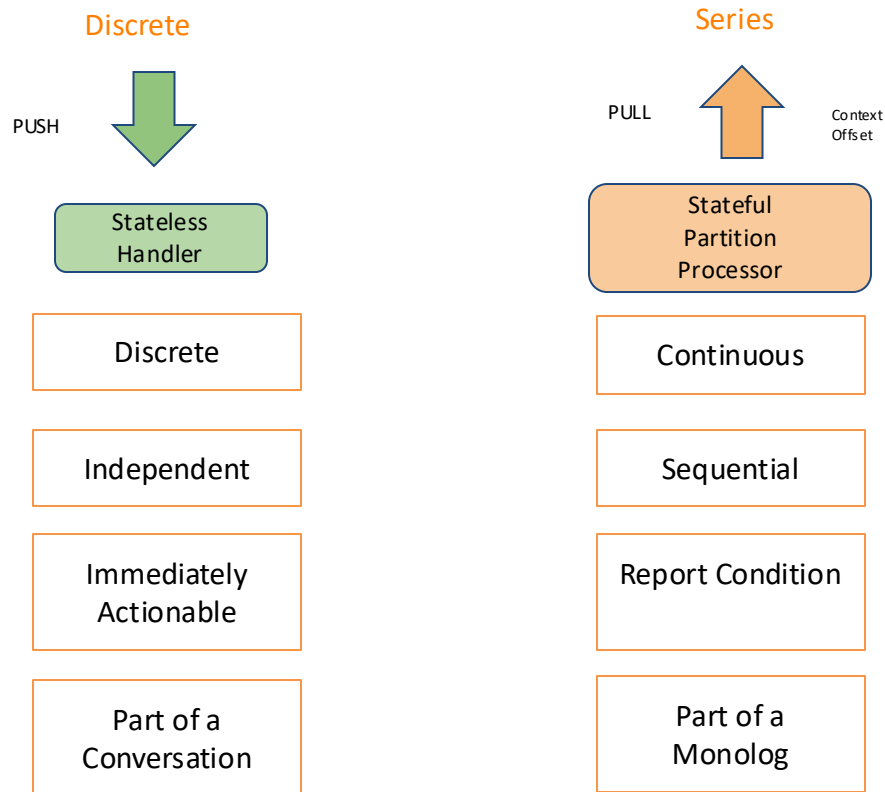
History

Context

Concerned with the Past

After Clemens Vasters <https://youtu.be/ITrILersqzY>

Eventing Types



After Clemens Vasters: <https://skillsmatter.com/skillscasts/10191-keynote-events-data-points-jobs-and-commands-the-rise-of-messaging>

See also: https://en.wikipedia.org/wiki/Discrete_time_and_continuous_time

Message Types

Messaging

Command

Eventing

Event (Notification)

Document

See Gregor Hohpe: <https://www.enterpriseintegrationpatterns.com/patterns/messaging/Message.html>

Command Message

Use a Command Message to reliably invoke a procedure in another application

Uses the well-established pattern for encapsulating a request as an object. The Command pattern [GoF] turns a request into an object that can be stored and passed around.

Document Message

Use a Document Message to reliably transfer a data structure between applications.

The receiver decides what, if anything, to do with the data

Event Message

Use an Event Message for reliable, asynchronous event notification between applications.

The difference between an Event Message and a Document Message is a matter of timing and content. An event's contents are typically less important.

Self-paced material

EXERCISES

EXERCISE MATERIAL

Introduction to Exercises

- Readme
- Videos
- Scripts & Slides

Introduction to RMQ



DON'T PANIC

CHANNELS

Channels

A virtual pipe that connects producer and consumer

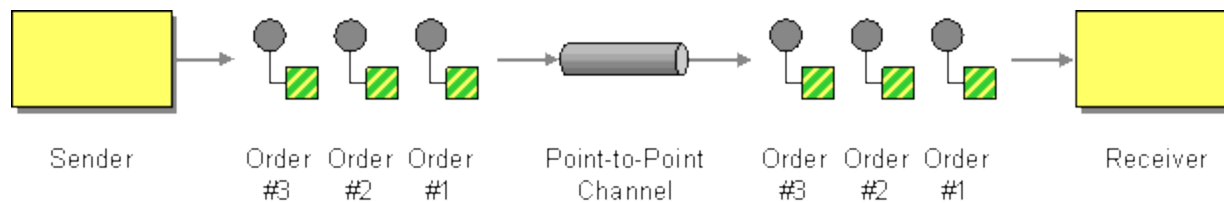
Logical Address (Topic or Routing Key)

Unidirectional

One-to-One or One-to-Many

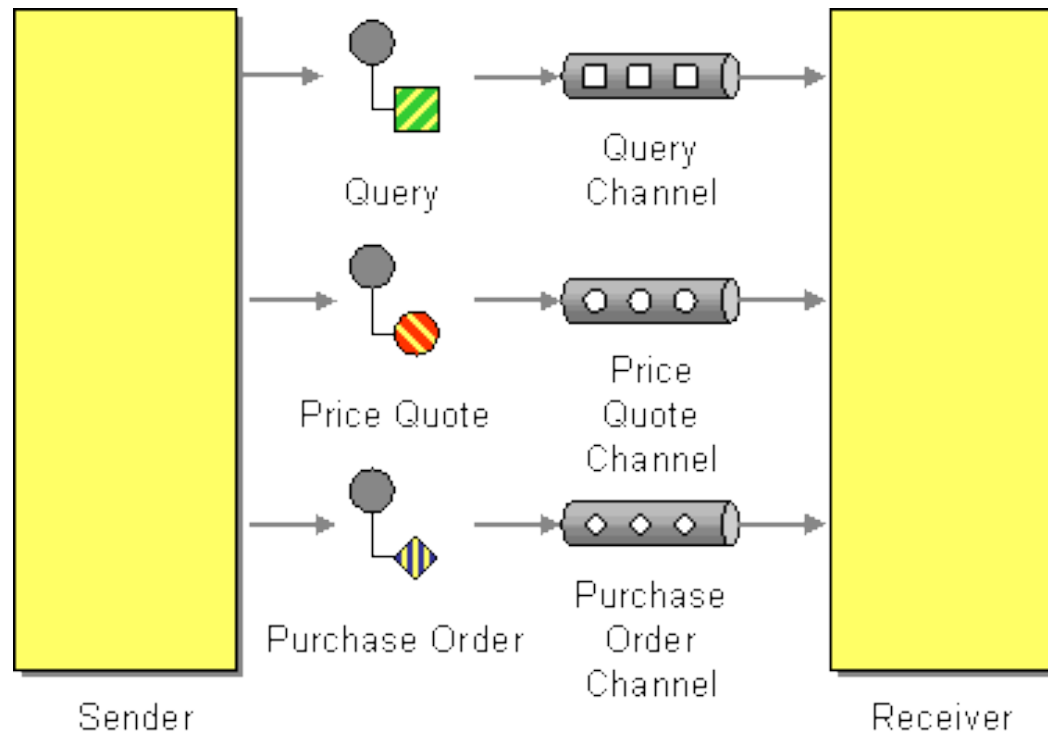
Messaging is a 'pipe' not a 'bucket'.

Point-to-Point Channel



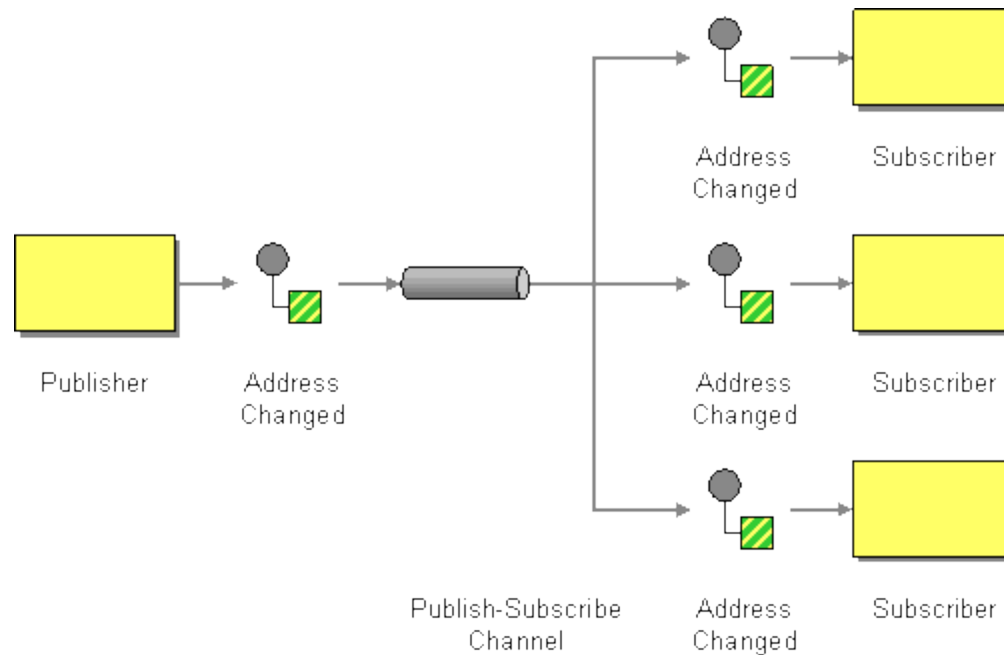
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/PointToPointChannel.html>

Datatype Channel



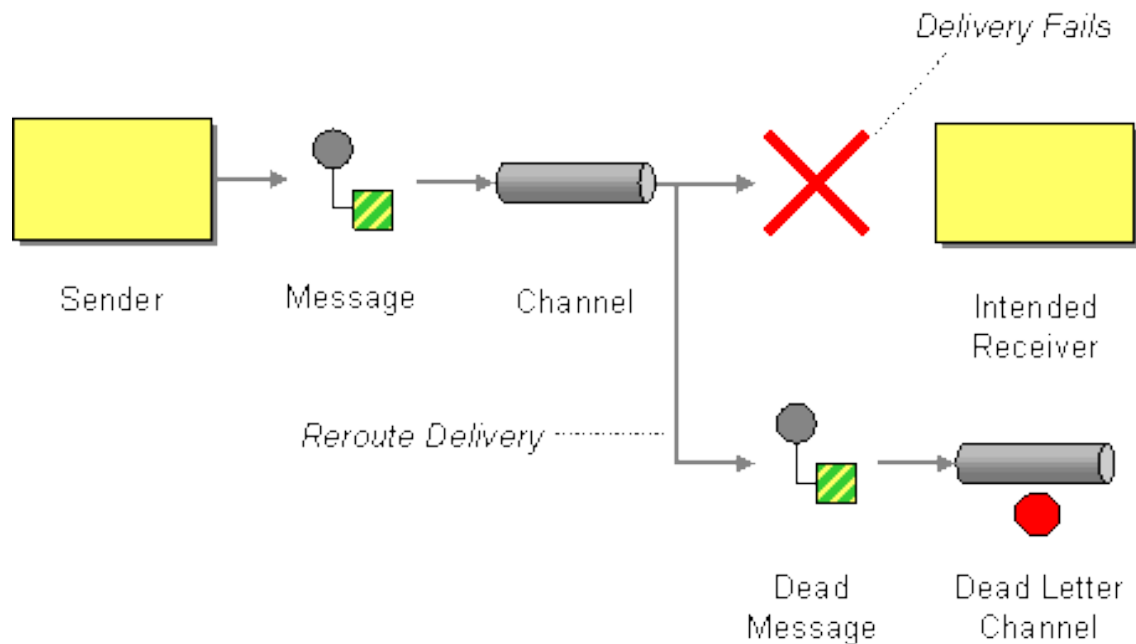
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/DatatypeChannel.html>

Publish-Subscribe Channel



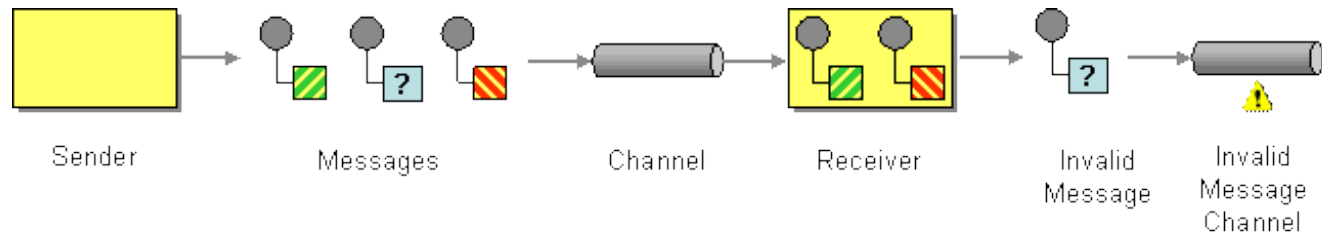
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/PublishSubscribeChannel.html>

Dead Letter Channel



<http://www.enterpriseintegrationpatterns.com/patterns/messaging/DeadLetterChannel.html>

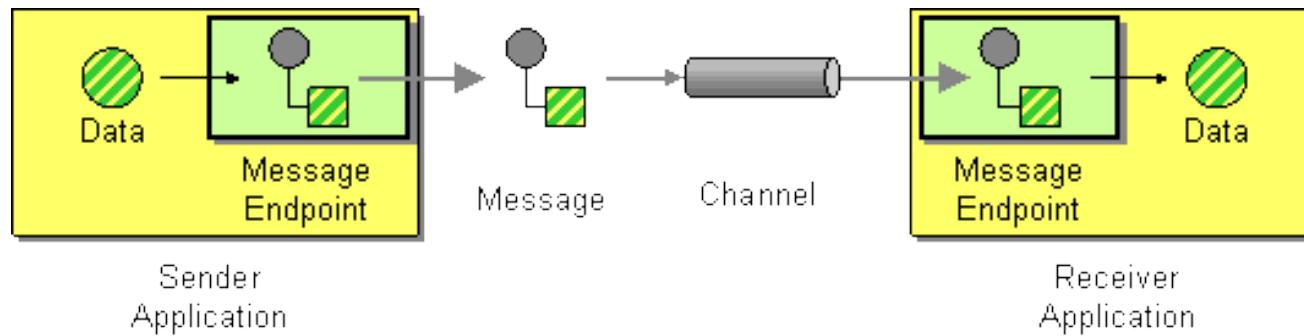
Invalid Message Channel



<http://www.enterpriseintegrationpatterns.com/patterns/messaging/InvalidMessageChannel.html>

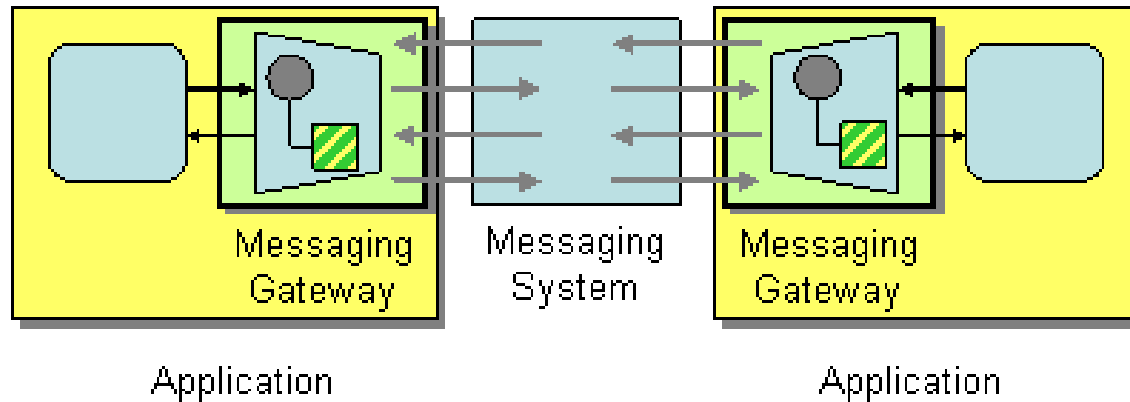
ENDPOINTS

Message Endpoint



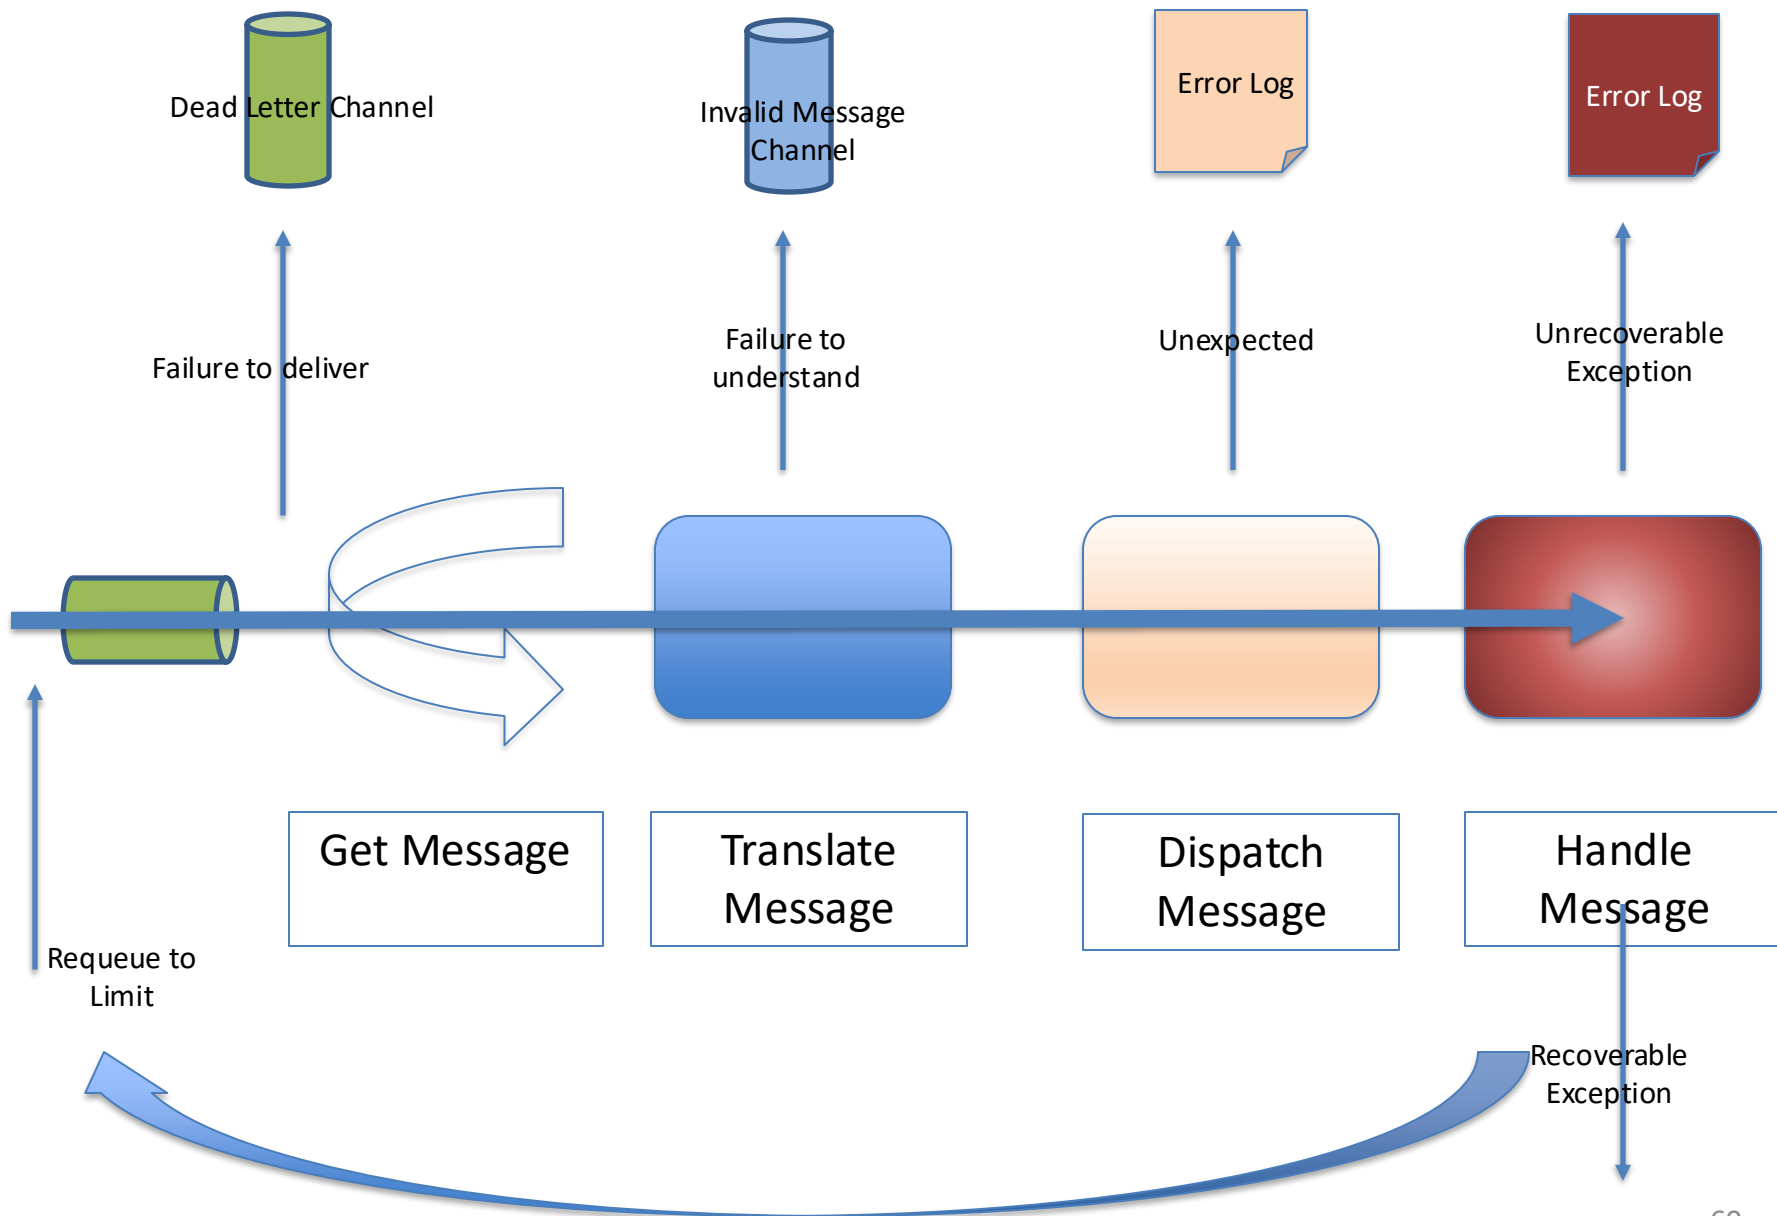
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/MessageEndpoint.html>

Messaging Gateway

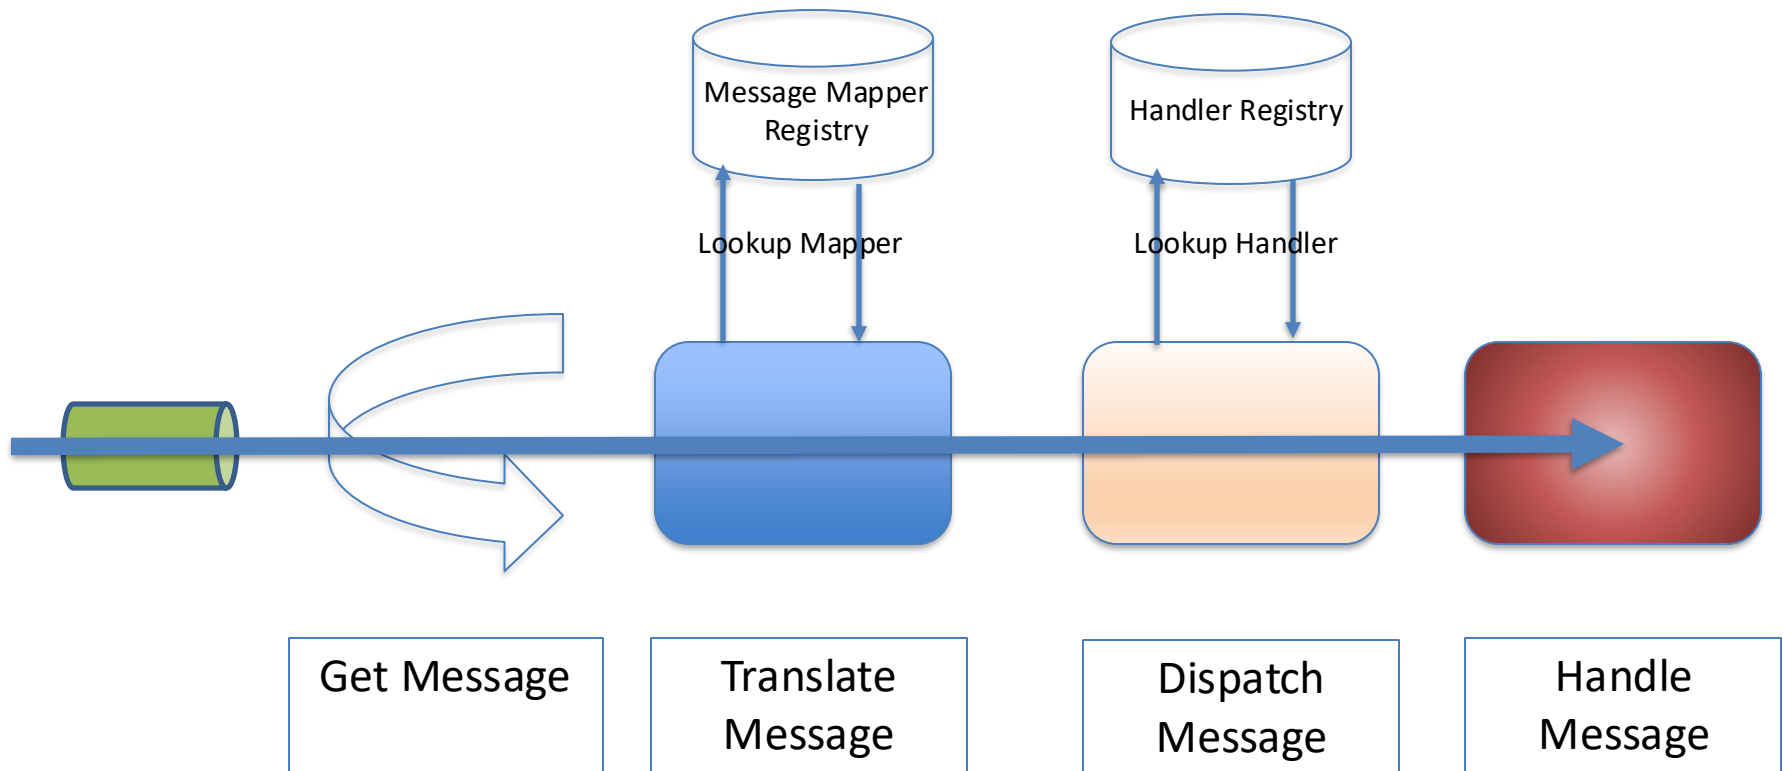


<http://www.enterpriseintegrationpatterns.com/patterns/messaging/MessagingGateway.html>

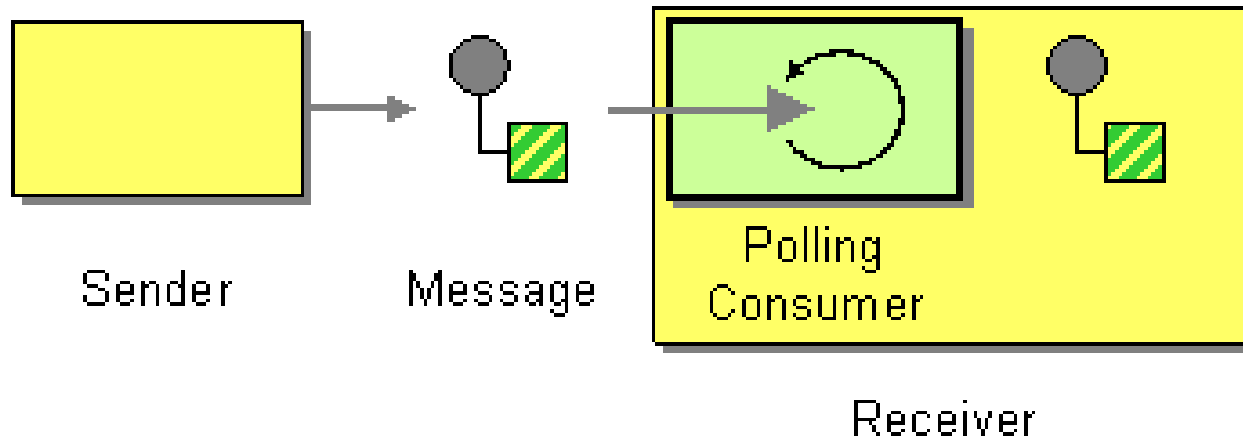
THE MESSAGE PUMP



Translate and Dispatch

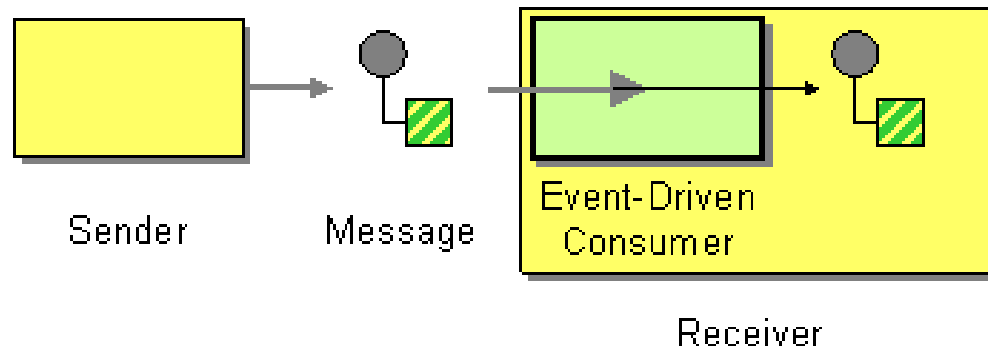


Polling Consumer



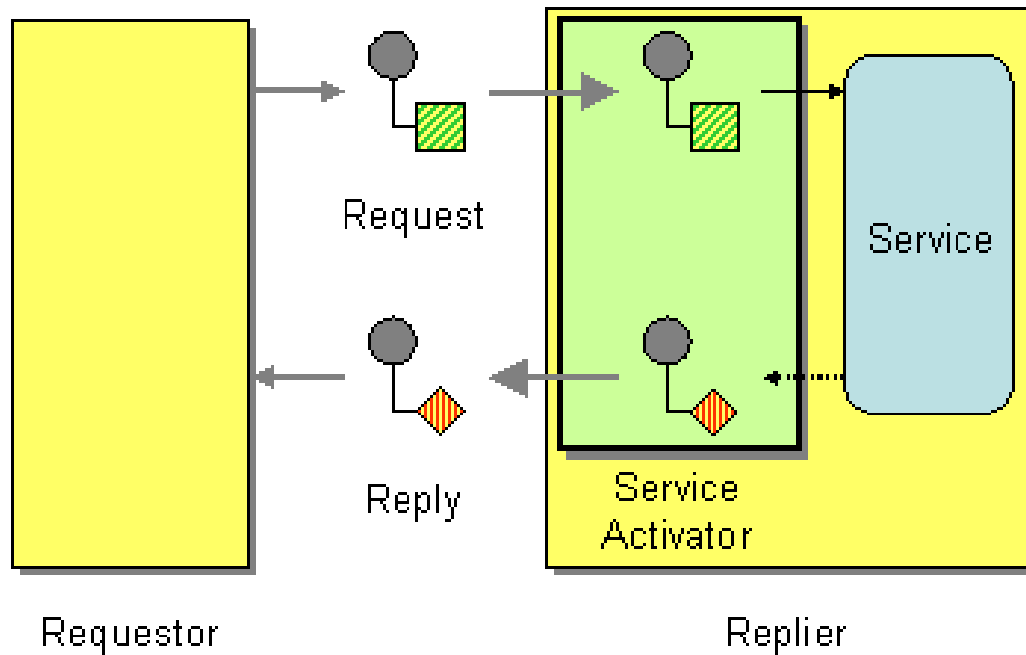
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/PollingConsumer.html>

Event Driven Consumer



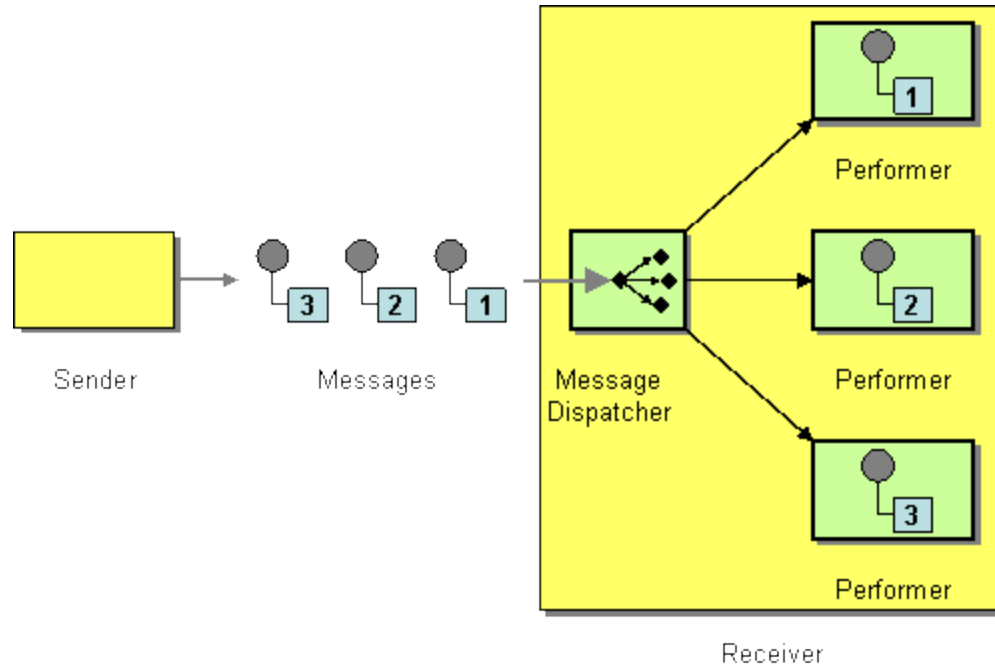
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/EventDrivenConsumer.html>

Service Activator



<http://www.enterpriseintegrationpatterns.com/patterns/messaging/MessagingAdapter.html>

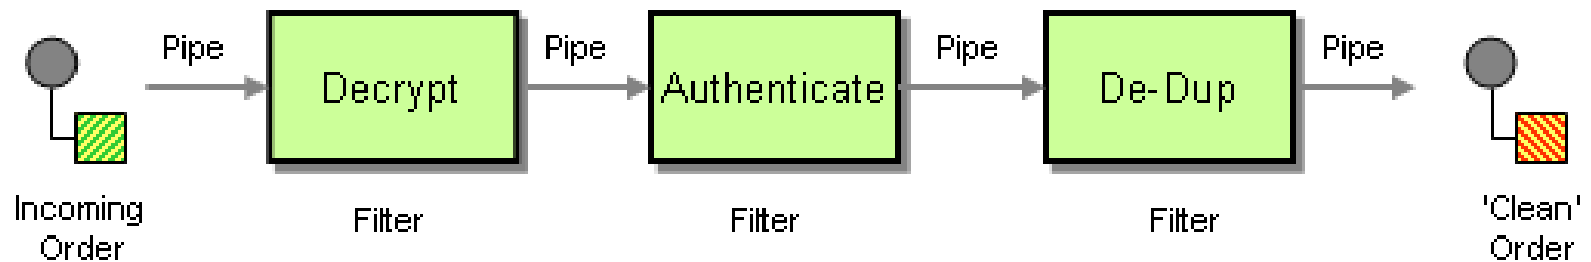
Competing Consumers



<http://www.enterpriseintegrationpatterns.com/patterns/messaging/MessageDispatcher.html>

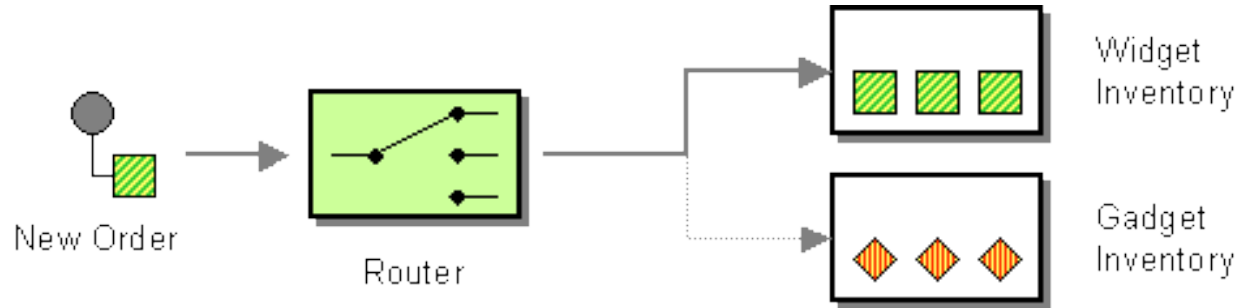
PIPELINES

Pipes and Filters



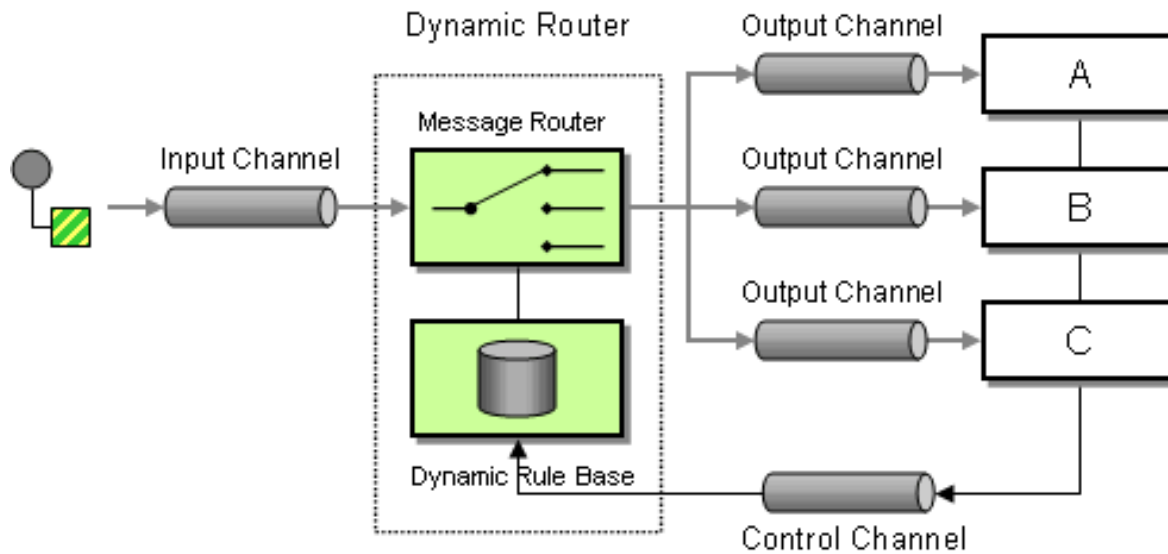
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/PipesAndFilters.html>

Content Based Router



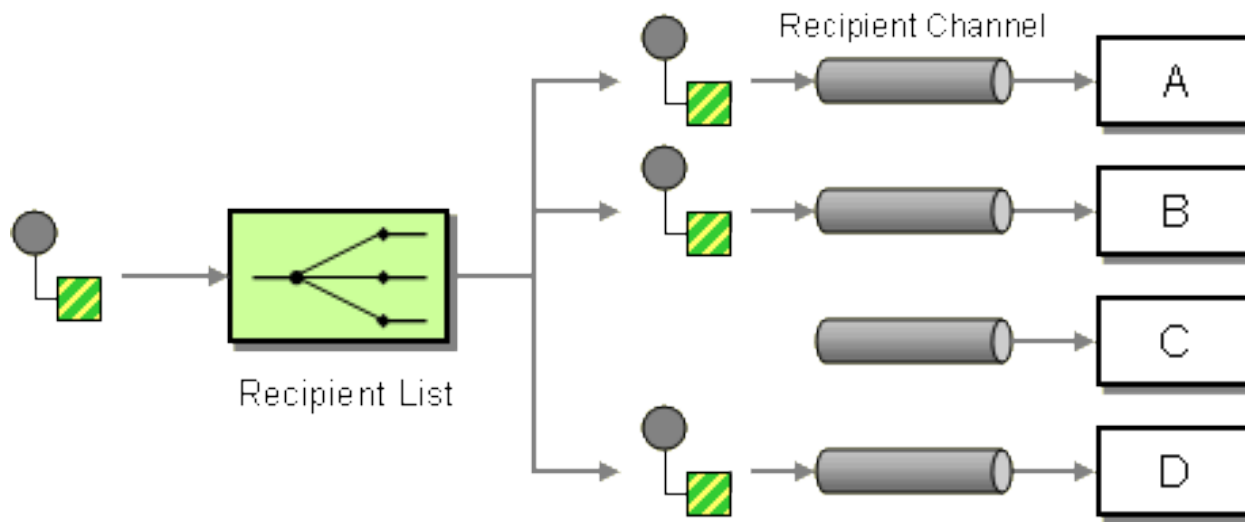
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/ContentBasedRouter.html>

Dynamic Router



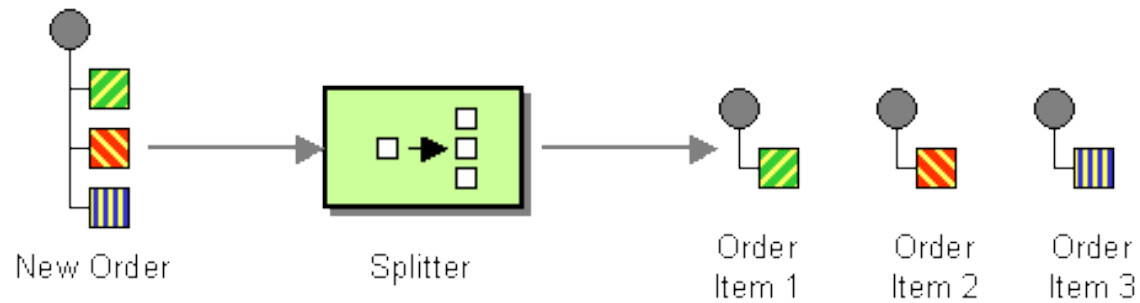
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/DynamicRouter.html>

Recipient List



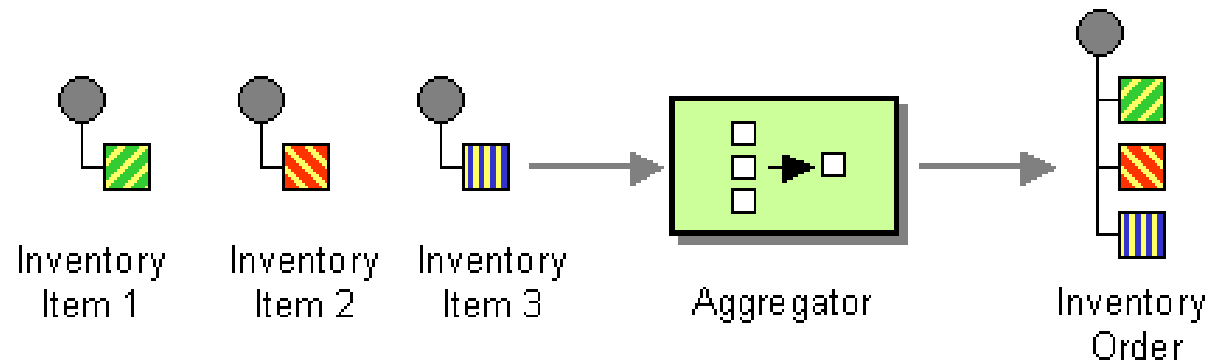
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/RecipientList.html>

Splitter



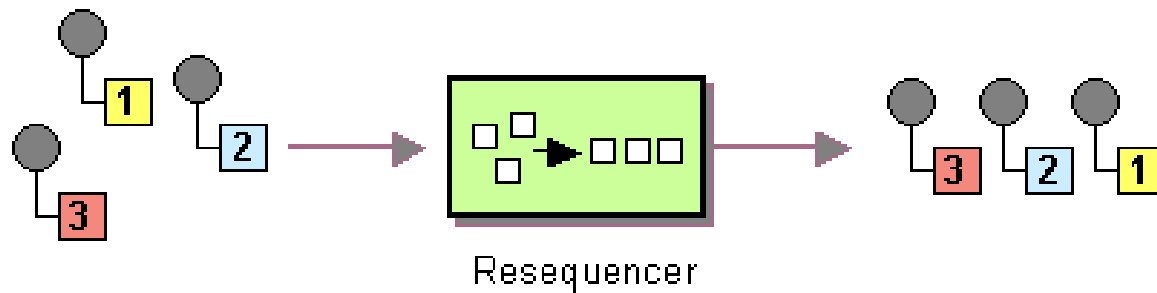
<http://www.enterpriseintegrationpatterns.com/patterns/messaging/Sequencer.html>

Aggregator



<http://www.enterpriseintegrationpatterns.com/patterns/messaging/Aggregator.html>

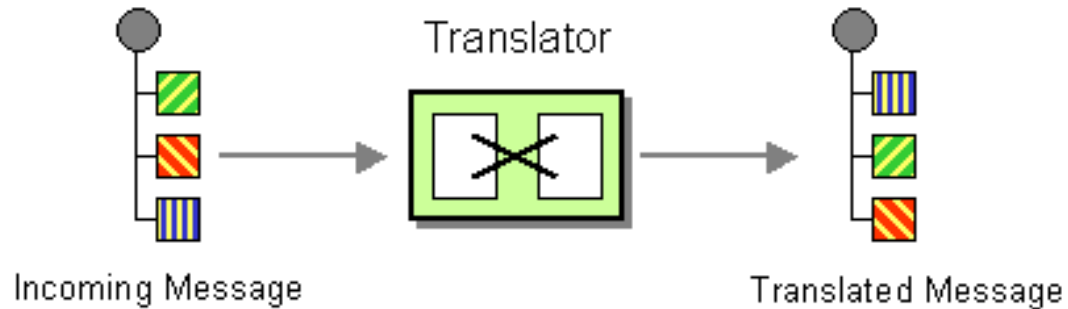
Resequencer



<http://www.enterpriseintegrationpatterns.com/patterns/messaging/Resequencer.html>

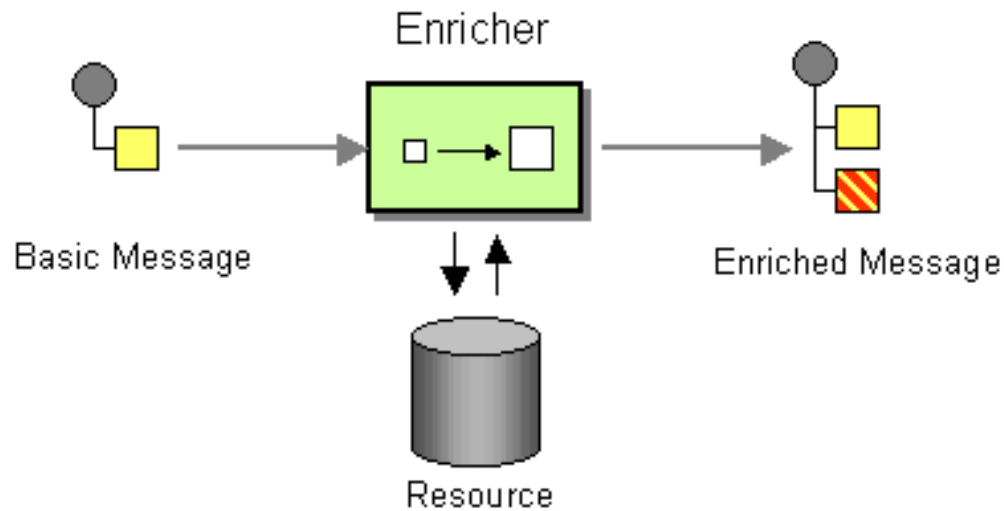
TRANSFORMATION

Message Translator



<http://www.enterpriseintegrationpatterns.com/patterns/messaging/MessageTranslator.html>

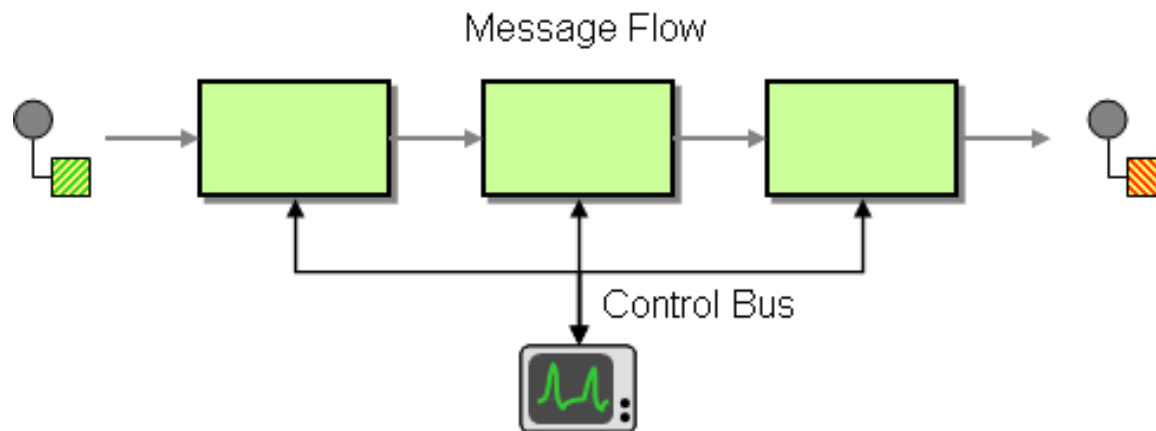
Content Enricher



<http://www.enterpriseintegrationpatterns.com/patterns/messaging/DataEnricher.html>

MANAGEMENT

Control Bus

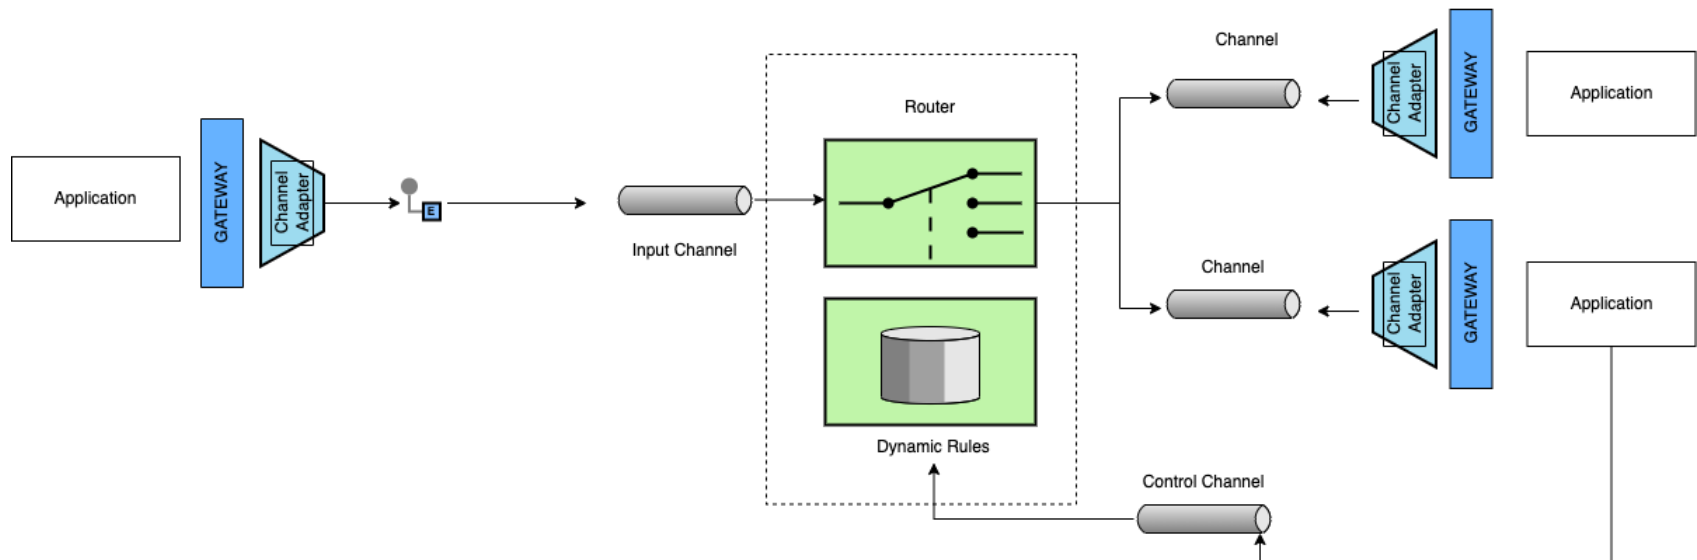


<http://www.enterpriseintegrationpatterns.com/patterns/messaging/ControlBus.html>

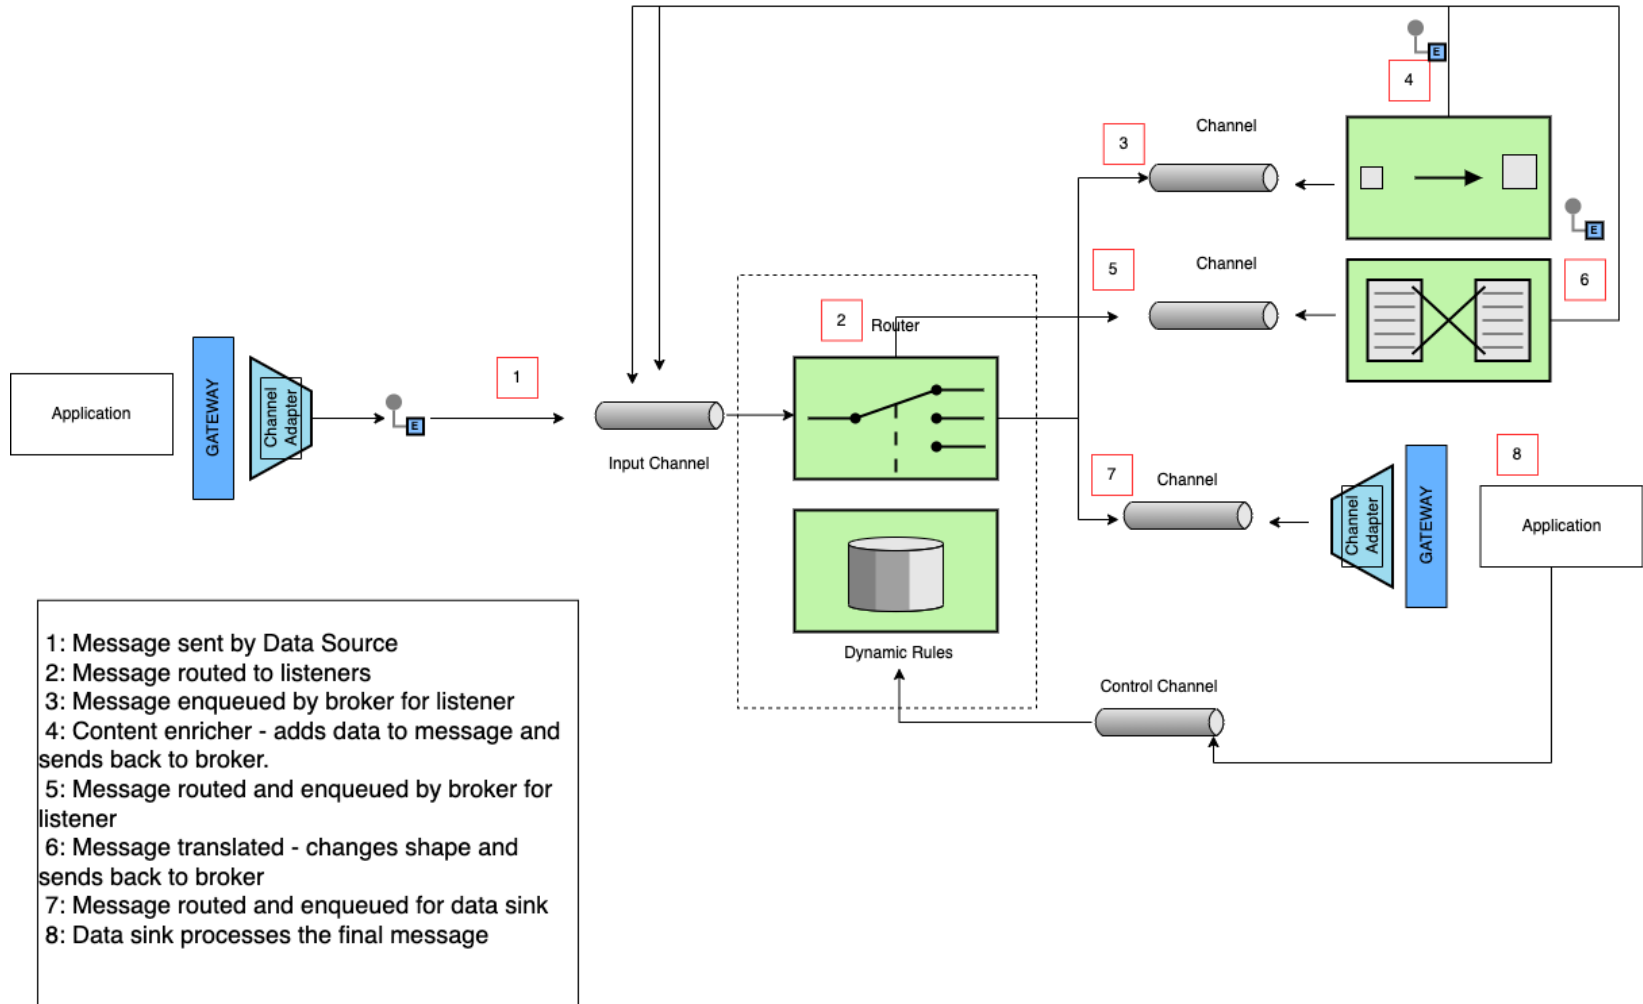
Integrating using events

BROKERS AND PIPELINES

Messaging with A Broker

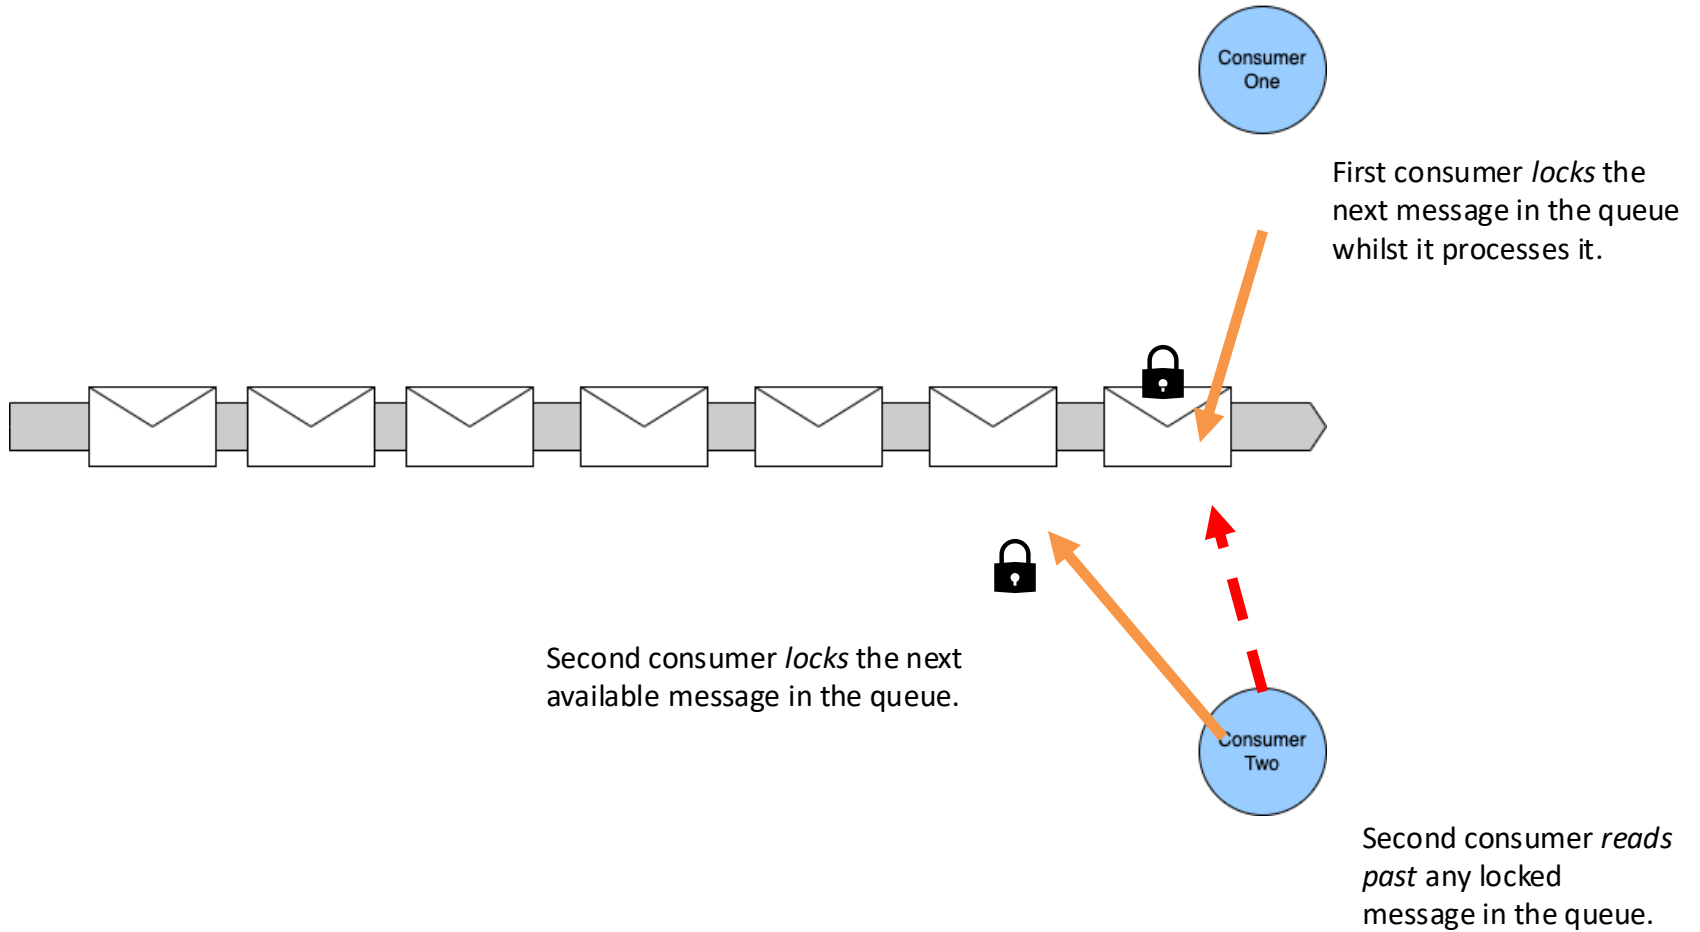


Pipes and Filters



QUEUES AND STREAMS

Queues

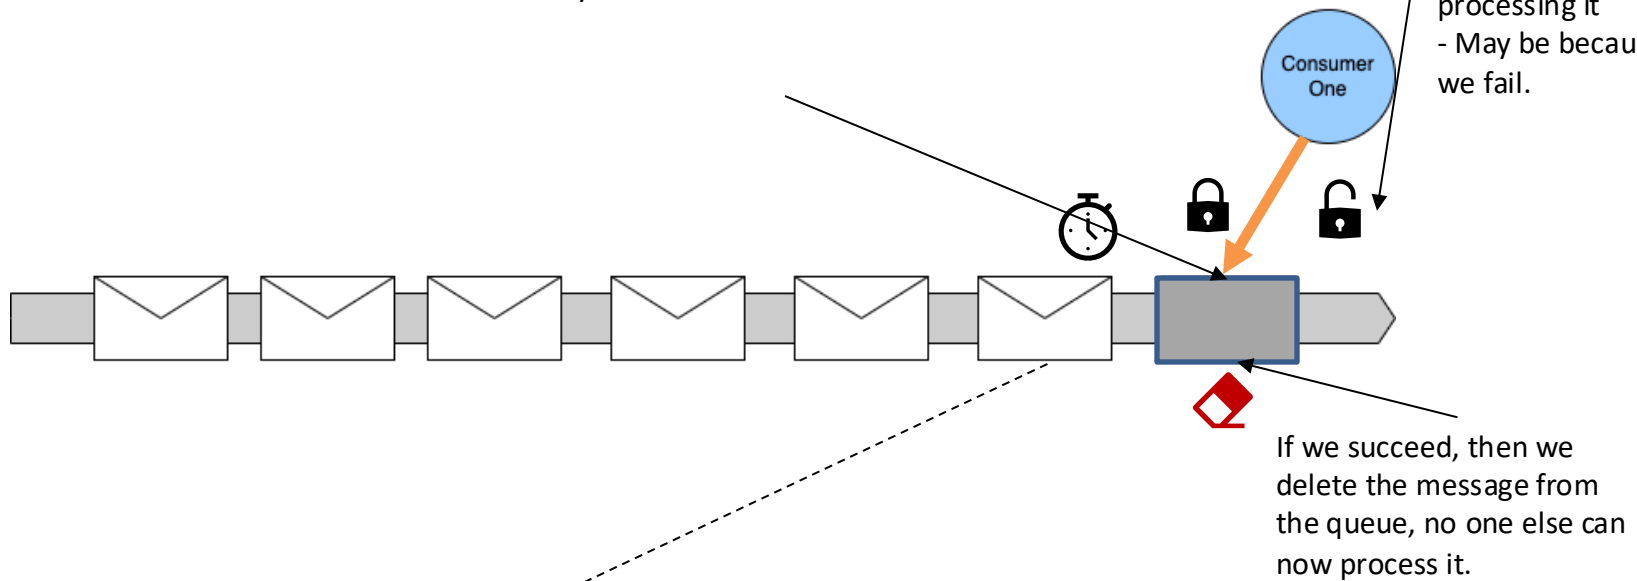


Queues

If we fail, then we may decide others could succeed later, so we let it become available to lock again, often with a delay.

When we are done processing, we unlock it.

- Usually because we finish processing it
- May be because we fail.



If we succeed, then we delete the message from the queue, no one else can now process it.



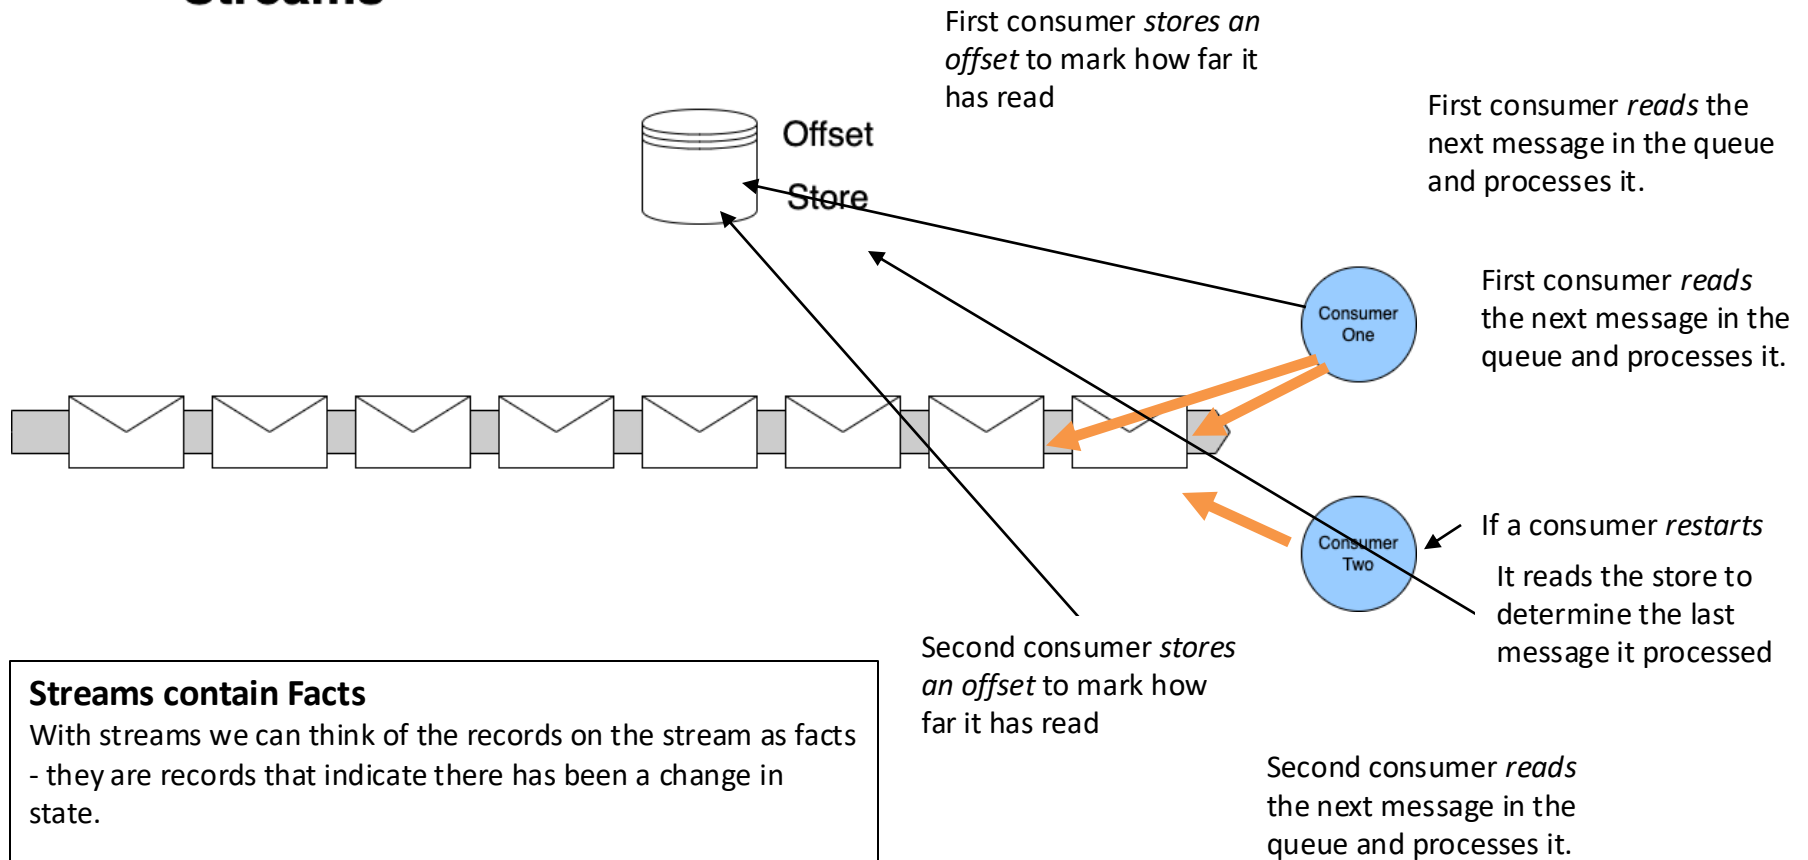
After a certain number of re-queues we may move the message to a dead-letter channel, it turns out that no one action the request within in a reasonable time frame

Queues contain Tasks

With queues we can think of the messages on a queue as tasks - they are a request for us to carry out an action. Once the action is done, we can delete the task.

- We don't anyone else to action it, it's already been done.
- Someone receiving a done task will have to discard it.
- If we can't action it, someone else will need to action it.

Streams



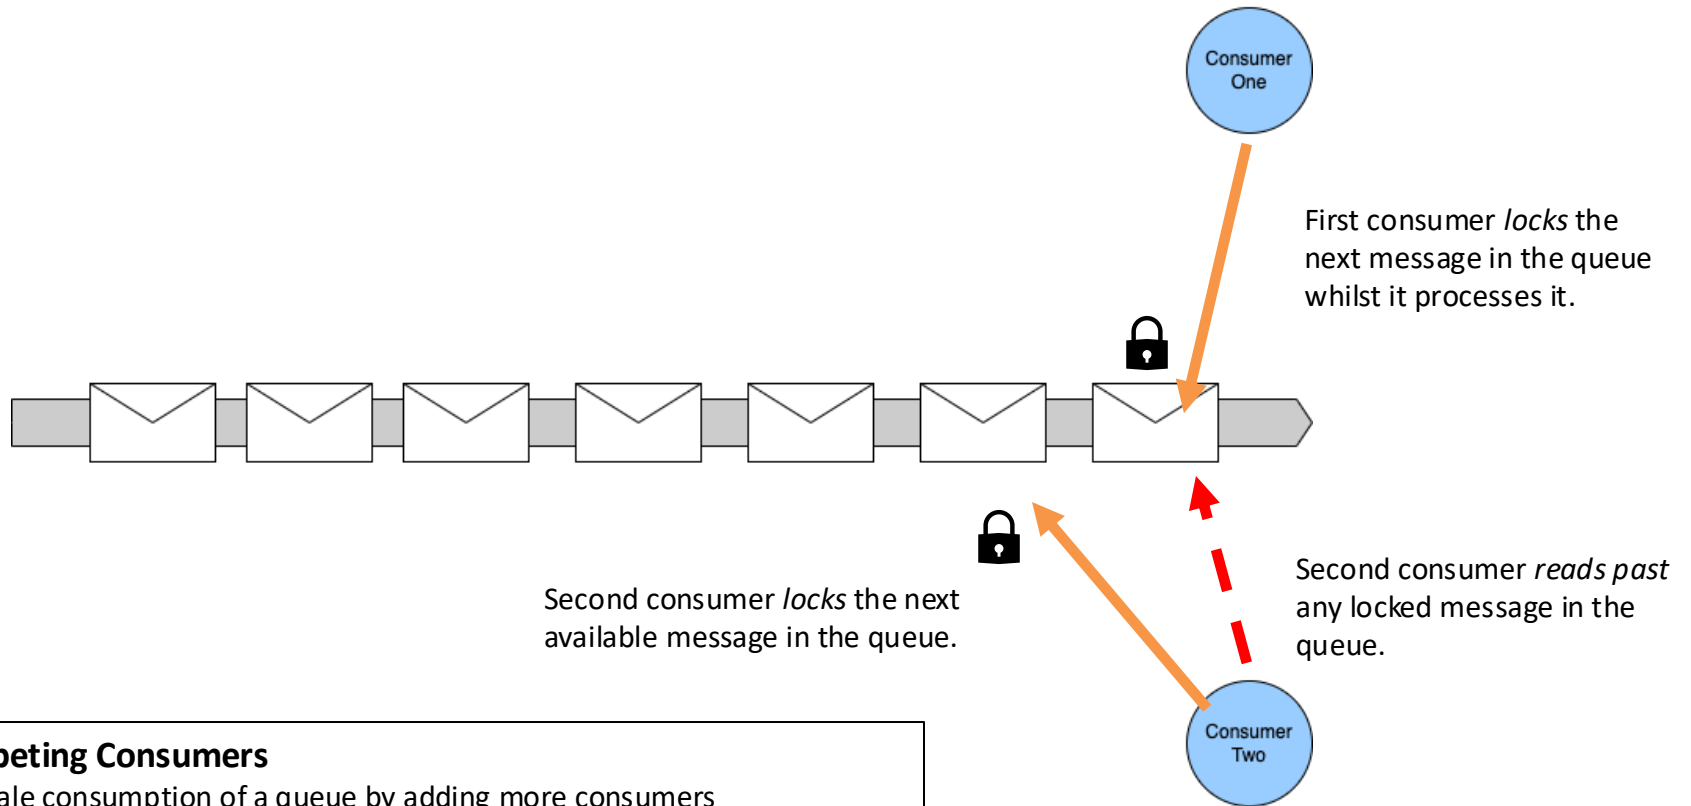
Streams contain Facts

With streams we can think of the records on the stream as facts - they are records that indicate there has been a change in state.

- We can view facts as an 'inverse database' they represent how current stat is arrived at
- We can navigate offsets to calculate a position at a 'point in time'
- We don't consume facts by reading them, they persist

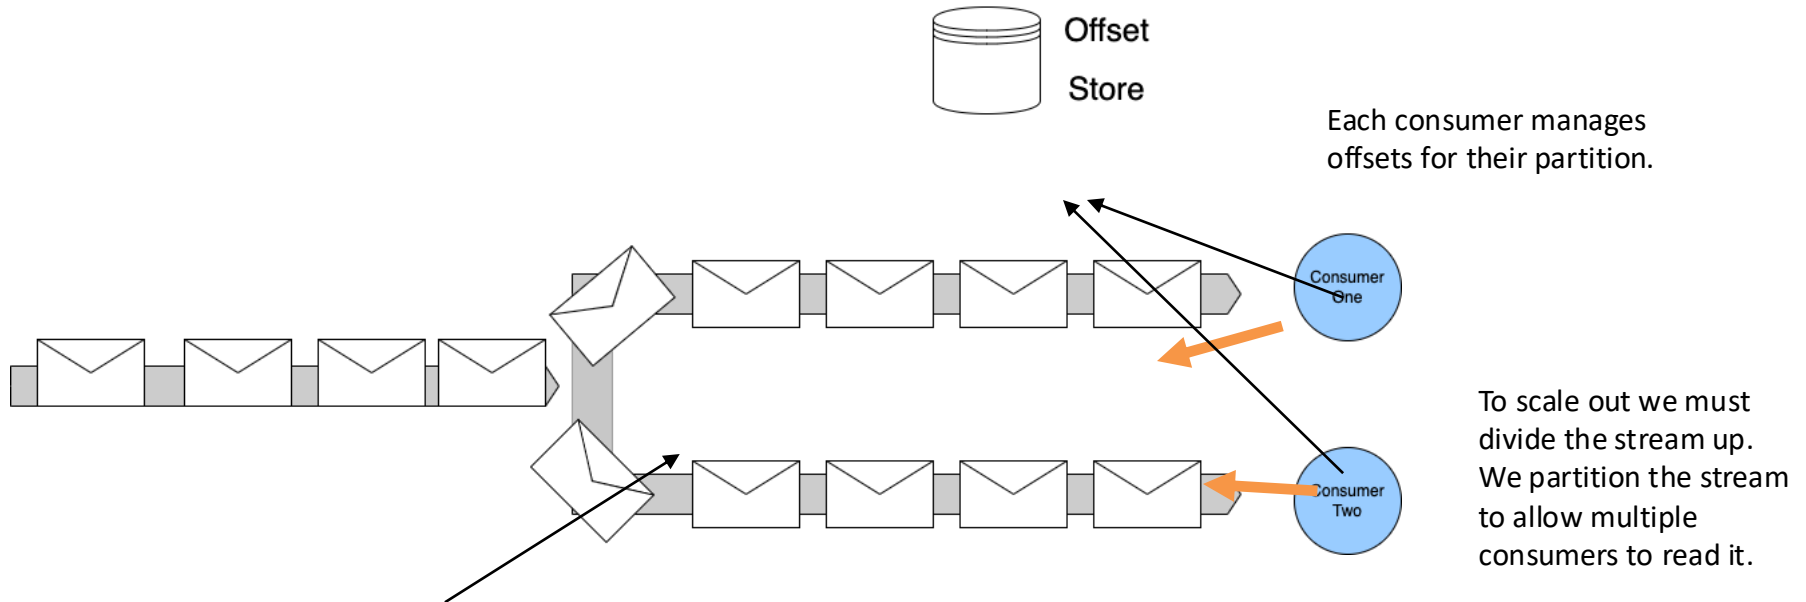
Scaling Queues and Streams

Queues



Streams

Partitions

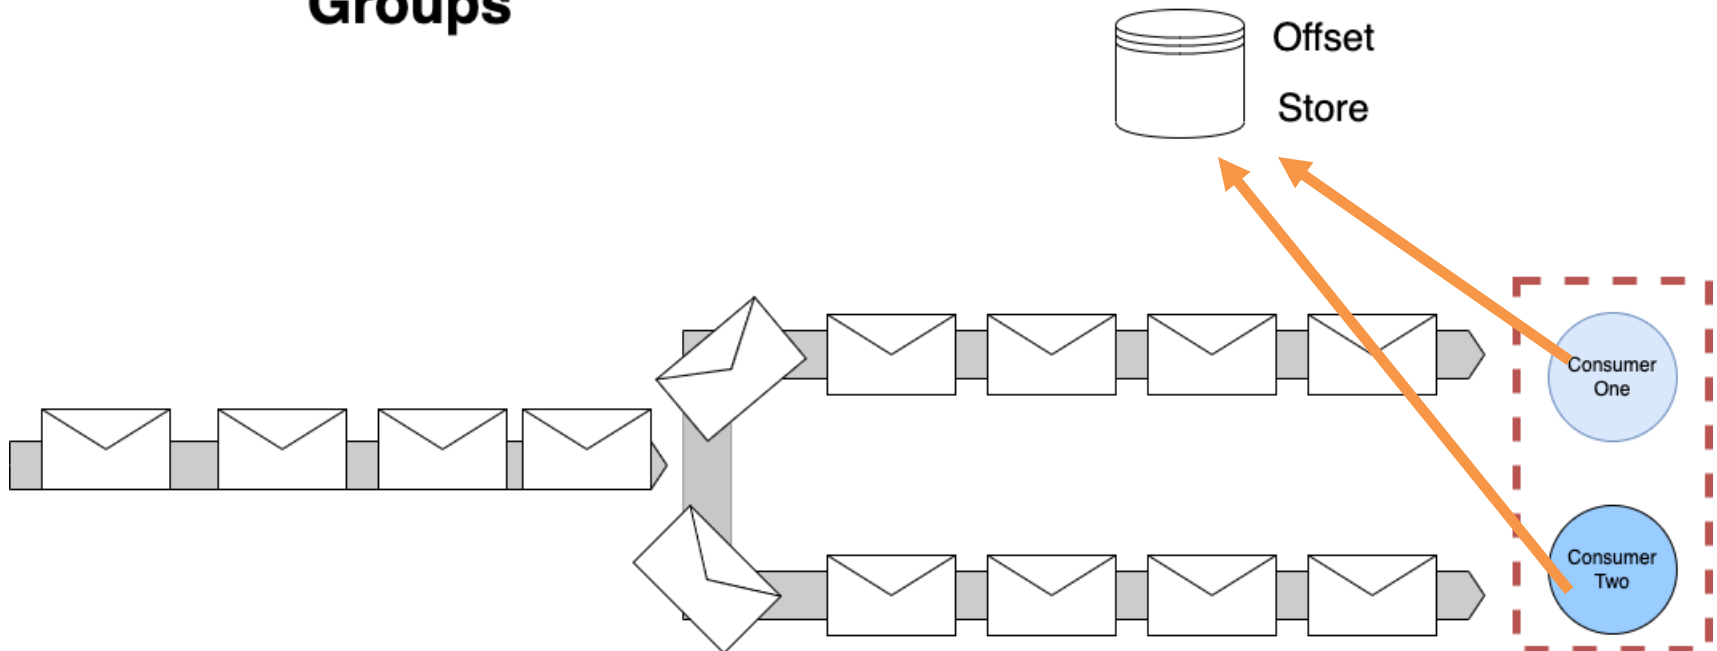


For any set of events that must be processed sequentially - all changes to one entity for example - we use consistent hashing to push messages with the same identifier to the same partition. This allows us to scale, whilst preserving our ordering.

Streams

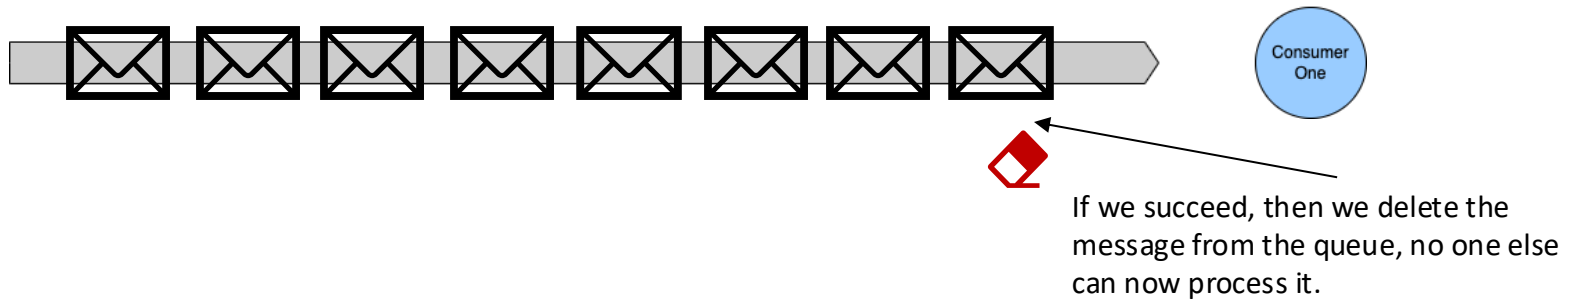
Consumer Groups

To provide availability – only one consumer in a group can read from a partition at a time – but a consumer in a group may read from more than one of the partitions owned by that group



Archive and Replay

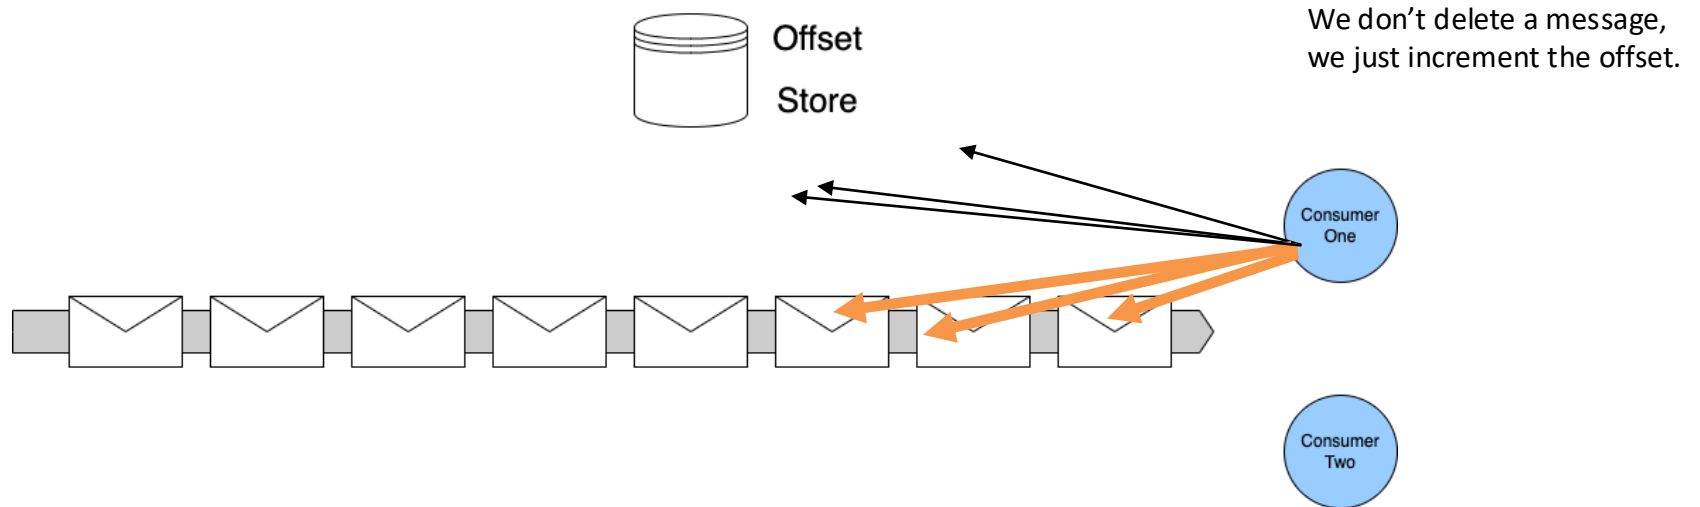
Queues



No Archive and Replay

With queues we delete a message once we have completed the associated action. That means we have no way to replay the request for work. Our only option is to ask the producer to resend their request.

Streams

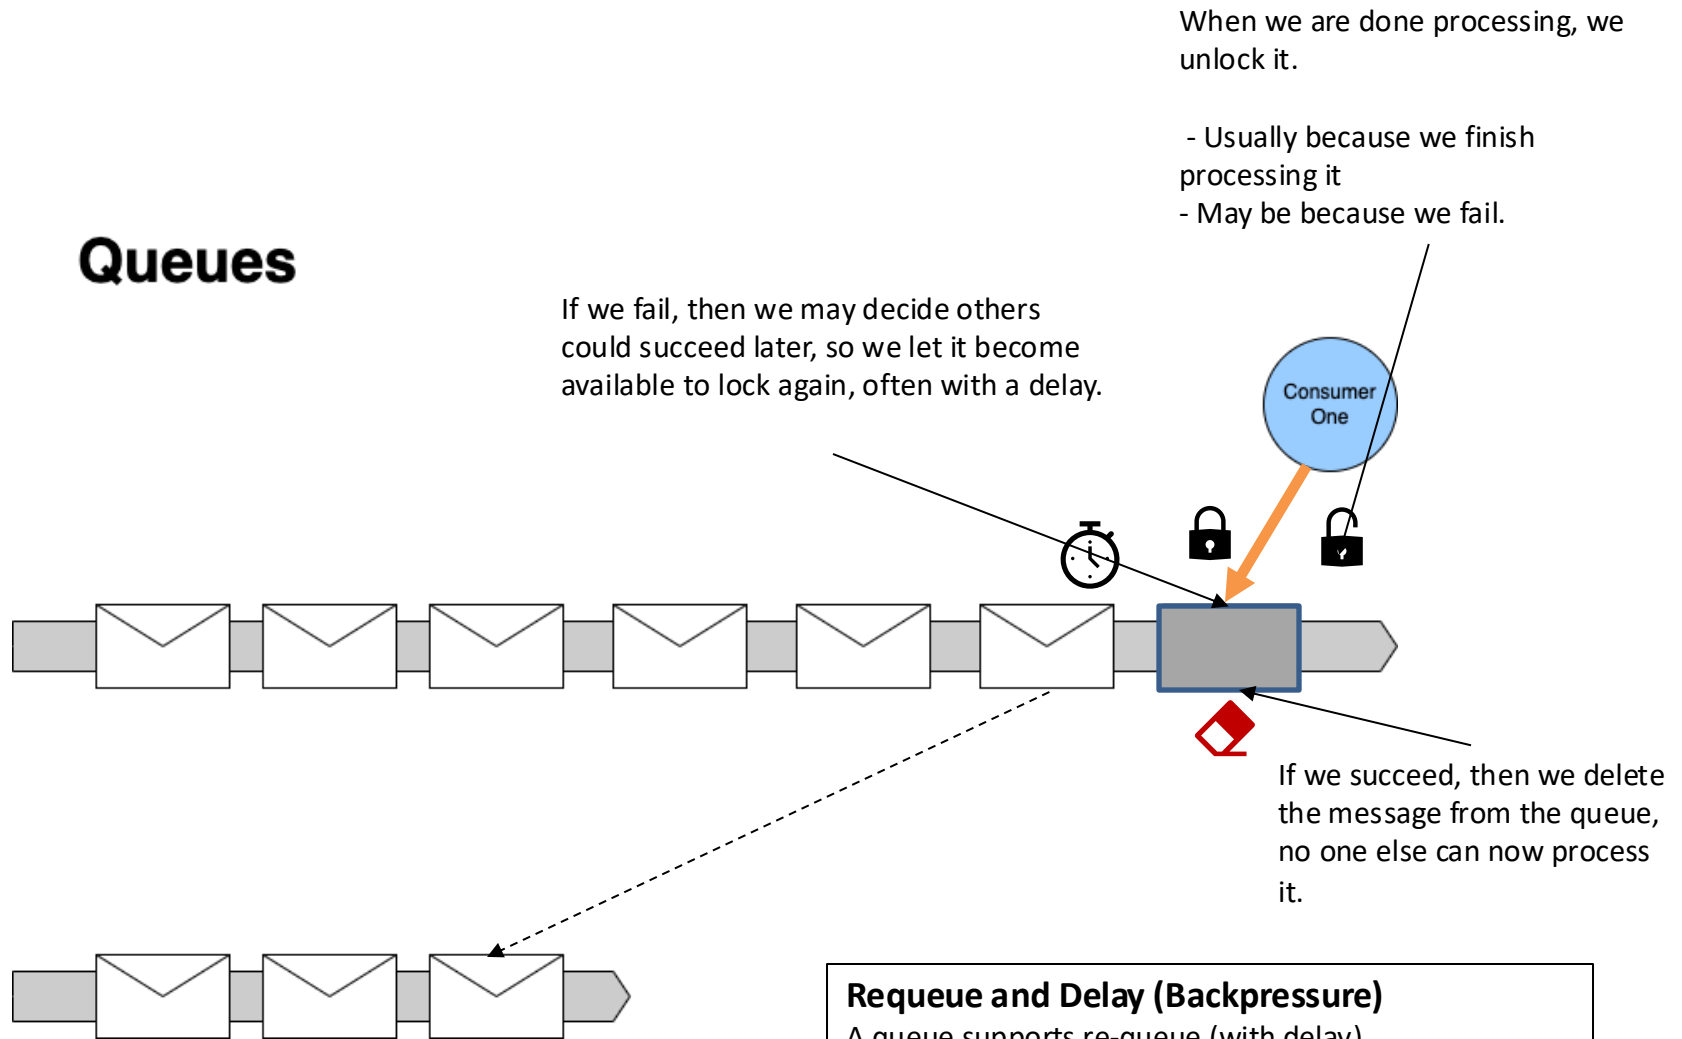


Archive and Replay

Archive and Replay is straightforward as nothing is deleted. We simply reset the consumer's offset to re-read the stream

Requeue and Delay (Backpressure)

Queues



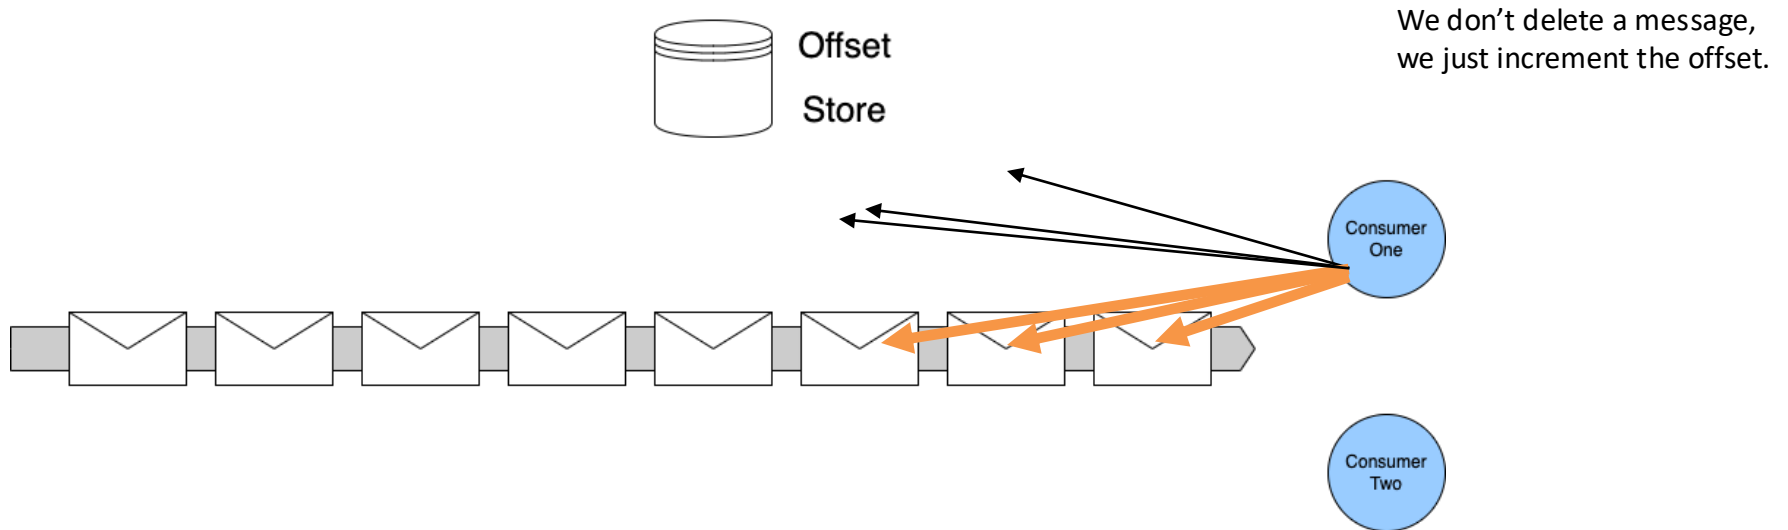
After a certain number of re-queues we may move the message to a dead-letter channel, it turns out that no one action the request within in a reasonable time frame

Requeue and Delay (Backpressure)

A queue supports re-queue (with delay)

- If the work is not done/acked, just make it available again to the next consumer
- If the work could not be done because of a transient issue, delay to let it pass

Streams



No Requeue or DLQ

Because we do not lock items, we do not requeue items, including requeue with delay. Your strategy is:

- Ignore and Continue (Load Shedding)
- Retry (Backpressure)
- Copy to another stream (a delay or DLQ stream)

	Messaging	Discrete Event	Series Event
Queue	✓	✓	✗
Stream	✗	✓	✓

	Ordering	Archive and Replay	Requeue with Delay
Queue	✓ ✗	✗	✓
Stream	✓	✓	✗

EXERCISE MATERIAL

Introduction to Kafka

- [Readme](#)
- [Slides](#)



DON'T PANIC

MANAGING ASYNCHRONOUS APIS

Versioning

Be strict when sending and tolerant when receiving.
Implementations must follow specifications precisely when sending to the network, and tolerate faulty input from the network.

Robustness Principal or Postel's Law – Jon Postel RFC 1958

Tolerant Reader

```
{
  "$schema": "http://json-schema.org/draft-07/schema#",
  "type": "object",
  "properties": {
    "orderid": {
      "type": "string"
    },
    "customerName": {
      "type": "string"
    },
    "addressLineOne": {
      "type": "string"
    },
    "postCode": {
      "type": "string"
    },
    "pinCode": {
      "type": "string",
      "pattern": "^[0-9]+$"
    }
  },
  "required": ["orderid", "customerName", "addressLineOne",
    "postCode", "pinCode"],
  "additionalProperties": false
}
```



```
{
  "$schema": "http://json-schema.org/draft-07/schema#",
  "type": "object",
  "properties": {
    "orderid": {
      "type": "string"
    },
    "customerName": {
      "type": "string"
    },
    "addressLineOne": {
      "type": "string"
    },
    "postCode": {
      "type": "string"
    },
    "pinCode": {
      "type": "string",
      "pattern": "^[0-9]+$"
    },
    "latitude": {
      "type": "number",
      "minimum": -90,
      "maximum": 90
    },
    "longitude": {
      "type": "number",
      "minimum": -180,
      "maximum": 180
    }
  },
  "required": ["orderid", "customerName", "addressLineOne",
    "postCode", "pinCode"],
  "additionalProperties": false
}
```

Ignore New Fields

Tolerant Reader

```
{
  "$schema": "http://json-schema.org/draft-07/schema#",
  "type": "object",
  "properties": {
    "orderid": {
      "type": "string"
    },
    "customerName": {
      "type": "string"
    },
    "addressLineOne": {
      "type": "string"
    },
    "postCode": {
      "type": "string"
    },
    "pinCode": {
      "type": "string",
      "pattern": "^[0-9]+$"
    },
    "latitude": {
      "type": "number",
      "minimum": -90,
      "maximum": 90
    },
    "longitude": {
      "type": "number",
      "minimum": -180,
      "maximum": 180
    }
  },
  "required": ["orderid", "customerName", "addressLineOne",
    "postCode", "pinCode"],
  "additionalProperties": false
}
```

Default Latitude: 0
Default Longitude: 0

Note: not required

Default Missing Fields

Breaking Change

```
{
  "$schema": "http://json-schema.org/draft-07/schema#",
  "type": "object",
  "properties": {
    "orderid": {
      "type": "string"
    },
    "firstName": {
      "type": "string"
    },
    "surName": {
      "type": "string"
    },
    "addressLineOne": {
      "type": "string"
    },
    "pinCode": {
      "type": "string",
      "pattern": "^[0-9]+$"
    },
    "latitude": {
      "type": "number",
      "minimum": -90,
      "maximum": 90
    },
    "longitude": {
      "type": "number",
      "minimum": -180,
      "maximum": 180
    }
  },
  "required": ["orderid", "firstName", "surName",
    "addressLineOne", "pinCode", "latitude", "longitude"],
  "additionalProperties": false
}
```

We might be able to write code to deal with this change, but we have to know that a required field is missing and we have new fields instead

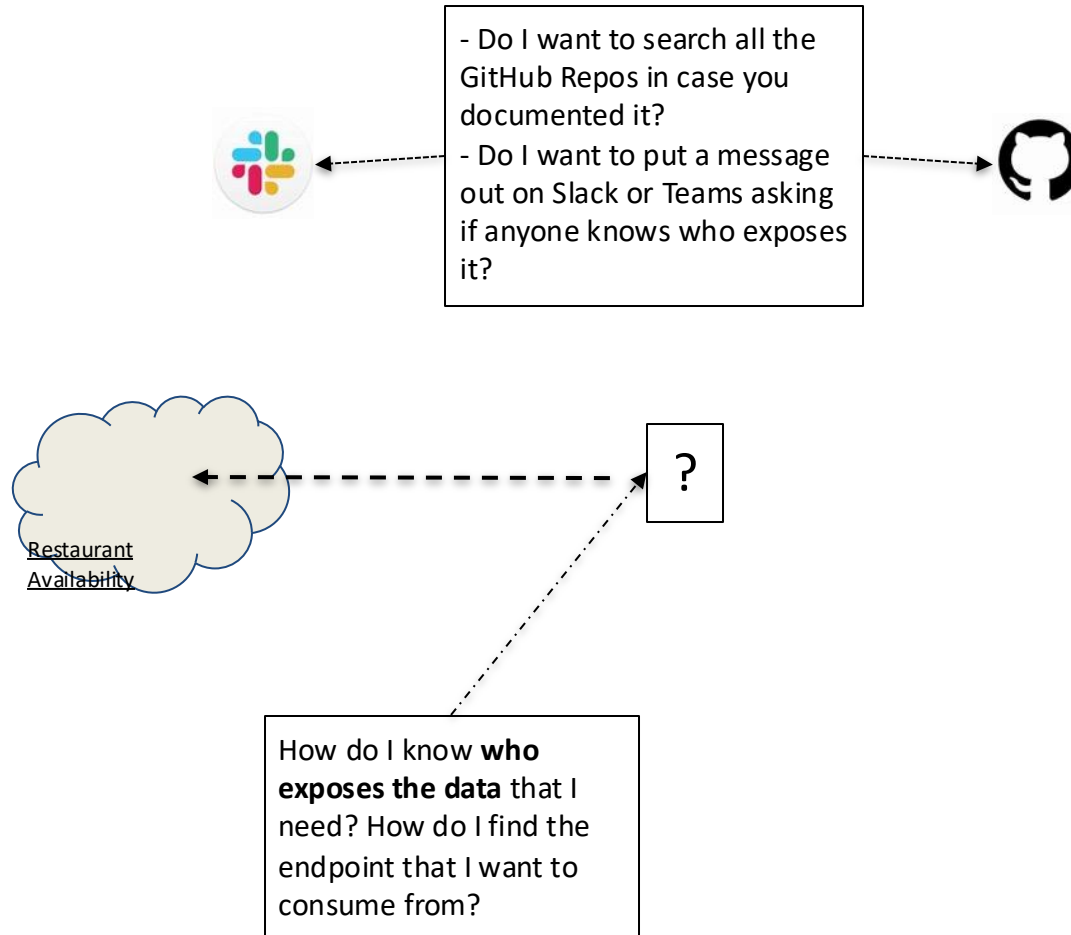
For this we need to rely on a version in the header, and the ability to process messages with this new version, alongside old ones to allow us to run out the old until new replaces it.

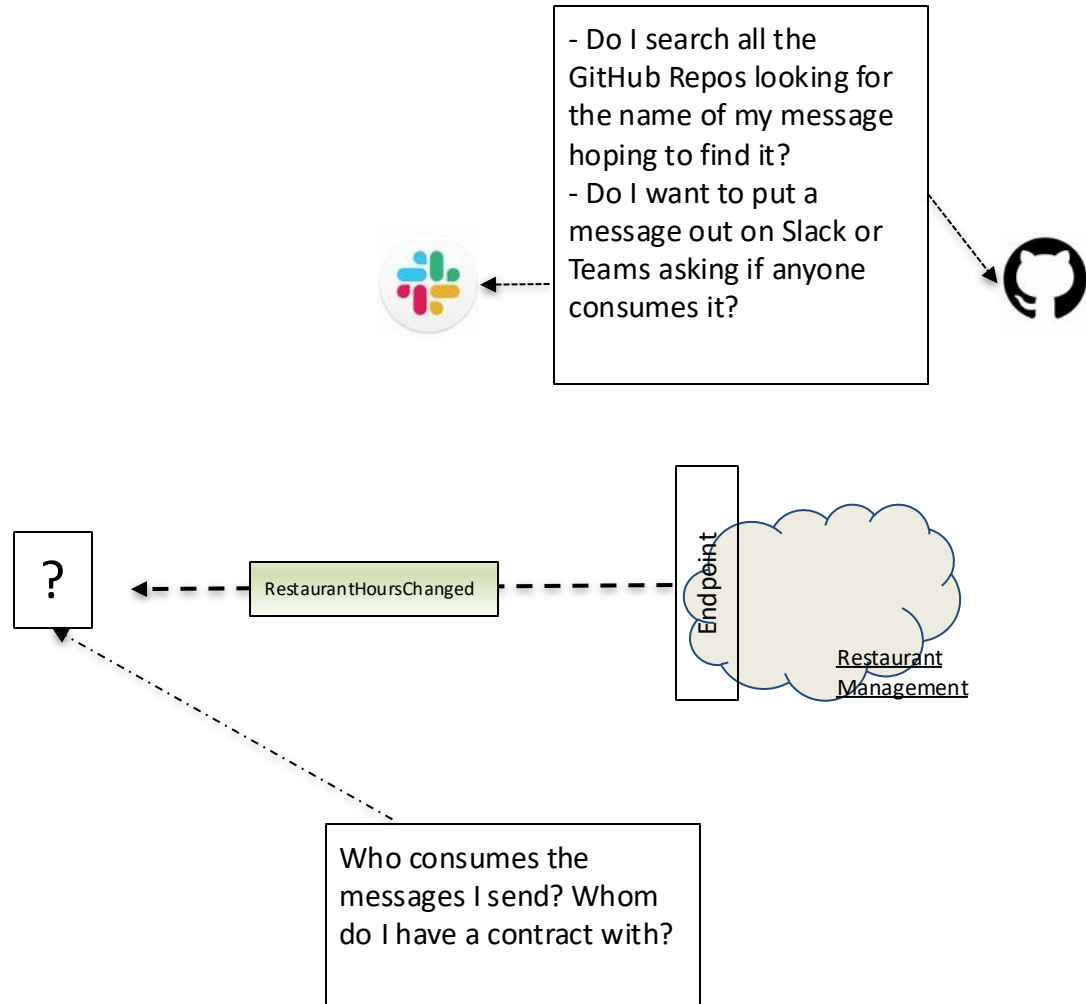
Documentation

Endpoints are **places where messages are sent or received** (or both), and they define all the information required for the message exchange.

An *endpoint* describes in a standard-based way **where messages should be sent, how they should be sent, and what the messages should look like.**

<https://docs.microsoft.com/en-us/dotnet/framework/wcf/fundamental-concepts>

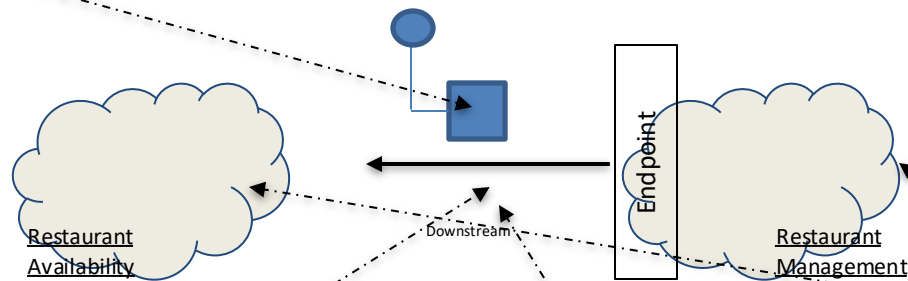




Asynchronous Endpoints

Our Asynchronous APIs need documenting just like any other API (HTTP etc).

We need to document the message, because it is the contract

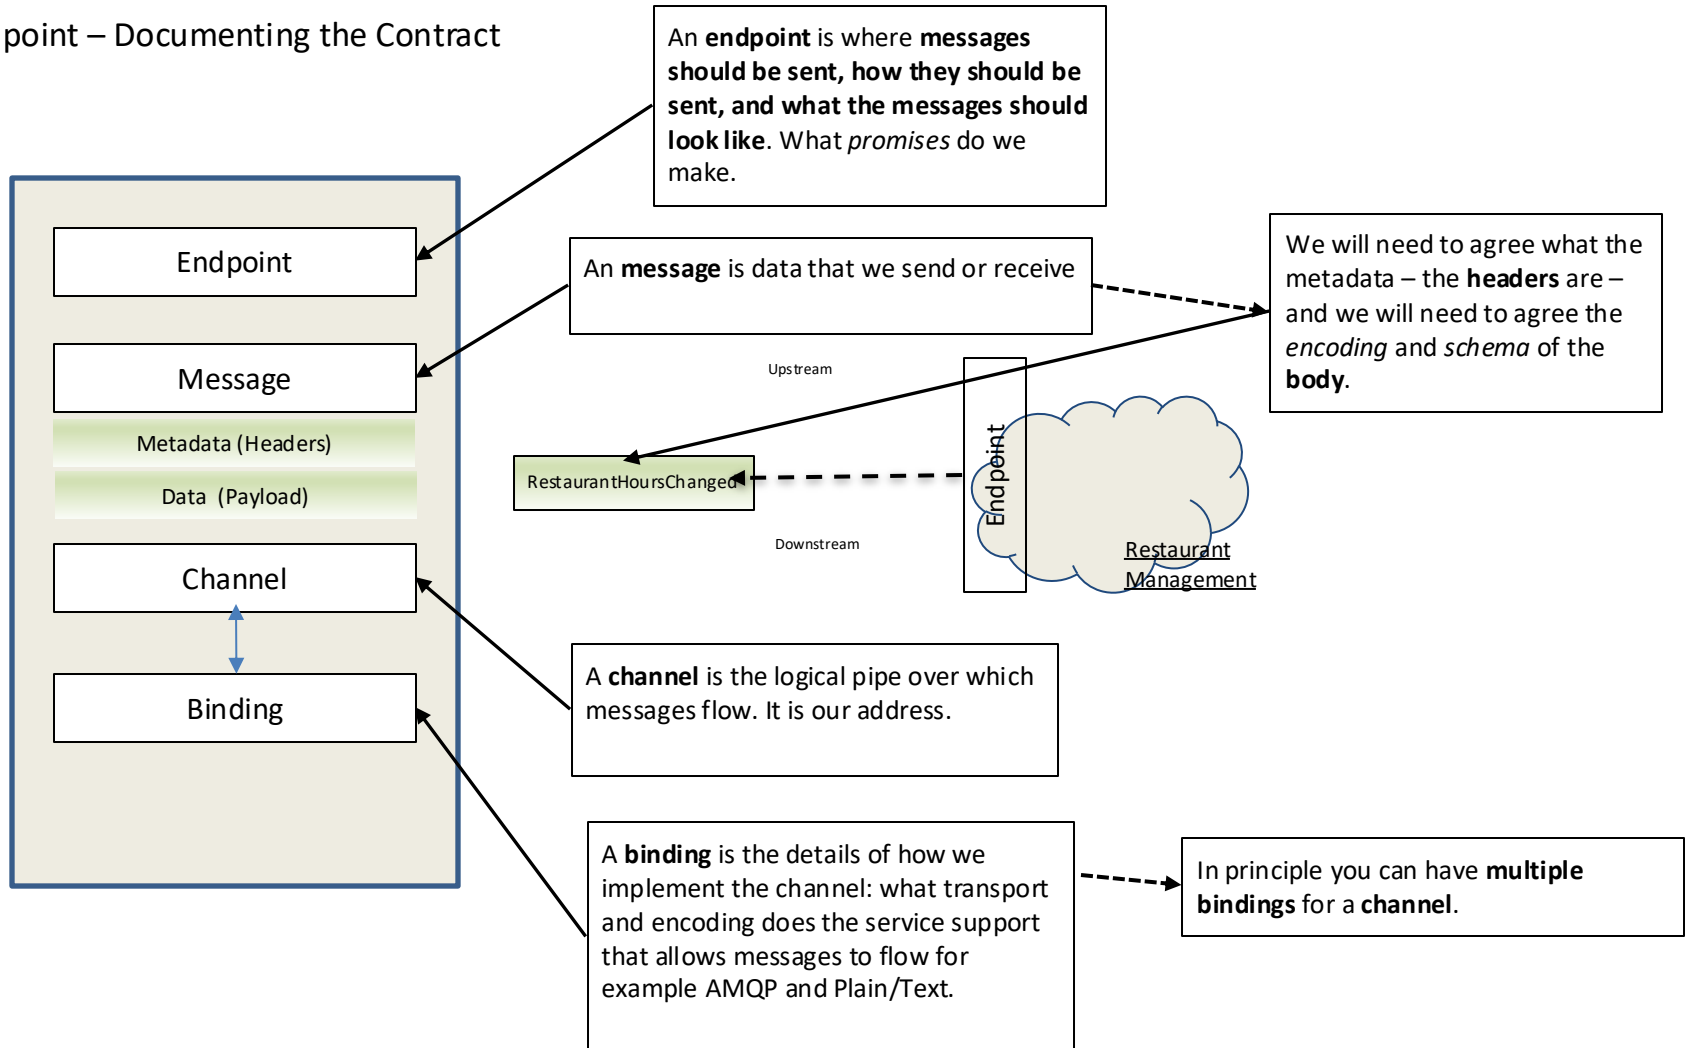


We need to document the channel so we know where the message is flowing

We need to document the protocol so we know how to send-receive

We need to document who sends and receives to understand flow

Endpoint – Documenting the Contract



Why AsyncAPI?

Improving the current state of Event-Driven Architectures (EDA)

Specification

Allows you to define the interfaces of asynchronous APIs and is protocol agnostic.

[Documentation](#)

Document APIs

Use our tools to generate documentation at the build level, on a server, and on a client.

[HTML Template](#)

[React Component](#)

Code Generation

Generate documentation, Code (TypeScript, Java, C#, etc), and more out of your AsyncAPI files.

[Generator](#)

[Modelina](#)

Community

We're a community of great people who are passionate about AsyncAPI and event-driven architectures.

[Join our Slack](#)

Open Governance

Our Open-Source project is part of Linux Foundation and works under an Open Governance model.

[Read more about Open Governance](#)

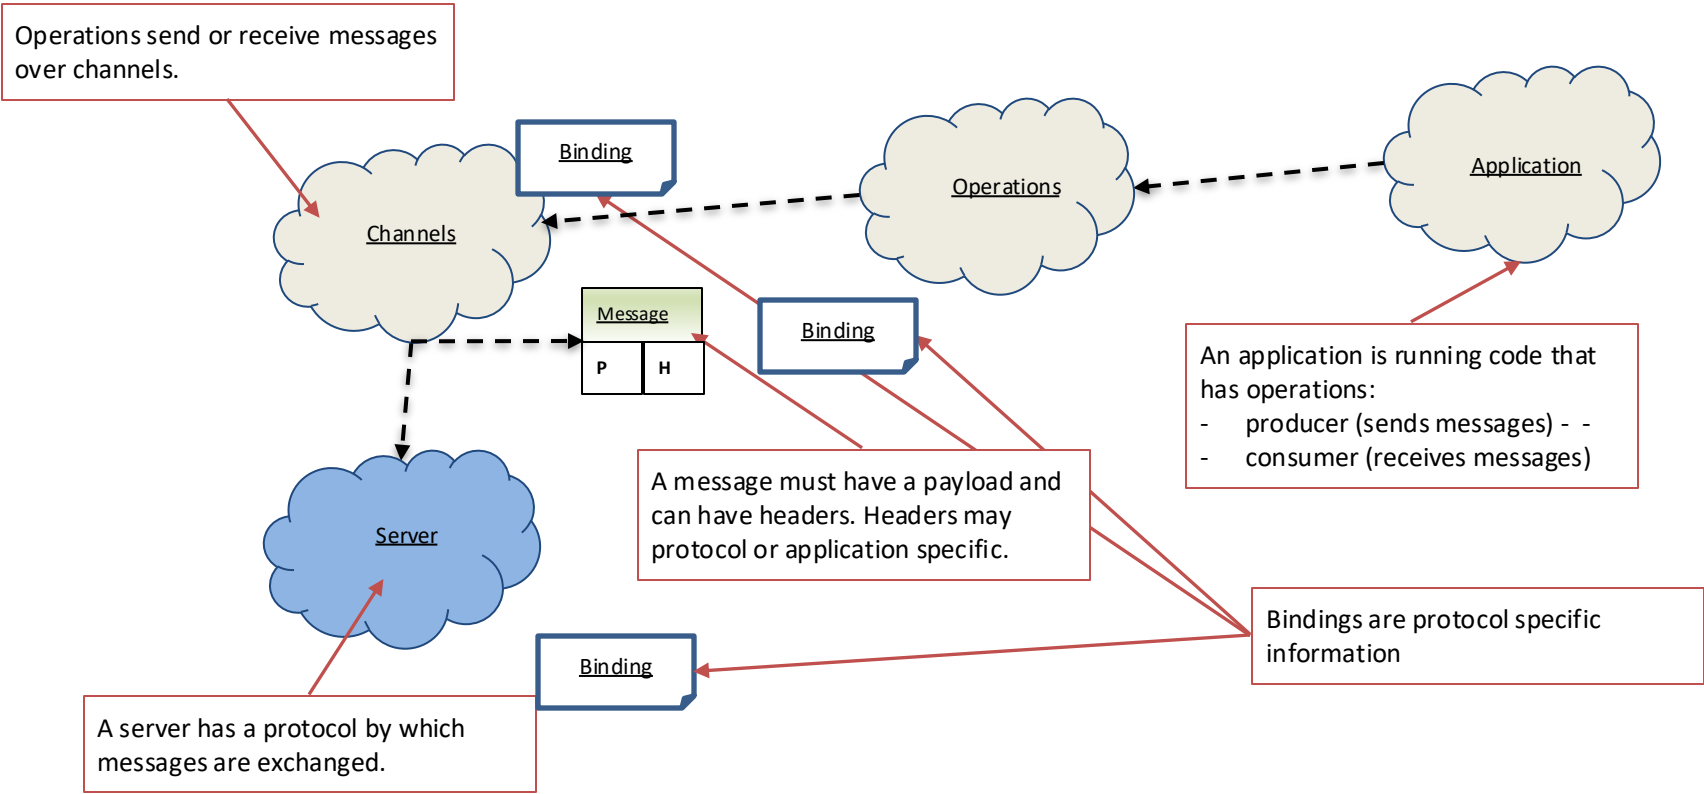
[TSC Members](#)

And much more...

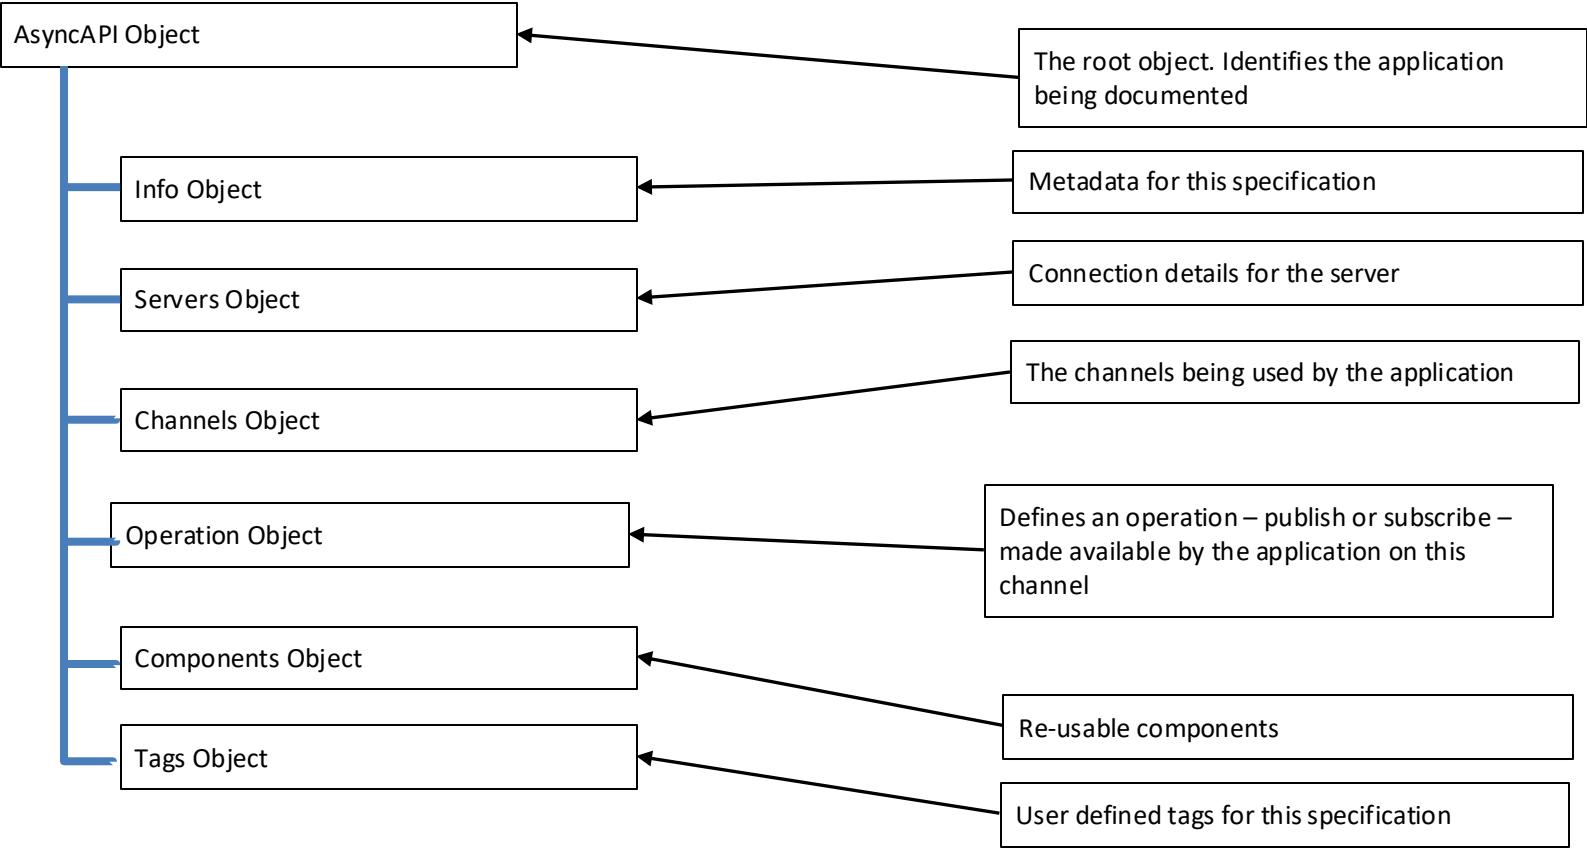
We have many different tools and welcome you to explore our ideas and propose new ideas to AsyncAPI.

[View GitHub Discussions](#)

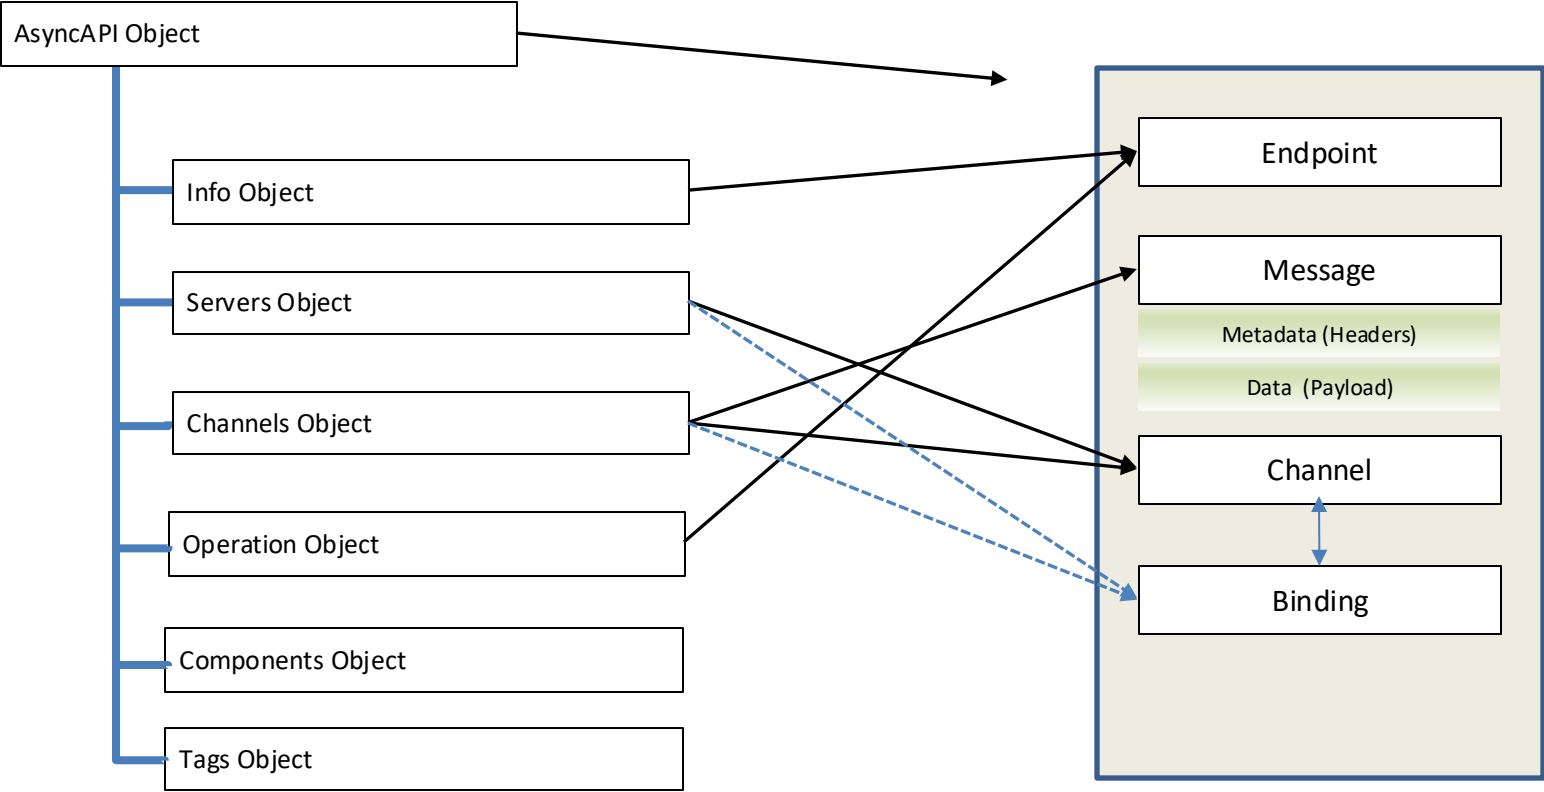
AsyncAPI Elements (V3)



Document Structure (V3)



AsyncAPI (V3) and Endpoint ABCs



id: <https://github.com/brightercommand/greetings/>

We use a specification file for an app, which is a producer or consumer, and identify them by id

info:

contact:

name: Paramore Brighter

url: <https://goparamore.io/support>

email: support@goparamore.io

license:

name: Apache 2.0

url: <https://www.apache.org/licenses/LICENSE-2.0.html>

description: Demonstrates sending a greeting over a messaging transport.

title: Brighter Sample App

version: 1.0.0

tags:

- name: brighter examples

development:

description: A Kafka broker for local development

url: localhost:9092

protocol: kafka

```
greeting:
  address: 'goparamore.io.greeting'
  summary: For sending greetings
  description: This channel contains greeting messages
  servers:
    - $ref: '#/servers/development'
  messages:
    greeting:
      $ref: "#/components/messages/greeting"
  bindings:
    kafka:
      partitions: 20
      replicas: 3
```

sendGreeting :

action: send

summary: sends a greeting

description: The application sends a greeting to a consumer.

channel:

\$ref: "#/channels/greeting"

bindings:

kafka:

partitions: 20

replicas: 3

Components

```
components:
  messages:
    greeting:
      name: greeting
      title: A salutation
      summary: This is how we send you a salutation
      contentType: application/json
      traits:
        - $ref: '#/components/messageTraits/commonHeaders'
      payload:
        $ref: "#/components/schemas/greetingContent"

  schemas:
    greetingContent:
      type: object
      properties:
        greeting:
          type: string
          description: The salutation you want to send
      ...
```

JSON Schema (AsyncAPI Schema Object)

\$schema

\$id

title

description

type

properties

```
{
  "$schema" : "https://json-schema.org/draft/2020-12/schema",
  "$id" : "https://goparamore.io/greeting.schema.json",
  "title" : "greeting",
  "description" : "A greeting message",
  "type" : "object",
  "properties" : {
    "greeting" : {
      "description" : "the salutation"
      "type" : string
    }
  }
}
```


Complex Types:

records

enums

arrays

maps

unions

fixed

Records:

name

namespace

doc

alias

fields

name

doc

type

default

```
{  
  "type": "record",  
  "name": "greeting",  
  "title": "greeting",  
  "fields": [  
    {"name": "greeting", "type": "string"}  
  ]  
}
```

Encodings:

JSON

Binary

Languages:

C

C++

C#

Java

Perl

Python

Ruby

Others...

```
syntax = "proto3";  
  
message Greeting {  
    string greeting = 1;  
}
```

Protobuf

Encodings:

Binary

Languages:

C++

C#

Dart

Go

Kotlin

Java

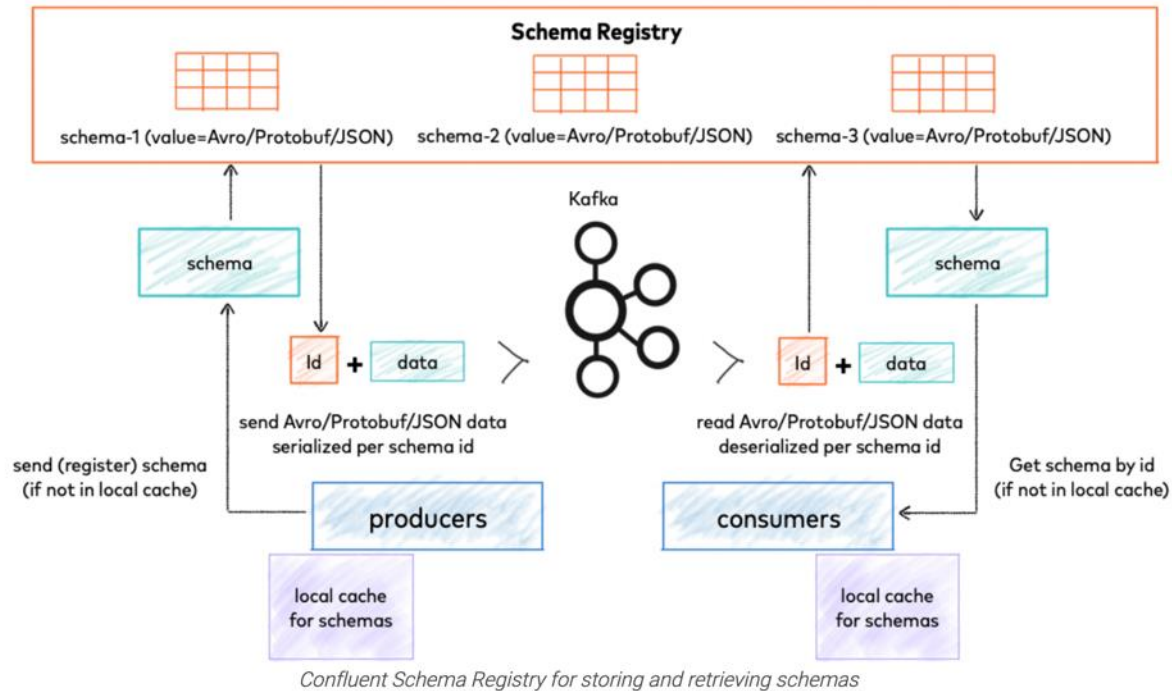
Objective-C

Python

Ruby

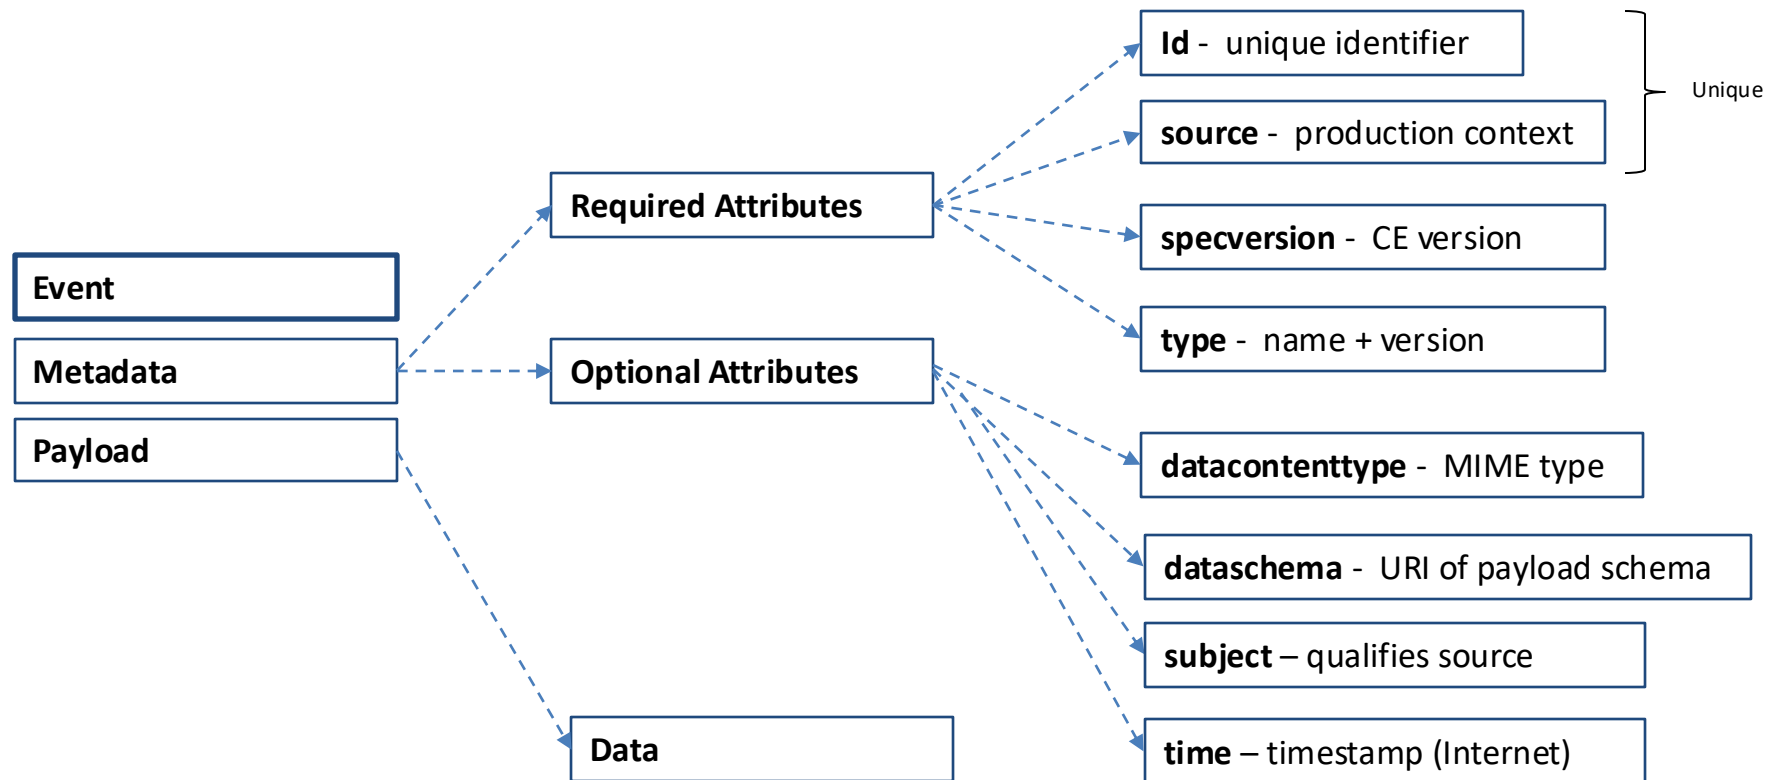
Others...

Schema Registry

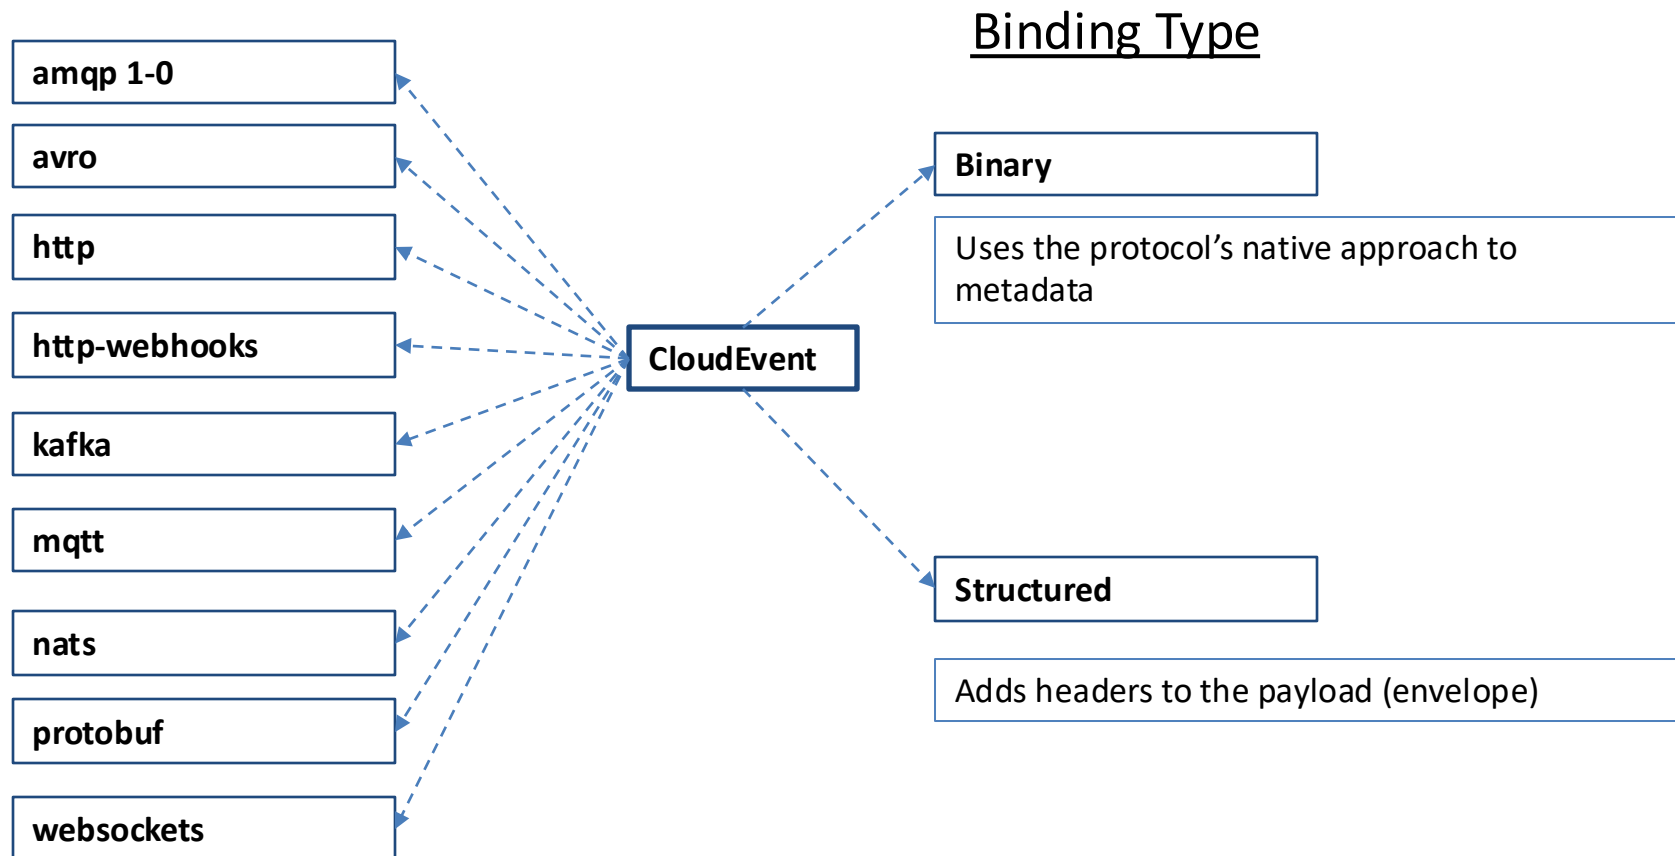


<https://docs.confluent.io/platform/current/schema-registry/index.html>

Cloud Events



Protocol Binding



Protocol Binding

Binary

```
----- Message -----
Topic Name: mytopic
----- key -----
Key: mykey
----- headers -----
ce_specversion: "1.0"
ce_type: "com.example.someevent"
ce_source: "/mycontext/subcontext"
ce_id: "1234-1234-1234"
ce_time: "2018-04-05T03:56:24Z"
content-type: application/avro
----- value -----
... application data encoded in Avro ...
-----
```

Structured

```
----- Message -----
Topic Name: mytopic
----- key -----
Key: mykey
----- headers -----
content-type: application/cloudevents+json; charset=UTF-8
----- value -----
{
  "specversion" : "1.0",
  "type" : "com.example.someevent",
  "source" : "/mycontext/subcontext",
  "id" : "1234-1234-1234",
  "time" : "2018-04-05T03:56:24Z",
  "datacontenttype" : "application/json",
  "data" : {
    ... application data encoded in JSON ...
  }
}
-----
```

AsyncAPI Studio

Information

Servers

Operations

Messages

Schemas

KAFKA localhost

SUB greeting

greeting

greetingContent

```
1  asyncapi: '2.0.0'
2  id: "https://github.com/brightercommand/greetings-sender/"
3  info:
4    contact:
5      name: Paramore Brighter
6      url: https://goparamore.io/support
7      email: support@goparamore.io
8    license:
9      name: Apache 2.0
10     url: https://www.apache.org/licenses/LICENSE-2.0.html
11    description: Demonstrates sending a greeting over a messaging
12     transport.
13    title: Brighter Sample App
14    version: 1.0.0
15    tags:
16      - name: brighter examples
17  servers:
18    localhost:
19      url: localhost:9092
20      protocol: kafka
21  channels:
22    greeting:
23      description: A channel for sending out greeting messages
24      subscribe:
25        description: This service lets you send the 'Hello World'
26        greeting to another service.
27        operationId: sendMessage
28        message:
29          $ref: '#/components/messages/greeting'
30  components:
31    messages:
32      greeting:
33        name: greeting
34        title: A salutation
35        summary: This is how we send you a salutation
36        contentType: application/json
37        traits:
38          - $ref: '#/components/messageTraits/commonHeaders'
39        payload:
40          $ref: "#/components/schemas/greetingContent"
41  schemas:
42    greetingContent:
43      type: string
```

Brighter Sample App 1.0.0

APACHE 2.0 PARAMORE BRIGHTER SUPPORT@GOPARAMORE.IO

ID: [HTTPS://GITHUB.COM/BRIGHTERCOMMAND/GREETINGS-SENDER/](https://github.com/brightercommand/greetings-sender/)

Demonstrates sending a greeting over a messaging transport.

#brighter examples

Servers

localhost:9092 KAFKA LOCALHOST

Security:

SECURITY.PROTOCOL: PLAINTEXT

Operations

SUB greeting

A channel for sending out greeting messages

This service lets you send the 'Hello World' greeting to another service.

Operation ID: sendMessage

Accepts the following message:

A salutation greeting

VS Code

The image shows a VS Code editor with a YAML file named `brighter-greetings-sender.yml` open. The file is a Brighter service definition. The right sidebar shows the rendered preview of the service.

```
! brighter-greetings-sender.yml > YAML > {} info > {} contact > url
1  asyncapi: '2.0.0'
2  id: "https://github.com/brightercommand/greetings-sender/"
3  info:
4    contact:
5      name: Paramore Brighter
6      url: https://goparamore.io/support
7      email: support@goparamore.io
8    license:
9      name: Apache 2.0
10     url: https://www.apache.org/licenses/LICENSE-2.0.html
11    description: Demonstrates sending a greeting over a messaging transport
12    title: Brighter Sample App
13    version: 1.0.0
14    tags:
15      - name: brighter examples
16
17  servers:
18    localhost:
19      url: localhost:9092
20      protocol: kafka
21
22  channels:
23    greeting:
24      description: A channel for sending out greeting messages
25      subscribe:
26        description: This service lets you send the 'Hello World' greeting
27        operationId: sendMessage
28        message:
29          $ref: '#/components/messages/greeting'
30
31  components:
32    messages:
33      greeting:
34        name: greeting
35        title: A salutation
36        summary: This is how we send you a salutation
37        contentType: application/json
38        traits:
39          - $ref: '#/components/messageTraits/commonHeaders'
40        payload:
41          $ref: '#/components/schemas/greetingContent'
42
43  schemas:
44    greetingContent:
```

Brighter Sample App 1.0.0

APACHE 2.0

Demonstrates sending a greeting over a messaging transport.

Contact link: [PARAMORE BRIGHTER](#) Contact email: [SUPPORT@GOPAMORE.IO](#)

Servers

localhost:9092 **KAFKA**

Operations

SUB greeting

A channel for sending out greeting messages

This service lets you send the 'Hello World' greeting to another service.

Accepts the following message:

A salutation **greeting**

This is how we send you a salutation

[Overview](#)[Getting Started](#)[Local Development](#)[Core Features](#)[Software Catalog](#)[Overview](#)[The Life of an Entity](#)[Catalog Configuration](#)[System Model](#)[YAML File Format](#)[Entity References](#)[Well-known Annotations](#)[Well-known Relations](#)

`spec.type` [required]

The type of the API definition as a string, e.g. `openapi`. This field is required.

The software catalog accepts any type value, but an organization should take great care to establish a proper taxonomy for these. Tools including Backstage itself may read this field and behave differently depending on its value. For example, an OpenAPI type API may be displayed using an OpenAPI viewer tooling in the Backstage interface.

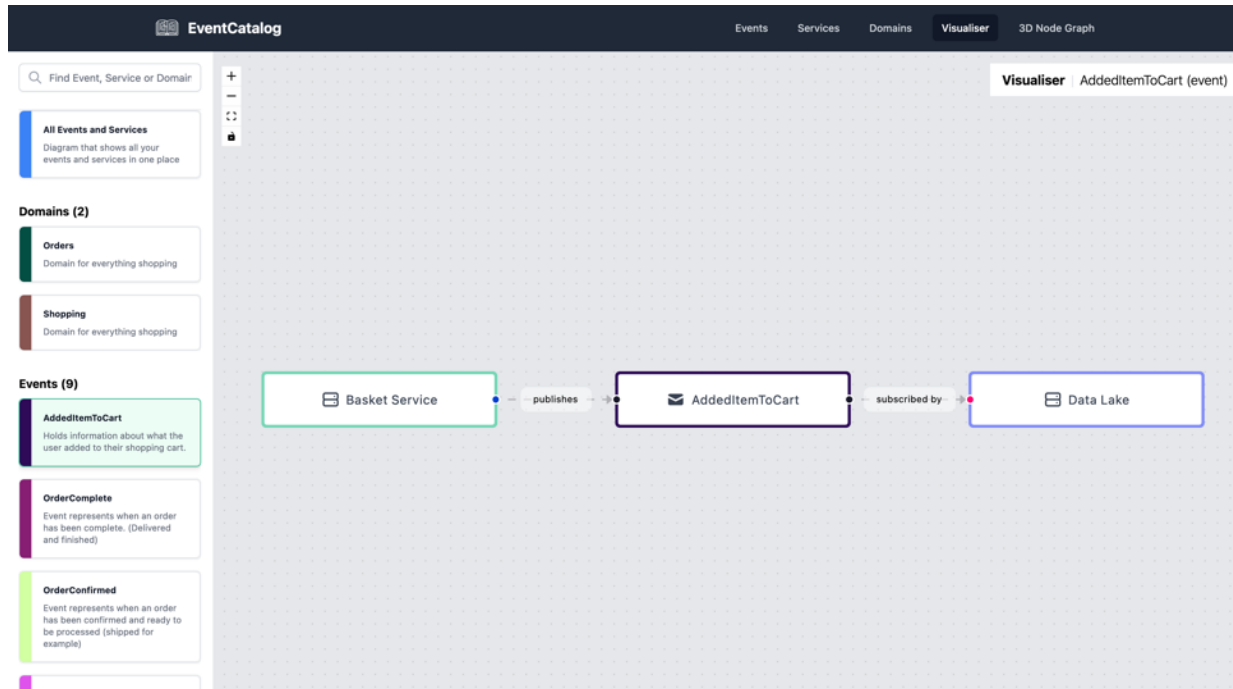
The current set of well-known and common values for this field is:

- `openapi` - An API definition in YAML or JSON format based on the [OpenAPI](#) version 2 or version 3 spec.
- `svncapi` - An API definition based on the [AsyncAPI](#) spec.
- `graphql` - An API definition based on [GraphQL schemas](#) for consuming [GraphQL](#) based APIs.
- `grpc` - An API definition based on [Protocol Buffers](#) to use with [gRPC](#).

Contents

[Overall Shape Of An Entity](#)[Substitutions In The Descriptor Format](#)[Common to All Kinds: The Envelope](#)`apiVersion` and `kind` [required]`metadata` [required]`spec` [varies][Common to All Kinds: The Metadata](#)`name` [required]`namespace` [optional]`title` [optional]`description` [optional]`labels` [optional]`annotations` [optional]`tags` [optional]`links` [optional]

Event Catalog



<https://github.com/boyney123/eventcatalog>

Observability



High-quality, ubiquitous, and portable telemetry to enable effective observability

[Learn more](#)

[Mission and vision](#)

Get started!

[Collector](#)

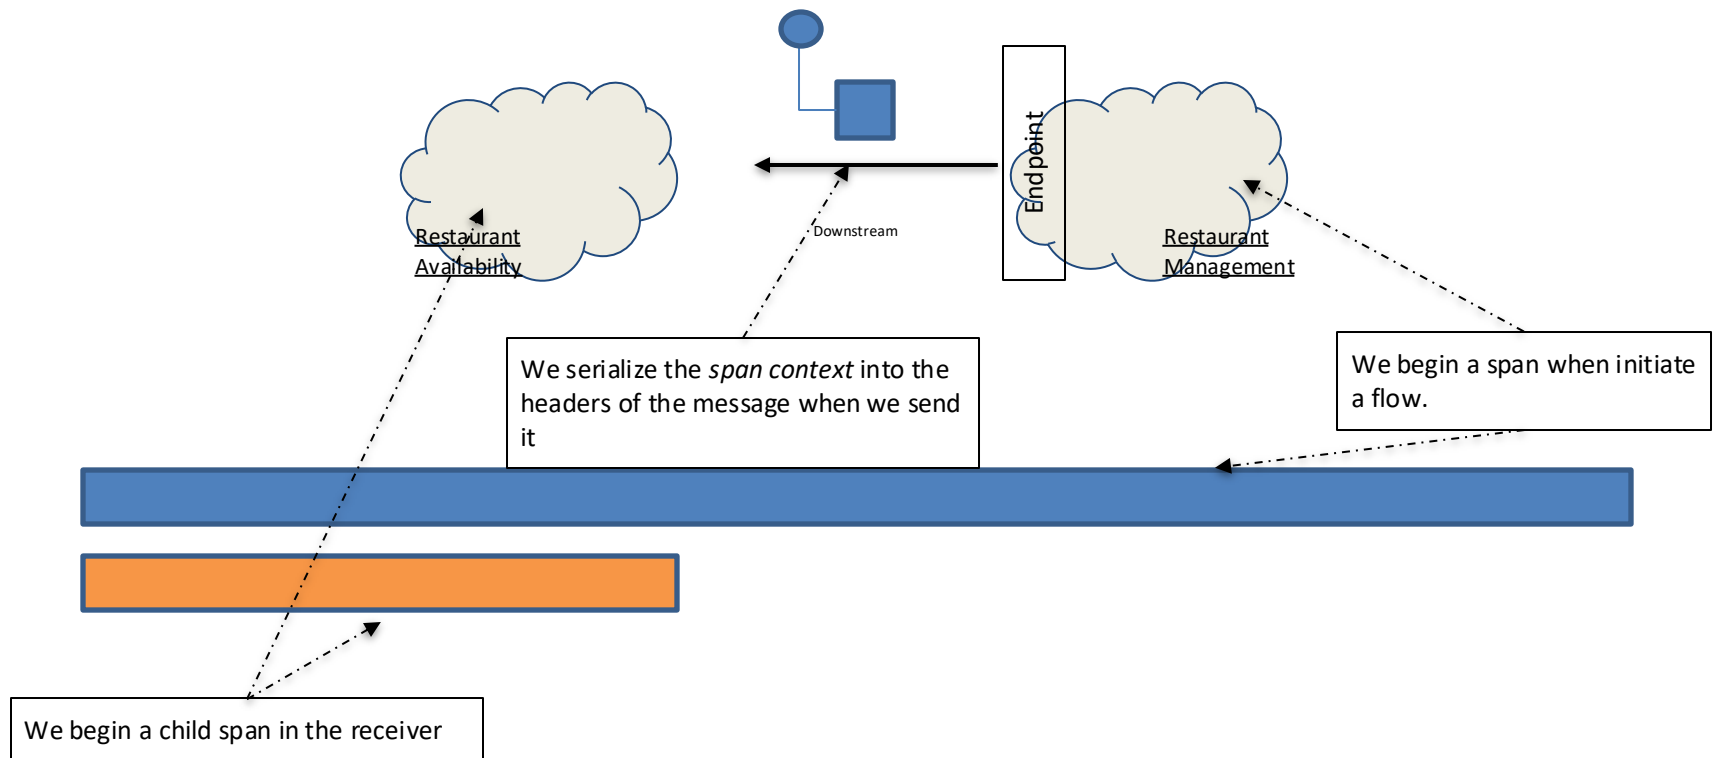
[Go](#)

[.NET](#)

[JavaScript](#)

[...](#)

OpenTelemetry Tracing



open-telemetry / semantic-conventions

<> Code Issues 275 Pull requests 42 Actions Projects 5 Security

main semantic-conventions / docs / messaging / messaging-spans.md

crossoverJie and joaopgrassi messaging.system support pulsar (#1099)

Preview Code Blame 609 lines (459 loc) · 35.7 KB

Semantic Conventions for Messaging Spans

Status: [Experimental](#)

- Definitions
 - [Message](#)
 - [Producer](#)
 - [Consumer](#)
 - [Intermediary](#)
 - [Destinations](#)
 - [Message consumption](#)
 - [Conversations](#)
 - [Temporary and anonymous destinations](#)
- Conventions
 - [Context propagation](#)
 - [Span name](#)
 - [Operation types](#)
 - [Span kind](#)
 - [Trace structure](#)
 - [Producer spans](#)
 - [Consumer spans](#)
- Messaging attributes
 - [Consumer attributes](#)
 - [Per-message attributes](#)
 - [Attributes specific to certain messaging systems](#)

name SHOULD only be used for the span name if it is known to be of low cardinality (cf. [generated](#) if it is statically derived from application code or configuration. Wherever possible, the or aliased names SHOULD be used. If the destination name is dynamic, such as a [conversation identifier](#), it SHOULD NOT be used for the span name. In these cases, an artificial destination name generic, static fallback like "(anonymous)" for [anonymous destinations](#) SHOULD be used

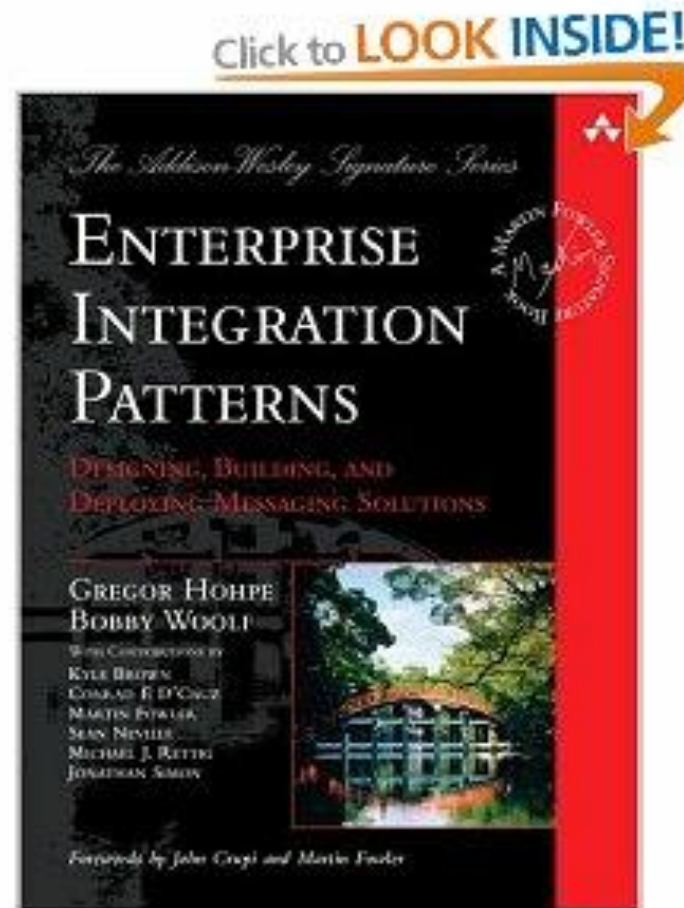
```
s publish
s subscribe
s settle
  publish
  nack
spaces process
tionRequest-Conversations settle
) send ( (anonymous) being a stable identifier for an unnamed destination)

m specific adaption to span naming MUST be documented in semantic conventions for spans
```

Operation types related to messages are defined for these semantic conventions:

Description
A message is created or passed to a client library for publishing. "Create" spans always include a "Create" span to provide a unique creation context for messages in batch publishing scenarios.
One or more messages are provided for publishing to an intermediary. If a single message is provided, the "Publish" span can be used as the creation context and no "Create" span needs to be added.
One or more messages are requested by a consumer. This operation refers to pull-based consumption where the consumer explicitly call methods of messaging SDKs to receive messages.
One or more messages are delivered to or processed by a consumer.
One or more messages are settled.

Further Reading



Q&A